

# The Fibonacci sequence

Eric Martin, CSE, UNSW

COMP9021 Principles of Programming

```
[1]: from functools import lru_cache
     from itertools import count, islice
     from math import sqrt
```

The Fibonacci sequence, say  $(F_n)_{n \in \mathbb{N}}$ , is defined as  $F_0 = 0$ ,  $F_1 = 1$  and for all  $n > 1$ ,  $F_n = F_{n-2} + F_{n-1}$ ; so it is 0, 1, 1, 2, 3, 5, 8, 13, 21, 34...

A generator function is the best option to generate the initial segment of the Fibonacci sequence of a given length:

```
[2]: def fibonacci_sequence():
     yield 0
     yield 1
     previous, current = 0, 1
     while True:
         previous, current = current, previous + current
         yield current
```

```
[3]: S = fibonacci_sequence()
     list(next(S) for _ in range(19))
```

```
[3]: [0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377, 610, 987, 1597, 2584]
```

It can also be used to generate the member of the Fibonacci sequence of a given index, which is facilitated by the `islice()` function from the `itertools` module; that function provides a counterpart to the slices of built sequences:

```
[4]: list(islice(count(), 10))
     list(islice(count(), 2, 10))
     list(islice(count(), 2, 10, 3))
```

```
[4]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
[4]: [2, 3, 4, 5, 6, 7, 8, 9]
```

```
[4]: [2, 5, 8]
```

$F_{18}$  can therefore be returned as follows:

```
[5]: next(islice(fibonacci_sequence(), 18, 19))
```

[5]: 2584

In case only one or a few specific members of the Fibonacci sequence are needed, a simple function is more appropriate:

```
[6]: def iterative_fibonacci(n):  
    if n < 2:  
        return n  
    previous, current = 0, 1  
    for _ in range(2, n + 1):  
        previous, current = current, previous + current  
    return current  
  
iterative_fibonacci(18)
```

[6]: 2584

A naive recursive implementation is elegant, but too inefficient, as we will see:

```
[7]: def recursive_fibonacci(n):  
    if n < 2:  
        return n  
    return recursive_fibonacci(n - 2) + recursive_fibonacci(n - 1)  
  
recursive_fibonacci(18)
```

[7]: 2584

Let an integer  $n$  greater than 1 be given. Then a call to `recursive_fibonacci(n)` involves:

- for all nonzero  $k \leq n$ ,  $F_{n-k+1}$  calls to `recursive_fibonacci(k)`;
- $F_{n-1}$  calls to `recursive_fibonacci(0)`.

In particular, `recursive_fibonacci(n)` calls `recursive_fibonacci(1)`  $F_n$  times. Proof is by induction on  $k \leq n$ :

- `recursive_fibonacci(n)` is called once indeed.
- `recursive_fibonacci(n)` directly calls `recursive_fibonacci(n - 1)` and does not call it indirectly, so calls it once indeed.
- For all  $k < n$  with  $k > 1$ , `recursive_fibonacci(n - k)` is directly called by `recursive_fibonacci(n - k + 1)` and by `recursive_fibonacci(n - k + 2)`. By inductive hypothesis, the latter two are called directly or indirectly by `recursive_fibonacci(n)`  $F_k$  and  $F_{k-1}$  times, respectively. Hence `recursive_fibonacci(n - k)` is called by `recursive_fibonacci(n)`  $F_{k+1}$  times.
- `recursive_fibonacci(0)` is directly called by `recursive_fibonacci(2)`, hence it is called by `recursive_fibonacci(n)`  $F_{n-1}$  times.

Let us illustrate this for  $n = 6$  with the following tracing function:

```
[8]: def trace_recursive_fibonacci(n, depth):  
    print('      ' * depth, 'Start of function call for n =', n)
```

```

if n >= 2:
    second_previous = trace_recursive_fibonacci(n - 2, depth + 1)
    previous = trace_recursive_fibonacci(n - 1, depth + 1)
    print('      ' * depth, f'End of function call for n = {n}, returning',
          second_previous + previous
          )
    return second_previous + previous
print('      ' * depth, f'End of function call for n = {n}, returning', n)
return n

```

```
trace_recursive_fibonacci(6, 0)
```

```

Start of function call for n = 6
  Start of function call for n = 4
    Start of function call for n = 2
      Start of function call for n = 0
        End of function call for n = 0, returning 0
      Start of function call for n = 1
        End of function call for n = 1, returning 1
      End of function call for n = 2, returning 1
    End of function call for n = 4, returning 3
  Start of function call for n = 3
    Start of function call for n = 1
      End of function call for n = 1, returning 1
    Start of function call for n = 2
      Start of function call for n = 0
        End of function call for n = 0, returning 0
      Start of function call for n = 1
        End of function call for n = 1, returning 1
      End of function call for n = 2, returning 1
    End of function call for n = 3, returning 2
  End of function call for n = 4, returning 3
Start of function call for n = 5
  Start of function call for n = 3
    Start of function call for n = 1
      End of function call for n = 1, returning 1
    Start of function call for n = 2
      Start of function call for n = 0
        End of function call for n = 0, returning 0
      Start of function call for n = 1
        End of function call for n = 1, returning 1
      End of function call for n = 2, returning 1
    End of function call for n = 3, returning 2
  End of function call for n = 5, returning 5
End of function call for n = 6, returning 8

```

```

End of function call for n = 2, returning 1
Start of function call for n = 3
  Start of function call for n = 1
  End of function call for n = 1, returning 1
  Start of function call for n = 2
  Start of function call for n = 0
  End of function call for n = 0, returning 0
  Start of function call for n = 1
  End of function call for n = 1, returning 1
  End of function call for n = 2, returning 1
End of function call for n = 3, returning 2
End of function call for n = 4, returning 3
End of function call for n = 5, returning 5
End of function call for n = 6, returning 8

```

[8]: 8

We can still save the recursive design by saving terms of the Fibonacci sequence as they get computed for the first time. As a result of processing the `def` statement below, a dictionary, `fibonacci`, is created and initialised with the values of the first two members of the Fibonacci sequence. Then the function `memoise_fibonacci()` is called, directly as `memoise_fibonacci(18)`, and indirectly as `memoise_fibonacci(18)` executes. For each of those calls, `memoise_fibonacci()` is given one argument only, so the second argument is set to its default value, namely, `fibonacci`, extended with a new key and associated value in case the condition of the `if` statement in the body of `memoise_fibonacci()` evaluates to `True`:

```

[9]: def memoise_fibonacci(n, fibonacci={0: 0, 1: 1}):
      if n not in fibonacci:
          fibonacci[n] = memoise_fibonacci(n - 2) + memoise_fibonacci(n - 1)
      return fibonacci[n]

memoise_fibonacci(18)

```

[9]: 2584

Let us illustrate the mechanism for  $n = 6$  with the following tracing function:

```

[10]: def trace_memoise_fibonacci(n, depth, fibonacci={0: 0, 1: 1}):
      print('    ' * depth, 'Start of function call for n =', n)
      print('    ' * (depth + 1), f'fibonacci now is {fibonacci}; ', end = '')
      if n not in fibonacci:
          print('compute value')
          fibonacci[n] = trace_memoise_fibonacci(n - 2, depth + 1)\
                          + trace_memoise_fibonacci(n - 1, depth + 1)
      else:
          print('retrieve value')
      print('    ' * depth, f'End of function call for n = {n}, returning',
            fibonacci[n])

```

```
return fibonacci[n]
```

```
trace_memoise_fibonacci(6, 0)
```

```
Start of function call for n = 6
  fibonacci now is {0: 0, 1: 1}; compute value
  Start of function call for n = 4
    fibonacci now is {0: 0, 1: 1}; compute value
    Start of function call for n = 2
      fibonacci now is {0: 0, 1: 1}; compute value
      Start of function call for n = 0
        fibonacci now is {0: 0, 1: 1}; retrieve value
        End of function call for n = 0, returning 0
      Start of function call for n = 1
        fibonacci now is {0: 0, 1: 1}; retrieve value
        End of function call for n = 1, returning 1
      End of function call for n = 2, returning 1
    End of function call for n = 4, returning 1
  Start of function call for n = 3
    fibonacci now is {0: 0, 1: 1, 2: 1}; compute value
    Start of function call for n = 1
      fibonacci now is {0: 0, 1: 1, 2: 1}; retrieve value
      End of function call for n = 1, returning 1
    Start of function call for n = 2
      fibonacci now is {0: 0, 1: 1, 2: 1}; retrieve value
      End of function call for n = 2, returning 1
    End of function call for n = 3, returning 2
  End of function call for n = 4, returning 3
Start of function call for n = 5
  fibonacci now is {0: 0, 1: 1, 2: 1, 3: 2, 4: 3}; compute value
  Start of function call for n = 3
    fibonacci now is {0: 0, 1: 1, 2: 1, 3: 2, 4: 3}; retrieve value
    End of function call for n = 3, returning 2
  Start of function call for n = 4
    fibonacci now is {0: 0, 1: 1, 2: 1, 3: 2, 4: 3}; retrieve value
    End of function call for n = 4, returning 3
  End of function call for n = 5, returning 5
End of function call for n = 6, returning 8
```

[10]: 8

`memoise_fibonacci()` illustrates the fact that when a function argument has a default value, that default value is not created at every function call, but when Python processes the function's `def` statement. This makes no difference for default values of a type such as `int`:

```
[11]: def f(x=0):
      x += 1
      return x
```

```

# Create the argument 0 before calling f(); let x denote it.
# From the value denoted by x and 1, create 1; let x denote it.
f(0)
f(1)
f(2)
# Let x denote the 0 created when def was processed.
# From the value denoted by x and 1, create 1; let x denote it.
f()
f()
f()

```

[11]: 1

[11]: 2

[11]: 3

[11]: 1

[11]: 1

[11]: 1

But it makes a difference for default values of a type such as `list`:

```

[12]: def g(x=[0]):
        x += [1]
        return x

# Create the argument [0] before calling g(); let x denote it.
# Then extend it to [0, 1]; let x denote the modified list.
g([0])
g([1])
g([2])
# Let x denote the list L created when def was processed, then and now
# equal to [0]. Then extend it to [0, 1]; let x denote the modified L.
g()
# Let x denote the list L created when def was processed, now equal to
# [0, 1]. Then extend it to [0, 1, 1], let x denote the modified L.
g()
g()

```

[12]: [0, 1]

[12]: [1, 1]

[12]: [2, 1]

[12]: [0, 1]

```
[12]: [0, 1, 1]
```

```
[12]: [0, 1, 1, 1]
```

What was good for `memoise_fibonacci()` might not be the intended behaviour for other functions, in other contexts: in case a function  $F$  is called without an argument for a parameter  $p$  that in  $F$ 's definition, receives a default value  $v$ , one might want  $p$  to always be assigned that default value, not the value currently denoted by  $p$  and possibly modified from the original value of  $v$  following previous calls to  $F$ . One should then opt for a different design:

```
[13]: def h(x=None):
      if x is None:
          x = [0]
      x += [1]
      return x

      # Create the argument [0] before calling h(); let x denote it.
      # Then extend it to [0, 1]; let x denote the modified list.
      h([0])
      h([1])
      h([2])
      # Let x denote None, then create [0]; let x denote it.
      # Then extend it to [0, 1]; let x denote the modified list.
      h()
      h()
      h()
```

```
[13]: [0, 1]
```

```
[13]: [1, 1]
```

```
[13]: [2, 1]
```

```
[13]: [0, 1]
```

```
[13]: [0, 1]
```

```
[13]: [0, 1]
```

The `lru_cache()` (“lru” is for *Least Recently Used*) function from the `functools` module returns a function that can be used as a **decorator** and applied to a function  $F$  to yield a memoised version of  $F$ . By default, the `maxsize` argument of `lru_cache()` is set to 128, to record up to the last 128 used values of the function, as witnessed by the `cache_info()` attribute of the memoised version of  $f$ :

```
[14]: @lru_cache()
      def lru_fibonacci(n):
          if n < 2:
              return n
```

```
    return lru_fibonacci(n - 2) + lru_fibonacci(n - 1)

lru_fibonacci.cache_info()
```

[14]: CacheInfo(hits=0, misses=0, maxsize=128, currsize=0)

Suppose that `lru_fibonacci()` is called for the first time with 2 as argument. Since `lru_fibonacci(2)` has not been computed yet, `lru_fibonacci(0)` and `lru_fibonacci(1)` are called, which both have not been computed yet either: a total of 3 values fail to be retrieved (3 misses). The last two values are computed and recorded, then the former value is computed and recorded, and the cache eventually stores those 3 values:

```
[15]: lru_fibonacci(2)
lru_fibonacci.cache_info()
```

[15]: 1

[15]: CacheInfo(hits=0, misses=3, maxsize=128, currsize=3)

Calling `lru_fibonacci(2)` again, the value is found in the cache (1 hit):

```
[16]: lru_fibonacci(2)
lru_fibonacci.cache_info()
```

[16]: 1

[16]: CacheInfo(hits=1, misses=3, maxsize=128, currsize=3)

When calling `lru_fibonacci(3)`, the value fails to be found in the cache (1 more miss), so `lru_fibonacci(1)` and `lru_fibonacci(2)` are called and retrieved from the cache (2 more hits), and the computed value of `lru_fibonacci(3)` is added to the cache.

```
[17]: lru_fibonacci(3)
lru_fibonacci.cache_info()
```

[17]: 2

[17]: CacheInfo(hits=3, misses=4, maxsize=128, currsize=4)

The cache can be cleared with the `cache_clear()` attribute of the memoised version of the function. Then calling `lru_fibonacci(3)` necessitates to call `lru_fibonacci(1)` and `lru_fibonacci(2)`, calling `lru_fibonacci(2)` necessitates to call `lru_fibonacci(0)` and `lru_fibonacci(1)`, for a total of 4 misses that are computed and all stored in the cache. The hit is for `lru_fibonacci(1)`, retrieved from the cache when computing `lru_fibonacci(2)`, assuming `lru_fibonacci(3)` first called `lru_fibonacci(1)`, or retrieved from the cache when computing `lru_fibonacci(1)`, assuming `lru_fibonacci(3)` first called `lru_fibonacci(2)`.

```
[18]: lru_fibonacci.cache_clear()
lru_fibonacci(3)
lru_fibonacci.cache_info()
```



[18]: 2

[18]: CacheInfo(hits=1, misses=4, maxsize=128, currsize=4)

Clearing the cache again, calling `lru_fibonacci(128)` necessitates to eventually call `lru_fibonacci(0)` and `lru_fibonacci(1)`, in that order or the other way around. Either way, the last time `lru_fibonacci(0)` is used is when `lru_fibonacci(2)` is computed, while `lru_fibonacci(1)`, `lru_fibonacci(3)`, `lru_fibonacci(4)`... `lru_fibonacci(128)` will all be computed or used afterwards. This explains the output obtained by executing the following cell.

```
[19]: lru_fibonacci.cache_clear()
lru_fibonacci(128)
lru_fibonacci.cache_info()
lru_fibonacci(1)
lru_fibonacci.cache_info()
lru_fibonacci(0)
lru_fibonacci.cache_info()
```

[19]: 251728825683549488150424261

[19]: CacheInfo(hits=126, misses=129, maxsize=128, currsize=128)

[19]: 1

[19]: CacheInfo(hits=127, misses=129, maxsize=128, currsize=128)

[19]: 0

[19]: CacheInfo(hits=127, misses=130, maxsize=128, currsize=128)

The capacity of the cache can be left unbounded by setting the value of the `maxsize` argument of `lru_cache()` to `None`:

```
[20]: @lru_cache(None)
def unbounded_lru_fibonacci(n):
    if n < 2:
        return n
    return unbounded_lru_fibonacci(n - 1) + unbounded_lru_fibonacci(n - 2)
```

```
[21]: unbounded_lru_fibonacci(150)
unbounded_lru_fibonacci.cache_info()
```

[21]: 9969216677189303386214405760200

[21]: CacheInfo(hits=148, misses=151, maxsize=None, currsize=151)

The argument `maxsize` of `lru_cache()` can also be set to any integer value. Let us set it to 4 and first call `bounded_lru_fibonacci(8)`. Then `bounded_lru_fibonacci(8)`, `bounded_lru_fibonacci(7)`, `bounded_lru_fibonacci(6)` and `bounded_lru_fibonacci(5)` are last called and recorded. If `bounded_lru_fibonacci(5)` is then called, its value

is retrieved (1 more hit). And if `bounded_lru_fibonacci(4)` is thereafter called, `bounded_lru_fibonacci(4)`, ..., `bounded_lru_fibonacci(0)` have to be recomputed (5 more misses), with `bounded_lru_fibonacci(1)` and `bounded_lru_fibonacci(2)` being retrieved in the process (2 more hits):

```
[22]: @lru_cache(4)
def bounded_lru_fibonacci(n):
    if n < 2:
        return n
    return bounded_lru_fibonacci(n - 1) + bounded_lru_fibonacci(n - 2)
```

```
[23]: bounded_lru_fibonacci(8)
bounded_lru_fibonacci.cache_info()
bounded_lru_fibonacci(5)
bounded_lru_fibonacci.cache_info()
bounded_lru_fibonacci(4)
bounded_lru_fibonacci.cache_info()
```

[23]: 21

[23]: CacheInfo(hits=6, misses=9, maxsize=4, cursize=4)

[23]: 5

[23]: CacheInfo(hits=7, misses=9, maxsize=4, cursize=4)

[23]: 3

[23]: CacheInfo(hits=9, misses=14, maxsize=4, cursize=4)

Set  $\varphi = \frac{1+\sqrt{5}}{2}$  and  $\psi = \frac{1-\sqrt{5}}{2}$ . For all  $n \in \mathbb{N}$ ,  $(\frac{1+\sqrt{5}}{2})^{n+2} = (\frac{1+\sqrt{5}}{2})^n \frac{1+2\sqrt{5}+5}{4} = (\frac{1+\sqrt{5}}{2})^n + (\frac{1+\sqrt{5}}{2})^n \frac{1+\sqrt{5}}{2}$ , hence for  $x = \varphi$  or  $x = \psi$ ,  $x^{n+2} = x^n + x^{n+1}$ . Let real numbers  $a$  and  $b$  be given, and for all  $n \in \mathbb{N}$ , let  $s_n$  denote  $a\varphi^n + b\psi^n$ . It follows from the previous equalities (one for  $x = \varphi$ , one for  $x = \psi$ ) that for all  $n \in \mathbb{N}$ ,  $s_{n+2} = s_n + s_{n+1}$ . So  $(s_n)_{n \in \mathbb{N}} = (F_n)_{n \in \mathbb{N}}$  iff  $s_0 = 0$  and  $s_1 = 1$ , which is equivalent to the two equalities  $a + b = 0$  and  $a\varphi + b\psi = 1$ , which is equivalent to  $a = \frac{1}{\varphi - \psi}$  and  $b = -a$ , so  $a = \frac{1}{\sqrt{5}}$  and  $b = -\frac{1}{\sqrt{5}}$ . Hence for all  $n \in \mathbb{N}$ ,

$$F_n = \frac{1}{\sqrt{5}} \left[ \left( \frac{1+\sqrt{5}}{2} \right)^n - \left( \frac{1-\sqrt{5}}{2} \right)^n \right]$$

This is the *closed-form expression* of the Fibonacci numbers. Note that  $\left| \frac{1-\sqrt{5}}{2} \right| < \frac{1}{2}$ , hence  $F_n$  can be computed as  $\frac{1}{\sqrt{5}} \left( \frac{1+\sqrt{5}}{2} \right)^n$  rounded to the closest integer, resulting in a simpler calculation:

```
[24]: def closed_form_fibonacci(n):
    sqrt_5 = sqrt(5)
    return round(1 / sqrt_5 * ((1 + sqrt_5) / 2) ** n)
```

But of course, due to the limited precision of floating point computation, it does not need a large input for `closed_form_fibonacci()` to fail and produce the correct result:

```
[25]: for n, correct in enumerate(fibonacci_sequence()):
      maybe_incorrect = closed_form_fibonacci(n)
      if maybe_incorrect != correct:
          print(f'{n}th term of the sequence is {correct}, '
                f'incorrectly computed as {maybe_incorrect}.')
      )
      break
```

71th term of the sequence is 308061521170129, incorrectly computed as 308061521170130.