

# Outline

---

- Introduction
- Background
- Distributed Database Design
- Database Integration
- Semantic Data Control
- Distributed Query Processing
  - ➔ Overview
  - ➔ Query decomposition and localization
  - ➔ Distributed query optimization
- Multidatabase Query Processing
- Distributed Transaction Management
- Data Replication
- Parallel Database Systems
- Distributed Object DBMS
- Peer-to-Peer Data Management
- Web Data Management
- Current Issues

# Step 3 – Global Query Optimization

---

**Input:** Fragment query

- Find the *best* (not necessarily optimal) global schedule
  - ➔ Minimize a cost function
  - ➔ Distributed join processing
    - ◆ Bushy vs. linear trees
    - ◆ Which relation to ship where?
    - ◆ Ship-whole vs ship-as-needed
  - ➔ Decide on the use of semijoins
    - ◆ Semijoin saves on communication at the expense of more local processing.
  - ➔ Join methods
    - ◆ nested loop vs ordered joins (merge join or hash join)

# Cost-Based Optimization

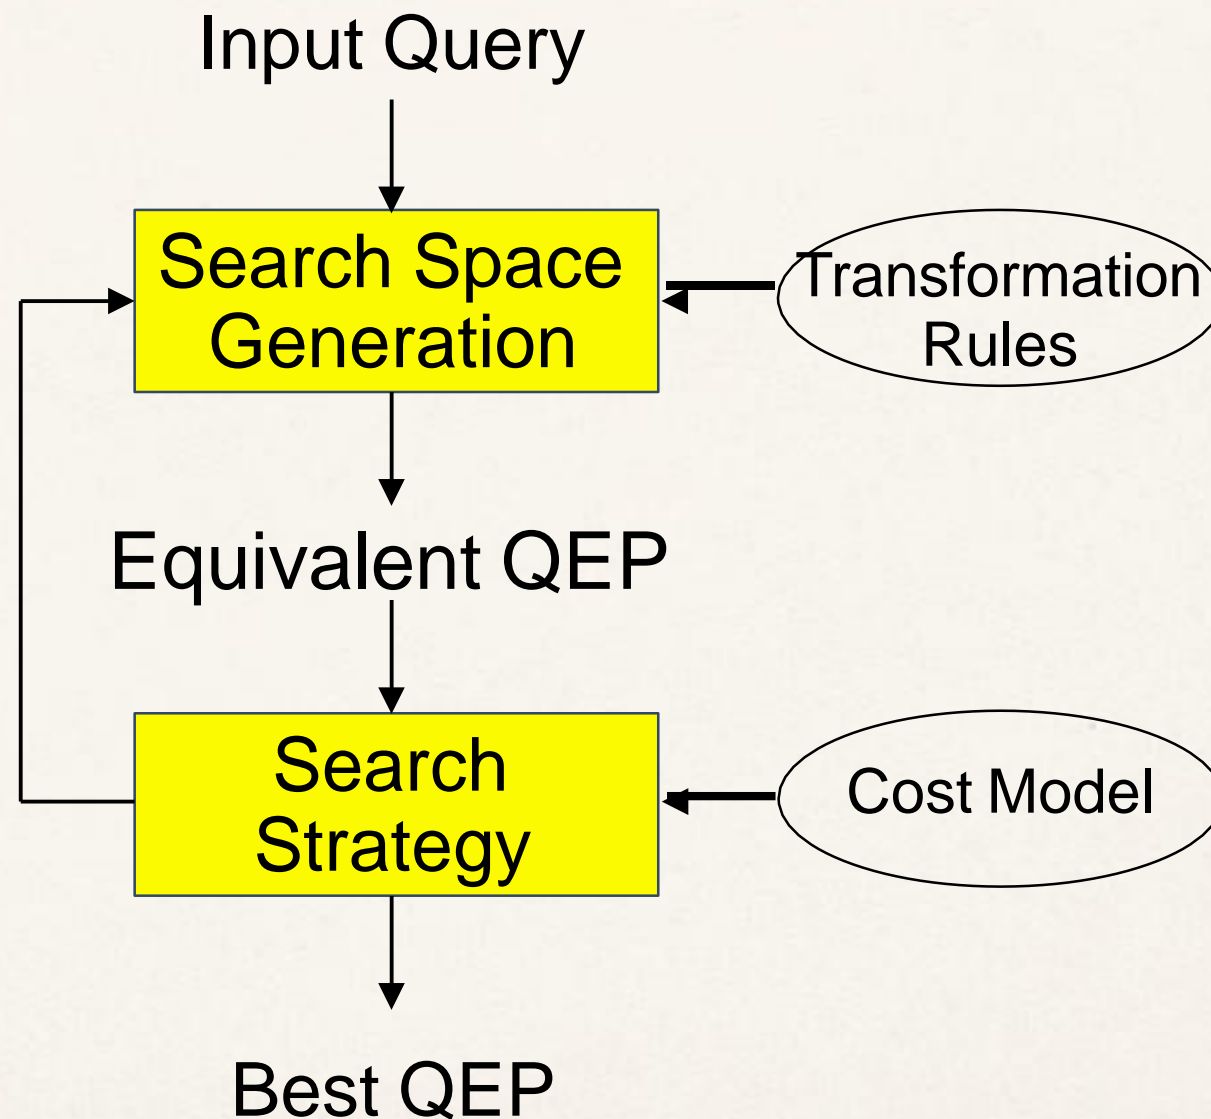
---

- Solution space
  - ➔ The set of equivalent algebra expressions (query trees).
- Cost function (in terms of time)
  - ➔ I/O cost + CPU cost + communication cost
  - ➔ These might have different weights in different distributed environments (LAN vs WAN).
  - ➔ Can also maximize throughput
- Search algorithm
  - ➔ How do we move inside the solution space?
  - ➔ Exhaustive search, heuristic algorithms (iterative improvement, simulated annealing, genetic,...)



# Query Optimization Process

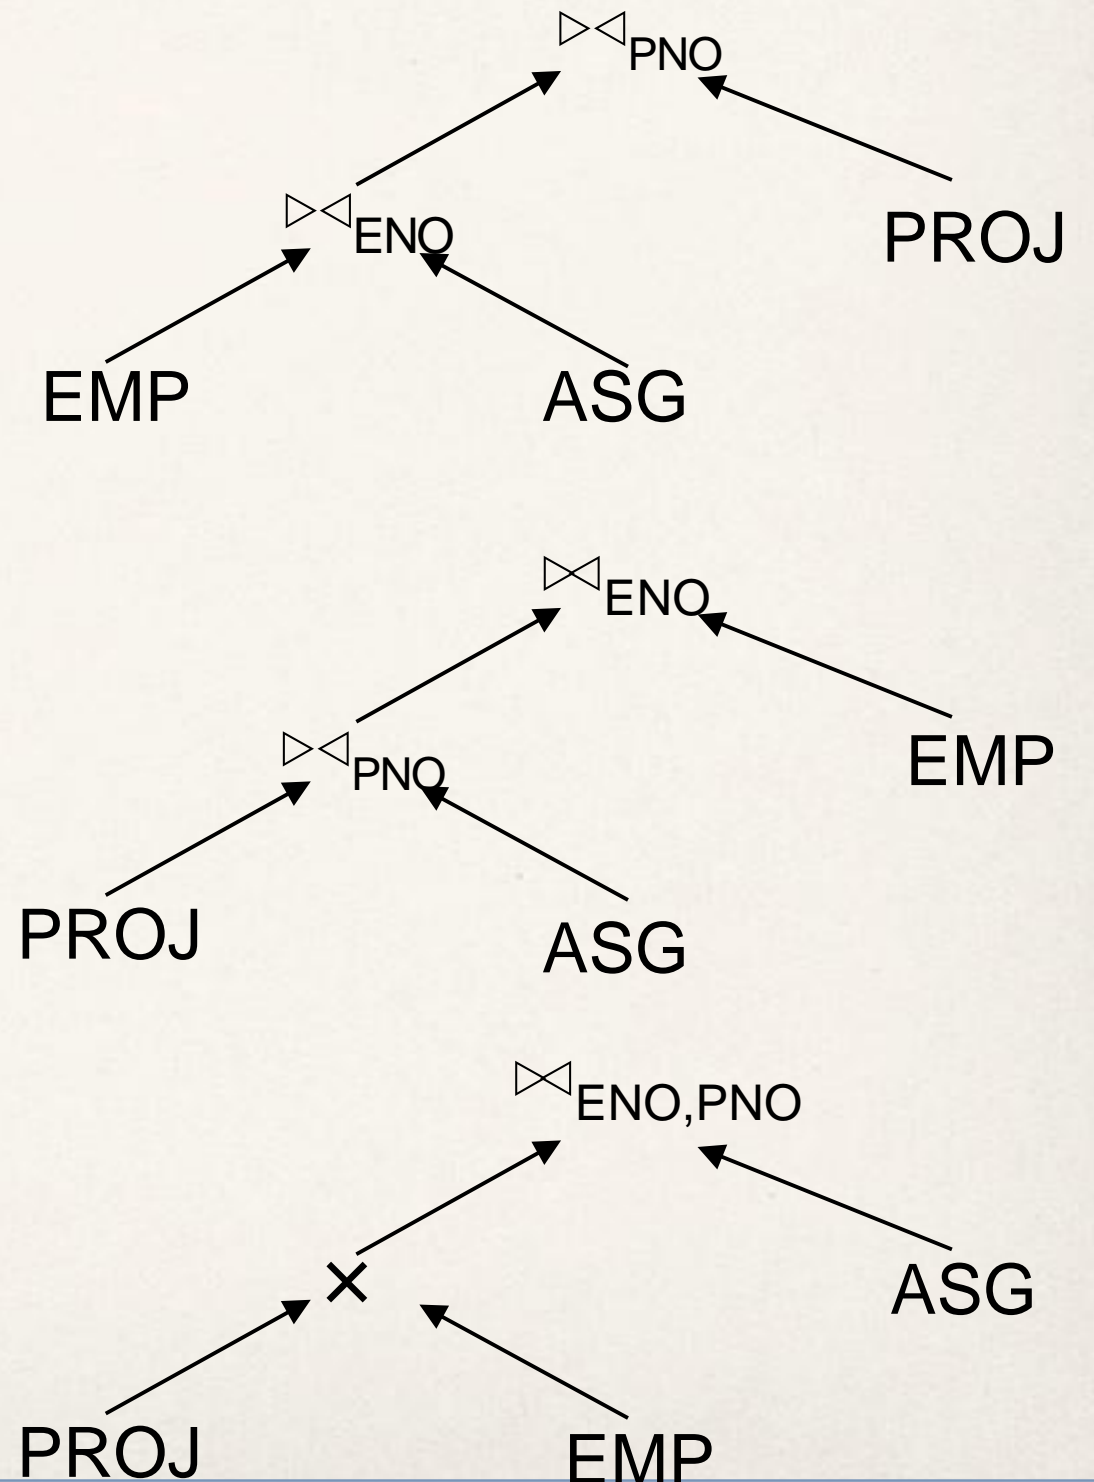
---



# Search Space

- Search space characterized by alternative execution
- Focus on join trees
- For  $N$  relations, there are  $O(N!)$  equivalent join trees that can be obtained by applying commutativity and associativity rules

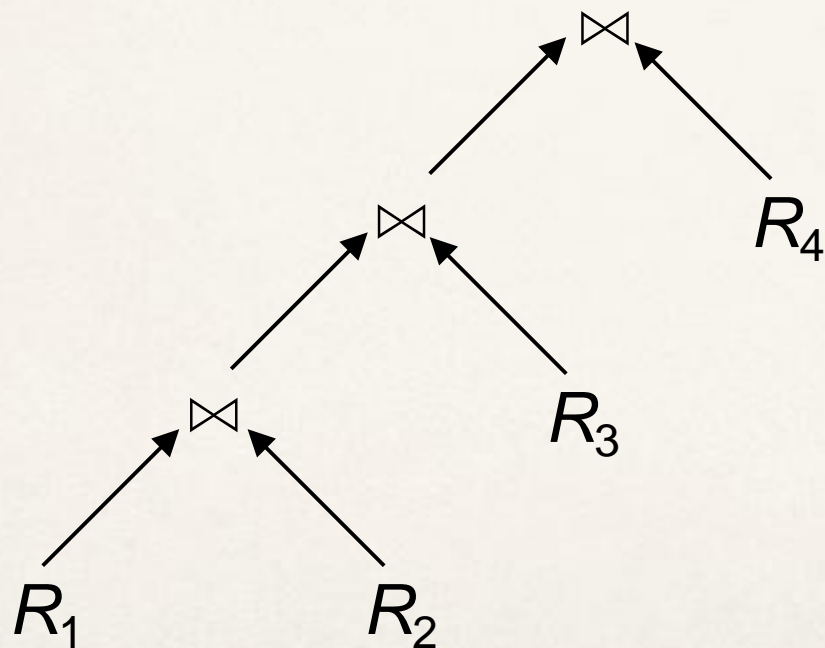
**SELECT**    ENAME, RESP  
**FROM**      EMP, ASG, PROJ  
**WHERE**     EMP.ENO=ASG.ENO  
**AND**        ASG.PNO=PROJ.PNO



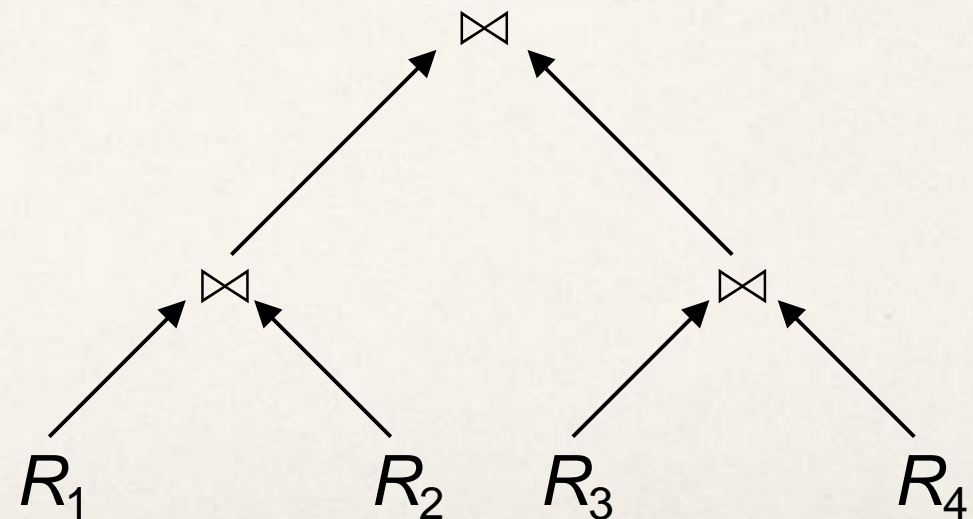
# Search Space

- Restrict by means of heuristics
  - ➔ Perform unary operations before binary operations
  - ➔ ...
- Restrict the shape of the join tree
  - ➔ Consider only linear trees, ignore bushy ones

Linear Join Tree



Bushy Join Tree





# Search Strategy

---

- How to “move” in the search space.

- Deterministic

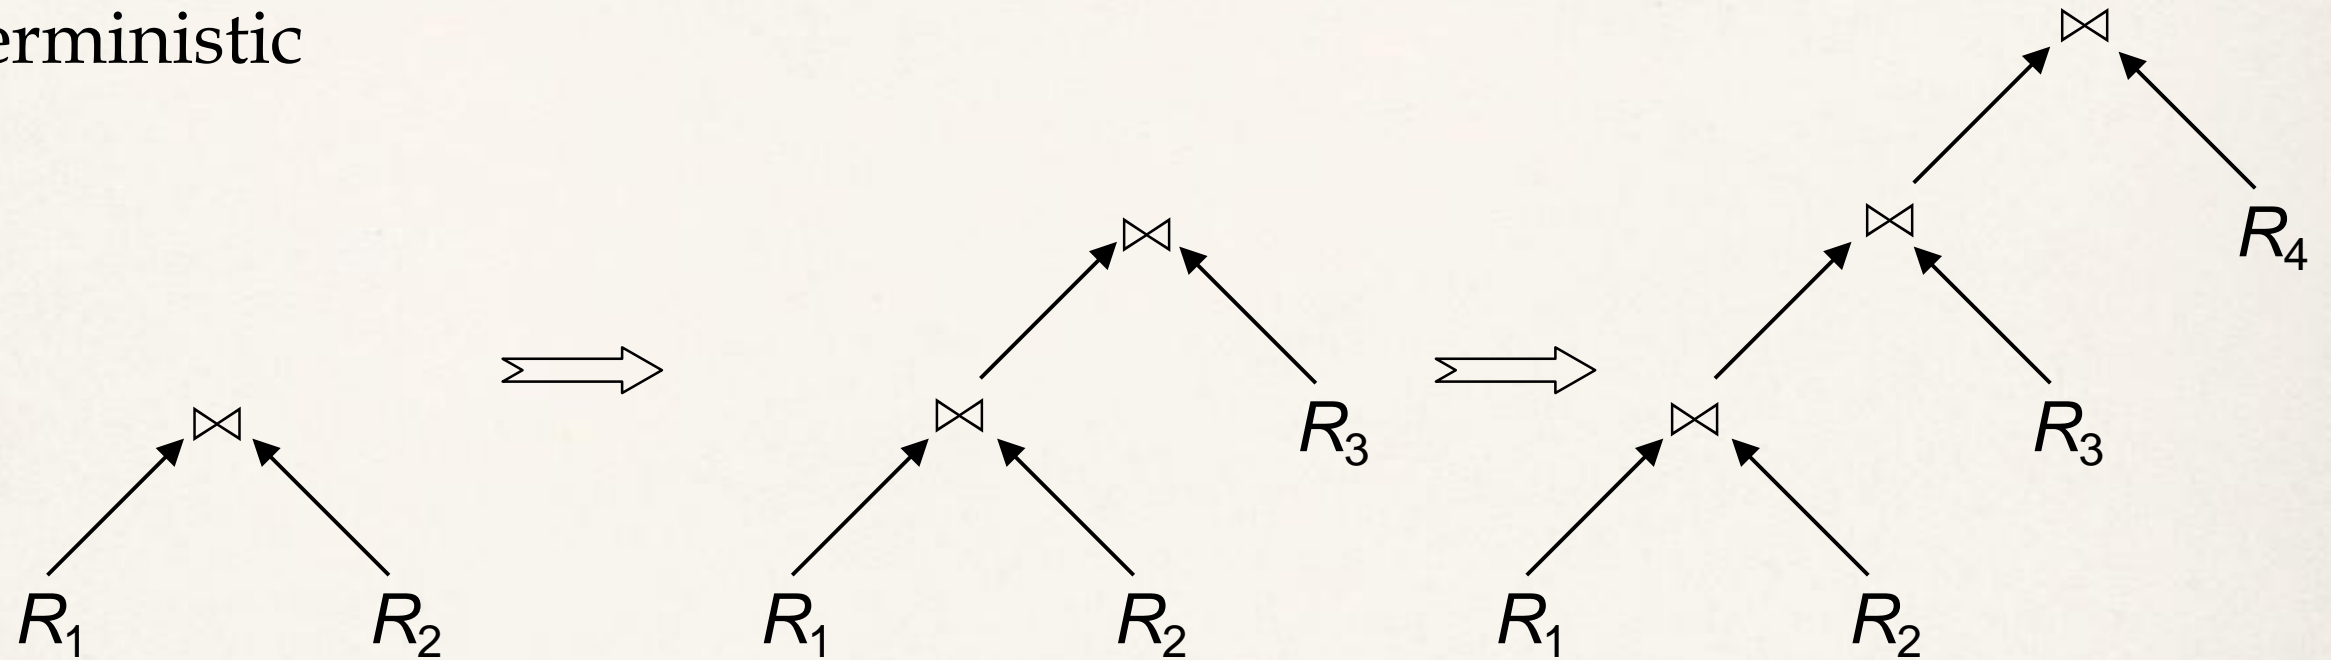
- ➔ Start from base relations and build plans by adding one relation at each step
- ➔ Dynamic programming: breadth-first
- ➔ Greedy: depth-first

- Randomized

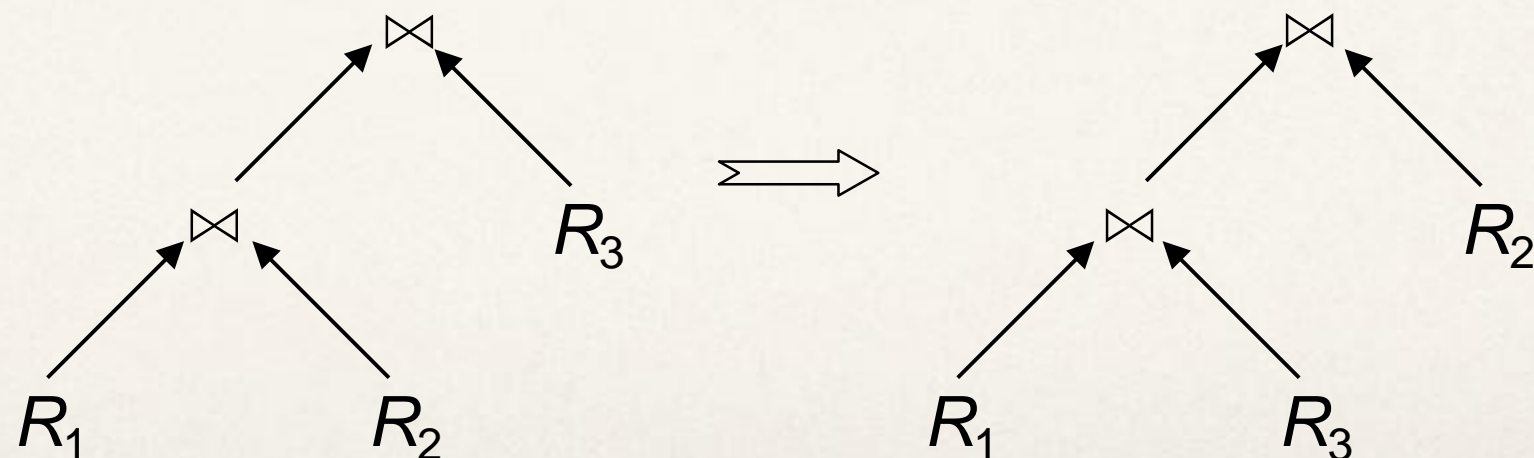
- ➔ Search for optimalities around a particular starting point
- ➔ Trade optimization time for execution time
- ➔ Better when  $> 10$  relations
- ➔ Simulated annealing
- ➔ Iterative improvement

# Search Strategies

- Deterministic



- Randomized





# Cost Functions

---

- Total Time (or Total Cost)
  - ➔ Reduce each cost (in terms of time) component individually
  - ➔ Do as little of each cost component as possible
  - ➔ Optimizes the utilization of the resources



Increases system throughput

- Response Time
  - ➔ Do as many things as possible in parallel
  - ➔ May increase total time because of increased total activity

# Total Cost

---

Summation of all cost factors

Total cost = CPU cost + I/O cost + communication cost

CPU cost = unit instruction cost \* no.of instructions

I/O cost = unit disk I/O cost \* no. of disk I/Os

communication cost = message initiation + transmission

# Total Cost Factors

---

- Wide area network
  - ➔ Message initiation and transmission costs high
  - ➔ Local processing cost is low (fast mainframes or minicomputers)
  - ➔ Ratio of communication to I/O costs = 20:1
- Local area networks
  - ➔ Communication and local processing costs are more or less equal
  - ➔ Ratio = 1:1.6



# Response Time

---

Elapsed time between the initiation and the completion of a query

Response time = CPU time + I/O time + communication time

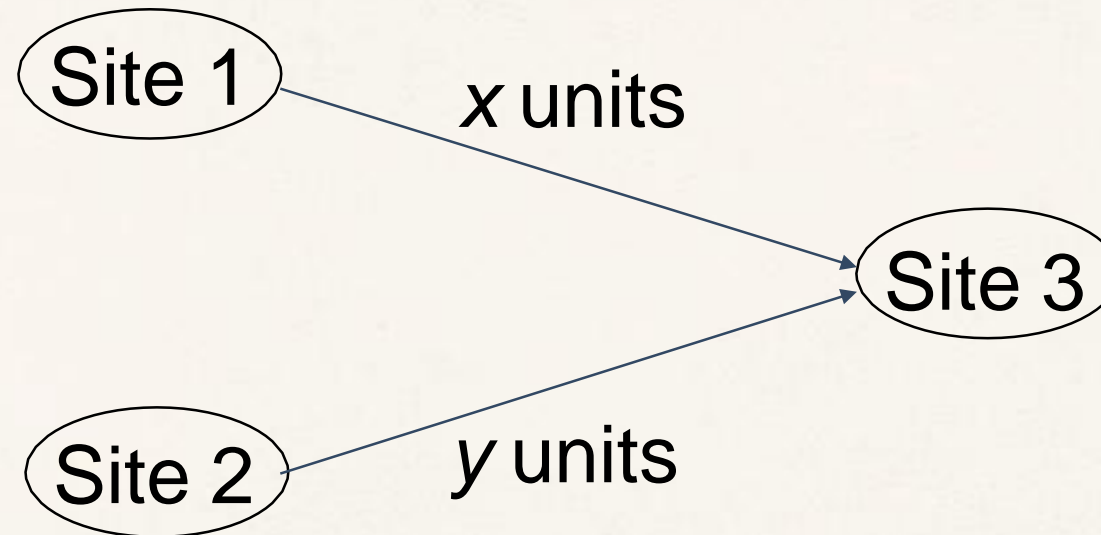
CPU time = unit instruction time \* no. of sequential instructions

I/O time = unit I/O time \* no. of sequential I/Os

communication time = unit msg initiation time \* no. of sequential msg  
+ unit transmission time \* no. of sequential bytes

# Example

---



Assume that only the communication cost is considered

Total time =  $2 \cdot \text{message initialization time} + \text{unit transmission time} * (x+y)$

Response time =  $\max \{ \text{time to send } x \text{ from 1 to 3, time to send } y \text{ from 2 to 3} \}$

time to send  $x$  from 1 to 3 = message initialization time  
+ unit transmission time \*  $x$

time to send  $y$  from 2 to 3 = message initialization time  
+ unit transmission time \*  $y$

# Optimization Statistics

---

- Primary cost factor: **size of intermediate relations**
  - ➔ Need to estimate their sizes
- Make them precise  $\Rightarrow$  more costly to maintain
- Simplifying assumption: uniform distribution of attribute values in a relation



# Statistics

---

- For each relation  $R[A_1, A_2, \dots, A_n]$  fragmented as  $R_1, \dots, R_r$ 
  - ➔ length of each attribute:  $length(A_i)$
  - ➔ The cardinalities of each fragment:  $card(R_j)$
  - ➔ the cardinalities of each domain:  $card(dom[A_i])$
  - ➔ the number of distinct values for  $A_i$  in fragment  $R_j$ :  $card(\Pi_{A_i} R_j)$
  - ➔ maximum and minimum values in the domain of each attribute:  $min(A_i)$ ,  $max(A_i)$
- Cardinality of each operation for relations

$$size(R) = card(R) \cdot length(R)$$

# Intermediate Relation Sizes

---

## Selection

$$card(\sigma_P(R)) = SF_\sigma(P) \cdot card(R)$$

where

$$SF_\sigma(A = value) = \frac{1}{card(\Pi_A(R))}$$

$$SF_\sigma(A > value) = \frac{max(A) - value}{max(A) - min(A)}$$

$$SF_\sigma(A < value) = \frac{value - min(A)}{max(A) - min(A)}$$

$$SF_\sigma(p(A_i) \wedge p(A_j)) = SF_\sigma(p(A_i)) \cdot SF_\sigma(p(A_j))$$

$$SF_\sigma(p(A_i) \vee p(A_j)) = SF_\sigma(p(A_i)) + SF_\sigma(p(A_j)) - (SF_\sigma(p(A_i)) \cdot SF_\sigma(p(A_j)))$$

$$SF_\sigma(A \in \{value\}) = SF_\sigma(A = value) * card(\{values\})$$

# Intermediate Relation Sizes

---

## Projection

$$\text{card}(\Pi_A(R)) = \text{card}(R)$$

## Cartesian Product

$$\text{card}(R \times S) = \text{card}(R) * \text{card}(S)$$

## Union

$$\text{upper bound: } \text{card}(R \cup S) = \text{card}(R) + \text{card}(S)$$

$$\text{lower bound: } \text{card}(R \cup S) = \max\{\text{card}(R), \text{card}(S)\}$$

## Set Difference

$$\text{upper bound: } \text{card}(R - S) = \text{card}(R)$$

$$\text{lower bound: } 0$$



# Intermediate Relation Size

---

## Join

- Special case:  $A$  is a key of  $R$  and  $B$  is a foreign key of  $S$

$$\text{card}(R \bowtie_{A=B} S) = \text{card}(S)$$

- Other cases:

$$\text{card}(R \bowtie S) = SF_{\bowtie} * \text{card}(R) \cdot \text{card}(S)$$

## Semi-join

$$\text{card}(R \ltimes_A S) = SF_{\ltimes}(R \ltimes_A S) * \text{card}(R)$$

where

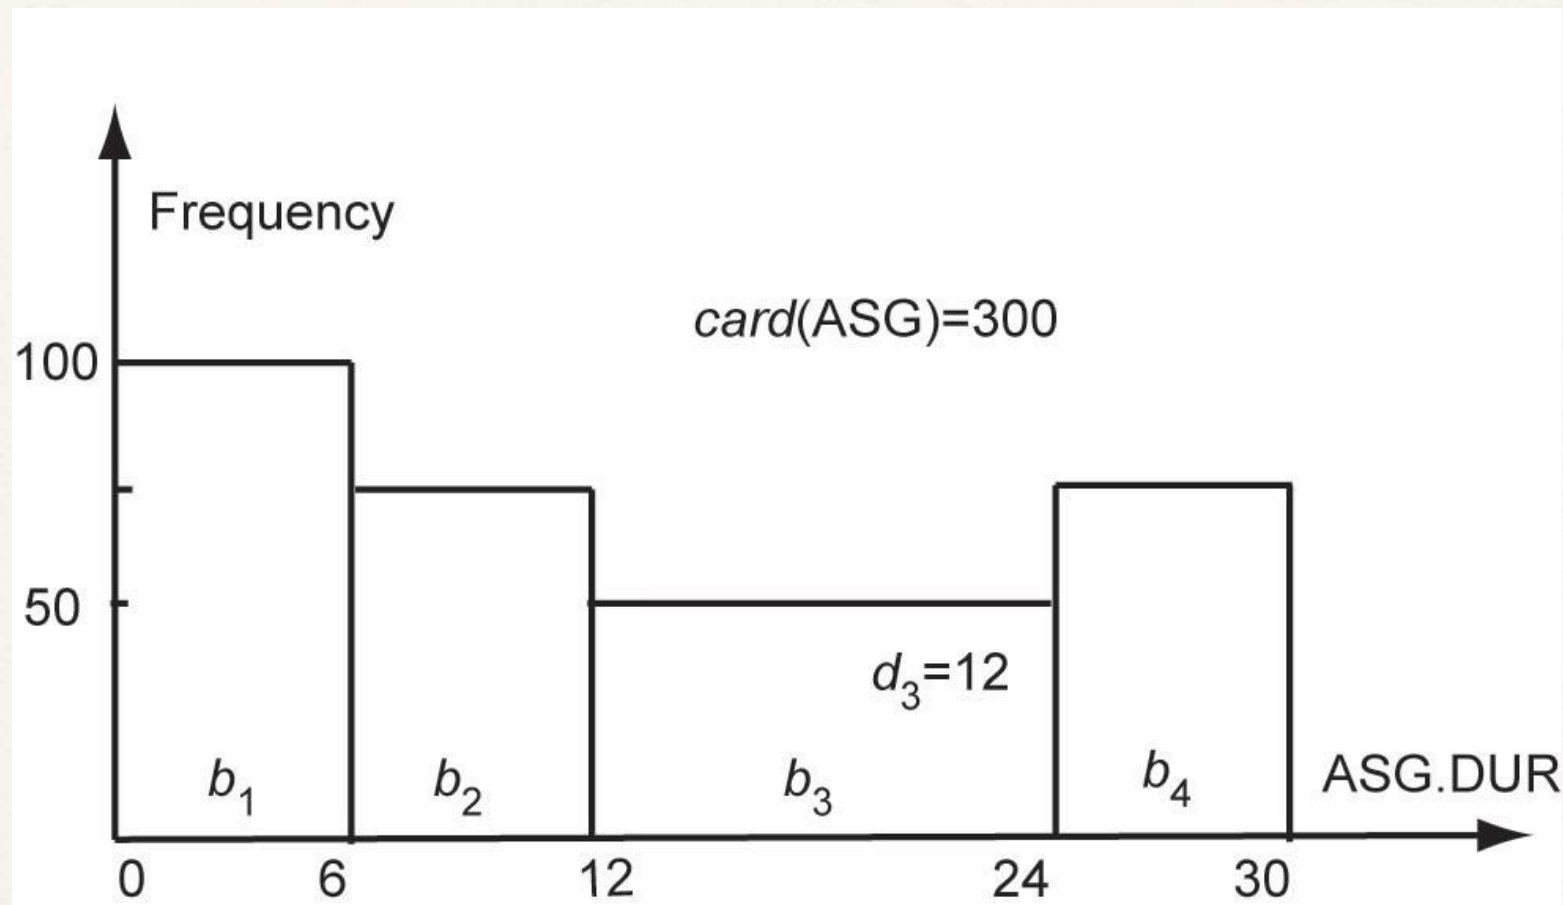
$$SF_{\ltimes}(R \ltimes_A S) = \frac{\text{card}(\Pi_A(S))}{\text{card}(\text{dom}[A])}$$

# Histograms for Selectivity Estimation

---

- For skewed data, the uniform distribution assumption of attribute values yields inaccurate estimations
- Use an histogram for each skewed attribute  $A$ 
  - ➔ Histogram = set of buckets
    - ◆ Each bucket describes a range of values of  $A$ , with its average frequency  $f$  (number of tuples with  $A$  in that range) and number of distinct values  $d$
    - ◆ Buckets can be adjusted to different ranges
- Examples
  - ➔ Equality predicate
    - ◆ With (value in  $\text{Range}_i$ ), we have:  $SF_{\sigma}(A = \text{value}) = 1/d_i$
  - ➔ Range predicate
    - ◆ Requires identifying relevant buckets and summing up their frequencies

# Histogram Example



For  $ASG.DUR=18$ : we have  $SF=1/12$  and  $card(b_3)=50$ , so the card of selection is  $50/12 \sim 4$  tuples

For  $ASG.DUR \leq 18$ : we have  $\min(range_3)=12$  and  $\max(range_3)=24$  so the card. of selection is  $100+75+(((18-12)/(24-12))*50) = 200$  tuples



# Centralized Query Optimization

---

- Dynamic (Ingres project at UCB)
- Static (System R project at IBM)
- Hybrid (Volcano project at OGI)

# Dynamic Algorithm

---

- ① Decompose each multi-relation query into a sequence of mono-relation queries
- ② Process each by a one relation query processor
  - ➔ Choose an initial execution plan (heuristics)
  - ➔ Order the rest by considering intermediate relation sizes



No statistical information is maintained

# Dynamic Algorithm— Decomposition

---

- Replace an  $n$  relation query  $q$  by a series of queries

$$q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_n$$

where  $q_i$  uses the result of  $q_{i-1}$ .

- Detachment

→ Query  $q$  decomposed into  $q' \rightarrow q''$  where  $q'$  and  $q''$  have a common relation which is the result of  $q'$

- Tuple substitution

→ Replace the value of each tuple with actual values and simplify the query

$$q(V_1, V_2, \dots, V_n) \rightarrow (q'(t_1, V_2, V_2, \dots, V_n), t_1 \in R)$$



# Detachment

---

$q$ :    **SELECT**     $R_2 \cdot A_2, R_3 \cdot A_3, \dots, R_n \cdot A_n$   
         **FROM**     $R_1, R_2, \dots, R_n$   
         **WHERE**     $P_1 (R_1 \cdot A_1')$  **AND**  $P_2 (R_1 \cdot A_1, R_2 \cdot A_2, \dots, R_n \cdot A_n)$



$q'$ :    **SELECT**     $R_1 \cdot A_1$  **INTO**  $R_1'$   
         **FROM**     $R_1$   
         **WHERE**     $P_1 (R_1 \cdot A_1')$

$q''$ :    **SELECT**     $R_2 \cdot A_2, \dots, R_n \cdot A_n$   
         **FROM**     $R_1', R_2, \dots, R_n$   
         **WHERE**     $P_2 (R_1' \cdot A_1, R_2 \cdot A_2, \dots, R_n \cdot A_n)$

# Detachment Example

---

Names of employees working on CAD/CAM project

$q_1$ :      **SELECT**            EMP.ENAME  
              **FROM**            EMP, ASG, PROJ  
              **WHERE**          EMP.ENO=ASG.ENO  
              **AND**            ASG.PNO=PROJ.PNO  
              **AND**            PROJ.PNAME="CAD/CAM"  
                                  $\Downarrow$

$q_{11}$ :      **SELECT**            PROJ.PNO **INTO** JVAR  
              **FROM**            PROJ  
              **WHERE**          PROJ.PNAME="CAD/CAM"

$q'$ :        **SELECT**            EMP.ENAME  
              **FROM**            EMP, ASG, JVAR  
              **WHERE**          EMP.ENO=ASG.ENO  
              **AND**            ASG.PNO=JVAR.PNO

# Detachment Example (cont'd)

---

$q'$ :      **SELECT**      EMP.ENAME  
             **FROM**        EMP, ASG, JVAR  
             **WHERE**      EMP.ENO=ASG.ENO  
             **AND**        ASG.PNO=JVAR.PNO



$q_{12}$ :      **SELECT**      ASG.ENO **INTO** GVAR  
             **FROM**        ASG, JVAR  
             **WHERE**      ASG.PNO=JVAR.PNO

$q_{13}$ :      **SELECT**      EMP.ENAME  
             **FROM**        EMP, GVAR  
             **WHERE**      EMP.ENO=GVAR.ENO



# Tuple Substitution

---

$q_{11}$  is a mono-relation query

$q_{12}$  and  $q_{13}$  are subject to tuple substitution

Assume GVAR has two tuples only:  $\langle E1 \rangle$  and  $\langle E2 \rangle$

Then  $q_{13}$  becomes

$q_{131}$ :     **SELECT**       EMP.ENAME  
              **FROM**       EMP  
              **WHERE**      EMP.ENO="E1"

$q_{13}$ :     **SELECT** EMP.ENAME  
              **FROM** EMP, GVAR  
              **WHERE** EMP.ENO=GVAR.ENO

$q_{132}$ :     **SELECT**       EMP.ENAME  
              **FROM**       EMP  
              **WHERE**      EMP.ENO="E2"

# Static Algorithm

---

- ① Simple (i.e., mono-relation) queries are executed according to the best access path
- ② Execute joins
  - ➔ Determine the possible ordering of joins
  - ➔ Determine the cost of each ordering
  - ➔ Choose the join ordering with minimal cost

# Static Algorithm

---

For joins, two alternative algorithms :

- Nested loops

```
for each tuple of external relation (cardinality  $n_1$ )  
  for each tuple of internal relation (cardinality  $n_2$ )  
    join two tuples if the join predicate is true  
  end  
end
```

→ Complexity:  $n_1 * n_2$

- Merge join

```
sort relations  
merge relations
```

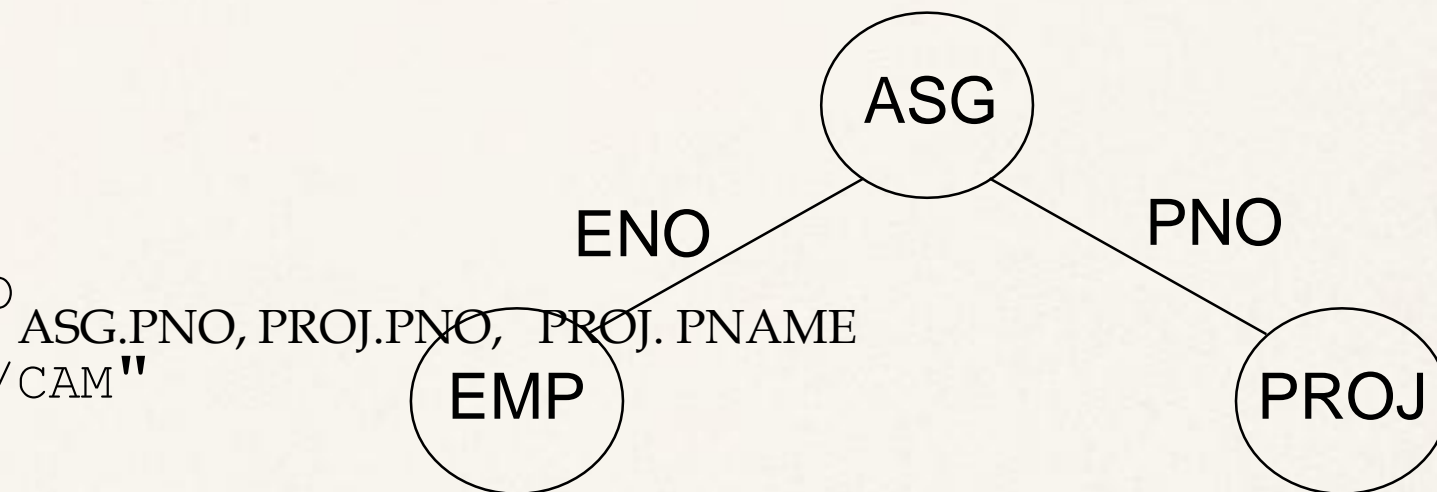
→ Complexity:  $n_1 + n_2$  if relations are previously sorted and equijoin



# Static Algorithm – Example

“Names of employees working on the CAD/CAM project”

```
SELECT EMP.ENAME  
FROM EMP, ASG, PROJ  
WHERE EMP.ENO=ASG.ENO  
AND ASG.PNO=PROJ.PNO  
AND PROJ.PNAME="CAD/CAM"
```



Assume

- EMP has an index on ENO,
- ASG has an index on PNO,
- PROJ has an index on PNO and an index on PNAME

# Example (cont'd)

## ① Choose the best access paths to each relation

- EMP: sequential scan (no predicate on EMP)
- ASG: sequential scan (no predicate on ASG)
- PROJ: index on PNAME (there is a predicate on PROJ based on PNAME)

## ② Determine the best join ordering

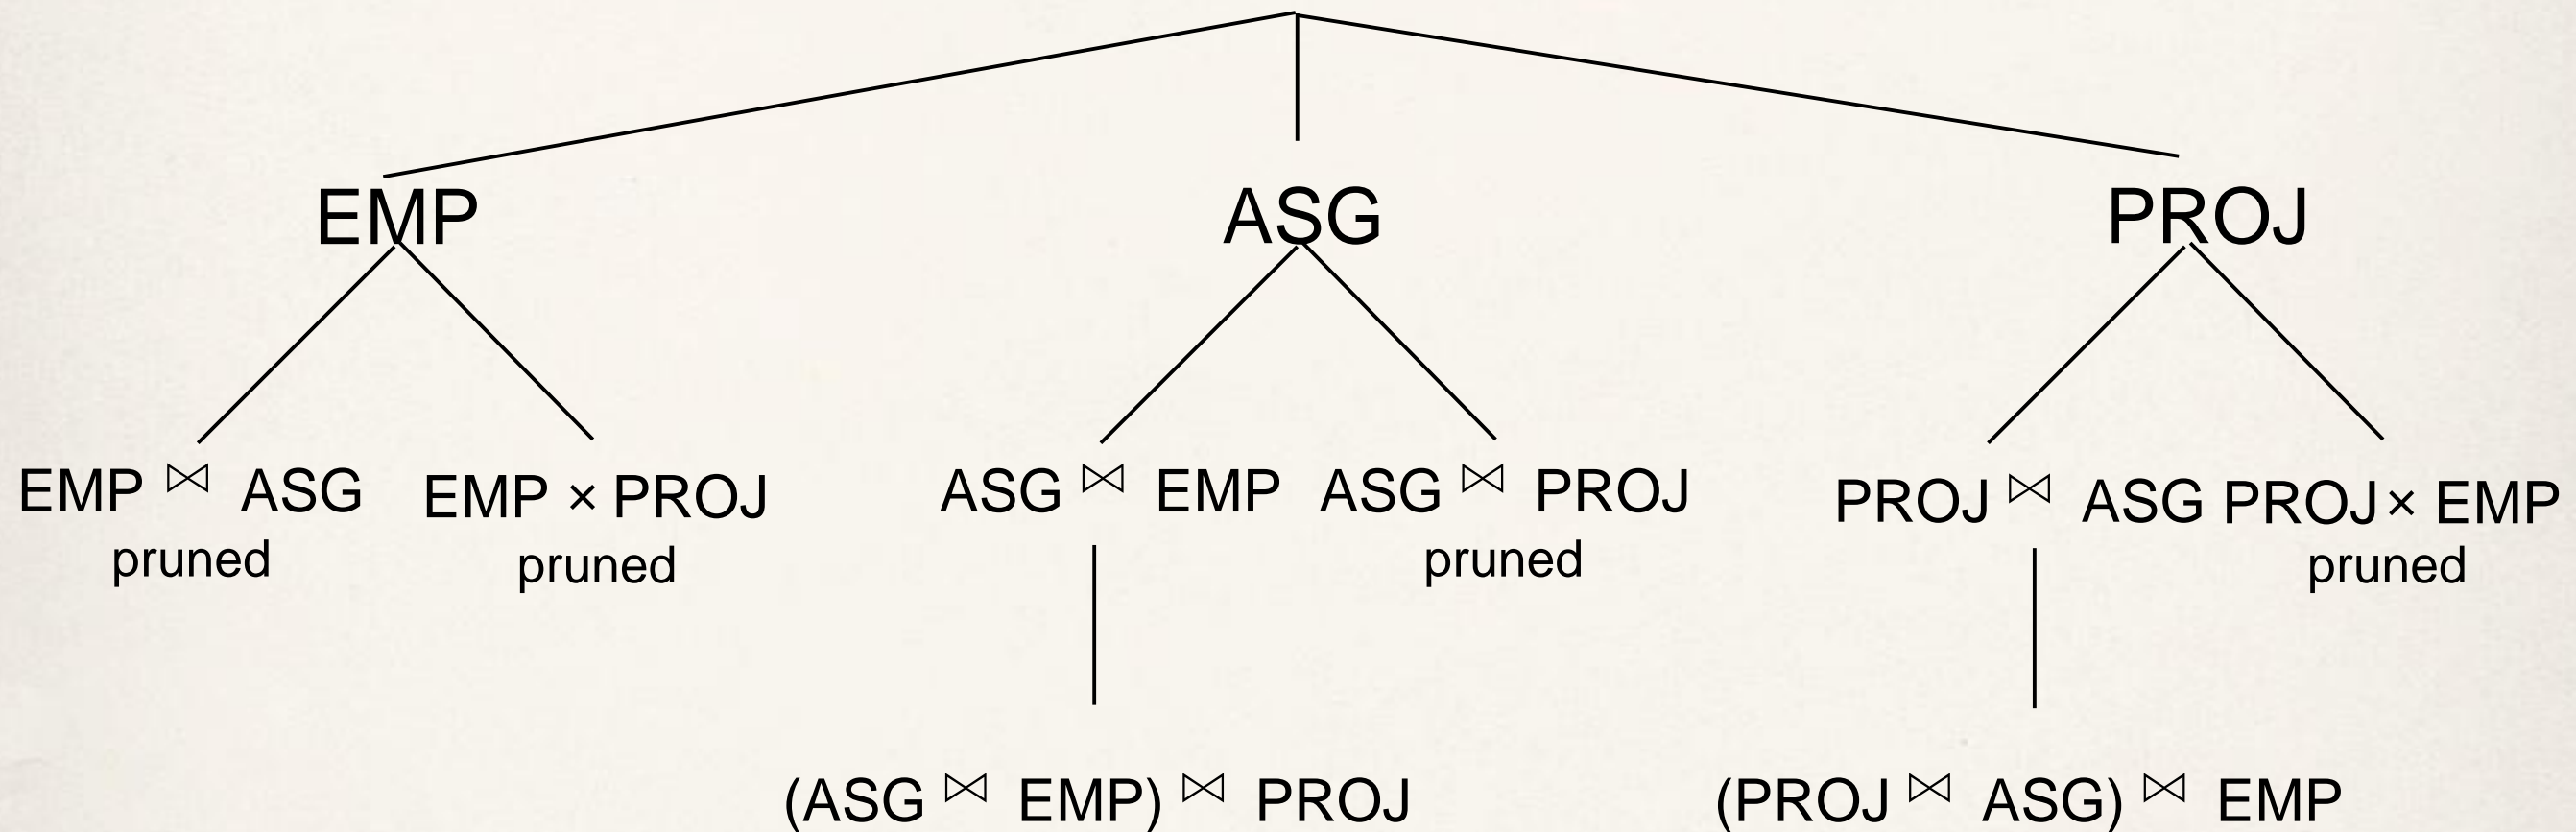
- EMP  $\bowtie$  ASG  $\bowtie$  PROJ
- ASG  $\bowtie$  PROJ  $\bowtie$  EMP
- PROJ  $\bowtie$  ASG  $\bowtie$  EMP
- ASG  $\bowtie$  EMP  $\bowtie$  PROJ
- EMP  $\times$  PROJ  $\bowtie$  ASG
- PROJ  $\times$  EMP  $\bowtie$  ASG

EMP. ENO=ASG. ENO  
ASG. PNO=PROJ. PNO

- Select the best ordering based on the join costs evaluated according to the two methods

# Static Algorithm

## Alternatives



Best total join order is one of

$((ASG \bowtie EMP) \bowtie PROJ)$

$((PROJ \bowtie ASG) \bowtie EMP)$

→ Index: EMP.ENO, ASG.PNO, PROJ.PNO,  
PROJ. PNAME

→ Join: EMP.ENO=ASG.ENO  
ASG.PNO=PROJ.PNO



# Static Algorithm

---

- $((\text{PROJ} \bowtie \text{ASG}) \bowtie \text{EMP})$  has a useful index on the select attribute and direct access to the join attributes of ASG and EMP
- Therefore, chose it with the following access methods:
  - ➔ select PROJ using index on PNAME
  - ➔ then join with ASG using index on PNO
  - ➔ then join with EMP using index on ENO

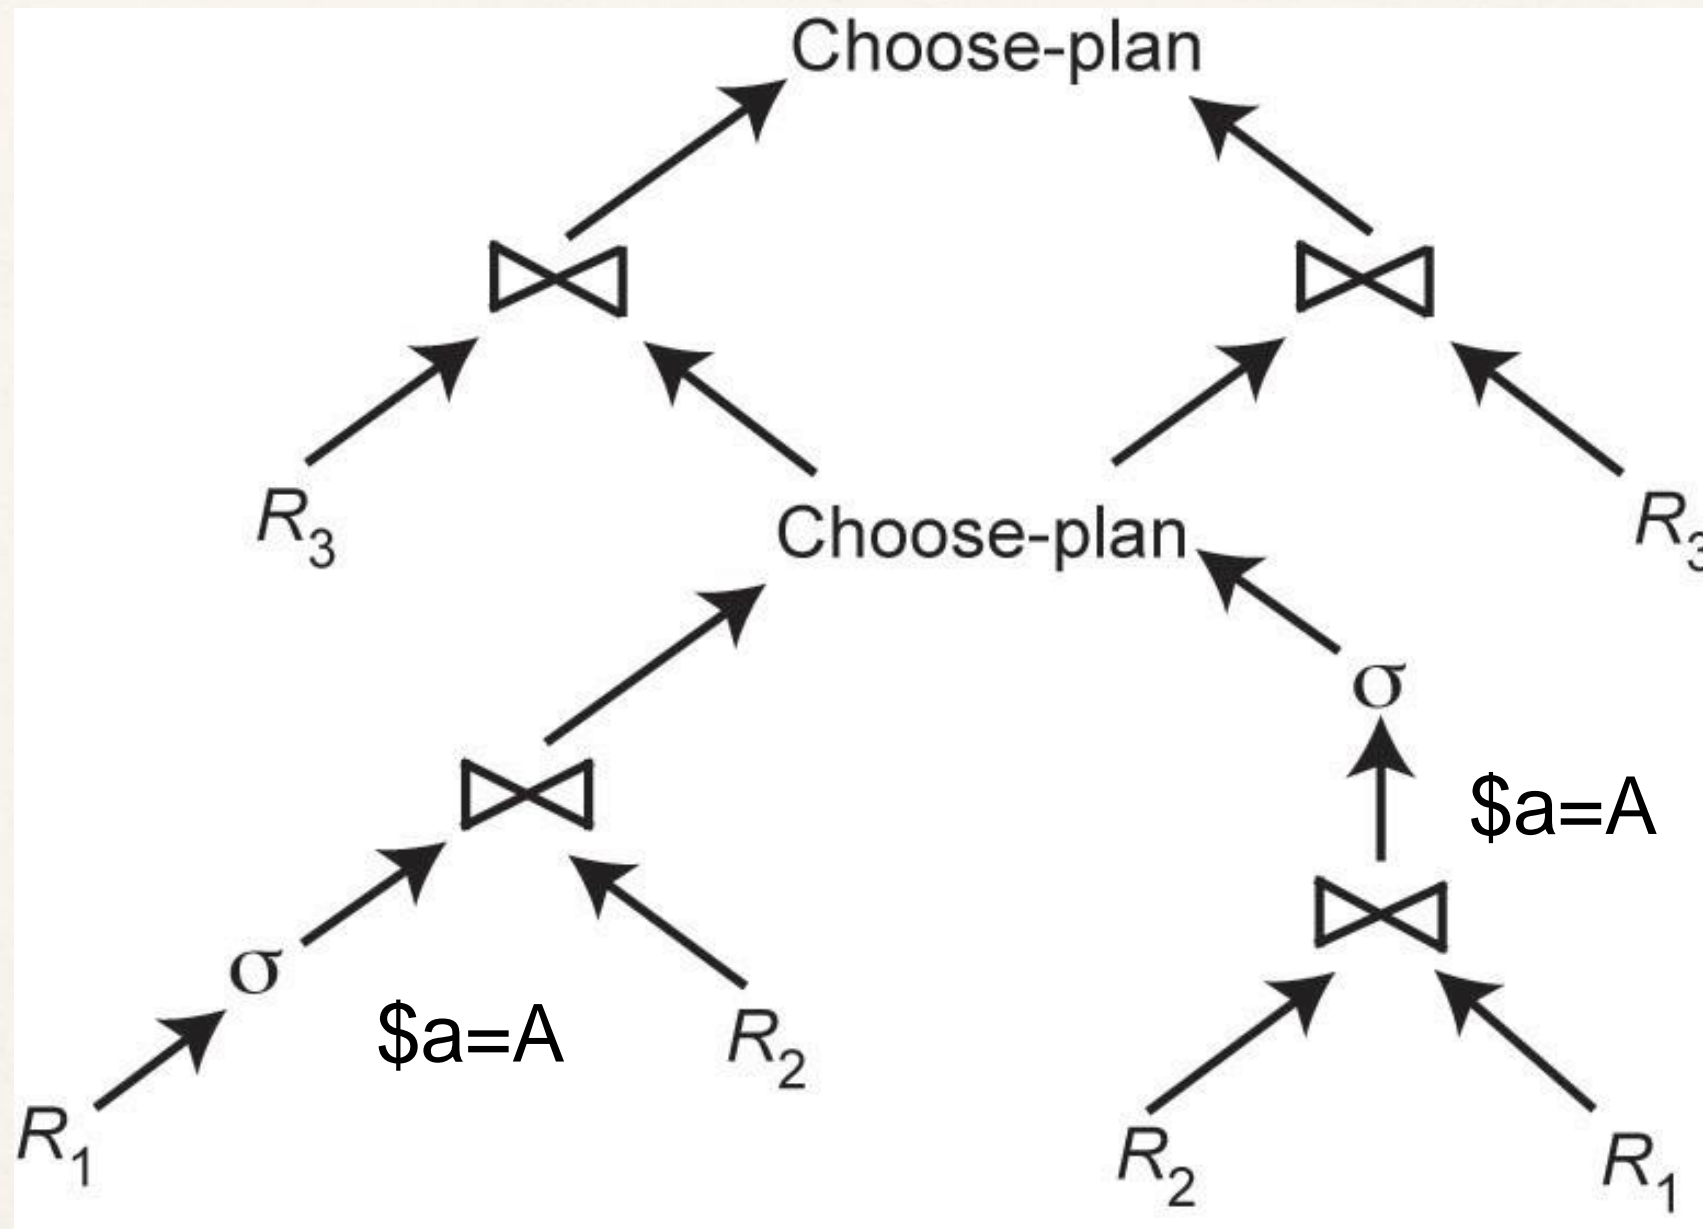
```
SELECT    EMP.ENAME
FROM      EMP, ASG, PROJ
WHERE     ASG.PNO=PROJ.PNO
AND       EMP.ENO=ASG.ENO
AND       PROJ.PNAME="CAD/CAM"
```

# Hybrid optimization

---

- In general, static optimization is more efficient than dynamic optimization
  - ➔ Adopted by all commercial DBMS
- But even with a sophisticated cost model (with histograms), accurate cost prediction is difficult
- Example: Consider a parametric query with predicate
$$\text{WHERE } R.A = \$a \quad // * \$a \text{ is a parameter}$$
$$\sigma_{R.A=\$a} (R_1) \bowtie R_2 \bowtie R_3$$
  - ➔ The only possible assumption at compile time is uniform distribution of values
- Solution: Hybrid optimization
  - ➔ Choose-plan done at runtime, based on the actual parameter binding

# Hybrid Optimization Example





# Join Ordering in Fragment Queries

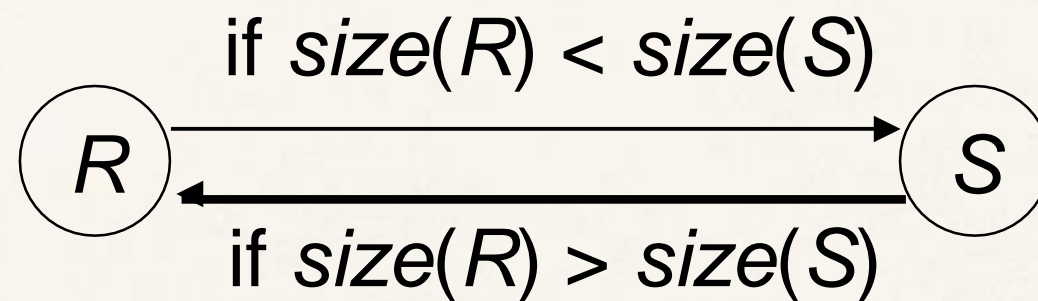
---

- Ordering joins
  - ➔ Distributed INGRES
  - ➔ System R\*
  - ➔ Two-step
- Semijoin ordering
  - ➔ SDD-1

# Join Ordering

---

- Consider two relations only



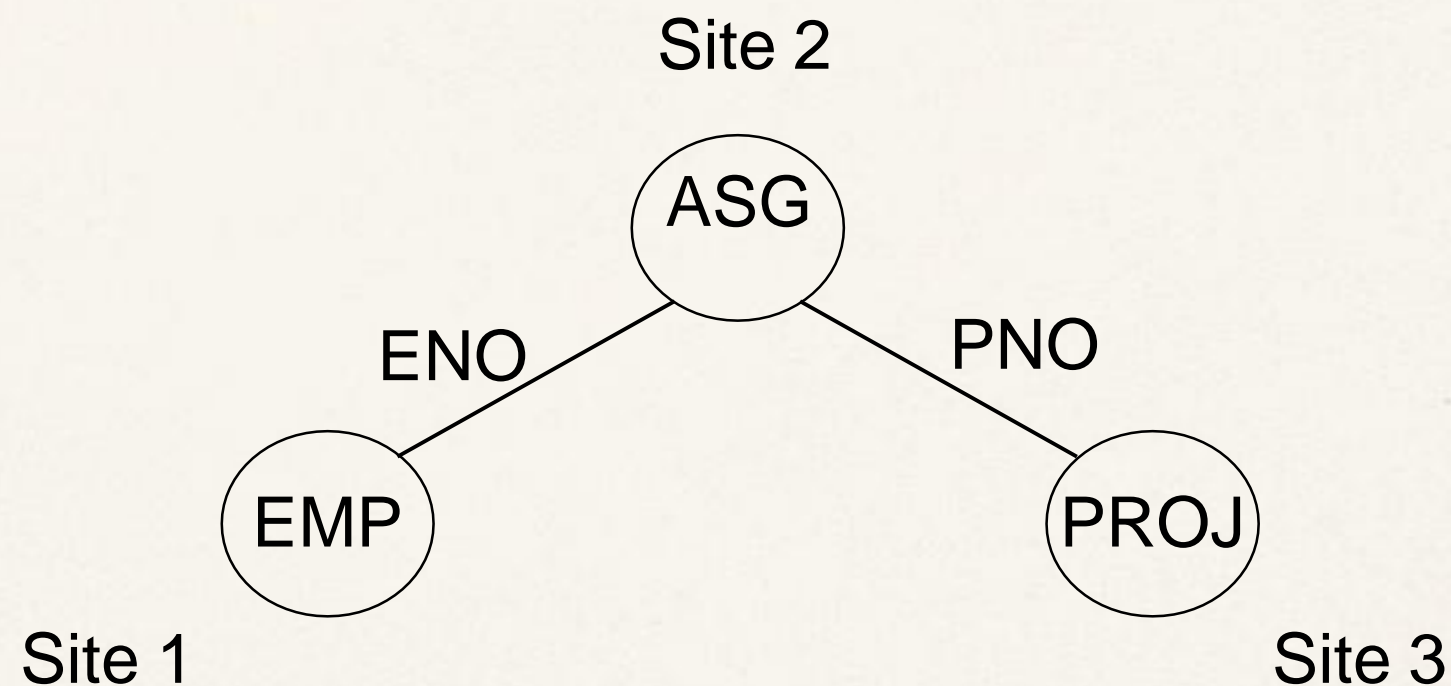
- Multiple relations more difficult because too many alternatives.
  - ➔ Compute the cost of all alternatives and select the best one.
    - ♦ Necessary to compute the size of intermediate relations which is difficult.
  - ➔ Use heuristics

# Join Ordering – Example

---

Consider

$\text{PROJ} \bowtie_{\text{PNO}} \text{ASG} \bowtie_{\text{ENO}} \text{EMP}$





# Join Ordering – Example

---

Execution alternatives:

1. EMP → Site 2

Site 2 computes  $EMP' = EMP \bowtie ASG$

$EMP' \rightarrow$  Site 3

Site 3 computes  $EMP' \bowtie PROJ$

2. ASG → Site 1

Site 1 computes  $EMP' = EMP \bowtie ASG$

$EMP' \rightarrow$  Site 3

Site 3 computes  $EMP' \bowtie PROJ$

3. ASG → Site 3

Site 3 computes  $ASG' = ASG \bowtie PROJ$

$ASG' \rightarrow$  Site 1

Site 1 computes  $ASG' \bowtie EMP$

4. PROJ → Site 2

Site 2 computes  $PROJ' = PROJ \bowtie ASG$

$PROJ' \rightarrow$  Site 1

Site 1 computes  $PROJ' \bowtie EMP$

5. EMP → Site 2

PROJ → Site 2

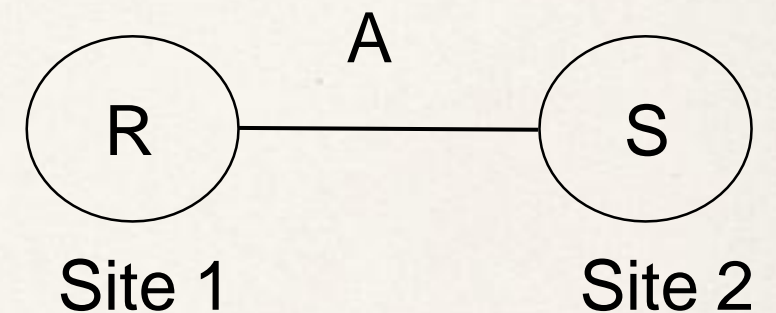
Site 2 computes  $EMP \bowtie PROJ \bowtie ASG$

# Semijoin Algorithms

---

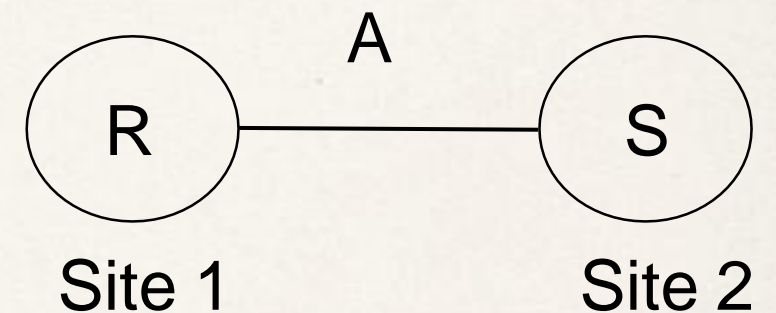
- Consider the join of two relations:
  - ➔  $R[A]$  (located at site 1)
  - ➔  $S[A]$  (located at site 2)
- Perform one of the semi-join equivalents :

$$\begin{aligned} R \bowtie_A S &\Leftrightarrow (R \bowtie_A S) \bowtie_A S \\ &\Leftrightarrow R \bowtie_A (S \bowtie_A R) \\ &\Leftrightarrow (R \bowtie_A S) \bowtie_A (S \bowtie_A R) \end{aligned}$$



# Semijoin Algorithms

- Consider semijoin  $(R \bowtie_A S) \bowtie_A S$ 
  - ➔  $S' = \Pi_A(S)$
  - ➔  $S' \rightarrow \text{Site 1}$
  - ➔ Site 1 computes  $R' = R \bowtie_A S'$
  - ➔  $R' \rightarrow \text{Site 2}$
  - ➔ Site 2 computes  $R' \bowtie_A S$
- Perform the join by sending  $R$  to site 2
  - ➔ send  $R$  to Site 2
  - ➔ Site 2 computes  $R \bowtie_A S$



Semijoin is better if

$$size(\Pi_A(S)) + size(R \bowtie_A S) < size(R)$$



# Distributed Dynamic Algorithm

---

1. Execute all monorelation queries (e.g., selection, projection)
2. Reduce the multirelation query to produce irreducible subqueries  $q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_n$  such that there is only one relation between  $q_i$  and  $q_{i+1}$
3. Choose  $q_i$  involving the smallest fragments to execute (call MRQ')
4. Find the best execution strategy for MRQ'
  - a) Determine processing site
  - b) Determine fragments to move
5. Repeat 3 and 4

# Static Approach

---

- Cost function includes local processing as well as transmission
- Considers only joins
- “Exhaustive” search
- Compilation
- Published papers provide solutions to handling horizontal and vertical fragmentations but the implemented prototype does not

# Static Approach – Performing Joins

---

- Ship whole
  - ➔ Larger data transfer
  - ➔ Smaller number of messages
  - ➔ Better if relations are small
- Fetch as needed
  - ➔ Number of messages =  $O(\text{cardinality of external relation})$
  - ➔ Data transfer per message is minimal
  - ➔ Better if relations are large and the selectivity is good



# Static Approach – Vertical Partitioning & Joins

---

1. Move outer relation tuples to the site of the inner relation
  - (a) Retrieve outer tuples
  - (b) Send them to the inner relation site
  - (c) Join them as they arrive

$$\begin{aligned} \text{Total Cost} = & \text{cost}(\text{retrieving qualified outer tuples}) \\ & + \text{no. of outer tuples fetched} * \text{cost}(\text{retrieving qualified} \\ & \quad \text{inner tuples}) \\ & + \text{msg. cost} * (\text{no. outer tuples fetched} * \text{avg. outer} \\ & \quad \text{tuple size}) / \text{msg. size} \end{aligned}$$

# Static Approach – Vertical Partitioning & Joins

---

## 2. Move inner relation to the site of outer relation

Cannot join as they arrive; they need to be stored

$$\begin{aligned} \text{Total cost} = & \text{cost}(\text{retrieving qualified outer tuples}) \\ & + \text{no. of outer tuples fetched} * \text{cost}(\text{retrieving} \\ & \quad \text{matching inner tuples from temporary storage}) \\ & + \text{cost}(\text{retrieving qualified inner tuples}) \\ & + \text{cost}(\text{storing all qualified inner tuples in temporary} \\ & \quad \text{storage}) \\ & + \text{msg. cost} * \text{no. of inner tuples fetched} * \text{avg. inner} \\ & \quad \text{tuple size/msg. size} \end{aligned}$$

# Static Approach – Vertical Partitioning & Joins

---

## 3. Move both inner and outer relations to another site

$$\begin{aligned}\text{Total cost} = & \text{cost}(\text{retrieving qualified outer tuples}) \\ & + \text{cost}(\text{retrieving qualified inner tuples}) \\ & + \text{cost}(\text{storing inner tuples in storage}) \\ & + \text{msg. cost} \cdot (\text{no. of outer tuples fetched} * \text{avg. outer} \\ & \quad \text{tuple size}) / \text{msg. size} \\ & + \text{msg. cost} * (\text{no. of inner tuples fetched} * \text{avg. inner} \\ & \quad \text{tuple size}) / \text{msg. size} \\ & + \text{no. of outer tuples fetched} * \text{cost}(\text{retrieving inner} \\ & \quad \text{tuples from temporary storage})\end{aligned}$$



# Static Approach – Vertical Partitioning & Joins

---

## 4. Fetch inner tuples as needed

- (a) Retrieve qualified tuples at outer relation site
- (b) Send request containing join column value(s) for outer tuples to inner relation site
- (c) Retrieve matching inner tuples at inner relation site
- (d) Send the matching inner tuples to outer relation site
- (e) Join as they arrive

$$\begin{aligned} \text{Total Cost} = & \text{cost}(\text{retrieving qualified outer tuples}) \\ & + \text{msg. cost} * (\text{no. of outer tuples fetched}) \\ & + \text{no. of outer tuples fetched} * \text{no. of} \\ & \quad \text{inner tuples fetched} * \text{avg. inner tuple} \\ & \quad \text{size} * \text{msg. cost} / \text{msg. size}) \\ & + \text{no. of outer tuples fetched} * \text{cost}(\text{retrieving} \\ & \quad \text{matching inner tuples for one outer value}) \end{aligned}$$

# Dynamic vs. Static vs Semijoin

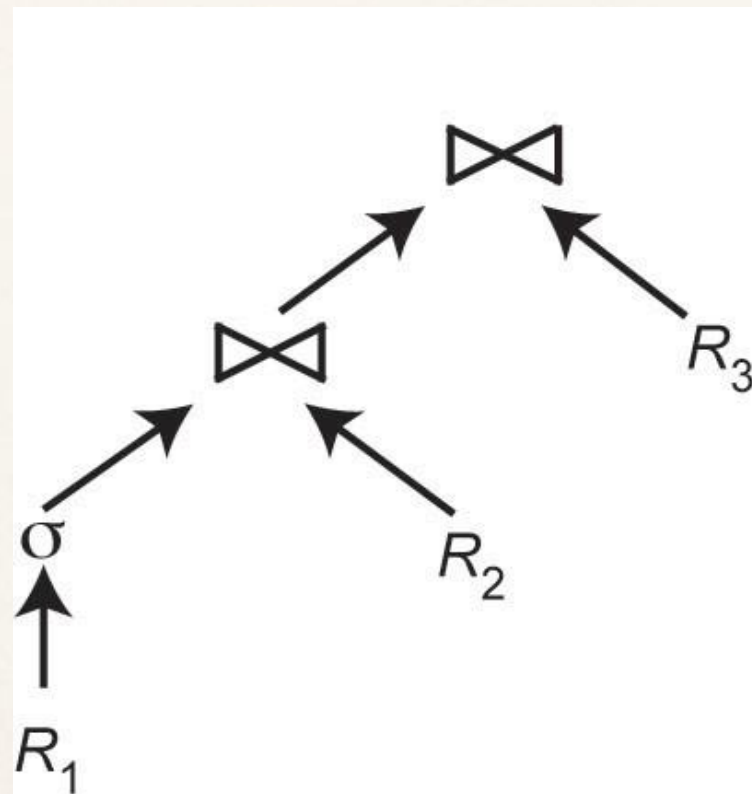
---

- Dynamic and static approaches have the same advantages and drawbacks as in centralized case
  - ➔ But the problems of accurate cost estimation at compile-time are more severe
    - ◆ More variations at runtime
    - ◆ Relations may be replicated, making site and copy selection important
- Semijoin
  - ➔ SDD1 selects only locally optimal schedules
- Hybrid optimization
  - ➔ Choose-plan approach can be used
  - ➔ 2-step approach simpler

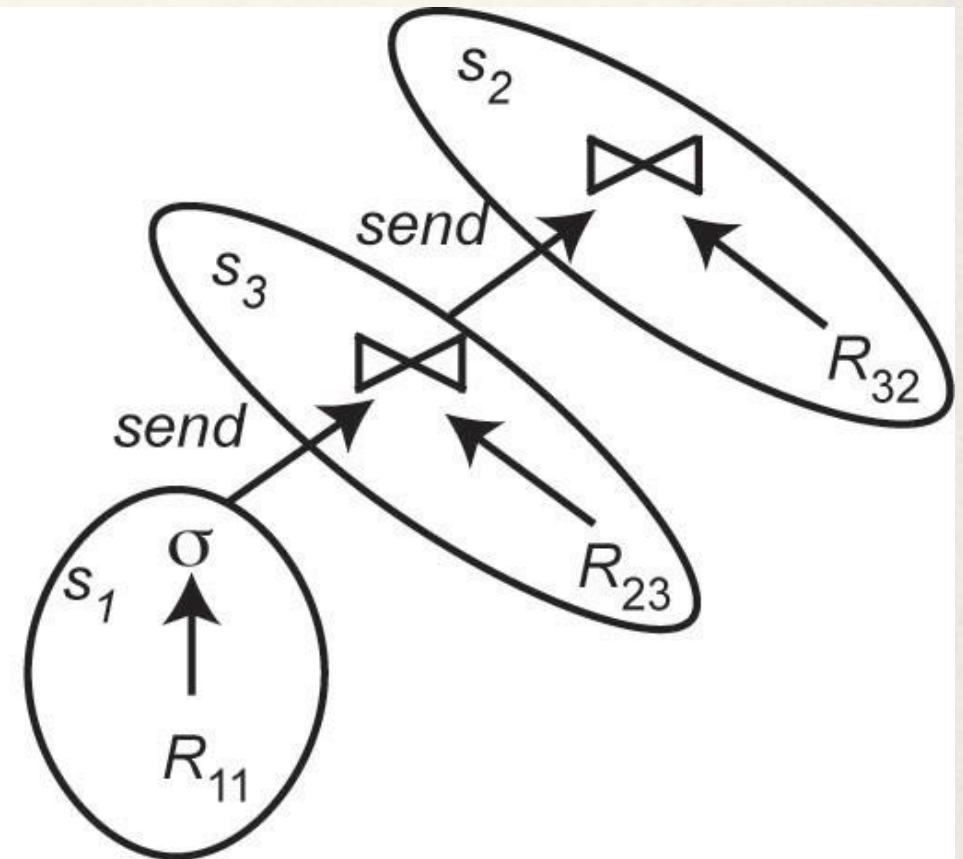
# Hybrid: 2-Step Optimization

1. At compile time, generate a static plan with operation ordering and access methods only
2. At startup time, carry out site and copy selection and allocate operations to sites

$$\sigma(R_1) \bowtie R_2 \bowtie R_3$$



(a) Static plan



(b) Run-time plan



# 2-Step – Problem Definition

---

- Given
  - ➔ A set of sites  $S = \{s_1, s_2, \dots, s_n\}$  with the load of each site
  - ➔ A query  $Q = \{q_1, q_2, \dots, q_m\}$  such that each subquery  $q_i$  is the maximum processing unit that accesses one relation and communicates with its neighboring queries
  - ➔ For each  $q_i$  in  $Q$ , a feasible allocation set of sites  $S_{q_i} = \{s_1, s_2, \dots, s_k\}$  where each site stores a copy of the relation in  $q_i$
- The objective is to find an optimal allocation of  $Q$  to  $S$  such that
  - ➔ the load unbalance of  $S$  is minimized
    - load( $s_i$ ): number of  $q_i$  submitted to  $s_i$
    - Ave\_load( $S$ ) =  $(1/n) \sum_{1 \leq i \leq n} \text{load}(s_i)$
    - UF( $S$ ) =  $(1/n) \sum_{1 \leq i \leq n} (\text{load}(s_i) - \text{Ave\_load}(S))^2$
  - ➔ The total communication cost is minimized

# 2-Step Algorithm

---

- For each  $q$  in  $Q$  compute load ( $S_q$ )
- While  $Q$  not empty do
  1. Select subquery  $a$  with least allocation flexibility
  2. Select best site  $b$  for  $a$  (with least load and best benefit)
  3. Remove  $a$  from  $Q$  and recompute loads if needed
- **allocation flexibility**: number of feasible allocation sites holding a copy of the relation involved in  $q$ .
- **load**: total number of subqueries in the site (existing + allocation)
- **benefit**: number of subqueries allocated to the site

# 2-Step Algorithm Example

- Let  $Q = \{q_1, q_2, q_3, q_4\}$  where  $q_1$  is associated with  $R_1$ ,  $q_2$  is associated with  $R_2$  joined with the result of  $q_1$ , etc.
- Iteration 1: select  $q_4$ , allocate to  $s_1$ , set  $\text{load}(s_1)=2$
- Iteration 2: select  $q_2$ , allocate to  $s_2$ , set  $\text{load}(s_2)=3$
- Iteration 3: select  $q_3$ , allocate to  $s_1$ , set  $\text{load}(s_1) = 3$
- Iteration 4: select  $q_1$ , allocate to  $s_3$  or  $s_4$

sites	load	$R_1$	$R_2$	$R_3$	$R_4$
$s_1$	1	$R_{11}$		$R_{31}$	$R_{41}$
$s_2$	2		$R_{22}$		
$s_3$	2	$R_{13}$		$R_{33}$	
$s_4$	2	$R_{14}$	$R_{24}$		

**Note:** if in iteration 2,  $q_2$ , were allocated to  $s_4$ , this would have produced a better plan. So hybrid optimization can still miss optimal plans