

# Tutorial 4

---

# Assignment 1

---

Remember that the deadline for assignment 1 is Friday at 5pm (i.e. the 29<sup>th</sup>)!!!

# This week's topic

---

We looked at locks and logs

More precisely, we saw 2PL, simple and more advanced locks and 3 logging schemes: Undo, Redo, Undo/redo and when and how they could be used

2PL: all locks are before every unlock in each transaction

Shared and exclusive locks: Shared is only for read and exclusive can do everything

Shared, exclusive and upgrade locks: Shared and upgrade are only for read, exclusive for everything, but you can't get an exclusive lock on an item if you have shared lock for it

Intention locks: Locks on different levels (like table, block and row level)

Undo is used to directly ensure atomicity (and if some requirements are satisfied also durability)

Redo is used to directly ensure durability (and if some requirements are satisfied also atomicity)

Undo/redo is used to ensure both

# Question 1

---

Consider the following two transactions:

$T_1$ : *read*( $X$ ); *read*( $Y$ );  $Y := X + 2Y$ ; *write*( $Y$ );

and

$T_2$ : *read*( $Y$ ); *read*( $X$ );  $X := Y + 2X$ ; *write*( $X$ );

There are multiple ways to add in shared lock, exclusive lock and unlock operations to the transactions so that they obey 2PL. Can you find a way to add in locks such that some schedules on them lead to a deadlock and another way so that no schedule on those leads to a deadlock?

# Solution for Question 1

---

A possible solution that could lead to a deadlock would be:

$T_1$  : s-lock(X);read(X);x-lock(Y);read(Y);Y := X + 2Y;write(Y);unlock(X);unlock(Y);

and

$T_2$  : s-lock(Y);read(Y);x-lock(X);read(X);X := Y + 2X;write(X);unlock(Y);unlock(X);

Those can lead to deadlocks, e.g. see the following part of a schedule:

$S$ :  $sl_1(Y)r_1(Y); sl_2(X); r_2(X); ?$

A possible solution that can't lead to a deadlock would be:

$T_1$  : x-lock(X);x-lock(Y);read(X);read(Y);Y := X + 2Y;write(Y);unlock(Y);unlock(X);

and

$T_2$  : x-lock(X);x-lock(Y);read(Y);read(X);X := Y + 2X;write(X);unlock(Y);unlock(X);

There are no deadlocks, since the first transaction to move can keep on moving (while the other transaction must wait until the first finishes).

# Question 2

---

(Exercise 17.2.4/17.2.5 in Database Systems: The complete book). The following is a sequence of undo-log records

written by two transactions,  $T_1$  and  $T_2$ :

<START  $T_1$ >

<  $T_1$ ,X,10>

<START  $T_2$ >

<  $T_2$ ,Y,20>

<  $T_1$ ,Z,30>

<  $T_2$ ,U,40>

<COMMIT  $T_2$ >

<  $T_1$ ,V,50>

<COMMIT  $T_1$  >

a) Describe the action of the recovery manager, including changes to both the database and the log on disk, if there is a system failure and the last log record to appear on disk is:

- i. <START  $T_2$ >
- ii. <COMMIT  $T_2$ >
- iii. <  $T_1$ ,V,50>
- iv. <COMMIT  $T_1$ >

b) For each of the situations i.–iv. in a), describe what values written by  $T_1$  and  $T_2$  must appear on disk? Which values might appear on disk?

# Solution for Question 2.a.i

---

In this case, the log on disk looks like this:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

Undo logging works in reverse order...

<START  $T_2$ >: No action is required by the recovery manager

< $T_1$ ,X,10>: Since the recovery manager hasn't seen a <COMMIT  $T_1$ > or an <ABORT  $T_1$ > record before it arrived at < $T_1$ ,X,10>, it changes the value of X on disk back to 10

<START  $T_1$ >: No action is required by the recovery manager

As a final step, the recovery manager appends to the log on disk the following two log records: <ABORT  $T_1$ > and <ABORT  $T_2$ >. The final log on disk is:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

<ABORT  $T_1$ >

<ABORT  $T_2$ >

This finishes the work of the recovery manager.

# Solution for Question 2.a.ii

---

In this case, the log on disk looks like this:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

< $T_2$ ,Y,20>

< $T_1$ ,Z,30>

< $T_2$ ,U,40>

<COMMIT  $T_2$ >

Undo logging works in reverse order...

<COMMIT  $T_2$ >: remember that  $T_2$  has finished

< $T_2$ ,U,40>:  $T_2$  has finished so skip

< $T_1$ ,Z,30>:  $T_1$  has **not** finished, so overwrite Z with 30

< $T_2$ ,Y,20>:  $T_2$  has finished so skip

<START  $T_1$ >: No action required

< $T_1$ ,X,10>:  $T_1$  has **not** finished, so overwrite X with 10

<START  $T_2$ >: No action required

As a final step, the recovery manager appends to the log on disk the following log record: <ABORT  $T_1$ >. We end with:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

< $T_2$ ,Y,20>

< $T_1$ ,Z,30>

< $T_2$ ,U,40>

<COMMIT  $T_2$ >

<ABORT  $T_1$ >



# Solution for Question 2.a.iii

---

In this case, the log on disk looks like this:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

< $T_2$ ,Y,20>

< $T_1$ ,Z,30>

< $T_2$ ,U,40>

<COMMIT  $T_2$ >

< $T_1$ ,V,50>

Undo logging works in reverse order...

< $T_1$ ,V,50>:  $T_1$  has **not** finished, so overwrite V with 50

<COMMIT  $T_2$ >: remember that  $T_2$  has finished

< $T_2$ ,U,40>:  $T_2$  has finished so skip

< $T_1$ ,Z,30>:  $T_1$  has **not** finished, so overwrite Z with 30

< $T_2$ ,Y,20>:  $T_2$  has finished so skip

<START  $T_1$ >: No action required

< $T_1$ ,X,10>:  $T_1$  has **not** finished, so overwrite X with 10

<START  $T_2$ >: No action required

As a final step, the recovery manager appends to the log on disk the following log record: <ABORT  $T_1$ >. We end with:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

< $T_2$ ,Y,20>

< $T_1$ ,Z,30>

< $T_2$ ,U,40>

<COMMIT  $T_2$ >

< $T_1$ ,V,50>

<ABORT  $T_1$ >

# Solution for Question 2.a.vi

---

In this case, the log on disk looks like this:

<START  $T_1$ >

< $T_1$ ,X,10>

<START  $T_2$ >

< $T_2$ ,Y,20>

< $T_1$ ,Z,30>

< $T_2$ ,U,40>

<COMMIT  $T_2$ >

< $T_1$ ,V,50>

<COMMIT  $T_1$ >

Undo logging works in reverse order...

<COMMIT  $T_1$ >: remember that  $T_1$  has finished

< $T_1$ ,V,50>:  $T_1$  has finished, so skip

<COMMIT  $T_2$ >: remember that  $T_2$  has finished

< $T_2$ ,U,40>:  $T_2$  has finished, so skip

< $T_1$ ,Z,30>:  $T_1$  has finished, so skip

< $T_2$ ,Y,20>:  $T_2$  has finished so skip

<START  $T_1$ >: No action required

< $T_1$ ,X,10>:  $T_1$  has finished, so skip

<START  $T_2$ >: No action required

# Solution for Question 2.b

---

**b.i:** X might have been written to disk, but without  $\langle \text{COMMIT } T_1 \rangle$  we can't be sure

**b.vi:** U, V, X, Y and Z have been written to disk, since we see  $\langle \text{COMMIT } T_1 \rangle$  and  $\langle \text{COMMIT } T_2 \rangle$

**b.ii:** X and Z might have been written to disk, but without  $\langle \text{COMMIT } T_1 \rangle$  we can't be sure. U and Y has been written to disk, since we see  $\langle \text{COMMIT } T_2 \rangle$

**b.iii:** V, X and Z might have been written to disk, but without  $\langle \text{COMMIT } T_1 \rangle$  we can't be sure. U and Y has been written to disk, since we see  $\langle \text{COMMIT } T_2 \rangle$

# Question 3

---

(Exercise 19.1.2/19.1.3 in Database Systems: the complete book). Consider the following schedules:

$S_1: r_1(X); r_2(Y); w_1(Y); w_2(Z); r_3(Y); r_3(Z); w_3(U)$

$S_2: r_1(X); w_1(Y); r_2(Y); w_2(Z); r_3(Z); w_3(U)$

$S_3: r_2(X); r_3(X); r_1(X); r_1(Y); r_2(Y); r_3(Y); w_2(Z); r_3(Z)$

$S_4: r_2(X); r_3(X); r_1(X); w_1(Y); r_3(Y); w_2(Z); r_3(Z)$

- a) Suppose that each of the schedules is followed by an abort operation for transaction  $T_1$ . Tell which transactions need to be rolled back.

- b) Now suppose that all three transactions commit and write their commit record on the log immediately after their last operation. However, a crash occurs, and a tail of the log was not written to disk before the crash and is therefore lost. Tell, depending on where the lost tail of the log begins:

- i. What transactions could be considered uncommitted?
- ii. Are any dirty reads created during the recovery process? If so, what transactions need to be rolled back?
- iii. What additional dirty reads could have been created if the portion of the log lost was not a tail, but rather some portions in the middle?

# Solution to question 3

---

**3.a:**  $S_1$ :  $T_1$  and  $T_3$  have to be rolled back:  $T_1$  because it might have written some values to disk, and  $T_3$  because it reads a value, Y, that was written before by  $T_1$ . No other transaction needs to be rolled back

$S_2$ : All three transactions have to be rolled back, because  $T_2$  reads a value, Y, written by  $T_1$ , and  $T_3$  reads a value, Z, written by  $T_2$

$S_3$ : No transaction needs to be rolled back. Note that  $T_1$  does not need to be rolled back, because it did not write any values.

$S_4$ : Same as for  $S_1$

**3.b:** No complete solution provided. Let us say transactions finishes so that transaction 1 finishes before 2, before 3. For each schedule, we distinguish the following possibilities for the time of the crash:

1. before the commit record of the first transaction in the schedule reaches disk;
2. after the commit record of the first transaction in the schedule reaches disk, but before the commit record of the second transaction in the schedule reaches disk;
3. after the commit record of the second transaction in the schedule reaches disk, but before the commit record of the third transaction in the schedule reaches disk;
4. after the commit record of the third transaction in the schedule reaches disk.

# Solution to question 3 continued

---

For part i. of the question: In case 1. “before the commit record of the first transaction in the schedule reaches disk;”, all transactions would be considered uncommitted.

In case 2, “after the commit record of the first transaction in the schedule reaches disk, but before the commit record of the second transaction in the schedule reaches disk;” only the second and the third transaction in the schedule would be considered uncommitted

In case 3, “after the commit record of the second transaction in the schedule reaches disk, but before the commit record of the third transaction in the schedule reaches disk;”, only the third transaction in the schedule would be considered uncommitted

In case 4, “after the commit record of the third transaction in the schedule reaches disk”, no transaction would be considered uncommitted

For part ii.: If a transaction reads an item that was written by a transaction that is considered uncommitted, then that is a dirty read. In this case, that transaction needs to be rolled back. All transactions who depend on that transaction also need to be rolled back and so on (cascading rollback)

For part iii.: No solution provided. For discussion.