

# Other kinds of logging

---

# Overview over this video

---

In this video, we will look at other kinds of logging (redo and undo/redo) and how to get atomicity and durability from logging

# Redo Logging

---

Logs activities with the goal of *restoring* committed transactions (ignores incomplete transactions).

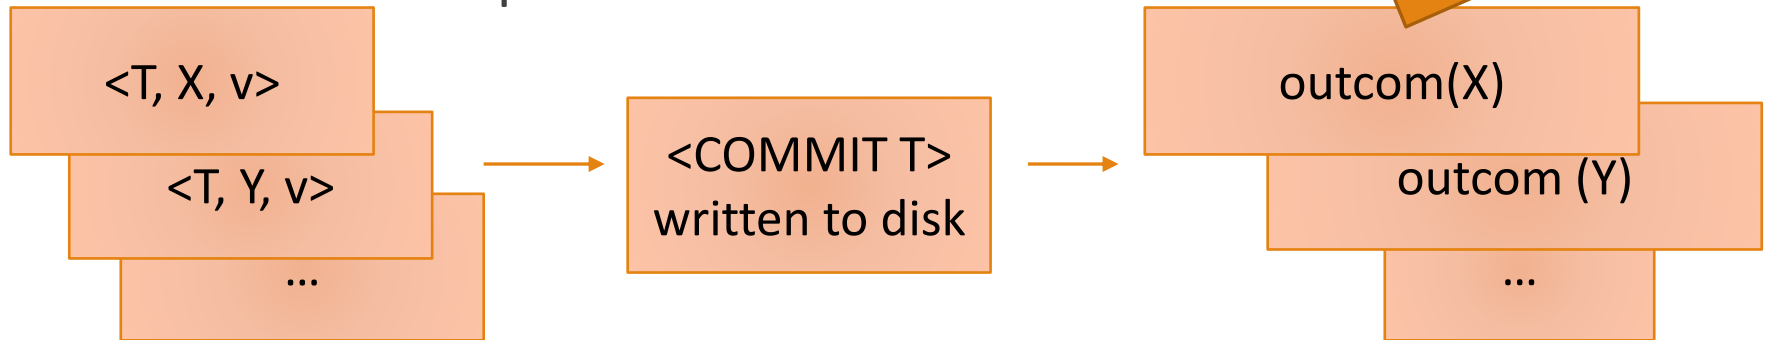
Log records:

- Same as before, but...
- New meaning of **<T, X, v>**: “Transaction T has updated the value of database item X & the new value of X is v.”
  - Direct response to **write(X)**
  - Haven’t changed X on disk yet!

Have to modify the logging procedure...

# Redo Logging: Procedure

1. T first writes all log records for all updates
2. T writes  $\langle \text{COMMIT } T \rangle$  to the log on disk
3. T writes all committed updates to disk

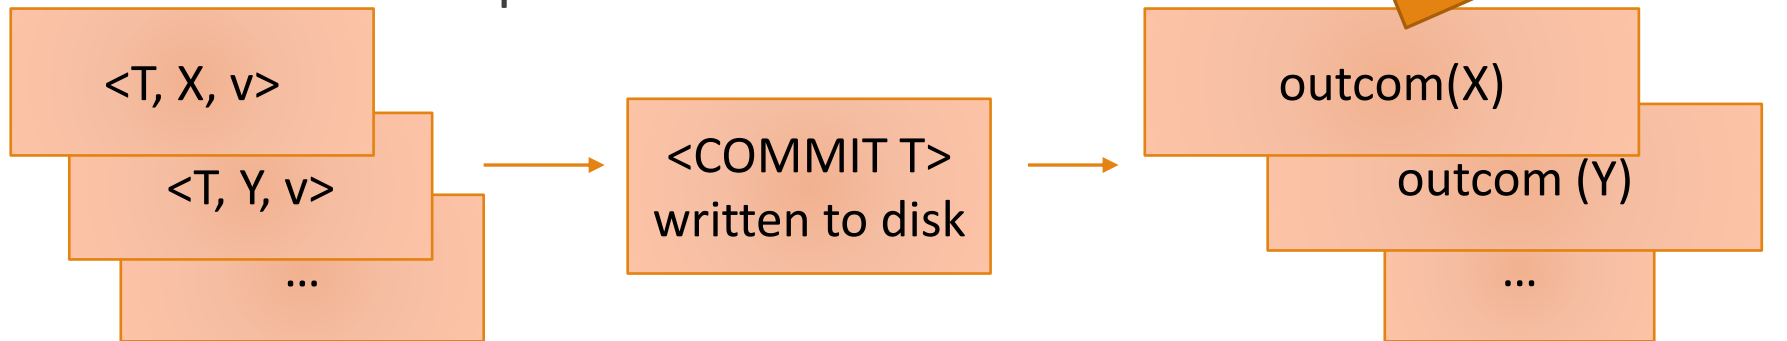


# Example

		Local		Buffer		Database			
Time	Transaction	X	Y	X	Y	X	Y	Log (buffer)	Log (disk)
0						1	10	<START T>	
1	read(X)	1		1		1	10		
2	X := X*2	2		1		1	10		
3	write(X)	2		2		1	10	<T, X, 2>	
4	read(Y)	2	10	2	10	1	10		
5	Y := Y*2	2	20	2	10	1	10		
6	write(Y)	2	20	2	20	1	10	<T, Y, 20>	
7								<COMMIT T>	
8	flush_log								
9	outcom(X)	2	20	2	20	2	10		
10	outcom(Y)	2	20	2	20	2	20		

# Redo Logging: Procedure

1. T first writes all log records for all updates
2. T writes  $\langle \text{COMMIT } T \rangle$  to the log on disk
3. T writes all committed updates to disk



## Fundamental property of redo logs:

- $\langle \text{COMMIT } T \rangle$  occurs in log  $\rightarrow$  log contains complete information on T
- $\langle \text{COMMIT } T \rangle$  doesn't occur in log  $\rightarrow$  T hasn't written anything to disk

# Recovery With Redo Logs

---

Essentially: reverse of undo logging

Procedure:

- Identify all the transactions with a **COMMIT** log record.
- Traverse the log from first to the last item.
- If we see  $\langle T, X, v \rangle$  and T has a COMMIT log record, then change the value of X on disk to v.
- For each *incomplete* transaction T, write **<ABORT T>** into the log on disk.

# Example with Redo Logging

Time	Transaction T <sub>1</sub>	Transaction T <sub>2</sub>
1	read(X)	
2	X := X * 2	
3	write(X)	
4		read(X)
5	read(Y)	
6		X := X * 3
7		write(X)
8	Y := X + Y	
9	write(Y)	

X = 1  
Y = 2

How does redo logging work on this schedule?

- Which **log entries** are written to buffer/disk & when?
- Which **other operations** must be executed & when?



Time	Transaction T <sub>1</sub>	Transaction T <sub>2</sub>	Log (buffer)	Log (disk)
0			<START T <sub>1</sub> >	
1	read(X)			
2	X := X * 2			
3	write(X)		<T <sub>1</sub> , X, 2>	
4			<START T <sub>2</sub> >	
5		read(X)		
6	read(Y)			
7		X := X * 3		
8		write(X)	<T <sub>2</sub> , X, 6>	
9	Y := X + Y			
10	write(Y)		<T <sub>1</sub> , Y, 4>	
11			<COMMIT T <sub>1</sub> >	
12	flush_log			
13	outcom(X)			
14	outcom(Y)			
15			<COMMIT T <sub>2</sub> >	
16		flush_log		
17		outcom(X)		
18				
19				

X = 1  
Y = 2

# Undo/Redo Logging

---

Good properties of undo logging and redo logging

Log records:

- Same as before, but replace  $\langle T, X, v \rangle$
- $\langle T, X, v, w \rangle$ : “Transaction T has updated the value of database item X, and the **old/new** value of X is v/w.”

Procedure:

- Write all log records for all updates to database items first
- Then write updates to disk
- $\langle \text{COMMIT } T \rangle$  can be written to disk before or after all changes have been written to disk

Recovery needs to process log in both directions

# Example with Undo/Redo Logging

Time	Transaction T <sub>1</sub>	Transaction T <sub>2</sub>
1	read(X)	
2	X := X * 2	
3	write(X)	
4		read(X)
5	read(Y)	
6		X := X * 3
7		write(X)
8	Y := X + Y	
9	write(Y)	

X = 1  
Y = 2

How does undo/redo logging work on this schedule?

- Which **log entries** are written to buffer/disk & when?
- Which **other operations** must be executed & when?

Time	Transaction T <sub>1</sub>	Transaction T <sub>2</sub>	Local T <sub>1</sub>		Local T <sub>2</sub>		Buffer		Disk		Buffer log
			X	Y	X	Y	X	Y	X	Y	
0									1	2	<START T <sub>1</sub> >
1	read(X)		1				1		1	2	
2	X := X * 2		2				1		1	2	
3	write(X)		2				2		1	2	<T <sub>1</sub> , X, 1, 2>
4			2				2		1	2	<START T <sub>2</sub> >
5		read(X)	2		2		2		1	2	
6	read(Y)		2	2	2		2	2	1	2	
7		X := X * 3	2	2	6		2	2	1	2	
8		write(X)	2	2	6		6	2	1	2	<T <sub>2</sub> , X, 2, 6>
9	Y := X + Y		2	4	6		6	2	1	2	
10	write(Y)		2	4	6		6	4	1	2	<T <sub>1</sub> , Y, 2, 4>
11			2	4	6		6	4	1	2	<COMMIT T <sub>1</sub> >
12	flush_log		2	4	6		6	4	1	2	
13	output(X)		2	4	6		6	4	6	2	
14	output(Y)		2	4	6		6	4	6	4	
15			2	4	6		6	4	6	4	<COMMIT T <sub>2</sub> >
16		flush_log	2	4	6		6	4	6	4	
17		output(X)	2	4	6		6	4	6	4	

Why are DBMS using Undo/Redo?

# Undo without Redo

---

Undo essentially ensures Atomicity

Can ensure durability using Force

- **Force** the writing of updates to disk before commit
- (**No Force** is not to require this)
- Force is expensive in disk operations

# Redo without Undo

---

Redo essentially ensures Durability

Can ensure atomicity using No Steal

- **No Steal** means that uncommitted data may not overwrite committed data on disk
- (**Steal** is not to require this)
- **No Steal** is expensive to ensure

# Ensuring atomicity and durability

---

Could ensure Atomicity and Durability without log using No Steal/Force

- Very hard and expensive to ensure
- (Must write every change to disk made by the transaction while performing the commit statement!)

In practice:

- Want Steal/No Force (cheapest in time) → Use Undo/Redo

# Summary

---

We covered Redo logging and Undo/Redo logging

- i.e. basically the kind of logging you should do to be able to do redos or undos and redos

Also covered why DBMS uses Undo/Redo logging (even if it is the more complex option)