

# (Undo-)Logging for your database

---

# Overview of this video

---

We will try to handle failures using a simple kind of logs, specifically using something called an undo-log

# Logging in DBMS

---

Idea: write important activities to a log so that a desired database state can be recovered later

Examples of important activities:

- Starts of transactions, commits, aborts
- Modification of database items


Should work *even in case of system failures!*

Techniques:


- Undo logging (for maintaining *Atomicity*)
- Redo logging (for maintaining *Durability*)
- Combinations thereof (specifically Undo/Redo for doing both *Atomicity* and *Durability*)

# Extension of Transaction Syntax

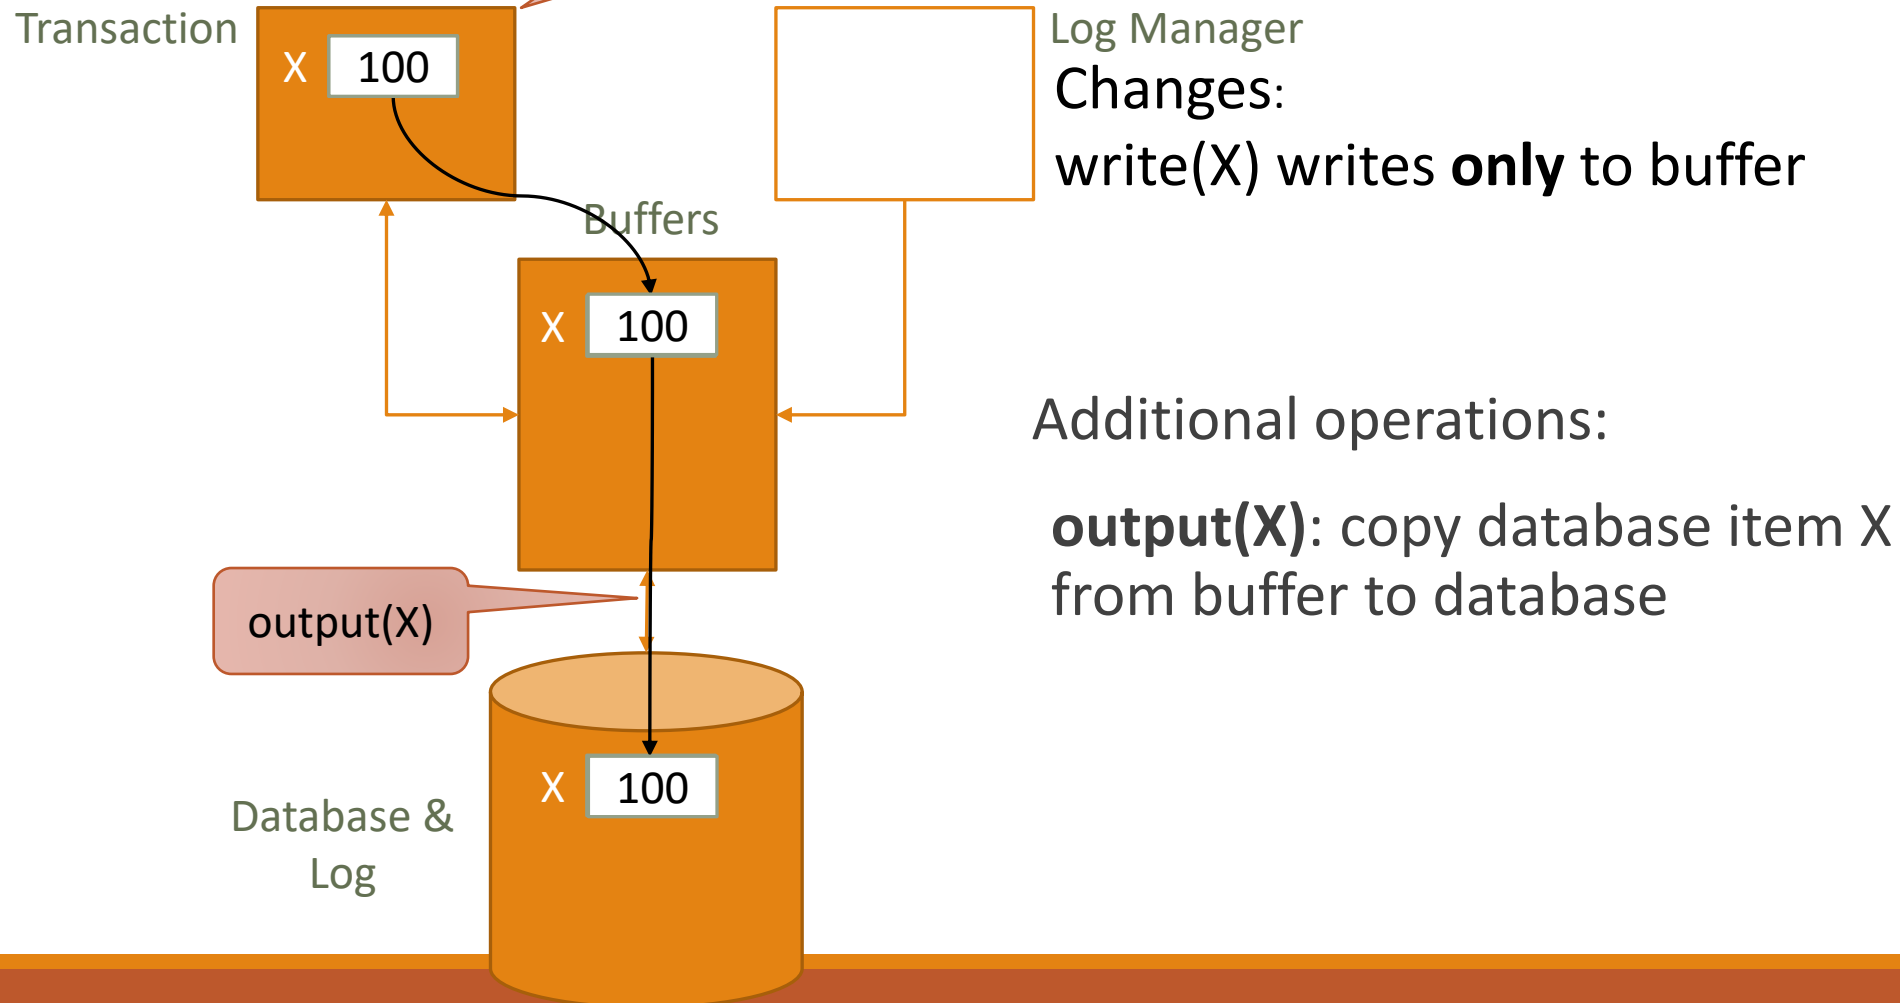
---



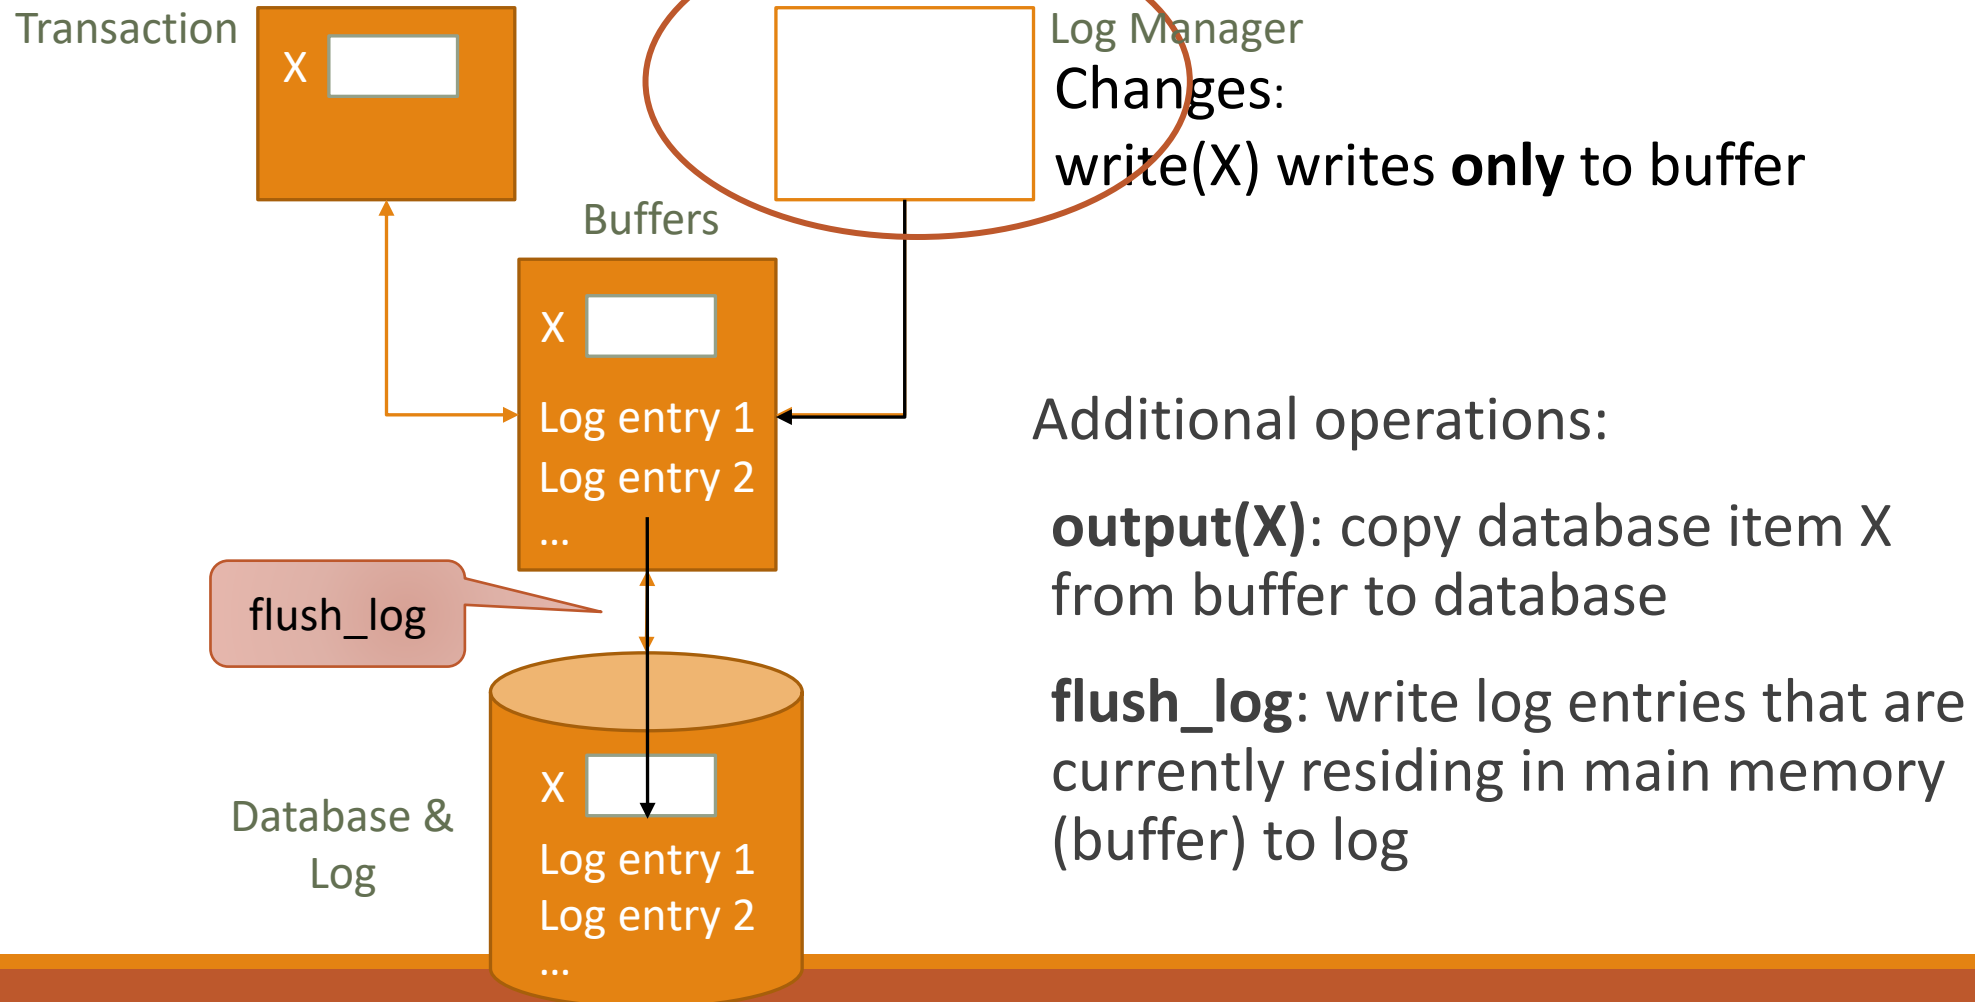
# Extension Transaction Syntax



# Extension Transaction Syntax



# Extension of Transaction Syntax



# Undo Logging

---

Logs activities with the goal of restoring (“undoing”) a previous database state.

Log records (or log entries):

- **<START T>**: Transaction T has started.
- **<COMMIT T>**: Transaction T has committed.
- **<ABORT T>**: Transaction T was aborted.
- **<T, X, v>**: Transaction T has updated the value of database item X, and the old value of X was v.
  - Response to **write(X)**
  - If this entry occurs in the log, then the new value of X might not have been written to the database yet.


Slightly different records for redo- and undo/redo logging (later video...)




# Undo Logging: Procedure

---

1. If transaction T updates database item X and the old value was v, then  $\langle T, X, v \rangle$  must be written to the log on disk **before** X is written to disk.



2. If transaction T commits, then  $\langle \text{COMMIT } T \rangle$  must be written to disk **as soon as** all database elements changed by T have been written to disk



# Example

---

Transaction
read(X)
$X := X * 2$
write(X)
read(Y)
$Y := Y * 2$
write(Y)
<b>flush_log</b>
<b>output(X)</b>
<b>output(Y)</b>
<b>flush_log</b>

Writes all log entries for updates to disk

Writes all updates to disk

Writes the <COMMIT T> record to disk

# Example

		Local		Buffer		Database		
Time	Transaction	X	Y	X	Y	X	Y	Log record written
0						1	10	<START T>
1	read(X)	1		1		1	10	
2	X := X*2	2		1		1	10	
3	write(X)	2		2		1	10	<T, X, 1>
4	read(Y)	2	10	2	10	1	10	
5	Y := Y*2	2	20	2	10	1	10	
6	write(Y)	2	20	2	20	1	10	<T, Y, 10>
7	flush_log	2	20	2	20	1	10	
8	output(X)	2	20	2	20	2	10	
9	output(Y)	2	20	2	20	2	20	
10		2	20	2	20	2	20	<COMMIT T>
11	flush_log	2	20	2	20	2	20	

---

What if a system failure occurs?  
(The Recovery Manager)

# Scenario 1

<COMMIT T> occurs in the log on disk

Time	Transaction	X	Y	X	Y	X	Y	
0						1	10	<START T>
1	read(X)	1		1		1	10	
2	X := X*2	2		1		1	10	
3	write(X)	2		2		1	10	<T, X, 1>
4	read(Y)	2	10	2	10	1	10	
5	Y := Y*2	2	20	2	10	1	10	
6	write(Y)	2	20	2	20	1	10	<T, Y, 10>
7	flush_log	2	20	2	20	1	10	
8	output(X)	2	20	2	20	2	10	
9	output(Y)	2	20	2	20	2	20	
10		2	20	2	20	2	20	<COMMIT T>
11	flush_log	2	20	2	20	2	20	

T has committed successfully  
(no recovery needed)

Error

# Scenario 2

<COMMIT T> **does not occur** in the log on disk

Time	Transaction	X	Y	X	Y			
0								
1	read(X)	1		1				
2	X := X*2	2		1				
3	write(X)	2		2		1	10	<T, X, 1>
4	read(Y)	2	10	2	10	1	10	
5	Y := Y*2	2	20	2	10	1	10	
6	write(Y)	2	20	2	20	1	10	<T, Y, 10>
7	flush_log	2	20	2	20	1	10	
8	output(X)	2	20	2	20	2	10	
9	output(Y)	2	20	2	20	2	20	
10		2	20	2	20	2	20	<COMMIT T>
11	flush_log	2	20	2	20	2	20	

Must undo all updates to database items that were written to disk.

For each log record <T, X, v> on disk, replace X on disk by v.

Error

# Scenario 3

<COMMIT T> does not occur in the log on disk

Time	Transaction	X	Y	X	Y	X	Y	Log
0						1	10	<START T>
1	read(X)	1		1		1	10	
2	X := X*2	2		1		1	10	
3	write(X)	2		2		1	10	<T, X, 1>
4	read(Y)	2	10	2	10	1	10	
5	Y := Y*2	2	20	2	10	1	10	
6	write(Y)	2	20	2	20	1	10	<T, Y, 10>
7	flush_log	2	20	2	20	1	10	
8	output(X)	2	20	2	20	2	10	
9	output(Y)	2	20	2	20	2	20	
10		2	20	2	20	2	20	<COMMIT T>
11	flush_log	2	20	2	20	2	20	

What to do in this case? Nothing!

Error

# Recovery With Undo Logs

(Simple Variant)

---

If an error occurs, the recovery manager restores the last consistent database state

Traverses the undo log backwards

If the current entry is...

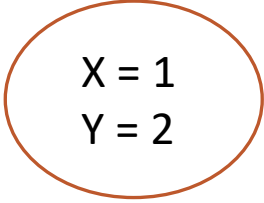
- **<COMMIT T>**: remember that T was committed successfully
- **<ABORT T>**: remember that T was aborted
- **<T, X, v>**: if T has not finished successfully (no COMMIT, no ABORT), change the value of X on disk to v

Write **<ABORT T>** for each *uncommitted* transaction T that was not previously aborted & call `flush_log`



# Example of Undo Logging

Time	Transaction T <sub>1</sub>	Transaction T <sub>2</sub>
1	read(X)	
2	X := X * 2	
3	write(X)	
4		read(X)
5	read(Y)	
6		X := X * 3
7		write(X)
8	Y := X + Y	
9	write(Y)	



X = 1  
Y = 2

How does undo logging work on this schedule?

- Which **log entries** are written to buffer/disk & when?
- Which **other operations** must be executed & when?

Time	Transaction T <sub>1</sub>	Transaction T <sub>2</sub>	Log (buffer)	Log (disk)
0			<START T <sub>1</sub> >	
1	read(X)			
2	X := X * 2			
3	write(X)		<T <sub>1</sub> , X, 1>	
4			<START T <sub>2</sub> >	
5		read(X)		
6	read(Y)			
7		X := X * 3		
8		write(X)	<T <sub>2</sub> , X, 2>	
9	Y := X + Y			
10	write(Y)		<T <sub>1</sub> , Y, 2>	
11	<b>flush_log</b>			
12	<b>output(X)</b>			
13	<b>output(Y)</b>			
14			<COMMIT T <sub>1</sub> >	
15	<b>flush_log</b>			
16		<b>flush_log</b>		
17		<b>output(X)</b>		
18			<COMMIT T <sub>2</sub> >	
19		<b>flush_log</b>		

X = 1  
Y = 2

What if a transaction aborts?

# If a Transaction Aborts...

---

Use the undo log to undo all changes made by the transaction.

- Similar to recovery with undo logs
- But focuses on a single transaction

Procedure:

- Assume T aborts
- Traverse the undo log from the last to the first item
- If we see  $\langle T, X, v \rangle$ , change the value of X on disk back to v.

# Summary

---

Aborted transactions & system failures can be dealt with using careful

Undo logging

- Maintains Atomicity
- Logs old value for each updated database item
- Recovery manager: use this information to restore the last consistent database state