

# Timestamp based schedules

---

# Overview of this video

---

In the last video, we saw how we could add deadlock-prevention into our Strict 2PL system, by among other approaches, use time stamps

This video, we try to make schedules that ensure no deadlocks directly, by using time-stamps!

# Deadlock prevention

---

Two approaches for deadlock prevention:

- **Detect deadlocks & fix them**
- **Enforce deadlock-free schedules**

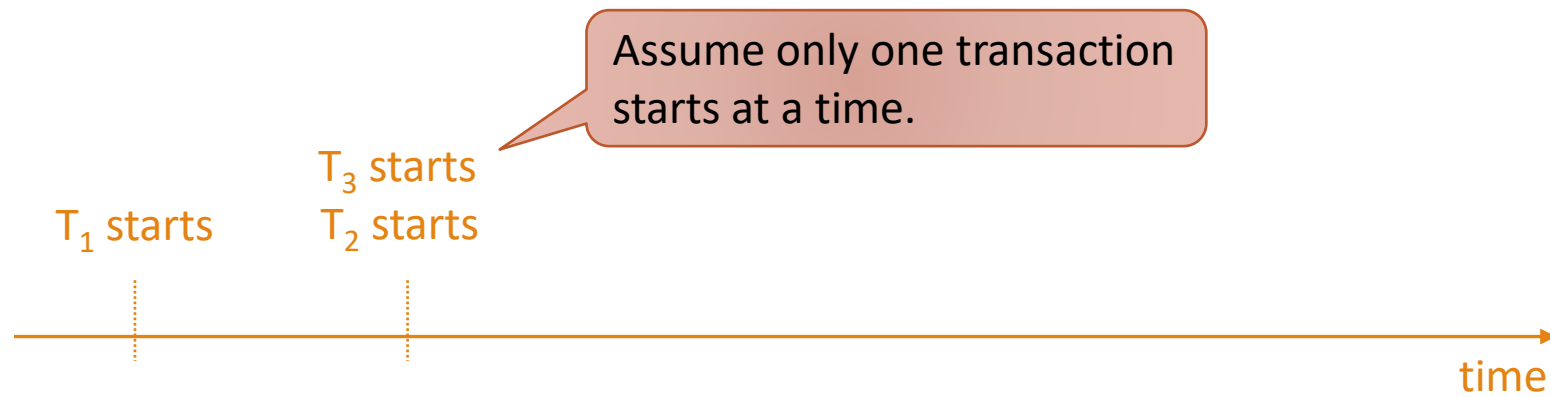


Uses timestamps too!

# Basic Idea

---

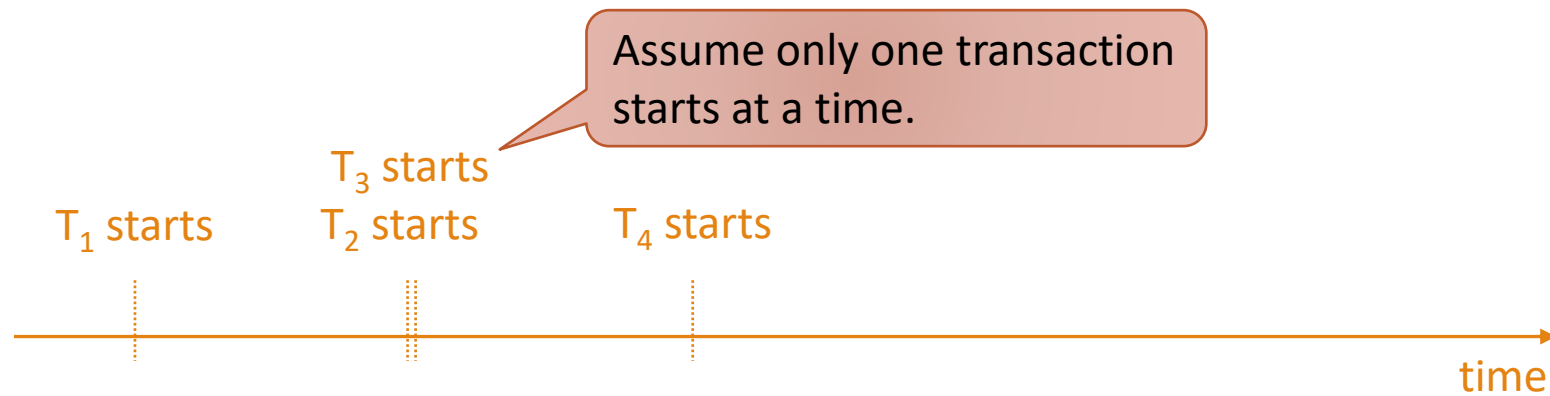
Schedule transactions so that the effect is the same as **executing each transaction instantaneously** when it is started.



# Basic Idea

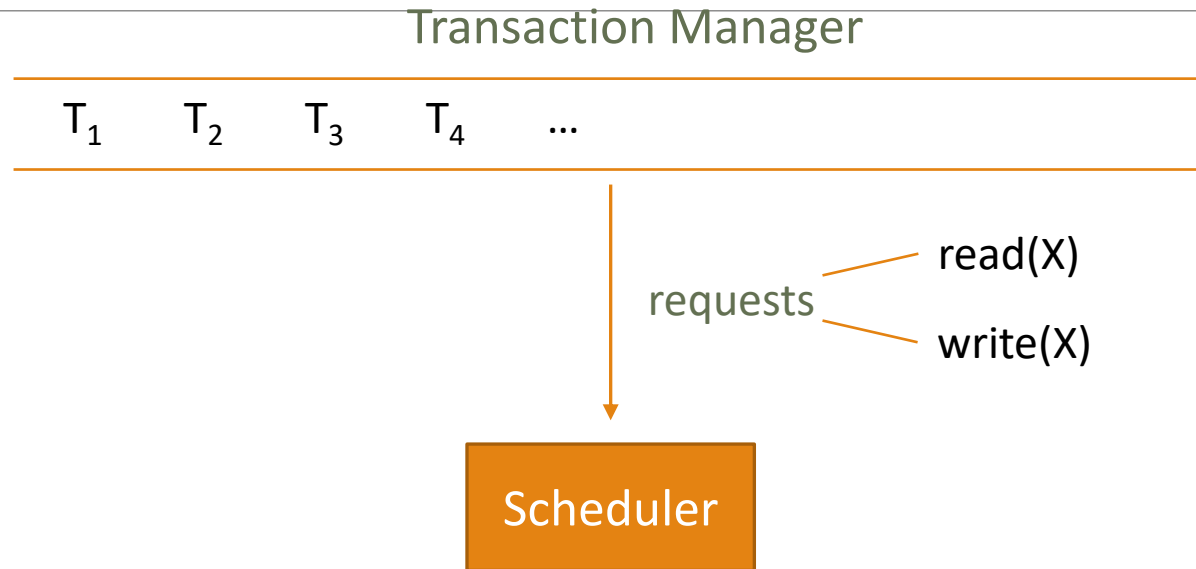
---

Schedule transactions so that the effect is the same as **executing each transaction instantaneously** when it is started.



Equivalent to serial schedule that has all transactions in the order of their start time.

# Timestamp-Based Schedulers



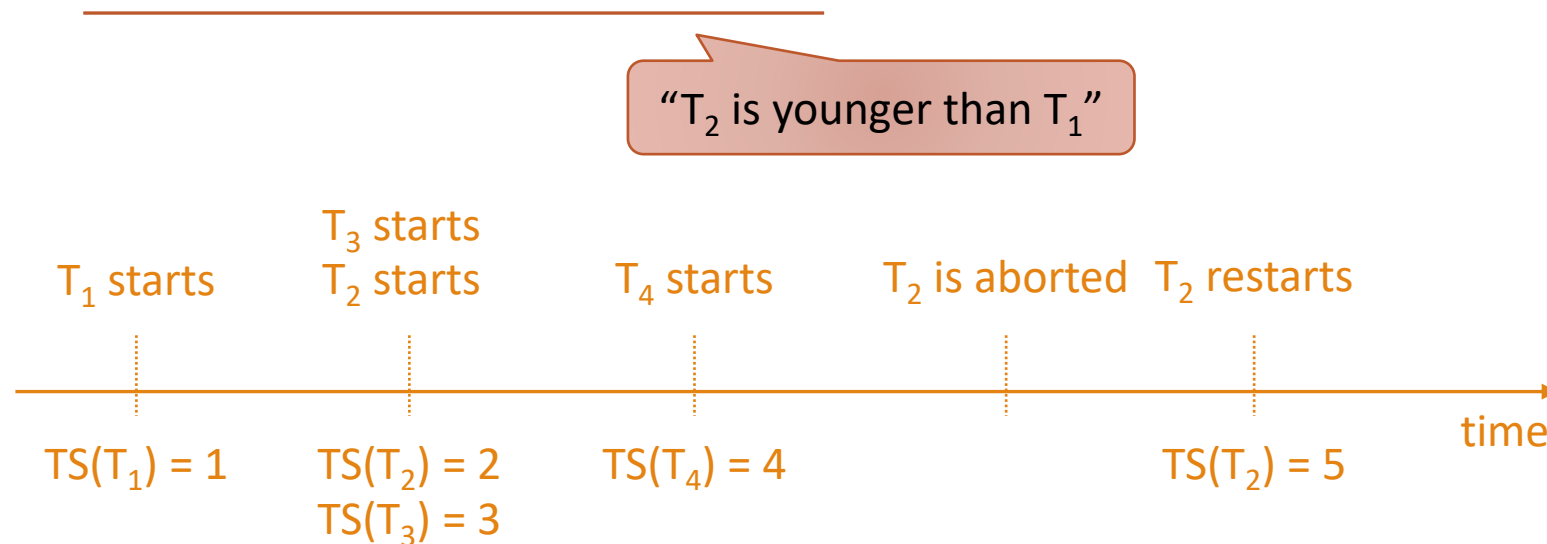
## Possible actions of the scheduler:

- Grant request
- Abort transaction
- Delay transaction

# Timestamps

Each transaction **T** is assigned a unique integer **TS(T)** when it starts (the **timestamp of T**).

If  $T_1$  started earlier than  $T_2$ , we require  $TS(T_1) < TS(T_2)$



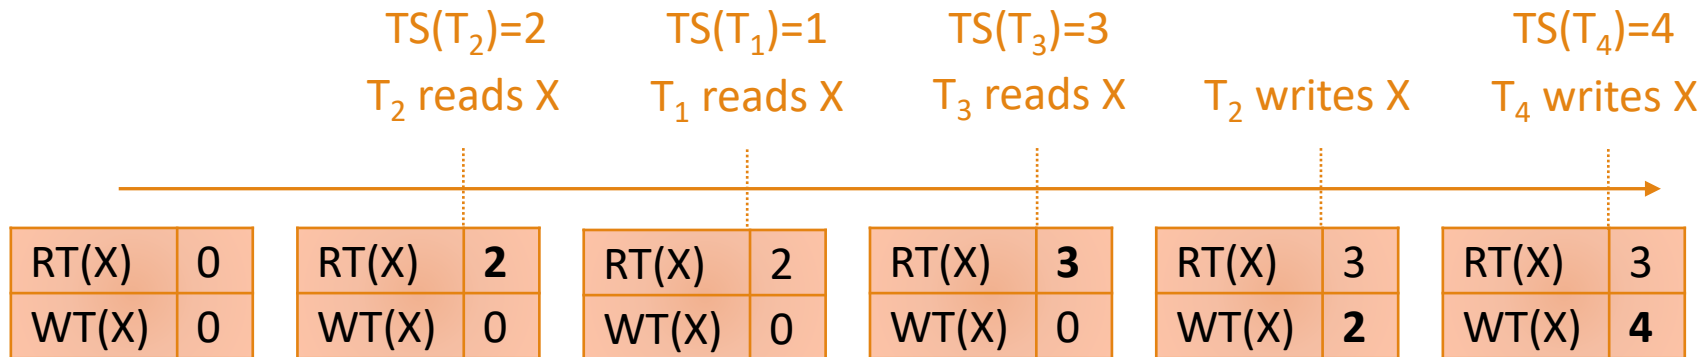
We assign a **new timestamp even after a restart!**

Recall: this is different from detection! There, restarts did not change timestamps!

# Additional Bookkeeping

For each database item X, maintain:

- **Read Time of X: RT(X)**  
Timestamp of youngest transaction that read X
- **Write Time of X: WT(X)**  
Timestamp of youngest transaction that wrote X

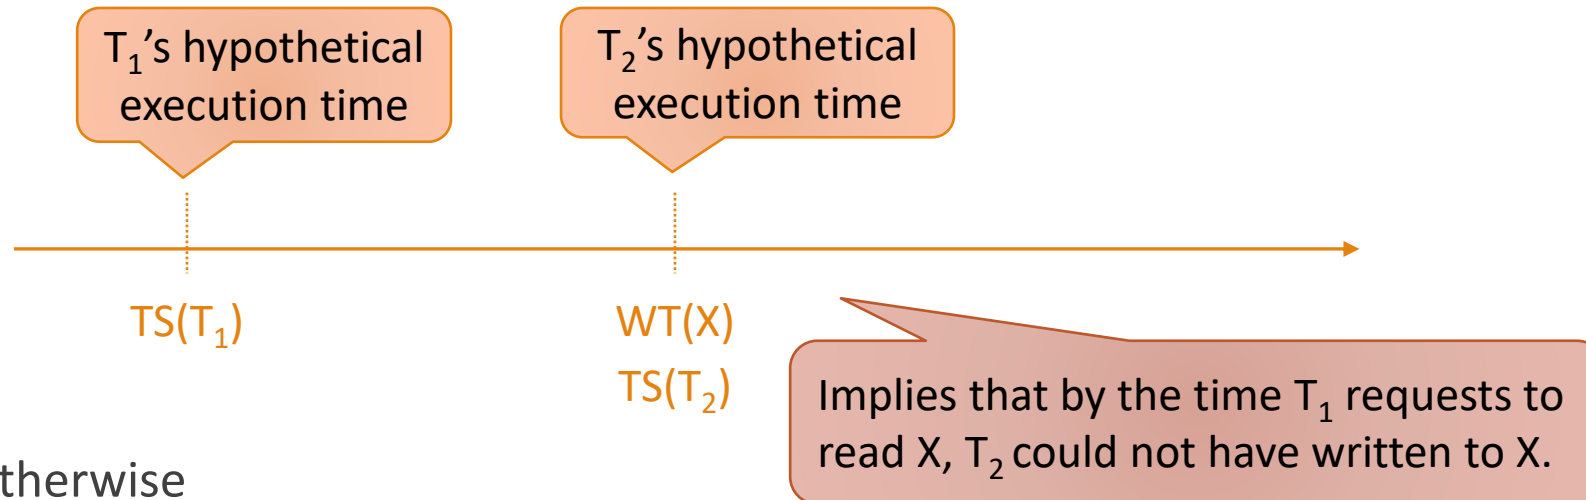




# Read Requests

If  $T_1$  requests to **read** X:

**Abort & restart  $T_1$**  if  $WT(X) > TS(T_1)$

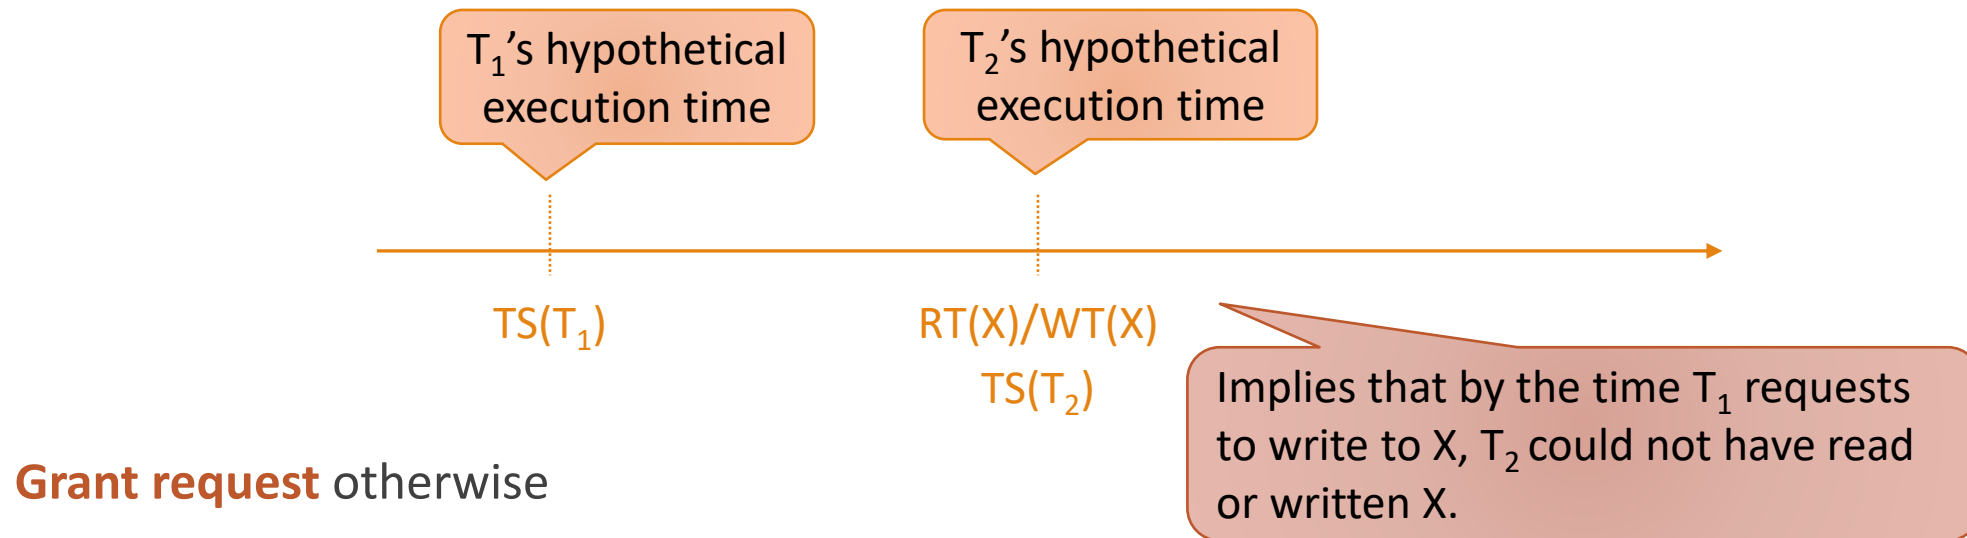


**Grant request** otherwise

# Write Requests

If  $T_1$  requests to **write** X:

**Abort & restart  $T_1$**  if  $RT(X) > TS(T_1)$  or  $WT(X) > TS(T_1)$



# Example 1

Time	$T_1$ (TS = 1)	$T_2$ (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1						
2						
3						
4						
5						
6						
7						
8						
9						
10						
11						

# Example 1

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2						
3						
4						
5						
6						
7						
8						
9						
10						
11						

X last  
read



# Example 1


Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3						
4						
5						
6						
7						
8						
9						
10						
11						

X last  
read




# Example 1

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4						
5						
6						
7						
8						
9						
10						
11						




# Example 1

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5						
6						
7						
8						
9						
10						
11						



# Example 1

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6						
7						
8						
9						
10						
11						






# Example 1

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)					
7						
8						
9						
10						
11						


Must abort  
& restart T<sub>1</sub>



# Exam

New timestamp


Time	$T_1$ (TS = 3)	$T_2$ (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	$X := X + 100$		1	0	0	0
3		read(Y)	1	0	2	0
4		$Y := Y * 2$	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		0	0	2	2
7						
8						
9						
10						
11						



# Exam


New timestamp

Time	$T_1$ (TS = 3)	$T_2$ (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	$X := X + 100$		1	0	0	0
3		read(Y)	1	0	2	0
4		$Y := Y * 2$	1	0	2	0
5		write(Y)	1	0	2	2
6	<del>read(Y)</del>		0	0	2	2
7	read(X)		3	0	2	2
8						
9						
10						
11						




# Example 1

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	<del>read(Y)</del>		0	0	2	2
7	read(X)		3	0	2	2
8	X := X + 100		3	0	2	2
9						
10						
11						




# Example 1

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	<del>read(Y)</del>		0	0	2	2
7	read(X)		3	0	2	2
8	X := X + 100		3	0	2	2
9	read(Y)		3	0	3	2
10						
11						




# Example 1

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	<del>read(Y)</del>		0	0	2	2
7	read(X)		3	0	2	2
8	X := X + 100		3	0	2	2
9	read(Y)		3	0	3	2
10	Y := Y * 3		3	0	3	2
11						



# Example 1

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	<del>read(Y)</del>		0	0	2	2
7	read(X)		3	0	2	2
8	X := X + 100		3	0	2	2
9	read(Y)		3	0	3	2
10	Y := Y * 3		3	0	3	2
11	write(Y)		3	0	3	3



# Example 2

---

T <sub>1</sub>	T <sub>2</sub>
read(X)	read(Y)
...	...
write(X)	write(Y)
read(Y)	read(X)
...	...
write(Y)	write(X)

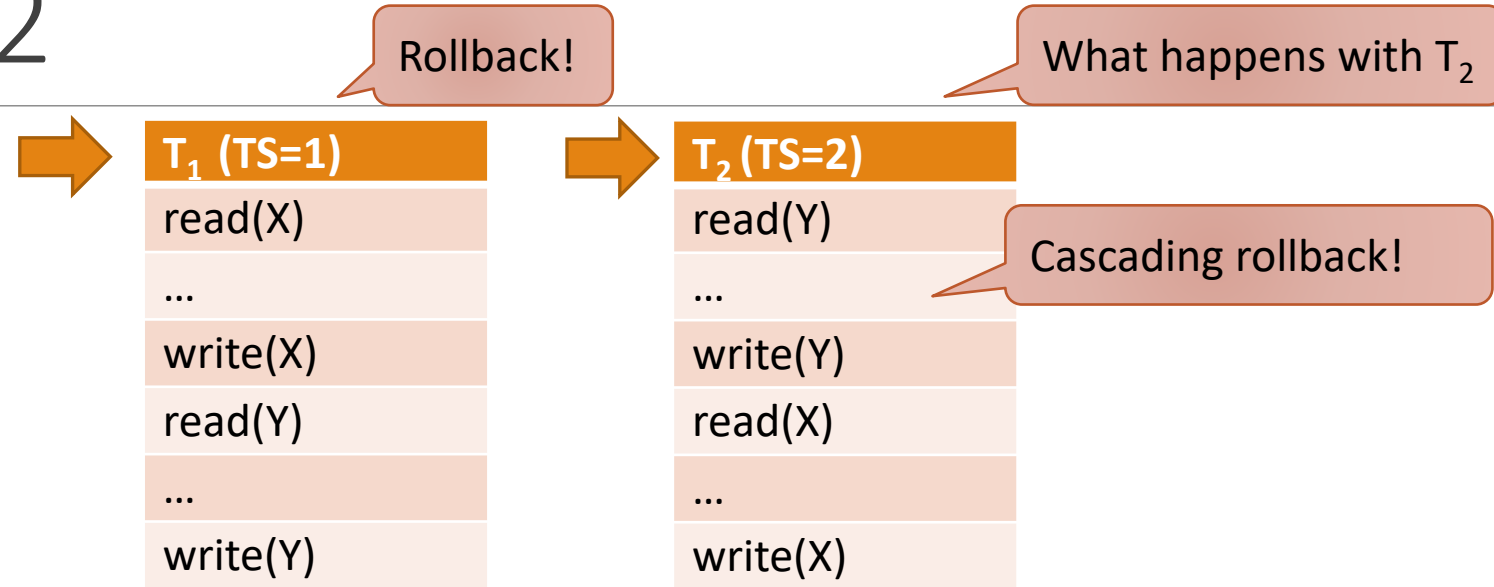
Assume a timestamp-based scheduler

- T<sub>1</sub> starts first
- Lines 1-3 of T<sub>1</sub> are executed first, then lines 1-3 of T<sub>2</sub>

Which operations could be executed next?



# Example 2

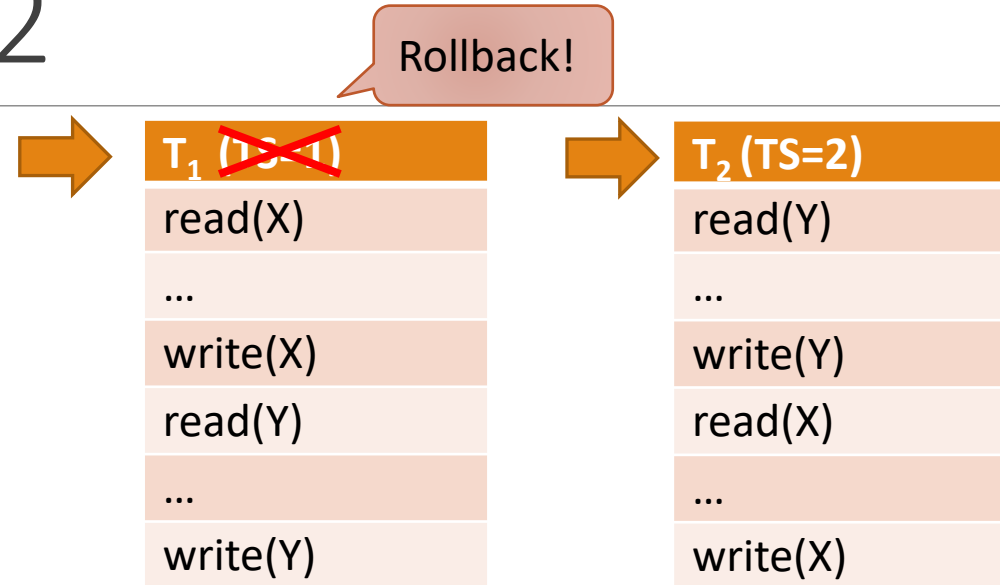


Assume a timestamp-based scheduler

- $T_1$  starts first
- Lines 1-3 of  $T_1$  are executed first, then lines 1-3 of  $T_2$

Which operations could be executed next?

# Example 2

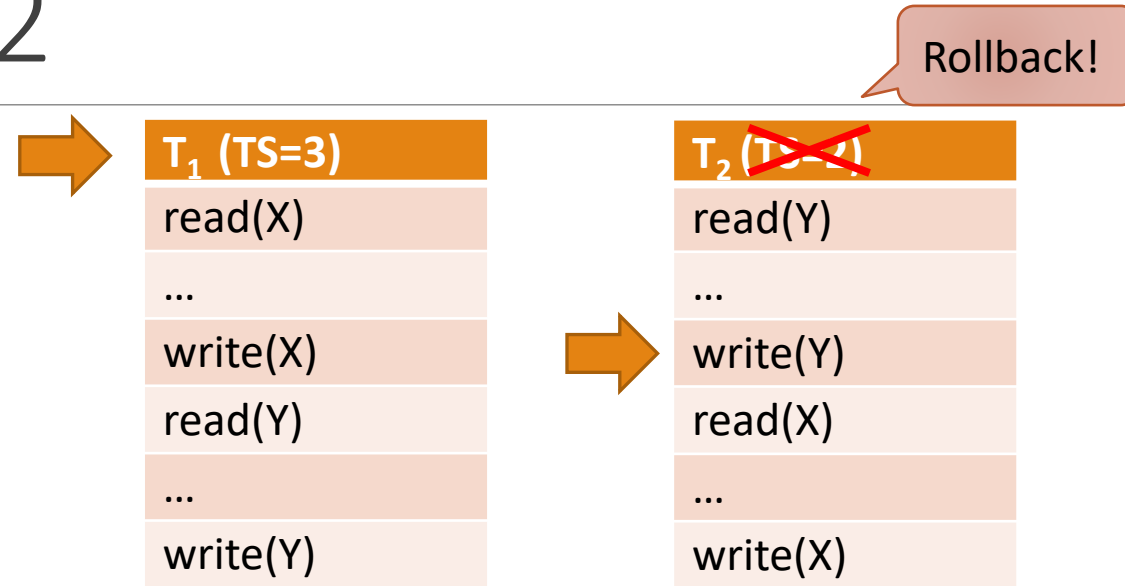


Assume a timestamp-based scheduler

- $T_1$  starts first
- Lines 1-3 of  $T_1$  are executed first, then lines 1-3 of  $T_2$

Which operations could be executed next?

# Example 2

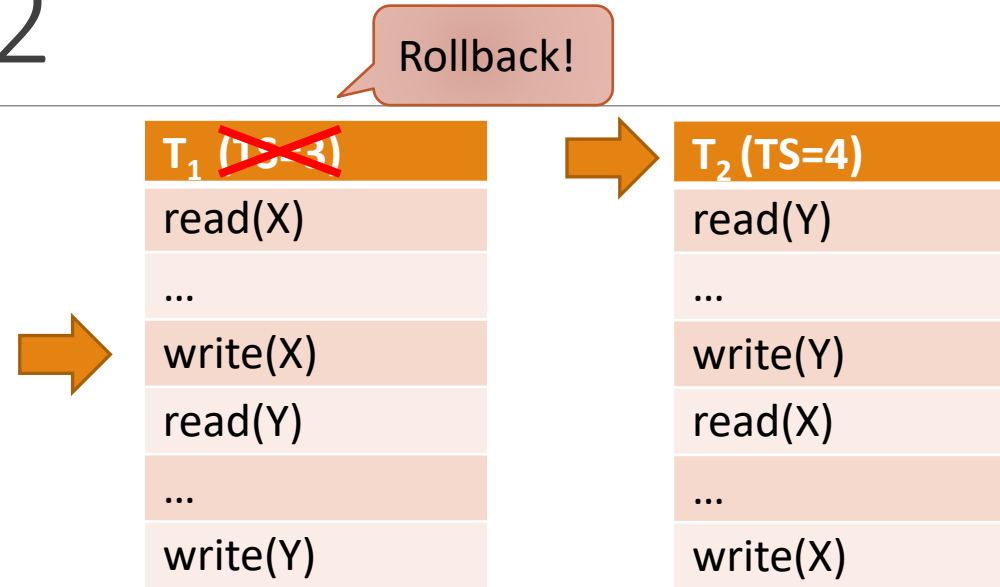


Assume a timestamp-based scheduler

- $T_1$  starts first
- Lines 1-3 of  $T_1$  are executed first, then lines 1-3 of  $T_2$

Which operations could be executed next?

# Example 2



Assume a timestamp-based scheduler

- $T_1$  starts first
- Lines 1-3 of  $T_1$  are executed first, then lines 1-3 of  $T_2$

Which operations could be executed next?

# Timestamp-based Scheduling

---

Nice properties:

- Enforces conflict-serialisable schedules
- Deadlocks don't occur

Bad properties:

- **Cascading rollbacks**
- **Starvation can occur** (cyclic aborts & restarts of transactions)

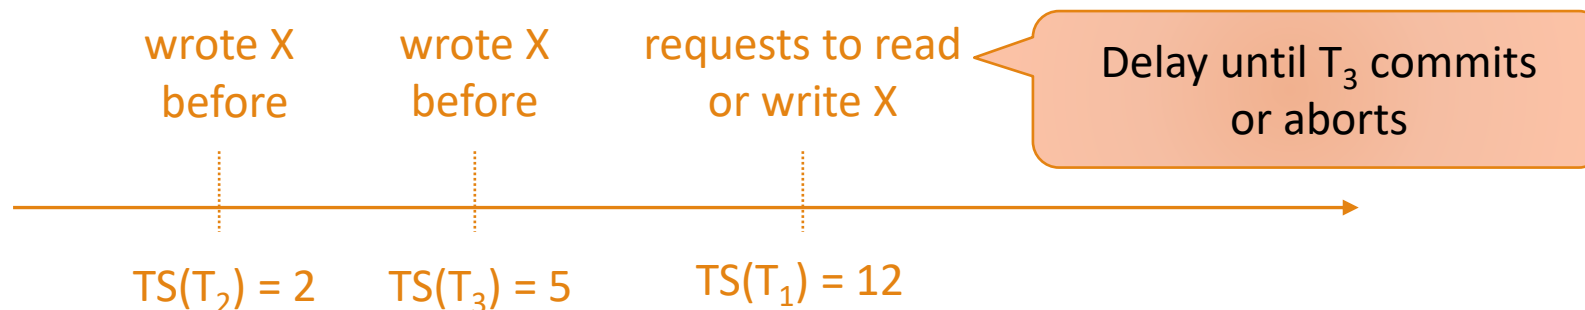
Starvation can be prevented using appropriate techniques (not in this module)

# Ensuring Strictness

Schedules enforced by timestamp-based schedulers are not strict.

Additional condition to enforce a **strict schedule**:

*Delay read or write requests until the youngest transaction who wrote X before has committed or aborted.*



# Multiversion concurrency control

---

Many DBMS are implementing a variant of time-stamp based approaches, called multiversion concurrency control

The idea is the same as time-stamp based approaches to scheduling, but you have multiple versions of the database!

- Meaning that write operations do not overwrite each other, but instead  $w_i(X)$  creates makes a new version of  $X$  at time  $TS(T_i)$
- Whenever you read, you read the latest version before your timestamp

This means that a transaction only need to restart if it tries to write AND the read timestamp is later than your timestamp

Or, written out in full, the rules are:

- For writes: **Abort & restart  $T_1$**  if  $RT(X) > TS(T_1)$  and **otherwise, grant request**
- For reads: **Always granted**

There is also a strict variant, where you delay reads until the transaction you read from commits

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1						
2						
3						
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10



# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2						
3						
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3						
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5						
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6						
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7						
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

No abort!

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8						
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20



# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 1)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(Y)					
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

Abort &  
restart

# Exam New timestamp MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	<del>write(Y)</del>		0	0	2	2
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

# Exam New timestamp MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	<del>write(Y)</del>		0	0	2	2
9						
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(X)		0	0	2	2
9	read(X)		3	0	2	2
10						
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(X)		0	0	2	2
9	read(X)		3	0	2	2
10	X := X + 100		3	0	2	2
11						
12						
13						

Timestamp	X	Y
0	10	10
2		20

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(X)		0	0	2	2
9	read(X)		3	0	2	2
10	X := X + 100		3	0	2	2
11	read(Y)		3	0	3	2
12						
13						

Timestamp	X	Y
0	10	10
2		20

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(X)		0	0	2	2
9	read(X)		3	0	2	2
10	X := X + 100		3	0	2	2
11	read(Y)		3	0	3	2
12	Y := Y * 3		3	0	3	2
13						

Timestamp	X	Y
0	10	10
2		20

# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(Y)		0	0	2	2
9	read(X)		3	0	2	2
10	X := X + 100		3	0	2	2
11	read(Y)		3	0	3	2
12	Y := Y * 3		3	0	3	2
13	write(Y)		3	0	3	3

Timestamp	X	Y
0	10	10
2		20



# Example 1 - MVCC

Time	T <sub>1</sub> (TS = 3)	T <sub>2</sub> (TS = 2)	X		Y	
			RT	WT	RT	WT
0			0	0	0	0
1	read(X)		1	0	0	0
2	X := X + 100		1	0	0	0
3		read(Y)	1	0	2	0
4		Y := Y * 2	1	0	2	0
5		write(Y)	1	0	2	2
6	read(Y)		1	0	2	2
7	Y := Y * 3		1	0	2	2
8	write(Y)		0	0	2	2
9	read(X)		3	0	2	2
10	X := X + 100		3	0	2	2
11	read(Y)		3	0	3	2
12	Y := Y * 3		3	0	3	2
13	write(Y)		3	0	3	3

Timestamp	X	Y
0	10	10
2		20
3		60

# Summary

---

One can use an alternate technique (instead of Strict 2PL) to ensure all the ACID properties and avoid deadlocks automatically: The idea is to just act as if the transaction was executed at its start time and restart any transaction that attempts to create a paradox with that!

- This only gives you conflict-serializability, but strict can be ensured by, on each read or write request (that would not be a paradox), wait until the transaction that last wrote to that item has committed or aborted

One can do the idea in two ways:

- “Normal” database with 1 version of each variable
- Multiversion concurrency control, where each write is saved with its own timestamp!