

# What would the DBMS implementations do?

---

THERE IS ALSO AN REVIEW OVER THIS PART, BECAUSE IT WAS SO LONG

# Overview over this video

---

In this video, we will look at how the different implementations implement ACID!

We will also see a review over this part, since it is fairly big

# Transaction Support in DBMS

---

# Transaction Support in DBMS

---

Widespread ACID support in major DBMSs

- Fully ACID compliant: PostgreSQL, Oracle DB, IBM DB2, ...
- Partly ACID compliant: MySQL (full compliance requires additional engines like InnoDB)

# Transaction Support in DBMS

---

Widespread ACID support in major DBMSs

- Fully ACID compliant: PostgreSQL, Oracle DB, IBM DB2, ...
- Partly ACID compliant: MySQL (full compliance requires additional engines like InnoDB)

Some does strict 2PL, e.g. PostgreSQL and MySQL/InnoDB (the latter only on the highest – i.e. serializable – isolation level)

# Transaction Support in DBMS

---

Widespread ACID support in major DBMSs

- Fully ACID compliant: PostgreSQL, Oracle DB, IBM DB2, ...
- Partly ACID compliant: MySQL (full compliance requires additional engines like InnoDB)

Some does strict 2PL, e.g. PostgreSQL and MySQL/InnoDB (the latter only on the highest – i.e. serializable – isolation level)

Most, however does MVCC (incl. MySQL/InnoDB on lower isolation levels)

- Database requires more storage, but relative little delay

# Deadlocks

# MySQL (with InnoDB)

---



# MySQL (with InnoDB)

---

Uses Wait-For graphs

# MySQL (with InnoDB)

---

Uses Wait-For graphs

Except: If transaction has line of length  $> 200$  it is rolled back


# MySQL (with InnoDB)

---

Uses Wait-For graphs

Except: If transaction has line of length  $> 200$  it is rolled back

E.g.  $T_1$  is rolled back in this case:




# MySQL (with InnoDB)

---

Uses Wait-For graphs

Except: If transaction has line of length  $> 200$  it is rolled back

E.g.  $T_1$  is rolled back in this case:



If cycle: rollback the smallest transaction


# MySQL (with InnoDB)

---

Uses Wait-For graphs

Except: If transaction has line of length  $> 200$  it is rolled back

E.g.  $T_1$  is rolled back in this case:



If cycle: rollback the smallest transaction


Deadlock detection can be switched off, in which case time-out is used (on locks)

# MySQL (with InnoDB)

Uses Wait-For graphs

Except: If transaction has line of length  $> 200$  it is rolled back

E.g.  $T_1$  is rolled back in this case:



Also uses a timestamp based approach to ensure that reads do not interfere with writes

If cycle: rollback the smallest transaction

Deadlock detection can be switched off, in which case time-out is used (on locks)

# Other DBMS

---

## **PostgreSQL**

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs



# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## **Oracle DB**

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## **Oracle DB**

Uses timeout directly or timeout followed by Wait-For graphs

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## **Oracle DB**

Uses timeout directly or timeout followed by Wait-For graphs

Does not use locks on read

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## **Oracle DB**

Uses timeout directly or timeout followed by Wait-For graphs

Does not use locks on read

## **IBM DB2**

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## **Oracle DB**

Uses timeout directly or timeout followed by Wait-For graphs

Does not use locks on read

## **IBM DB2**

Uses lock timeout or global time-out followed by Wait-For graphs

# Other DBMS

---

## **PostgreSQL**

Uses timeout on locks followed by Wait-For graphs

Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## **Oracle DB**

Uses timeout directly or timeout followed by Wait-For graphs

Does not use locks on read

## **IBM DB2**

Uses lock timeout or global time-out followed by Wait-For graphs

Uses update-locks

# Transactions Beyond DBMS

---



# Transactions Beyond DBMS

---

The techniques covered in this chapter are not confined to DBMS

# Transactions Beyond DBMS

---

The techniques covered in this chapter are not confined to DBMS

Similar issues whenever systems share resources

# Transactions Beyond DBMS

---

The techniques covered in this chapter are not confined to DBMS

Similar issues whenever systems share resources

Some example scenarios:

- Processes in an operating system that access the same files, network resources, etc.
- Users editing the same document online
- Document versioning systems like subversion, git, etc.

# Try it out...

---

```
CREATE TABLE Student (id INT NOT NULL, name ...);
```

```
INSERT INTO Student VALUES (1, 'Anna', ...);
```

```
SELECT * FROM Student;
```

```
START TRANSACTION;
```

```
INSERT INTO Student VALUES (2, 'Ben', ...);
```

```
INSERT INTO Student VALUES (3, 'Chloe', ...);
```

```
ROLLBACK;
```

```
SELECT * FROM Student;
```

Try out reads, writes, different isolation levels, dirty reads, look up the documentation,  
...

Experiment with more complex scenarios...

# Try it out...

---

```
CREATE TABLE Student (id INT NOT NULL, name ...);
```

```
INSERT INTO Student VALUES (1, 'Anna', ...);
```

```
SELECT * FROM Student;
```

```
START TRANSACTION;
```

```
INSERT INTO Student VALUES (2, 'Ben', ...);
```

```
INSERT INTO Student VALUES (3, 'Chloe', ...);
```

```
ROLLBACK;
```

```
SELECT * FROM Student;
```

Try out reads, writes, different isolation levels, dirty reads, look up the documentation,  
...


Experiment with more complex scenarios...

I would suggest to use some programming language for interacting with it, if you want to try with multiple transactions at a time...

# Review

# Relational DBMS Components

(Simplified from Content video)



# Transaction Support

---



# Transaction Support

---

Part of SQL:

- Begin/end transactions, isolation levels, auto commit, ...
- Need to understand the consequences of these commands to make effective use of DBMS

# Transaction Support

---

## Part of SQL:

- Begin/end transactions, isolation levels, auto commit, ...
- Need to understand the consequences of these commands to make effective use of DBMS
  - When to combine different SQL statements into a transaction?
  - When do we need (conflict) serialisability? When is a weaker isolation level fine?

# Transaction Management Review

---

Dealing with **transactions** is a core task of DBMS

- Many things can go wrong when processing transactions, even when executing single SQL statements.
- Need to ensure **ACID properties**

# Transaction Management Review

---

Dealing with **transactions** is a core task of DBMS

- Many things can go wrong when processing transactions, even when executing single SQL statements.
- Need to ensure **ACID properties**

Requires careful **scheduling** of transactions and **logging** of relevant information

- Schedules should be **conflict-serialisable**
- Schedules should be **strict**

# Transaction Management Review

---

Dealing with **transactions** is a core task of DBMS

- Many things can go wrong when processing transactions, even when executing single SQL statements.
- Need to ensure **ACID properties**

Requires careful **scheduling** of transactions and **logging** of relevant information

- Schedules should be **conflict-serialisable**
- Schedules should be **strict**

Methods for enforcing conflict-serialisability & strictness:

- **Strict two-phase locking & deadlock prevention** methods
- **Timestamping**

# ACID

---

## Atomicity

- Transactions are fully executed or not at all
- Ensured by Undo logging, Undo/Redo logging or Force

# ACID

---

## Atomicity

- Transactions are fully executed or not at all
- Ensured by Undo logging, Undo/Redo logging or Force

## Consistency

- Schedule executes transactions equivalent to a serial schedule
- (needs two assumptions for this: non-database operations can be ignored and if a schedule is serial, then it is consistent)
- Ensured by Serializability, Conflict-Serializability, 2PL and Timestamp-based Scheduling (also Strict versions of the last two)

# ACID

---

## Atomicity

- Transactions are fully executed or not at all
- Ensured by Undo logging, Undo/Redo logging or Force

This is intuition.  
Exact def. is different

## Consistency

- Schedule executes transactions equivalent to a serial schedule
- (needs two assumptions for this: non-database operations can be ignored and if a schedule is serial, then it is consistent)
- Ensured by Serializability, Conflict-Serializability, 2PL and Timestamp-based Scheduling (also Strict versions of the last two)



# ACID continued

---

## Iisolation

- Transactions are isolated from each other (how well depends on level!)
- Ensured by Cascadeless and Strict schedules (incl. Strict 2PL and Strict Timestamp-based schedules)

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

# ACID continued

---

## I Isolation

- Transactions are isolated from each other (how well depends on level!)
- Ensured by Cascadeless and Strict schedules (incl. Strict 2PL and Strict Timestamp-based schedules)

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

Alternately:  
READ COMMITTED,  
REPEATABLE READ,  
SERIALIZABLE

# ACID continued

---

## I Isolation

- Transactions are isolated from each other (how well depends on level!)
- Ensured by Cascadeless and Strict schedules (incl. Strict 2PL and Strict Timestamp-based schedules)

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

Can also be:  
READ ONLY

Alternately:  
READ COMMITTED,  
REPEATABLE READ,  
SERIALIZABLE

# ACID continued

---

## I Isolation

- Transactions are isolated from each other (how well depends on level!)
- Ensured by Cascadeless and Strict schedules (incl. Strict 2PL and Strict Timestamp-based schedules)

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

Can also be:  
READ ONLY

Alternately:  
READ COMMITTED,  
REPEATABLE READ,  
SERIALIZABLE

## Durability

- If a transaction is committed, it does not disappear
- Ensured by Redo logging, Undo/Redo logging or No Steal
- Recoverable schedules are also required

# Summary

---

The different database implementations do things differently in regards to transactions

# Summary

---

The different database implementations do things differently in regards to transactions

In this part we saw how DBMS ensures the ACID properties (far too much to summarize)