

Faster joins using sorting

Overview

In the last video, we saw how to do (Naïve) loop join to handle joins, e.g. natural join

Here, we will see a more advanced technique called sort join

Can We Go Faster?

Yes, we can!

Equijoins

Equijoin $R \bowtie_{A=B} S$ is defined as $\sigma_{A=B}(R \times S)$

A, B are the **join attributes**

Stores

code	city
12345	1
678910	2

Employees

name	depart
Oscar	12345
Janice	678910
David	678910

Stores $\bowtie_{\text{code}=\text{depart}}$ **Employees**

code	city	name	depart
12345	1	Oscar	12345
678910	2	Janice	678910
678910	2	David	678910


If R is sorted on A and S is sorted on B , then $R \bowtie_{A=B} S$ can be computed with one pass over R and S + run time equal to the size of the output

Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B



as in merge sort

Assume: **R** is sorted on **A** and **S** is sorted on **B**

sorted

sorted


A	B	C	D	...
2	2	99	50	...
2	2	99	94	...
2	2	99	11	...
4	4	52	74	...

Merging

as in merge sort

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B



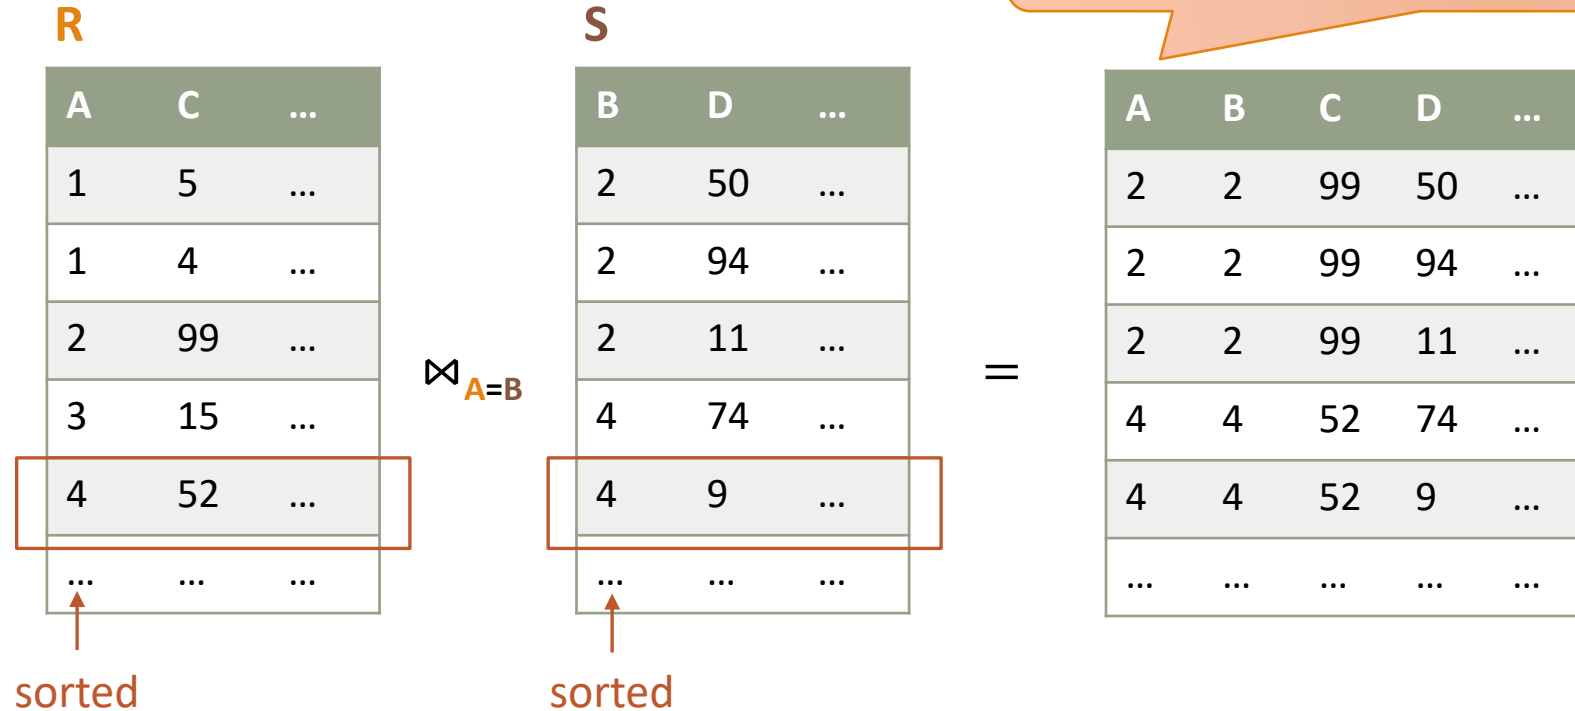
Merging

as in merge sort

Goal: compute $\mathbf{R} \bowtie_{A=B} \mathbf{S}$

Assume: \mathbf{R} is sorted on \mathbf{A} and \mathbf{S} is sorted on \mathbf{B}


Running time: $O(|\mathbf{R}| + |\mathbf{S}| + \text{size of output})$



Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B




Remember tuples in S that match with the current value of A

Merging With Duplicates in Column A

Runtime: $O(|R| + |S| + \text{size of output})$

Goal: compute $R \bowtie_{A=B} S$

Assume: R is sorted on A and S is sorted on B



Remember tuples in S that match with the current value of A

Faster Joins With Sorting

Sort Join Algorithm:

Compute $R \bowtie_{A=B} S$:

1. Sort **R** on **A**

Running time: $O(|R| \times \log_2 |R|)$

2. Sort **S** on **B**

Running time: $O(|S| \times \log_2 |S|)$

3. Merge the sorted **R** and **S**

Running time: $O(|R| + |S| + \text{size of output})$

Typical running time: $O(|R| \log_2 |R| + |S| \log_2 |S|)$

- If not “too many” values in **A** occur multiple times
- E.g., this is the case if **A** is a key

Having a run time depending on the size of output is called output sensitive

Typically much faster than Nested Loop Join

- Same time in the worst case, because output can have size up to $|R| \times |S|$

Remarks


Various **join algorithms** in practice:

- Index joins
- Hash joins
- Multiway joins: join more than two relations at once

Can compute **other operations of relational algebra** using similar methods as those in this lecture

We've neglected that relations are stored on disk

Running Time vs Disk Accesses



Relevant parameters:

- **B** = size of a disk block (typically 512→4096 bytes)
- **M** = number of disk blocks that fit into available RAM

Algorithm	No. of elementary operations	No. of disk accesses
Reading a relation R	$O(R)$	$O\left(\frac{ R }{B}\right)$
Sorting R on attribute A	$O(R \log_2 R)$	$O\left(\frac{ R }{B} \log_M \frac{ R }{B}\right)$

External memory merge sort

Summary

We saw one approach to do faster joins, namely sort joins

Sort join has run time close to linear (i.e. like sort) + the size of the output