Introduction to query processing

Overview over this video

This video is meant to serve as an introduction to query processing AND also how and what a query plan is (basically, a way to look at relational algebra)

Where we are…

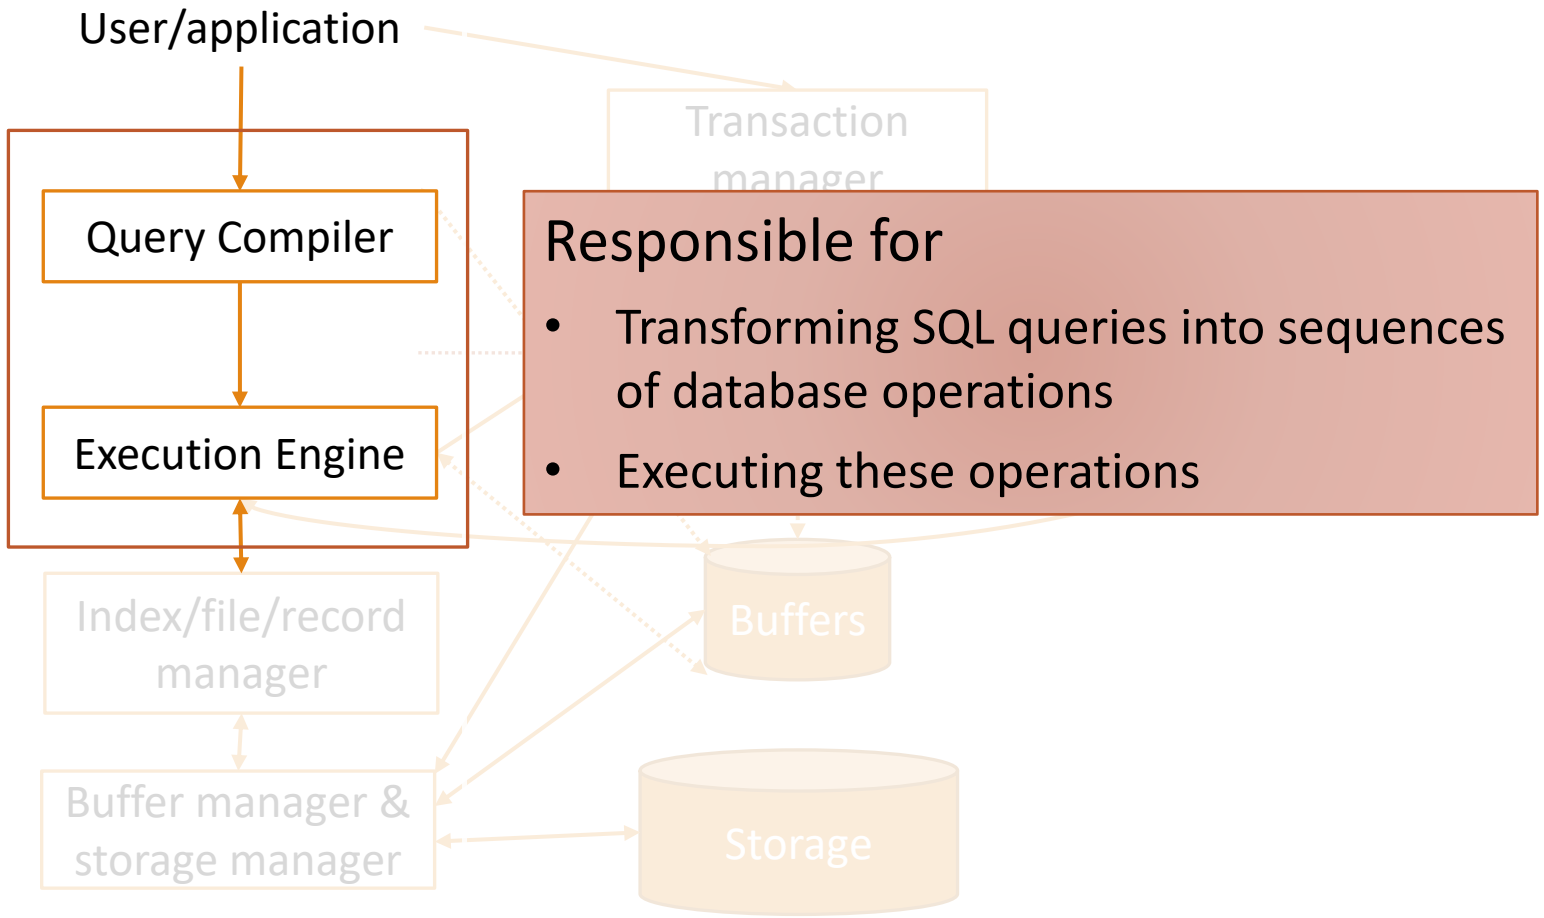

Responsible for

- Transforming SQL queries into sequences of database operations
- Executing these operations

SQL Queries

Here typically SELECT statements

SQL query: SQL SELECT/INSERT/UPDATE/DELETE statement

```
SELECT id AS employee_id
FROM Employees
WHERE department = 'Liv';
```

Employees(name, id, departmnet)

Salary(e_id, month, amount)

```
SELECT name, sum(amount)
FROM Employees, Salary
WHERE id = e_id AND department = 'Liv'
GROUP BY name;
```

Declarative: tells the DBMS **what we want**, ***not*** **how to get it**

DBMS selects a good sequence of database operations
to execute the query

How Does Query Processing Work?

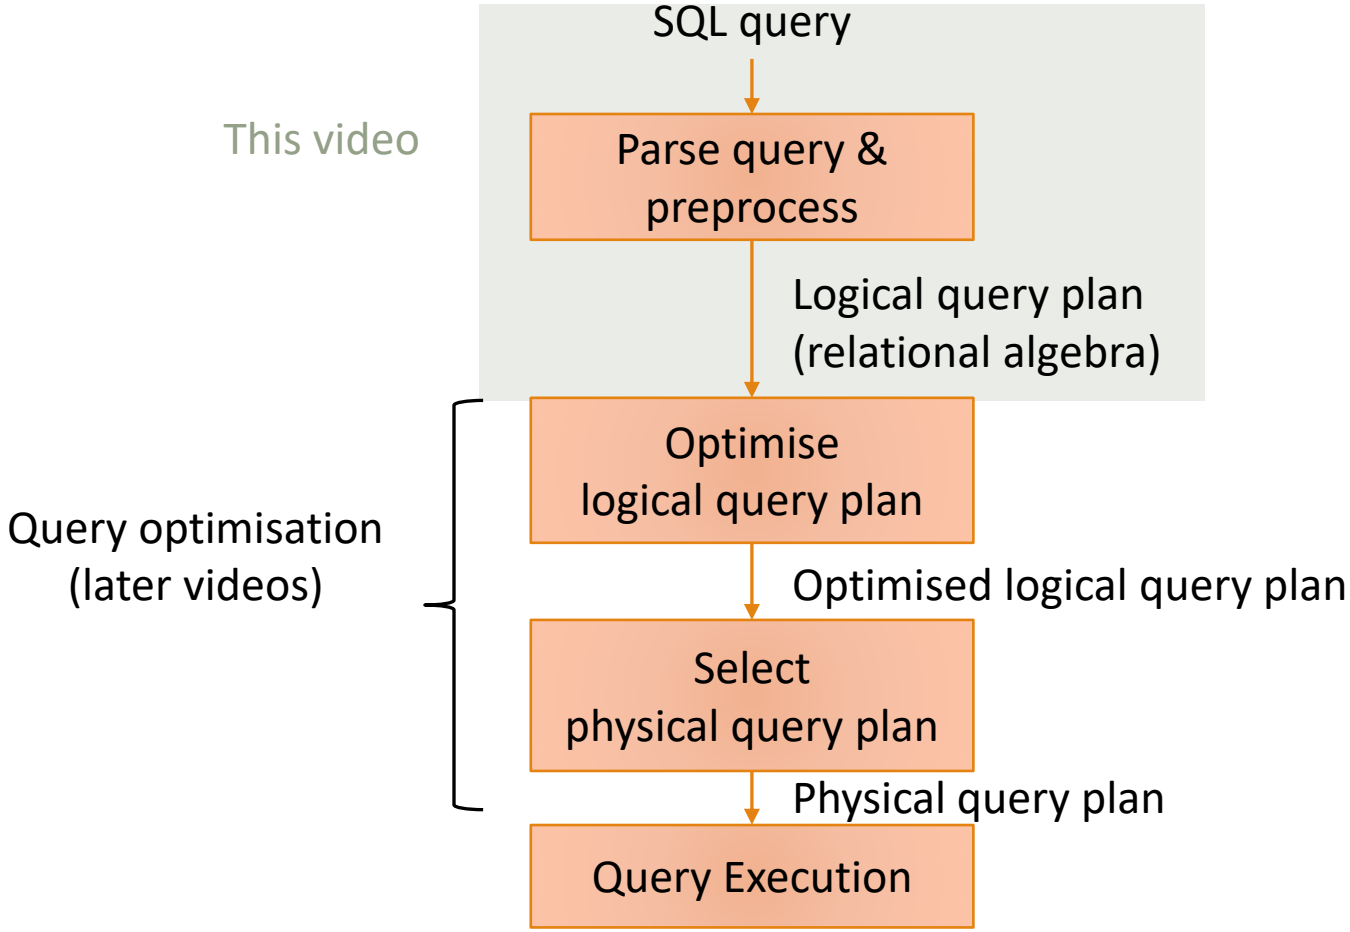

Relational Algebra

Set of **operations** that can be applied to **relations** to **compute new relations**

Basic relational algebra:

- Selection (σ)
- Projection (π)
- Cartesian product ($\times$)
- Union ($\cup$)
- Renaming (ρ)
- Natural join ($\bowtie$)
- Semijoin ($\ltimes$)
- Many others can be defined

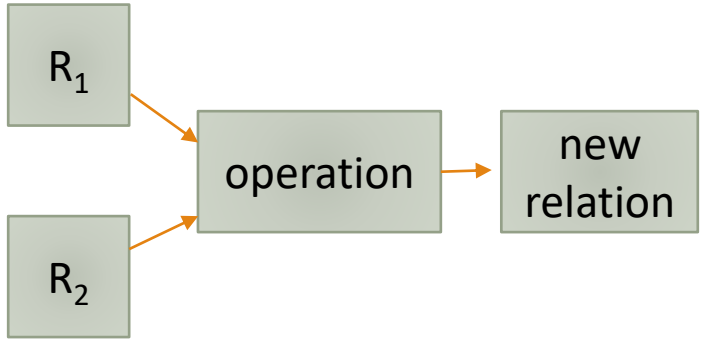

If you are unsure what any of these means, go back and watch the SQL query – Misc video (and the second tutorial for Union ($\cup$))

Query Plans

A **relational algebra expression** that is obtained from an SQL query is also called a **(logical) query plan**

```
SELECT department, name
FROM Stores, Employees
WHERE department=worksAt AND city='Liverpool';
```

SQL query

$$\pi_{\text{department,name}}(\sigma_{\text{department=worksAt AND city='Liverpool'}}(\text{Stores} \times \text{Employees}))$$

Query plan for the query

Query plans are typically **represented as trees**

Query Plans As Trees

$$\pi_{\text{department,name}}(\sigma_{\text{department=worksAt AND city='Liverpool'}}(\text{Stores} \times \text{Employees}))$$

Tree representation:

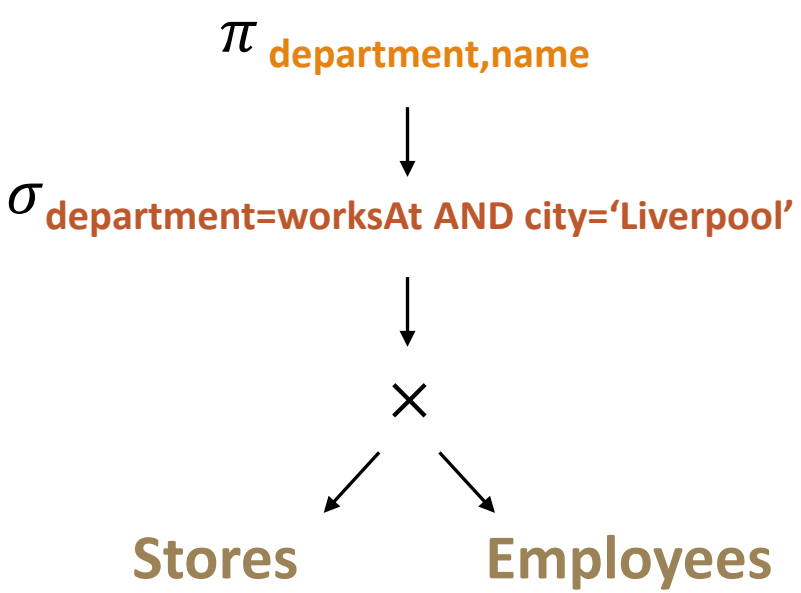

- Leaves = input relations
- Inner nodes = operators

- Such trees are evaluated from the leaves to the root

Another example

Represent the following query plan as a tree:

$$\sigma_{\text{department=worksAt}}(\pi_{\text{department}}(\text{Stores}) \times \pi_{\text{worksAt,name}}(\text{Employees}))$$

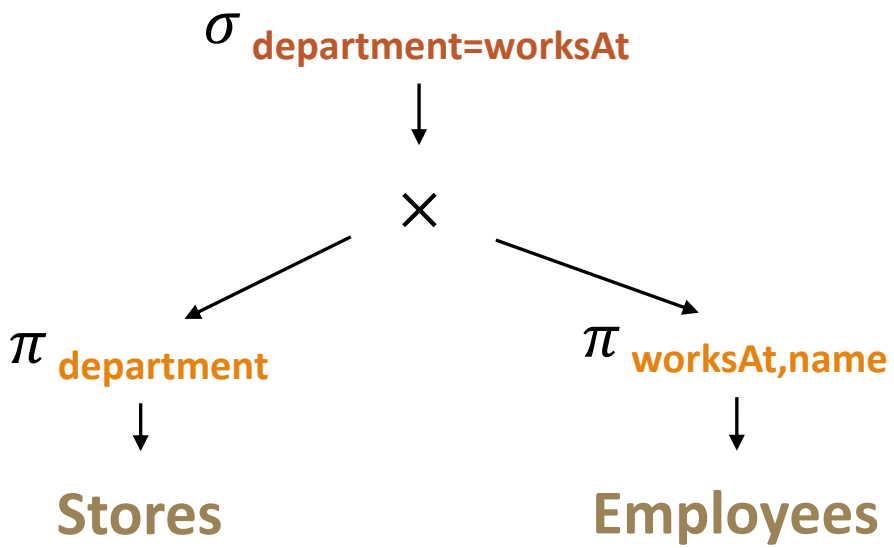

Equivalent Query Plans

There are typically **many different query plans**

DBMSs aim to **select a best possible query plan**

Relational algebra is better suited than SQL for this

- Can use **equivalence laws** of relational algebra to generate a query plan for the same query that can be executed faster!
- Example:
  - $\sigma_{condition}(R_1 \bowtie R_2) = \sigma_{condition}(R_1) \bowtie \sigma_{condition}(R_2)$

Details will come later…

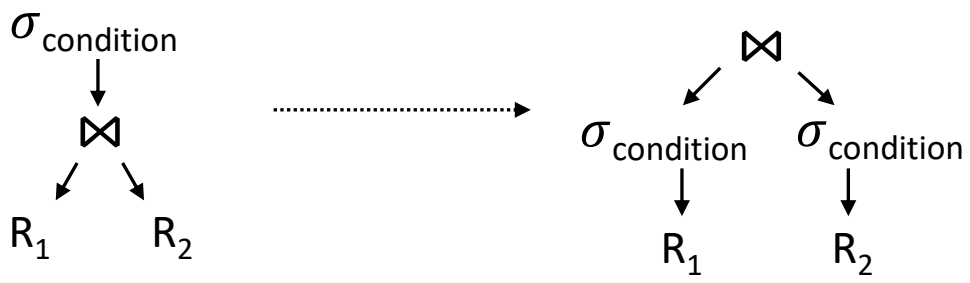

Summary

DBMS translate SQL queries into relational algebra expressions, also called **(logical) query plans**

The DBMS will then

- Optimise the logical query plan by using equivalence laws (later…)
- Select suitable algorithms for computing each operator in the logical query plan (later…)