

Executing a query plan  
naively

---

# Overview over this video

---

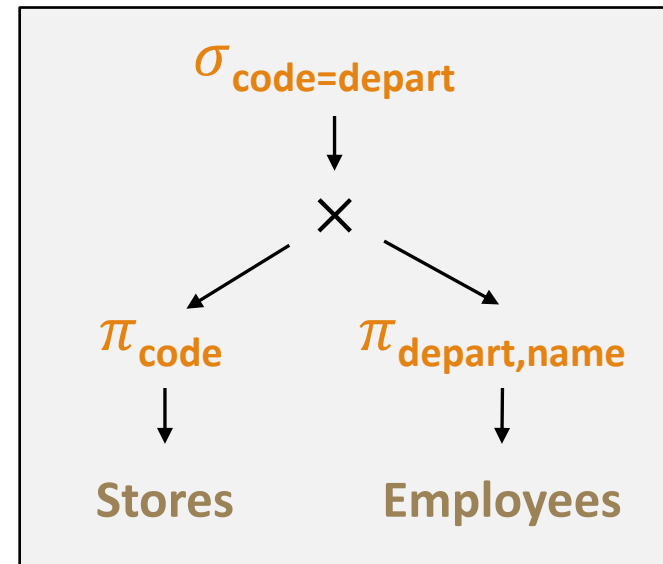
We will see a fairly naïve approach to executing a query plan – the next video will show a more advanced approach to some of it

# Executing Query Plans

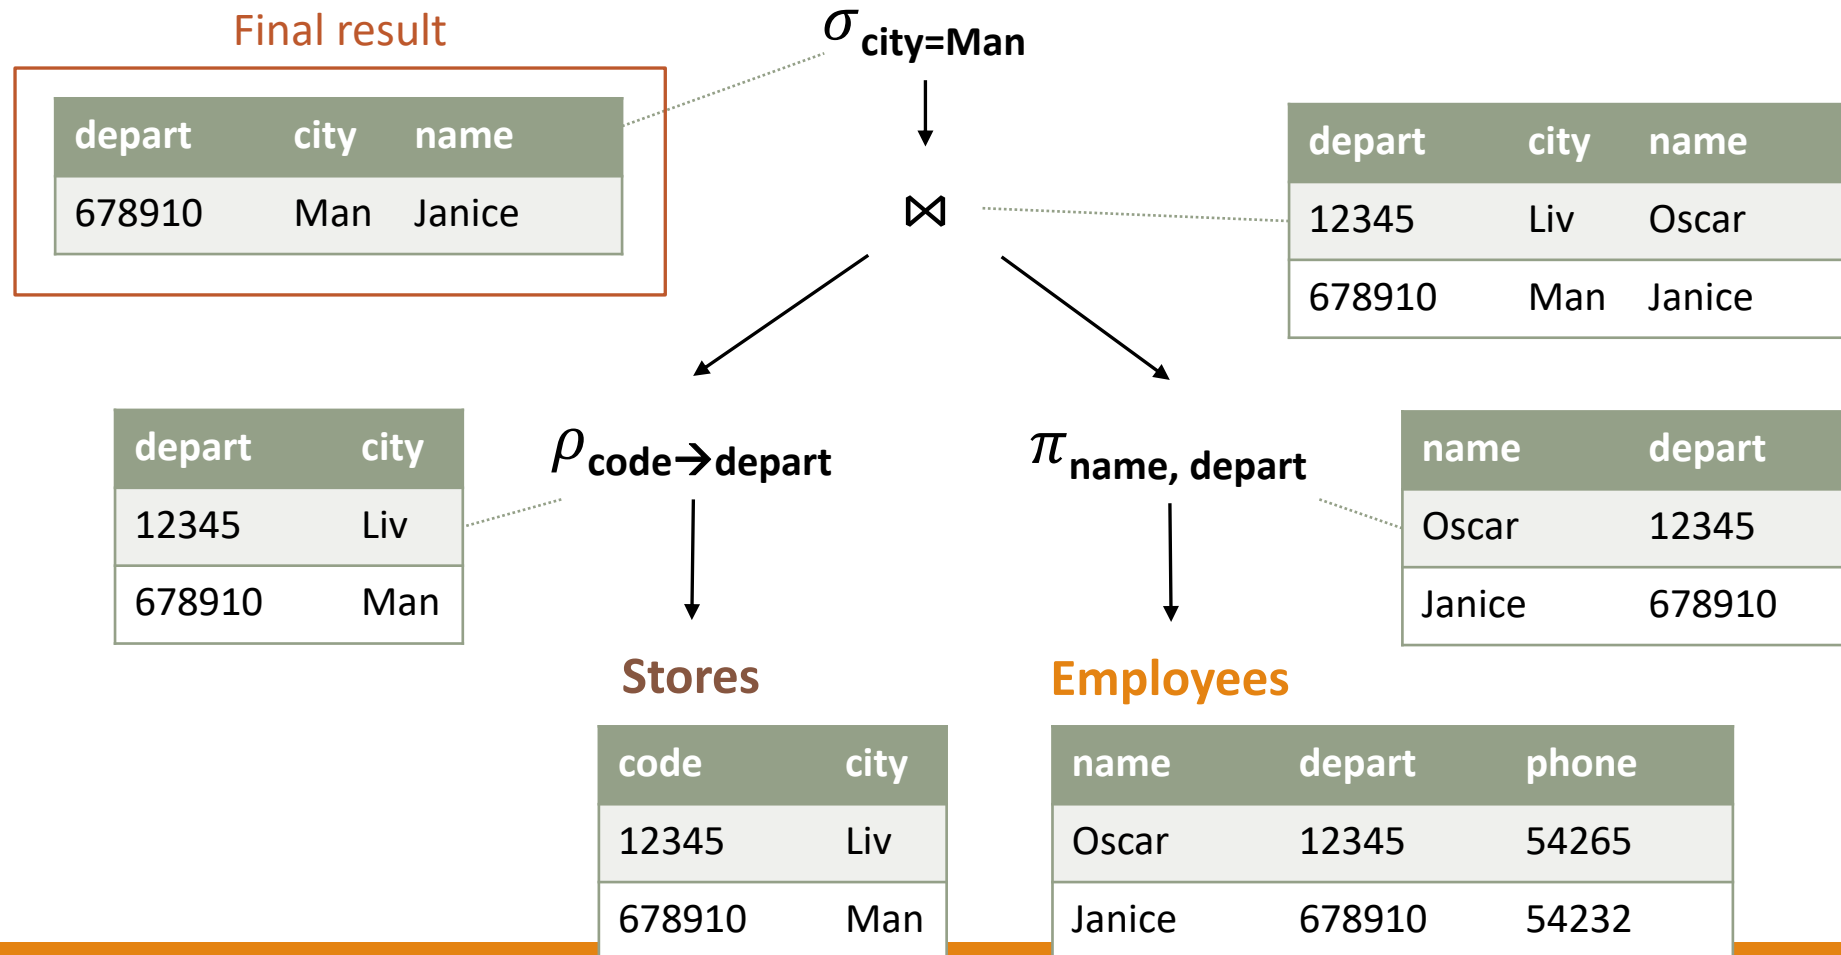
Query plans tell us exactly how to compute the result to a query

Proceed from bottom to top:

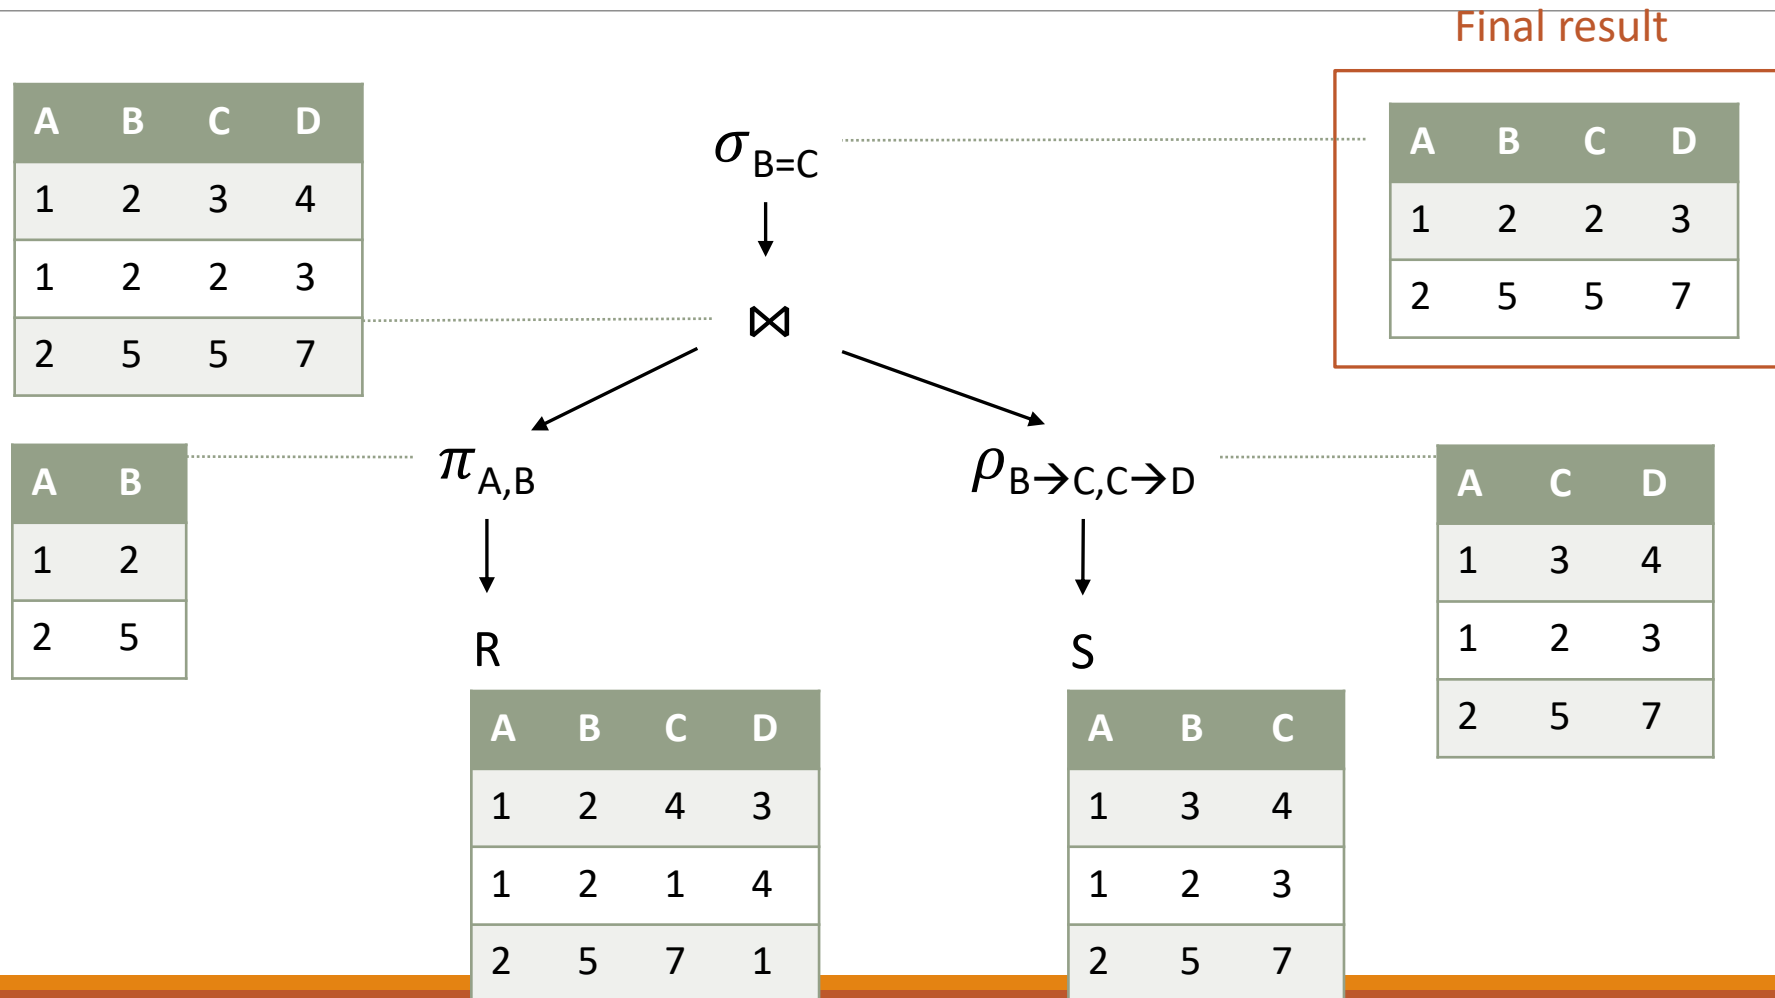
- Compute an **intermediate result** for each node
- For a **leaf** labeled with relation **R**, the intermediate result is **R**.
- For an **inner node** labeled with operator **op**, get the intermediate result by applying **op** to the children's intermediate results.
- **Result of the query** = intermediate result of the root



# Example



# Example 2



# How to “Apply” An Operator?

How to compute  $\sigma_{\text{condition}}(R)$ ?



have to read entire file

```
for each tuple t in R:  
  if t satisfies condition:  
    output t
```

Is there a faster way? Sometimes, e.g. if it is sorted and the condition is “easy” or we have an index on it

How to compute  $\pi_{\text{attribute list}}(R)$ ?

Similar! Read **R** only once...

```
for each tuple t in R:  
  output the restriction  
    of t to the attributes  
    in attribute list
```

How to compute  $R \bowtie S$ ?

Many different ways...

Is there a faster way? Not really, if you actually do it – can do it lazily though

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**

the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

$r$

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

$s$

depart	city	name
12345	Liv	Oscar

$r \bowtie s$

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**

the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

$r$

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

$\bowtie$

$=$

$s$

depart	city	name
12345	Liv	Oscar



# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**

the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

**r**

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

**s**

$\bowtie$

$=$

depart	city	name
12345	Liv	Oscar

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**

the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

**r**

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

**s**

depart	city	name
12345	Liv	Oscar

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**  the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

**r**

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

**=**

depart	city	name
12345	Liv	Oscar

**s**

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**

the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

$r$

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

=

depart	city	name
12345	Liv	Oscar
678910	Man	Janice

$s$

$r \bowtie s$

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**  the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

**r**

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

=

depart	city	name
12345	Liv	Oscar
678910	Man	Janice

**s**

# Naïve Computation of Joins

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**  the tuple obtained by joining  $r$  and  $s$

**Stores**

depart	city
12345	Liv
678910	Man

**r**

**Employees**

name	depart
Oscar	12345
Janice	678910
David	678910

=

depart	city	name
12345	Liv	Oscar
678910	Man	Janice
678910	Man	David

**s**

**$r \bowtie s$**

# Naïve Computation of Joins

---

Nested Loop Join Algorithm:

**Compute  $R \bowtie S$ :**

**for each tuple  $r$  in  $R$ :**

**for each tuple  $s$  in  $S$ :**

**if  $r$  and  $s$  have the same values for all common attributes:**

**output  $r \bowtie s$**

Slow: for each tuple  $r$  in  $R$  reads entire relation  $S$

Running time:  $O(|R| \times |S|)$

Number of tuples in  $R$

Number of tuples in  $S$

# Summary

---

Query plans are evaluated bottom-up – from the leaves to the root

Each operator can be computed in different ways

- Selection: e.g., linear scans (reading the relation once)
- Projection: linear scans
- Joins: Only Nested Loop Join in this video...