

# External memory merge-sort

---

THIS ONE IS NOT REQUIRED FOR THE EXAM!

# Overview over this video

---

This video (which is not required for the exam), we will cover external memory merge sort


# Reminder: Computing $\sigma_{\text{condition}}(R)$

---

Basic procedure:

**R:**

tuple 1
tuple 2
tuple 3
tuple 4
...



```
for each tuple t in R:  
    if t satisfies condition:  
        output t
```


# Reminder: Computing $\sigma_{\text{condition}}(R)$

---

Basic procedure:

**R:**

tuple 1
tuple 2
tuple 3
tuple 4
...



```
for each tuple t in R:  
    if t satisfies condition:  
        output t
```

Needs to read the entire relation

# Reminder: Computing $\sigma_{\text{condition}}(R)$

Basic procedure:

R:

tuple 1
tuple 2
tuple 3
tuple 4
...

```
for each tuple t in R:  
    if t satisfies condition:  
        output t
```

Needs to read the entire relation

Can this be done faster?

# Reminder: Computing $\sigma_{\text{condition}}(R)$

Basic procedure:

R:

tuple 1
tuple 2
tuple 3
tuple 4
...

```
for each tuple t in R:  
    if t satisfies condition:  
        output t
```

Needs to read the entire relation

Yes, sometimes!

Can this be done faster?

# Example

---

$\sigma_{\text{programme}='G401'}(\text{Students})$

Students

id	name	programme
...	...	...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...	...	...

# Example

---

$\sigma_{\text{programme}='G401'}(\text{Students})$

Students

id	name	programme
...	...	...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...	...	...

Selection can be performed faster if we know  
**where to find the rows for 'G401'**



# Example

---

$\sigma_{\text{programme}='G401'}(\text{Students})$

Students

id	name	programme
...	...	...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...	...	...

Selection can be performed faster if we know  
**where to find the rows for 'G401'**

Two solutions: **sorting & index**

# Example

---

$\sigma_{\text{programme}='G401'}(\text{Students})$

Students

id	name	programme
...	...	...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...	...	...

Selection can be performed faster if we know  
**where to find the rows for 'G401'**

Two solutions: **sorting & index**

This video

# Faster Joins With Sorting

(from video “Faster Joins With Sorting”)

Sort Join Algorithm:

**Compute  $R \bowtie_{A=B} S$ :**

1. Sort **R** on **A**

Running time:  $O(|R| \times \log_2 |R|)$

2. Sort **S** on **B**

Running time:  $O(|S| \times \log_2 |S|)$

3. Merge the sorted **R** and **S**

Running time:  $O(\text{size of output})$

Typical running time:  $O(|R| \log_2 |R| + |S| \log_2 |S|)$

- If not “too many” values in **A** occur multiple times
- E.g., this is the case if **A** is a key

Having a run time depending on the size of output is called output sensitive

Typically much faster than Nested Loop Join

- Same time in the worst case, because output can have size up to  $|R| \times |S|$

# Faster Joins With Sorting

(from video “Faster Joins With Sorting”)

Sort Join Algorithm:

**Compute  $R \bowtie_{A=B} S$ :**

1. Sort **R** on **A**

Running time:  $O(|R| \times \log_2 |R|)$

2. Sort **S** on **B**

Running time:  $O(|S| \times \log_2 |S|)$

3. Merge the sorted **R** and **S**

Running time:  $O(\text{size of output})$

Want to do that faster!

Typical running time:  $O(|R| \log_2 |R| + |S| \log_2 |S|)$

- If not “too many” values in **A** occur multiple times
- E.g., this is the case if **A** is a key

Having a run time depending on the size of output is called output sensitive

Typically much faster than Nested Loop Join

- Same time in the worst case, because output can have size up to  $|R| \times |S|$

# Outside databases

---

Many problems are much easier then the data is sorted and in general, a significant fraction of all computing time is spend on sorting

- Was estimated to be around 25% back in the 60s
  - Found some claims that it is still in the range 25%-50%, but it is hard to verify

# Outside databases

---

Many problems are much easier then the data is sorted and in general, a significant fraction of all computing time is spend on sorting

- Was estimated to be around 25% back in the 60s
  - Found some claims that it is still in the range 25%-50%, but it is hard to verify

Some problems were you need sorting are “small” in that they fit in main memory

# Outside databases

---

Many problems are much easier then the data is sorted and in general, a significant fraction of all computing time is spend on sorting

- Was estimated to be around 25% back in the 60s
  - Found some claims that it is still in the range 25%-50%, but it is hard to verify

Some problems were you need sorting are “small” in that they fit in main memory

Some ain't

# Understanding the issue solved

---

To find one record on hard disk (with very fast “normal” drive of 15,000 rpm):

- 2.00 **milli-seconds** or  $2 \cdot 10^6$  **nano-seconds**



# Understanding the issue solved

---

To find one record on hard disk (with very fast “normal” drive of 15,000 rpm):

- 2.00 **milli-seconds** or  $2 \cdot 10^6$  **nano-seconds**

To find one record on SSD:

- 0.031 **milli-seconds** or  $3.1 \cdot 10^4$  **nano-seconds**

# Understanding the issue solved

---

To find one record on hard disk (with very fast “normal” drive of 15,000 rpm):

- 2.00 **milli-seconds** or  $2 \cdot 10^6$  **nano-seconds**

To find one record on SSD:

- 0.031 **milli-seconds** or  $3.1 \cdot 10^4$  **nano-seconds**

To find one record in RAM (with relative normal DDR4-2666 RAM):

- 13 **nano-seconds**

# Understanding the issue solved

---

To find one record on hard disk (with very fast “normal” drive of 15,000 rpm):

- 2.00 **milli-seconds** or  $2 \cdot 10^6$  **nano-seconds**

SLOW!!!

To find one record on SSD:

- 0.031 **milli-seconds** or  $3.1 \cdot 10^4$  **nano-seconds**

slow

To find one record in RAM (with relative normal DDR4-2666 RAM):

- 13 **nano-seconds**

fast

# Merge Sort

---

(Internal memory) merge sort:

**Divide** input in two parts  $P_1, P_2$   
**Sort**  $P_1$  &  $P_2$  **recursively**  
**While**  $P_1$  or  $P_2$  is not empty:  
    Add smallest remaining  
    element from  $P_1$  or  $P_2$   
    to output

# Merge Sort

---

(Internal memory) merge sort:

Merge



**Divide** input in two parts  $P_1, P_2$

**Sort**  $P_1$  &  $P_2$  **recursively**

**While**  $P_1$  or  $P_2$  is not empty:

Add smallest remaining  
element from  $P_1$  or  $P_2$   
to output

# External Merge Sort

---

External merge sort:

**Divide** input in **M** parts  $P_1, P_2, \dots, P_M$   
**Sort**  $P_1, P_2, \dots, P_M$  **recursively**  
**While** not all  $P_i$  are empty:  
    Add smallest remaining  
    element from any part  $P_i$  to  
    output

# External Merge Sort

---

External merge sort:

Merge

**Divide** input in **M** parts  $P_1, P_2, \dots, P_M$

**Sort**  $P_1, P_2, \dots, P_M$  **recursively**

**While** not all  $P_i$  are empty:

Add smallest remaining  
element from any part  $P_i$  to  
output

# External Merge Sort

External merge sort:

Number of disk blocks that fit in RAM

Merge

**Divide** input in **M** parts  $P_1, P_2, \dots, P_M$

**Sort**  $P_1, P_2, \dots, P_M$  **recursively**

**While** not all  $P_i$  are empty:

Add smallest remaining  
element from any part  $P_i$  to  
output



# Merge in more details


---

Goal: Must merge  $M$  (perhaps large) parts into one

# Merge in more details

---

Goal: Must merge  $M$  (perhaps large) parts into one




# Merge in more details

---

Goal: Must merge  $M$  (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet

$P_1$ :



$P_2$ :



...

$P_M$ :



# Merge in more details

Goal: Must merge  $M$  (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet

$P_1$ :



$P_2$ :



...

$P_M$ :



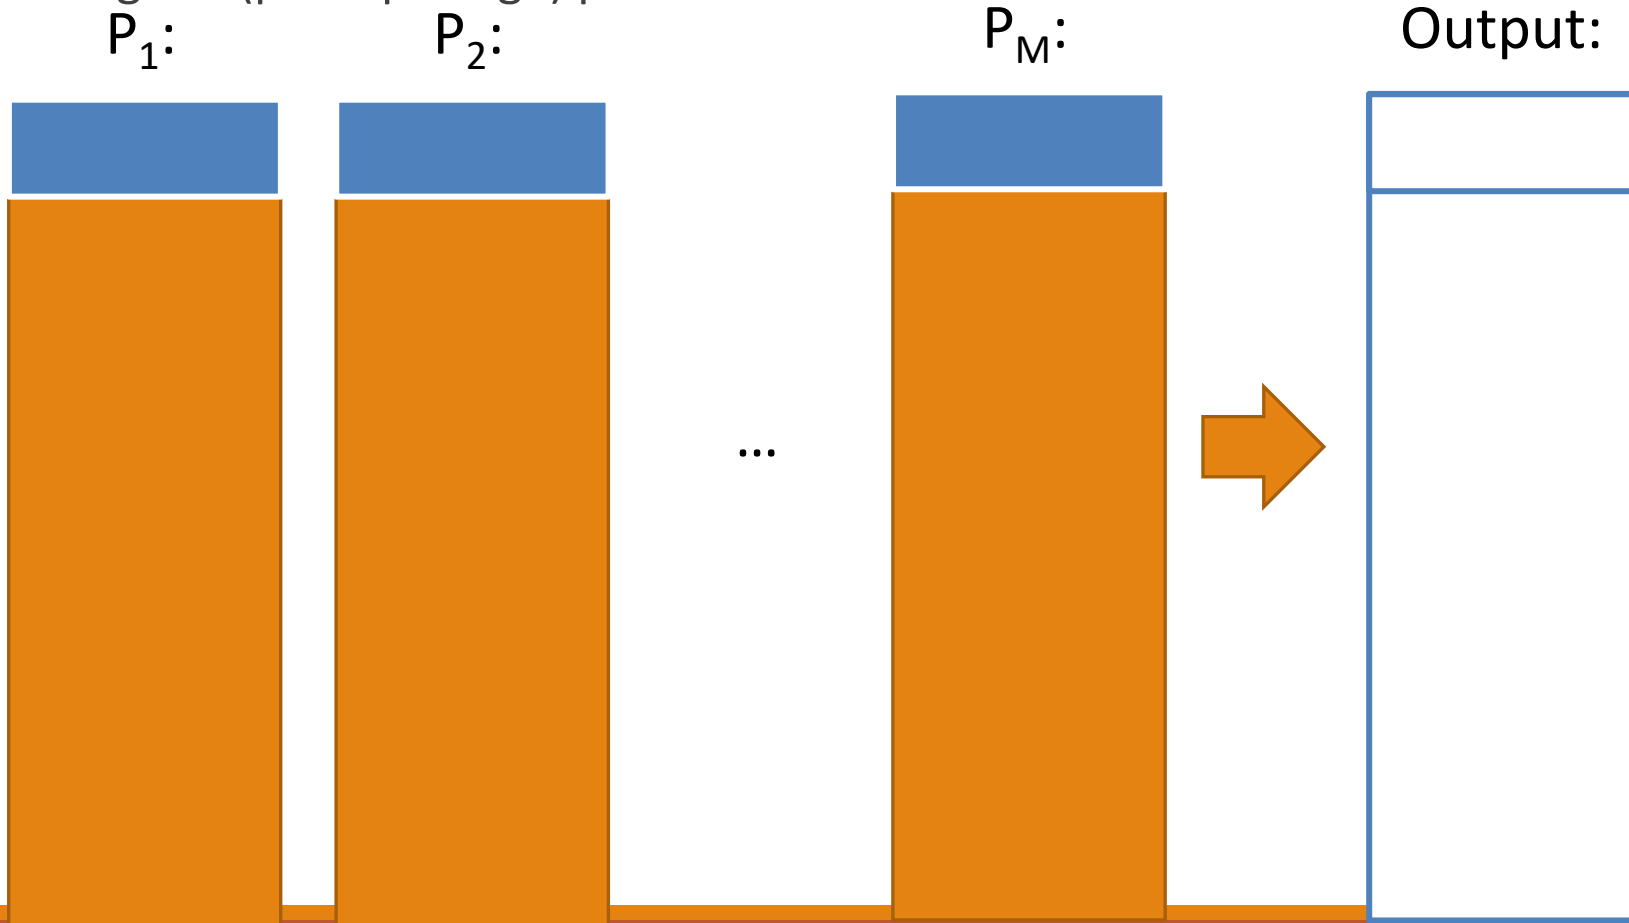
Output:



# Merge in more details

Goal: Must merge  $M$  (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge  $M$  (perhaps large) parts into one


Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet

$P_1$ :

$P_2$ :

$P_M$ :


Output:



# Merge in more details

Goal: Must merge M (perhaps large) parts into one


Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet





# Merge in more details

Goal: Must merge M (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet

1. Move to disk, ...

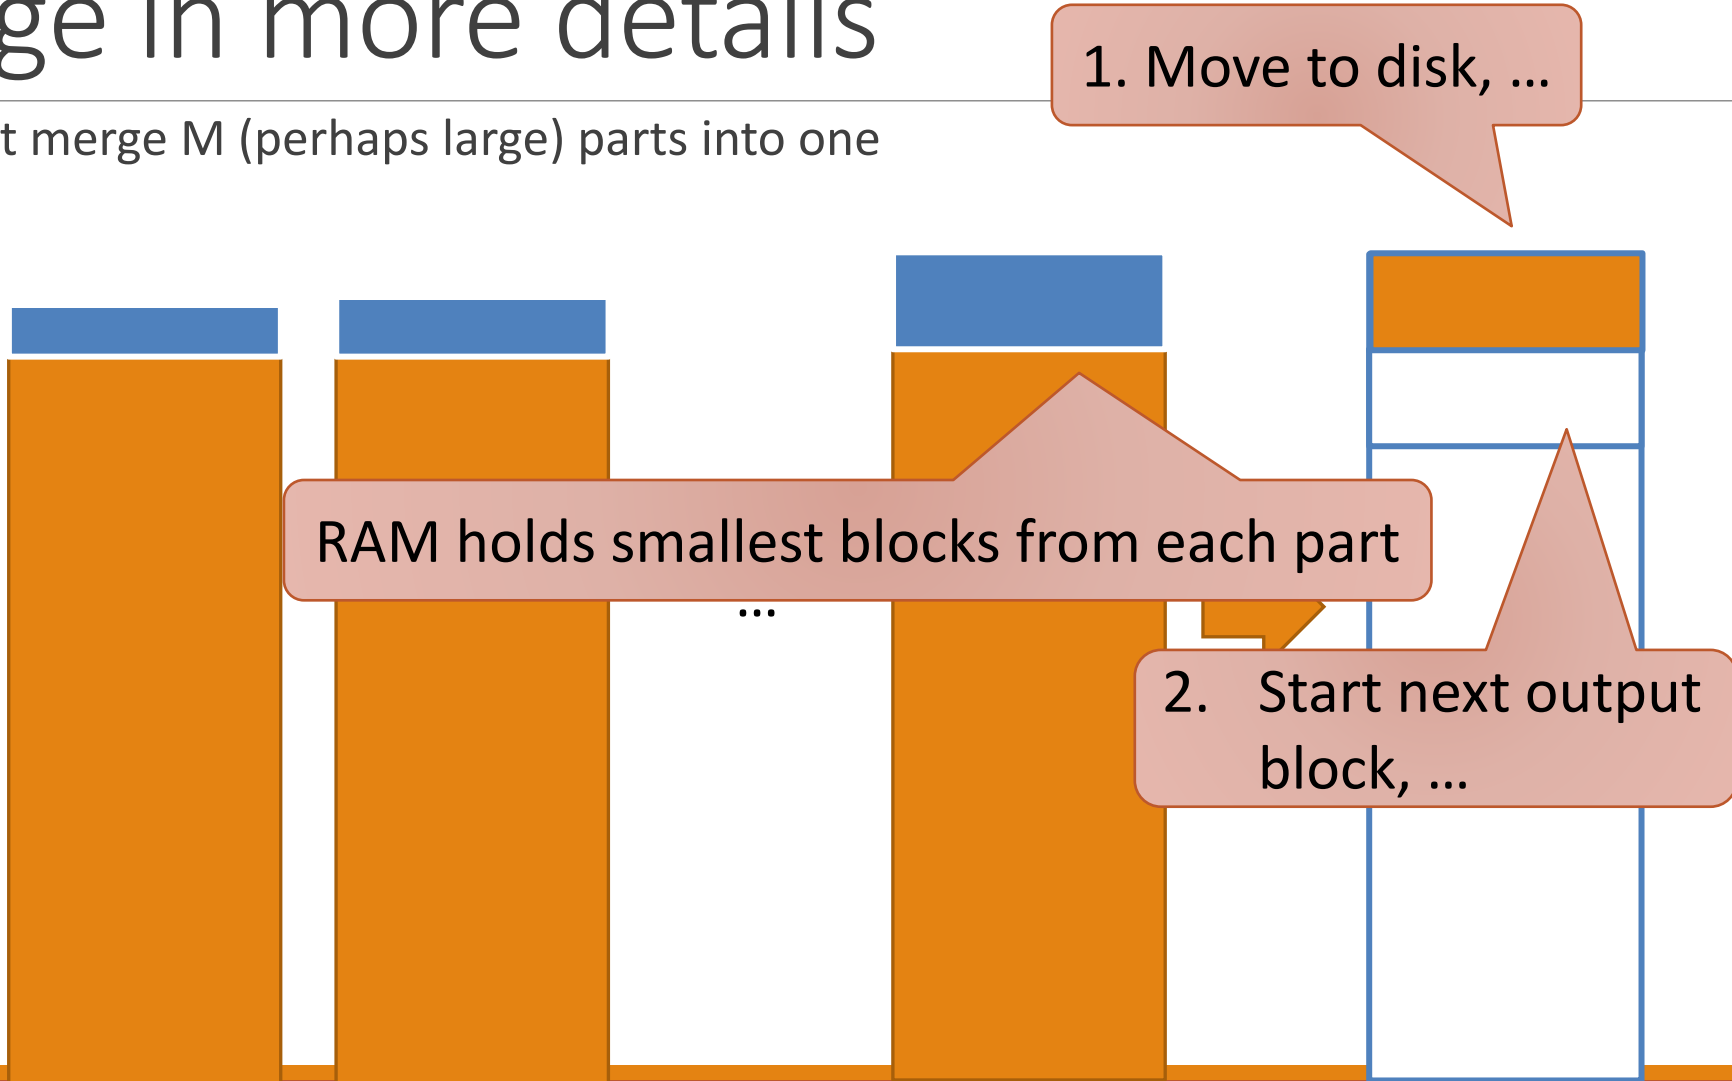
RAM holds smallest blocks from each part

...

# Merge in more details

Goal: Must merge M (perhaps large) parts into one

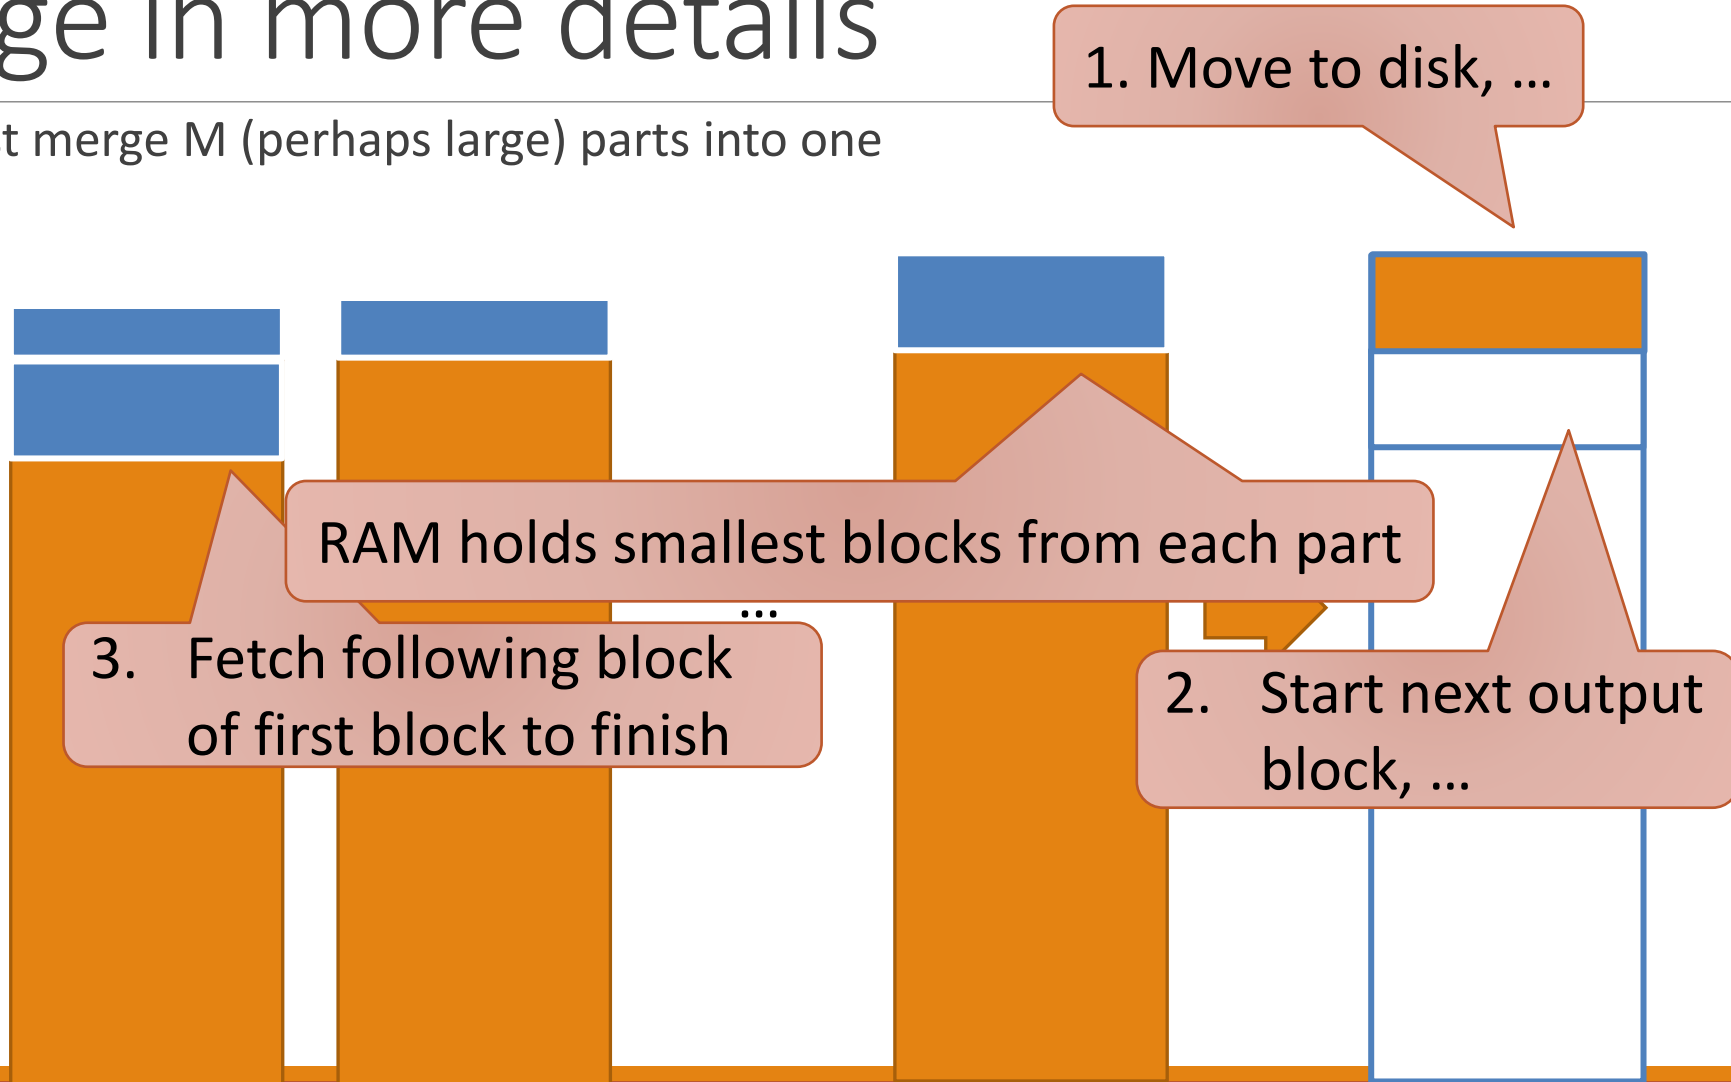
Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one

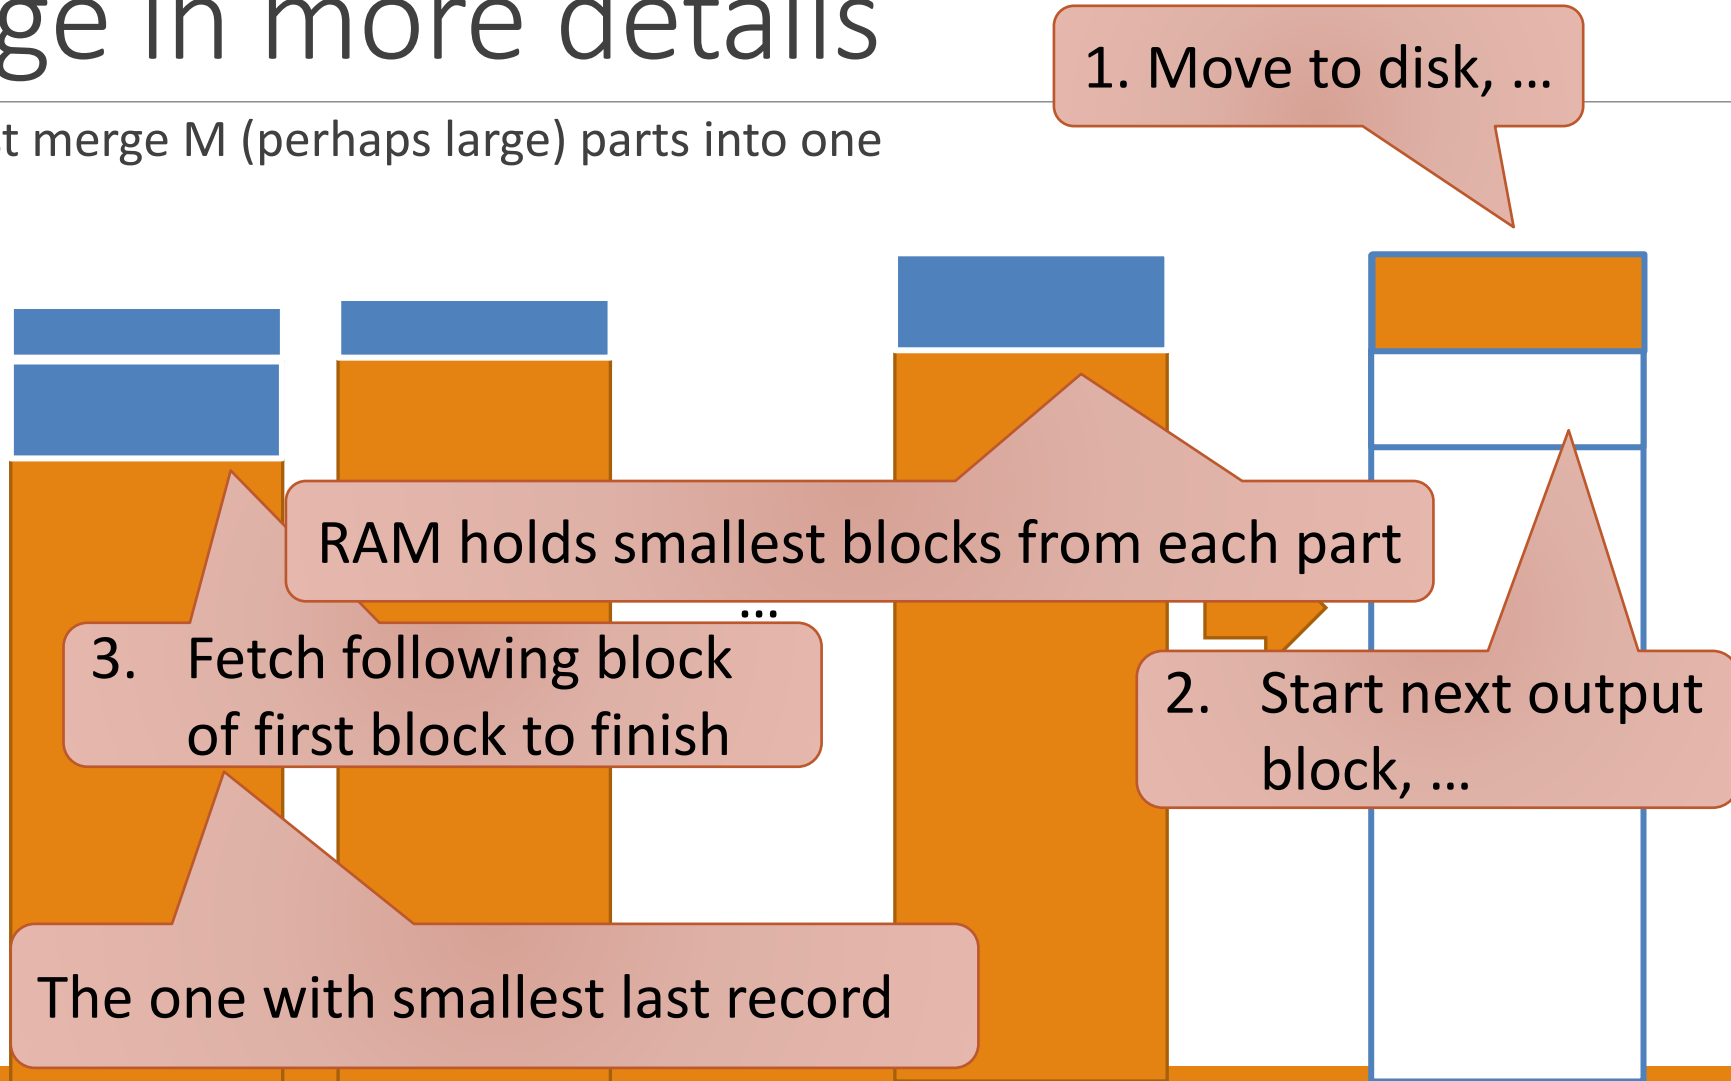
Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one


Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one


Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one


Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one

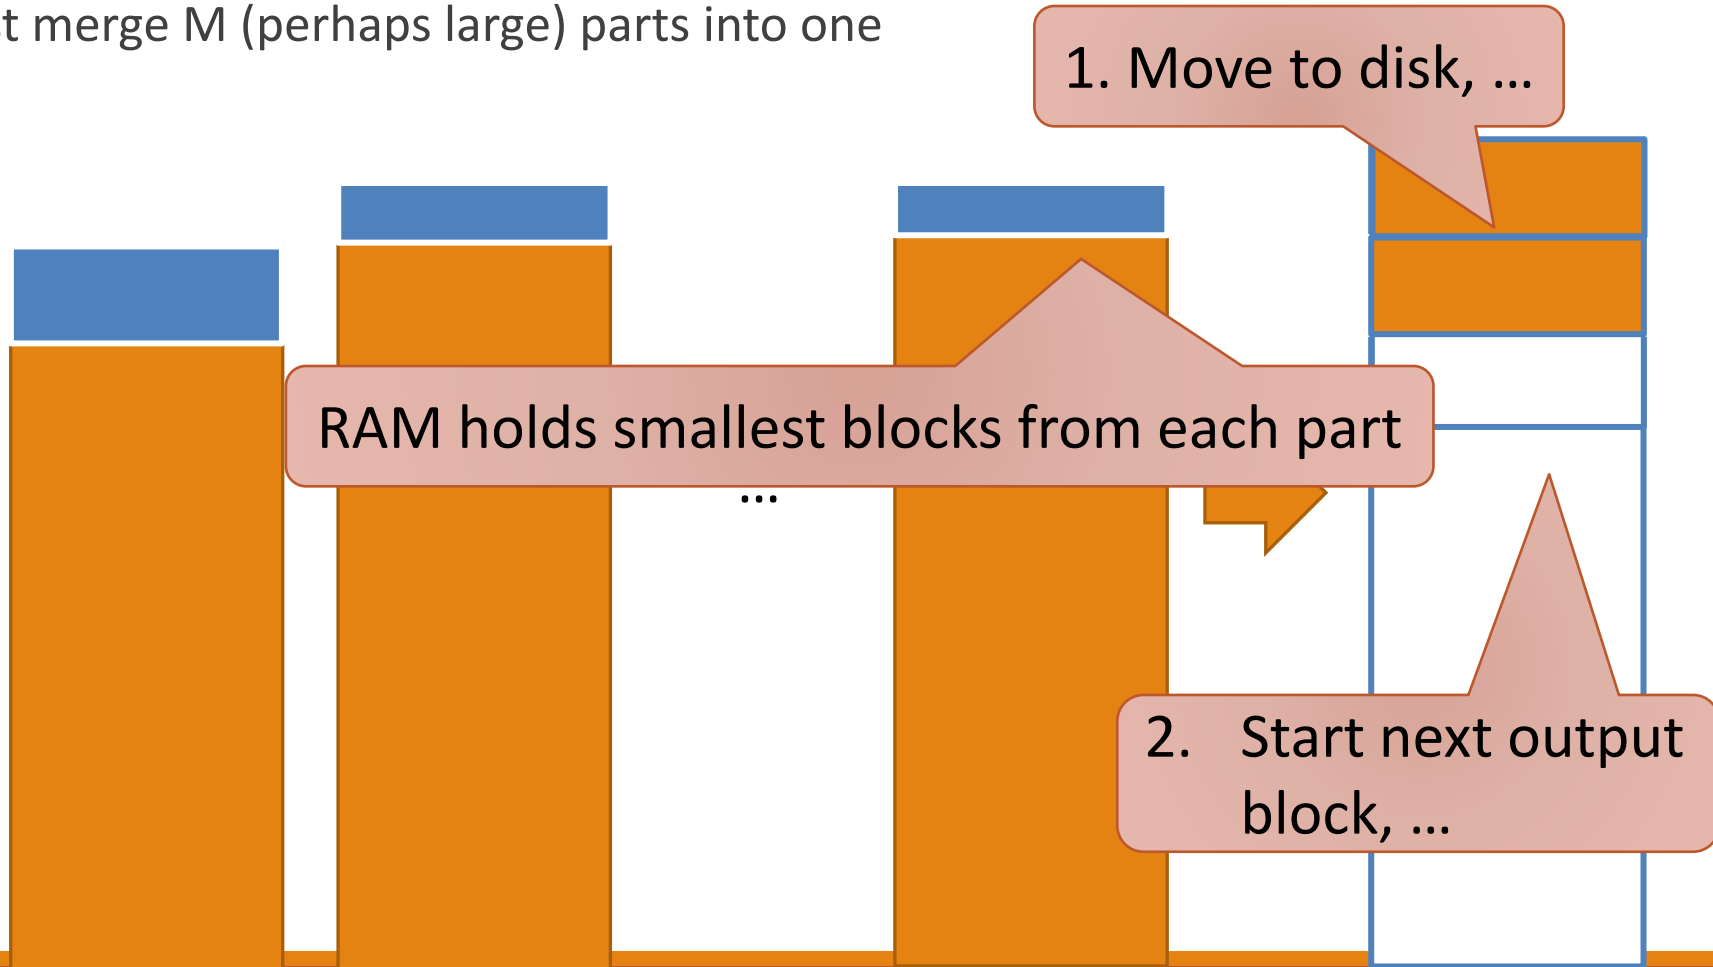
Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet

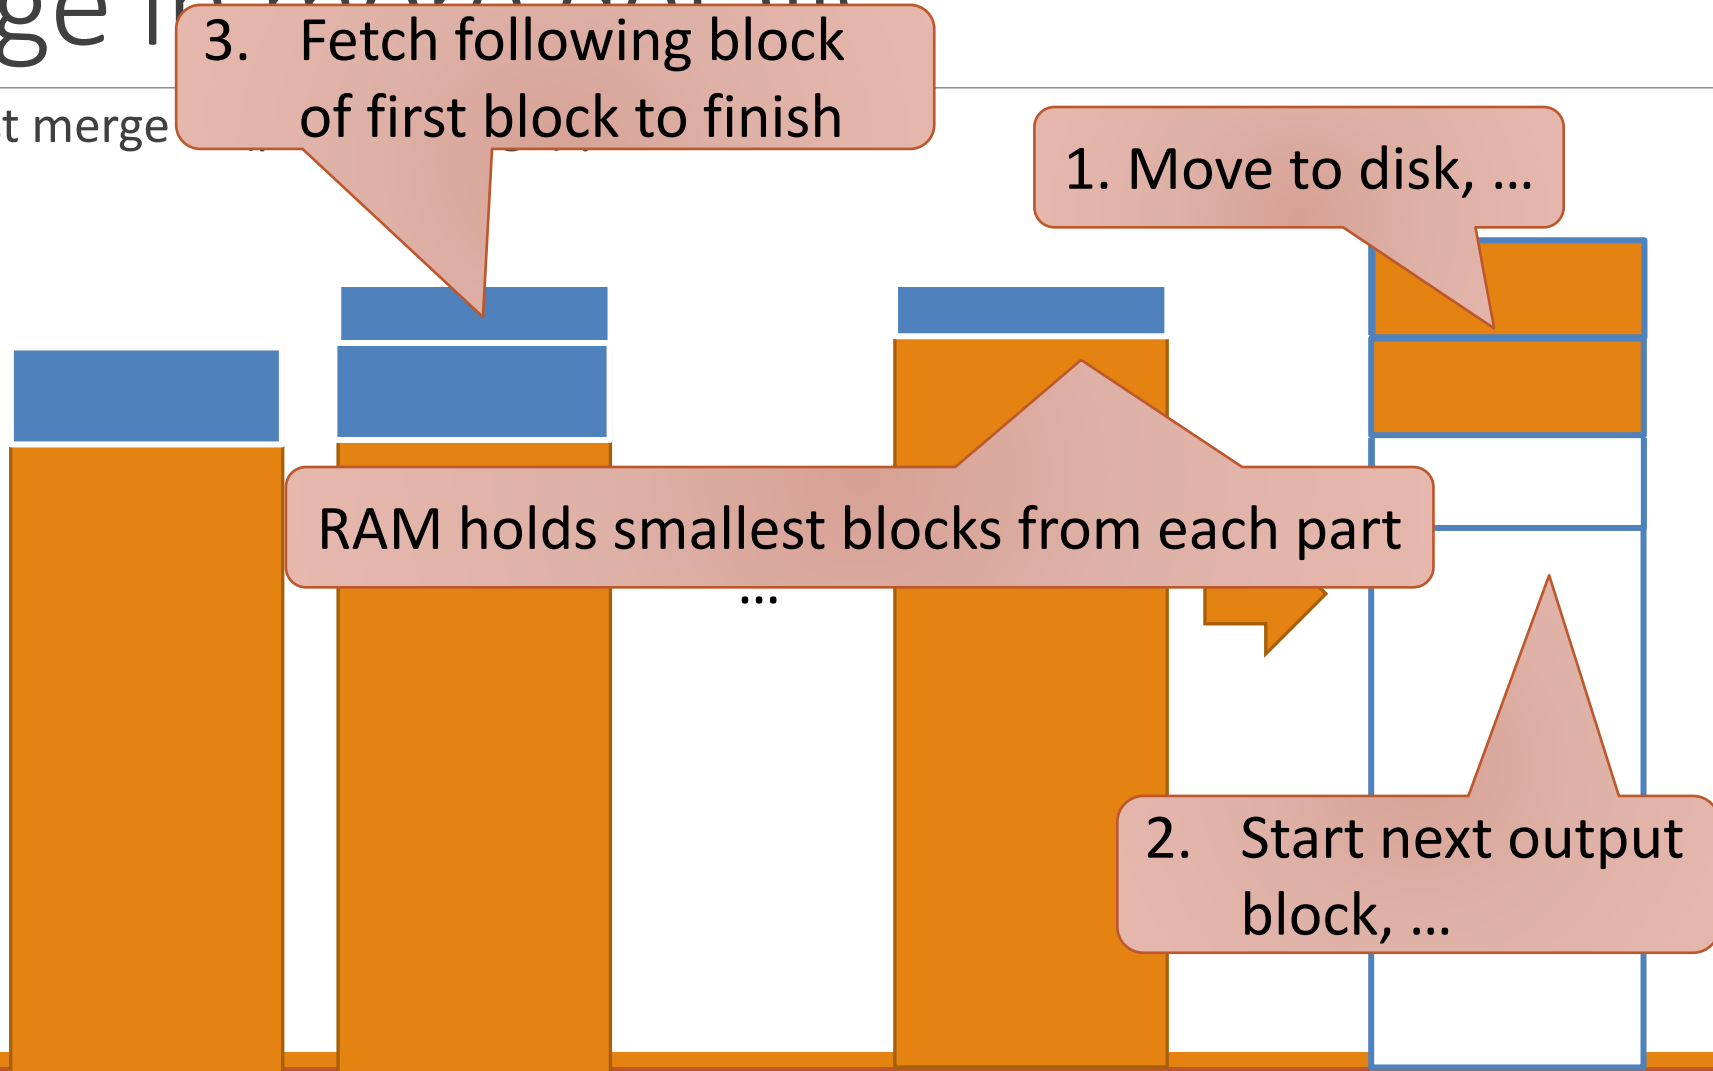




# Merge in more details

Goal: Must merge


Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Merge in more details

Goal: Must merge M (perhaps large) parts into one

Colors:  
Blue in RAM  
Orange in disk  
White does not  
exists yet



# Analysis of External Merge Sort

---

External merge sort:

Merge

**Divide** input in **M** parts  $P_1, P_2, \dots, P_M$

**Sort**  $P_1, P_2, \dots, P_M$  **recursively**

**While** not all  $P_i$  are empty:

Add smallest remaining  
element from any part  $P_i$  to  
output

# Analysis of External Merge Sort

External merge sort:

Merge

**Divide** input in **M** parts  $P_1, P_2, \dots, P_M$   
**Sort**  $P_1, P_2, \dots, P_M$  **recursively**  
**While** not all  $P_i$  are empty:  
    Add smallest remaining  
    element from any part  $P_i$  to  
    output

On each level of recursion we  
scan through all blocks once  
=  $O(N/B)$  disk operations,  
where  $B$  is block size

# Analysis

We split in  $M$  buckets in each level of recursion until remainder is below  $MB$  (where it fits in RAM and can be sorted directly):  $\log_M(N/(MB))$  levels

# Sort

External merge sort:

Merge

**Divide** input into  $M$  parts  $P_1, P_2, \dots, P_M$   
**Sort**  $P_1, P_2, \dots, P_M$  **recursively**  
**While** not all  $P_i$  are empty:  
    Add smallest remaining element from any part  $P_i$  to output

On each level of recursion we scan through all blocks once  
=  $O(N/B)$  disk operations,  
where  $B$  is block size

# Analysis

We split in  $M$  buckets in each level of recursion until remainder is below  $MB$  (where it fits in RAM and can be sorted directly):  $\log_M(N/(MB))$  levels

# Sort

External merge sort:

Total disk operations:  
 $O(N/B \log_M(N/(MB)))$

**Divide** input into  $M$  parts  $P_1, P_2, \dots, P_M$   
**Sort**  $P_1, P_2, \dots, P_M$  **recursively**  
**While** not all  $P_i$  are empty:  
    Add smallest remaining element from any part  $P_i$  to output

On each level of recursion we scan through all blocks once  
=  $O(N/B)$  disk operations,  
where  $B$  is block size

# Comparison

---

In practice, since hard disks are slow:

- $D$ )
- $D$  is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for  $D$ ,  $B$  and  $M$ :

- $D$  was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- $B$  is typically around 512-4096 bytes
- $M$  is the size of your RAM divided by  $B$ , so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- $O(N/B \log_2(N) \times D)$
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$



# Comparison

---

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- Thus, time is  $O(N/B \log_M(N/(MB)) \times D)$
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- D is time for a disk operation
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

D) time!

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

D) time!

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

D) time!

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

D) time!

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Comparison

---

D) time!

D)

In practice, since hard disks are slow:

- Running time is basically time spend on disk operations
- D is time for a disk operation

Quicksort (or other internal memory sorting algorithms) uses  $O(N/B \log_2(N) \times D)$  time!

Common values for D, B and M:

- D was around  $3.1 \cdot 10^4$  **nano-seconds** on SSDs and  $2 \cdot 10^6$  **nano-seconds** on fast normal hard disks
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$
- B is typically around 512-4096 bytes
- M is the size of your RAM divided by B, so say 8-32 GB RAM / 512-4096 bytes:  $2.0 \cdot 10^6$  to  $6.5 \cdot 10^7$

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$



# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

We are thus using 2 recursive calls, for 3 scans in total

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

We are thus using 2 recursive calls, for 3 scans in total

If we used a normal amount of RAM, we would be done in 1 recursive call...

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

We are thus using 2 recursive calls, for 3 scans in total

## (Internal) Quick Sort:

If we used a normal amount of RAM, we would be done in 1 recursive call...

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

We are thus using 2 recursive calls, for 3 scans in total

## (Internal) Quick Sort:

- $\log_2(N)=42$

If we used a normal amount of RAM, we would be done in 1 recursive call...

# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

We are thus using 2 recursive calls, for 3 scans in total

## (Internal) Quick Sort:

- $\log_2(N)=42$

We are thus using 42 recursive calls, for 43 scans in total

If we used a normal amount of RAM, we would be done in 1 recursive call...



# Numeric example

---

Want to sort  $N=4$  TB in 4 MB of RAM with  $B=4096$

## External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$  is then  $\approx \log_{1024}(1024^2) = 2$

We are thus using 2 recursive calls, for 3 scans in total

If we used a normal amount of RAM, we would be done in 1 recursive call...

## (Internal) Quick Sort:

- $\log_2(N)=42$

We are thus using 42 recursive calls, for 43 scans in total

Assuming we are lucky about pivot elements

# Summary

---

External memory sorting is very much faster on inputs that are much larger than main memory