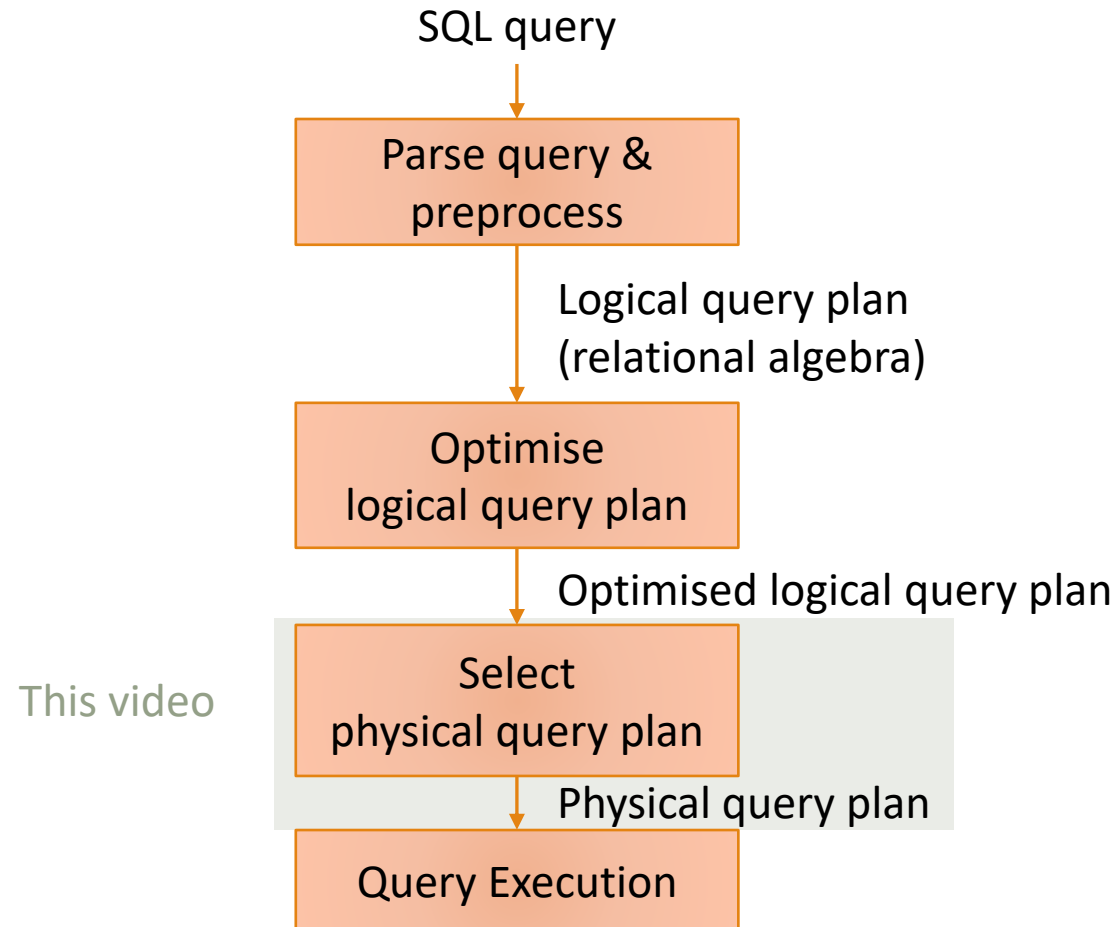


# Physical query plans

---

# Overview of this video

---



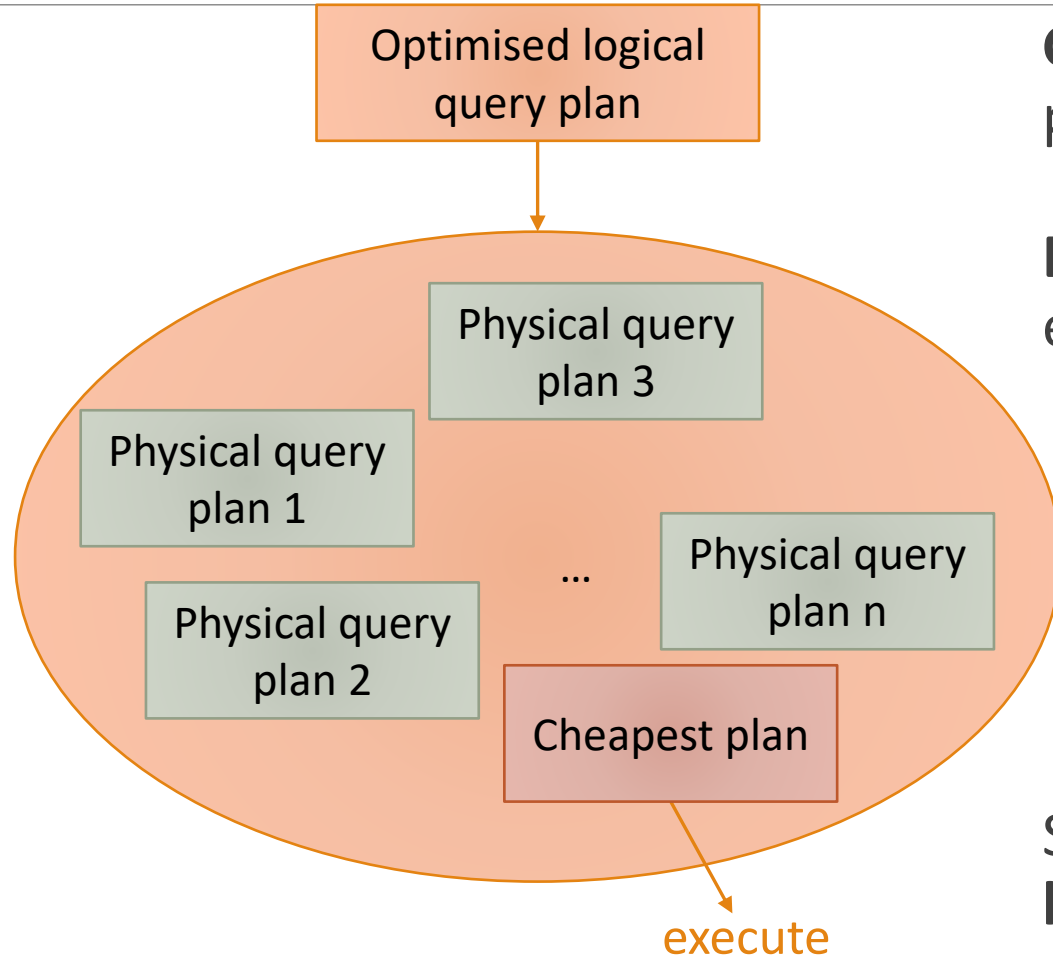
# What is the purpose of a physical query plan?

---

A **physical query plan** adds information required to execute the optimised query plan

- Which **algorithm** to use for execution of operators?
  - Naïve selection or selection with an index?
  - Nested Block Join or Sort Join or Hash Join etc.?
- How to **pass information** from one operator to the other?
  - Write to disk, keep in memory, pipelining operators, etc.?
- Good **order** for computing joins, unions, etc.?
- Additional operations such as sorting

# From Logical Plans To Physical Plans



**Generate** many different physical query plans

**Estimate cost** of execution for each plan

- Time
- Disk accesses
- Memory
- Communication
- ...

Select physical plan with **lowest cost estimate**

# Estimating the Cost of Execution

---

Here: **number of disk access operations**

Number of disk accesses influenced by many factors:

- Selection of algorithms for the individual operators
- Method for passing information
- **Size of intermediate results**



One of the most critical factors

Estimated from parameters of the database

- Important parameters:
  - **Size of relations**
  - **Number of distinct items per attribute per relation**
- Computed exactly from the database or are estimated (“statistics gathering”)


# Estimating Intermediate Result Sizes

One of the most challenging tasks of a DBMS

- Difficult even for nodes close to the leaves
- Cannot afford executing the query
- Rely on statistics gathered from data

Many different approaches

- Some easier than others
- With join size estimation, we enter active research...



# Estimating the Size of a Selection

---

Estimate for the size of  $\sigma_{A=a}(R)$ :

Recall:  $|R|$  = number of tuples in  $R$

$$\frac{|R|}{\text{number of distinct values in column } A \text{ of relation } R}$$

Estimate for the # of blocks required to store  $\sigma_{A=a}(R)$ :

$$\frac{\text{number of blocks for } R}{\text{number of distinct values in column } A \text{ of relation } R}$$

Good if values in column  $A$  of  $R$  occur equally often, but can be bad

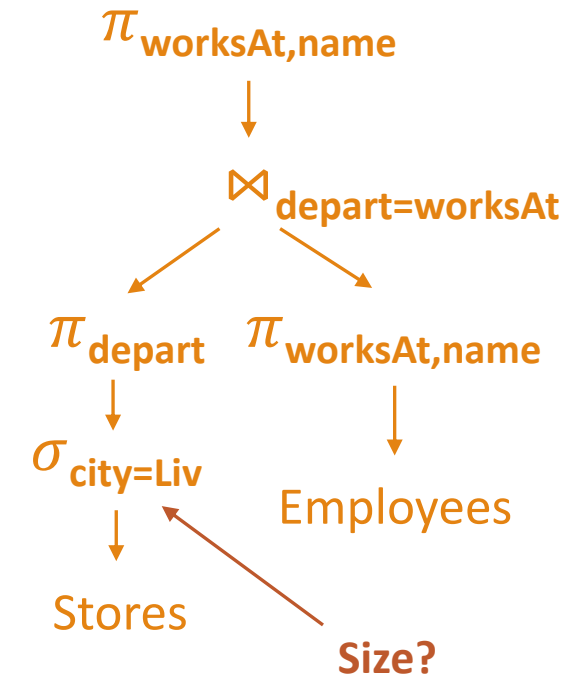
# Example

Assume:

- Stores contains **80 tuples**, stored in **20 blocks**
- There are exactly **4 distinct values** for the city attribute of Stores

Estimate for size of  $\sigma_{\text{city=Liv}}(\text{Stores})$ :  
 $80/4 = 20$  tuples

Estimate for number of blocks that are required to store  $\sigma_{\text{city=Liv}}(\text{Stores})$ :  $20/4 = 5$  blocks





# Joins

Assume A is the only common attribute.

How to estimate  $R \bowtie S$ ?

Simple estimate based on size of  $R$  &  $S$  and number of distinct values in common attributes

$$\frac{|R| \times |S|}{\text{max. number of distinct values for } A \text{ in } R \text{ or } S}$$

As for selection, based on assumptions that might not always lead to good estimates

More sophisticated methods:

- Still a topic of active research
- See, e.g., SIGMOD/PODS/VLDB conferences

# Other Issues

---

## How to generate physical query plans?

- Explore all?
- More sensible approaches: top-down/bottom-up

## Selection of a **suitable algorithm** for each operator


- based on size of intermediate result

## Selection of a **good join order**

- also based on size of intermediate results


# Example where join order matters

```
SELECT *  
FROM Lecturers NATURAL JOIN Modules NATURAL JOIN Students  
WHERE Lecturers.name = Rasmus Ibsen-Jensen
```



# Example where join order matters

```
SELECT *  
FROM Lecturers NATURAL JOIN Modules NATURAL JOIN Students  
WHERE Lecturers.name = Rasmus Ibsen-Jensen
```



# Other Issues

---

**How to generate** physical query plans?

- Explore all?
- More sensible approaches: top-down/bottom-up

Selection of a **suitable algorithm** for each operator

- based on size of intermediate result

Selection of a **good join order**

- also based on size of intermediate results

How to **pass information** from one operator to another?

# Passing Information

---

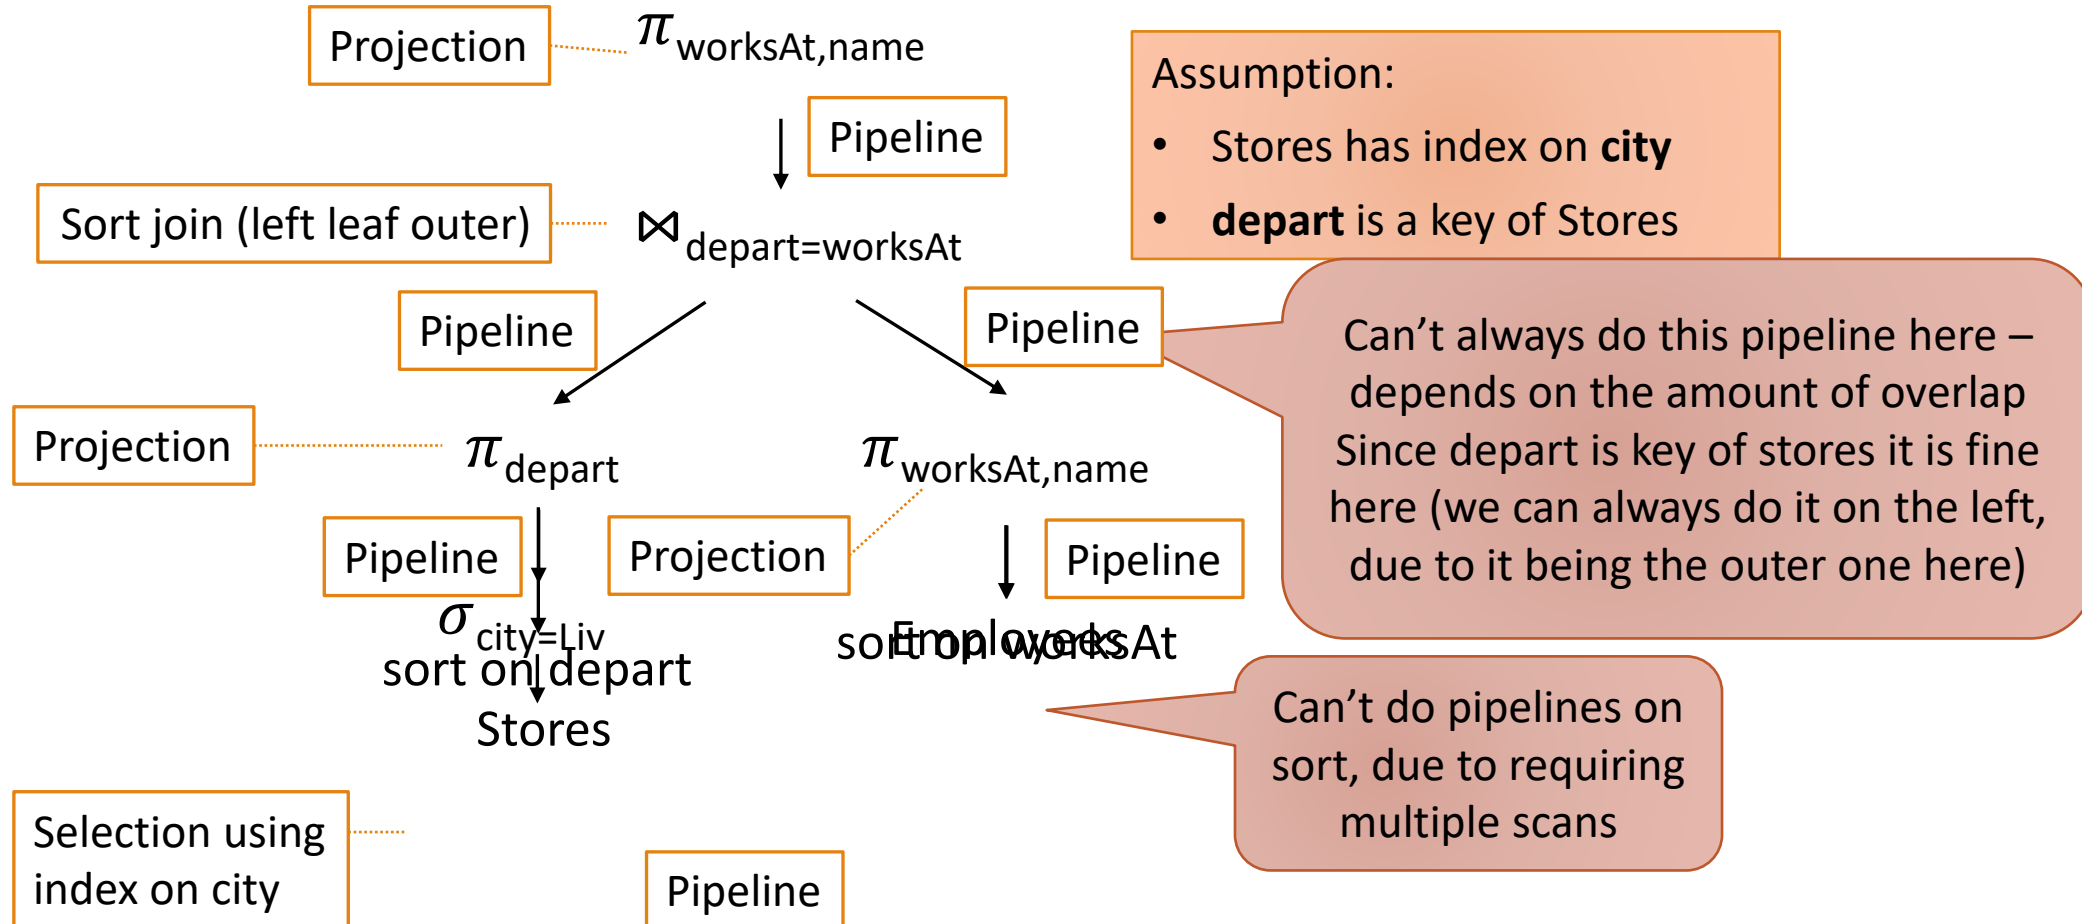
Materialisation: write intermediate results to disk

Pipelining (“stream-based processing”)

- Passes the tuples of one operation directly to the next operation without using disk
- Extra buffer for each pair of adjacent operations to hold tuples passing from one relation to the other
- Example:
  - $\pi_{\text{title,year}}(\sigma_{\text{length} \geq 100 \text{ AND studio} = \text{'Fox'}}(\text{Film}))$
  - With pipelining, the intermediate result of the selection will be written into a buffer in memory, from which the projection operator will read and process these tuples directly



# From Logical to Physical Query Plan



# Summary

---

In this part, we have seen how to go from a SQL query to how exactly we are going to solve it!

Specifically, in this video, we saw how to go from a optimised logical query plan to a physical ditto