

# Query processing for DDBBMS

---


# Overview over this video

---

We discuss some problems that comes up in DDBMS for query processing and an example of how to do better than the naïve approach in some cases

# Query Processing in Distributed DBMS

---




Try to answer query at site where query is raised

If not possible: request information from other sites

- Slow → design database to reduce this as much as possible
- Most expensive: joins

# Joins

Goal: compute join at site **B**



Obvious approach:

- Site **B** asks site **A** to send **R**
- Site **B** computes  $R \bowtie S$

R might be very large – do we have to send all tuples?

- Better: only send data that is actually required

# Semijoins ( $\bowtie$ )

$$R \bowtie S := R \bowtie \pi_{\text{common attributes of R and S}}(S)$$

**Modules**

module	year
COMP105	1
COMP201	2
COMP207	2

**Lecturers**

name	module
J. Fearnley	COMP105
S. Coope	COMP201

**Modules  $\bowtie$  Lecturers**


=

module	year
COMP105	1
COMP201	2

Intuition:  $R \bowtie S$  = set of all tuples in  $R$  that NATURAL JOIN  
at least one tuple in  $S$

# Semijoin Reduction

Goal: compute join at site **B**



With semijoins:

- Site **B** sends  $S' := \pi_{\text{common attributes of R and S}}(S)$  to site **A**
- Site **A** sends  $R' := R \bowtie S (= R \bowtie S')$  to site **B**
- Site **B** outputs  $R' \bowtie S$

Communication costs  $\approx$

$$|S'| \times (\text{size of tuple in } S') + |R'| \times (\text{size tuple in } R')$$

# Efficiency

---

Is this more efficient than computing the join in the obvious approach way (exchanging relations)?

Depends:

- Is the projection much smaller than the full relation?
  - Many duplicates to be eliminated? I.e., do many tuples of S share values of the common attributes? Not the case if one of the join attributes is a key...
- Is the size of the semijoin much smaller?
- In general:  $|\pi_{\text{common attributes}}(S)| + |R \bowtie S|$  should be smaller than  $|R|$
- Note that: You can calculate if sending  $\pi_{\text{common attributes}}(S)$  is less than sending R, if so, you can try. In the worst case you end sending twice as much as needed (e.g. in case  $\pi_{\text{common attributes}}(S)$  is close to the size of R and  $R \bowtie S$  is R)

# Example

$\text{Films} \bowtie \sigma_{\text{city}='Liverpool' \text{ AND } \text{date}=2020-12-3}(\text{InTheatres})$



**A**

Films(film\_title, genre, ...)



**B**

InTheatres(film\_title, theatre, city, date, ...)

## Procedure:

- At **B**, send  $S' := \pi_{\text{film\_title}}(\sigma_{\text{city}='Liverpool' \text{ AND } \text{date}=2020-12-3}(\text{InTheatres}))$  to **A**
- At **A**, send  $R' := \text{Films} \bowtie S'$  to **B**
- At **B**, output  $R' \bowtie \sigma_{\text{city}='Liverpool' \text{ AND } \text{date}=2020-12-3}(\text{InTheatres})$

## Communication costs:

- Assume  $|\text{Films}| = 10,000$ ,  $|S'| = |R'| = 20$ , and 1000 bytes per tuple
- Communication cost =  $(20+20) \times 1000 = 40,000$  bytes

10,000,000 bytes for obvious approach



# Summary

---

Query processing have new challenges in DDBMS, because of the distributed nature of the database

- We saw how to use Semijoins to in some cases speed up joins

Also, if you are watching these in order, this was it for the required part on distributed databases

- There are some more videos on MapReduce – special programming framework that fits somewhere between distributed databases and the next topic of NoSQL databases (i.e. not only SQL databases)
- Those are not required for the exam though