

Transaction management in DDBMS

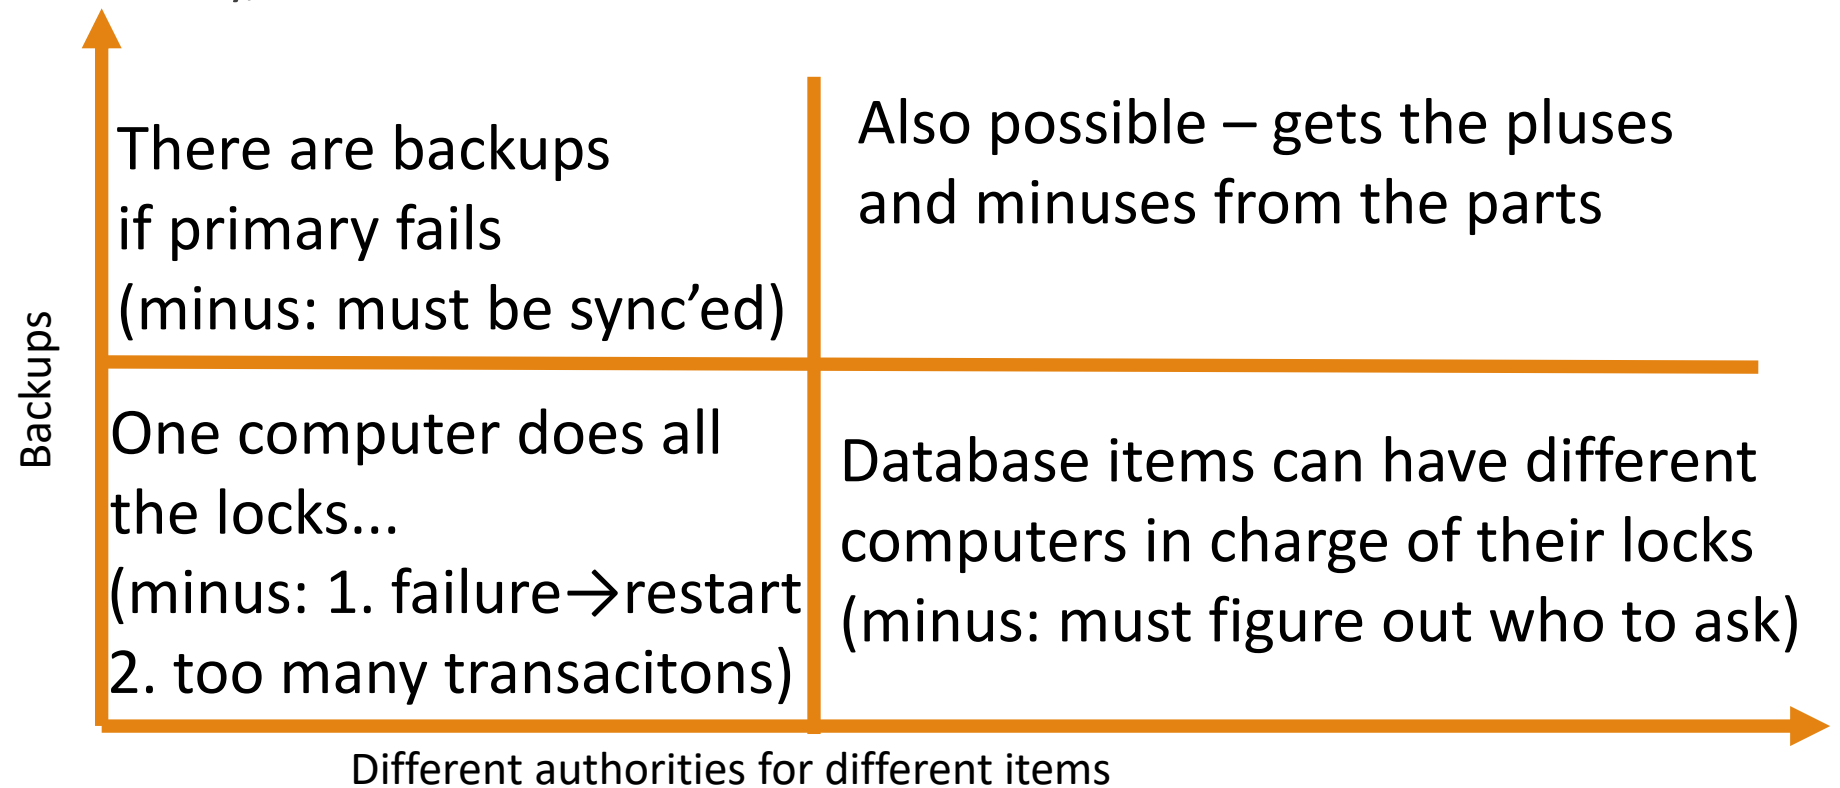
Overview over this video

Distributed databases creates new issues for transaction management

In this video we see what some of them are and ways to deal with these

Concurrency control in DDBMS

For full isolation/consistency, often based on locks:



Concurrency control based on voting

Another approach to locks (instead of having a single designated computer granting them):

Voting!

Idea:

- Each site with a copy of an item has a local lock that it can grant transactions for that item
- If a transaction gets over half the local locks for an item, it has a global lock on the item
 - If so, it must tell the sites with a copy that it has the lock
 - If it takes too long, it must stop trying to get the lock
- Plus: Much more distributed than the non-voting approach
- Minus: Requires more communication


Recovery in DDBMS

Transactions in Distributed Databases

Let's revisit our CS_Store chain...

At central office:

- Determine inventory for product X at each site
- Move product X between stores to balance inventory



Global transaction T

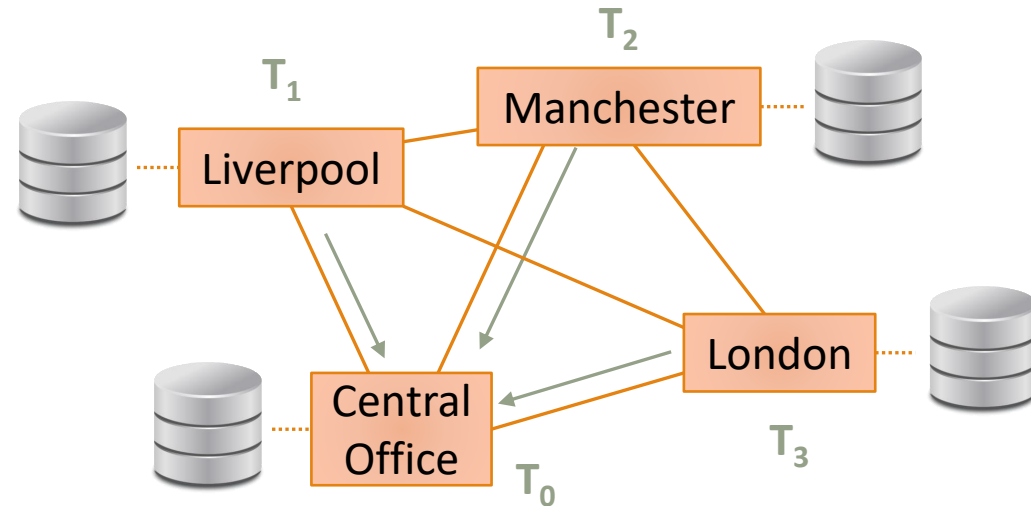
- Starts **local transaction** T_0 at central office
- T_0 instructs other sites to start **local transactions** T_1, T_2, T_3

Transactions in Distributed Databases

Let's revisit our CS_Store chain...

At central office:

- Determine inventory for product X at each site
- Move product X between stores to balance inventory



Global transaction T

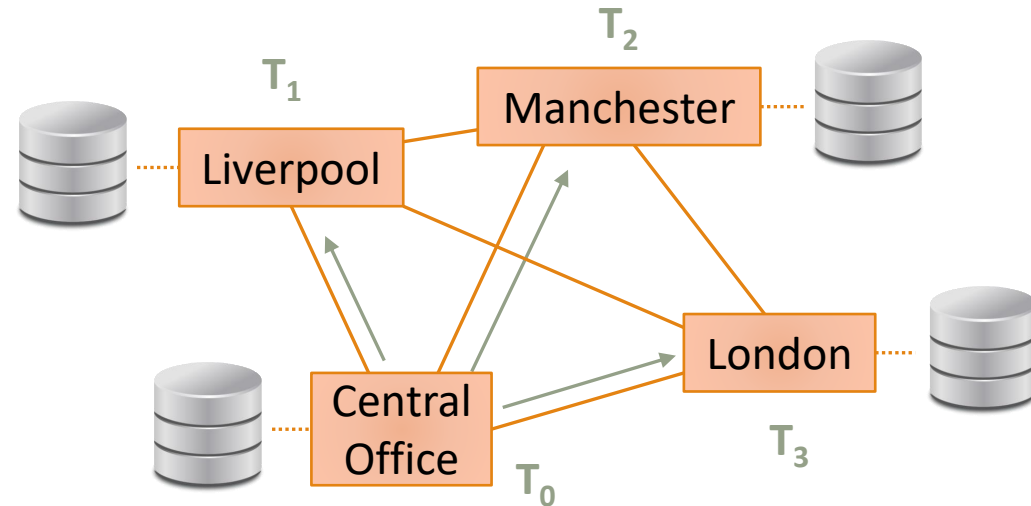
- Starts **local transaction** T_0 at central office
- T_0 instructs other sites to start **local transactions** T_1, T_2, T_3
- T_1, T_2, T_3 find out inventory for product X at sites & send it back to T_0

Transactions in Distributed Databases

Let's revisit our CS_Store chain...

At central office:

- Determine inventory for product X at each site
- Move product X between stores to balance inventory




Global transaction T

- Starts **local transaction** T_0 at central office
- T_0 instructs other sites to start **local transactions** T_1, T_2, T_3
- T_1, T_2, T_3 find out inventory for product X at sites & send it back to T_0
- T_0 determines how to move product X between sites
- T_0 instructs T_1, T_2, T_3 to move product X accordingly

Violation of Atomicity

Global transaction T

- Start T_0 at central office
- T_0 instructs other sites to start T_1, T_2, T_3
- T_1, T_2, T_3 report inventory for product X
- T_0 determines how to move product X
- T_0 instructs T_1, T_2, T_3 to move product X accordingly



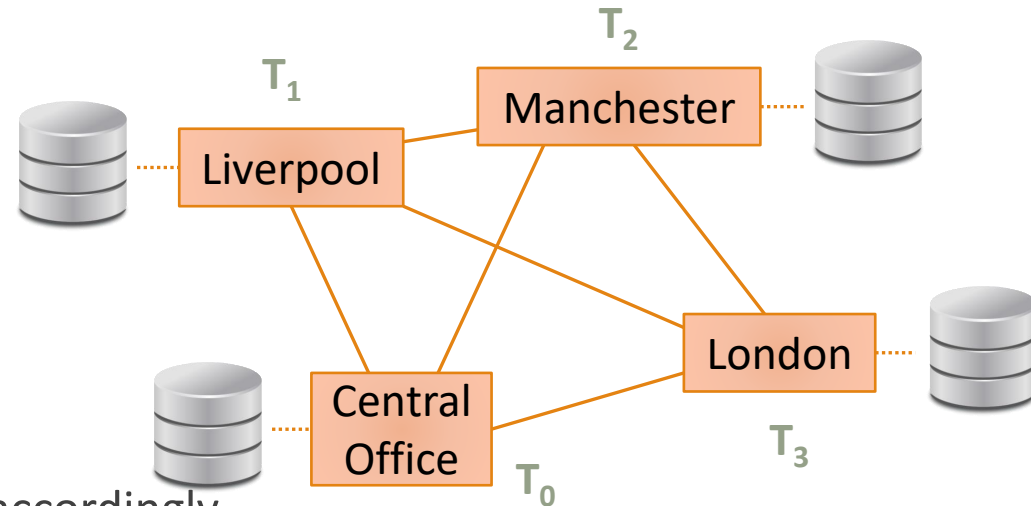
Atomicity:

- Can assume to be enforced at each node *locally*
- Could be violated *globally*

Problems With Failing Nodes

Global transaction T

- Start T_0 at central office
- T_0 instructs other sites to start T_1, T_2, T_3
- T_1, T_2, T_3 report inventory for product X
- T_0 determines how to move product X
- T_0 instructs T_1, T_2, T_3 to move product X accordingly



Nodes can fail during execution of T


- Should we abort or wait?
- What about the failing node after recovery?

Distributed Commit

Not related to 2PL!

Two-Phase Commit Protocol

Coordinates commit actions globally



Coordinator: executed at some node & decides if and when local transactions can commit

Logging: at each node locally

- Messages sent to & received from other nodes are logged, too!

The Two Phases


Phase 1: Decide when to commit or abort

Phase 2: Commit or abort

Phase 1: When To Commit?

Coordinator: ask nodes if they want to commit

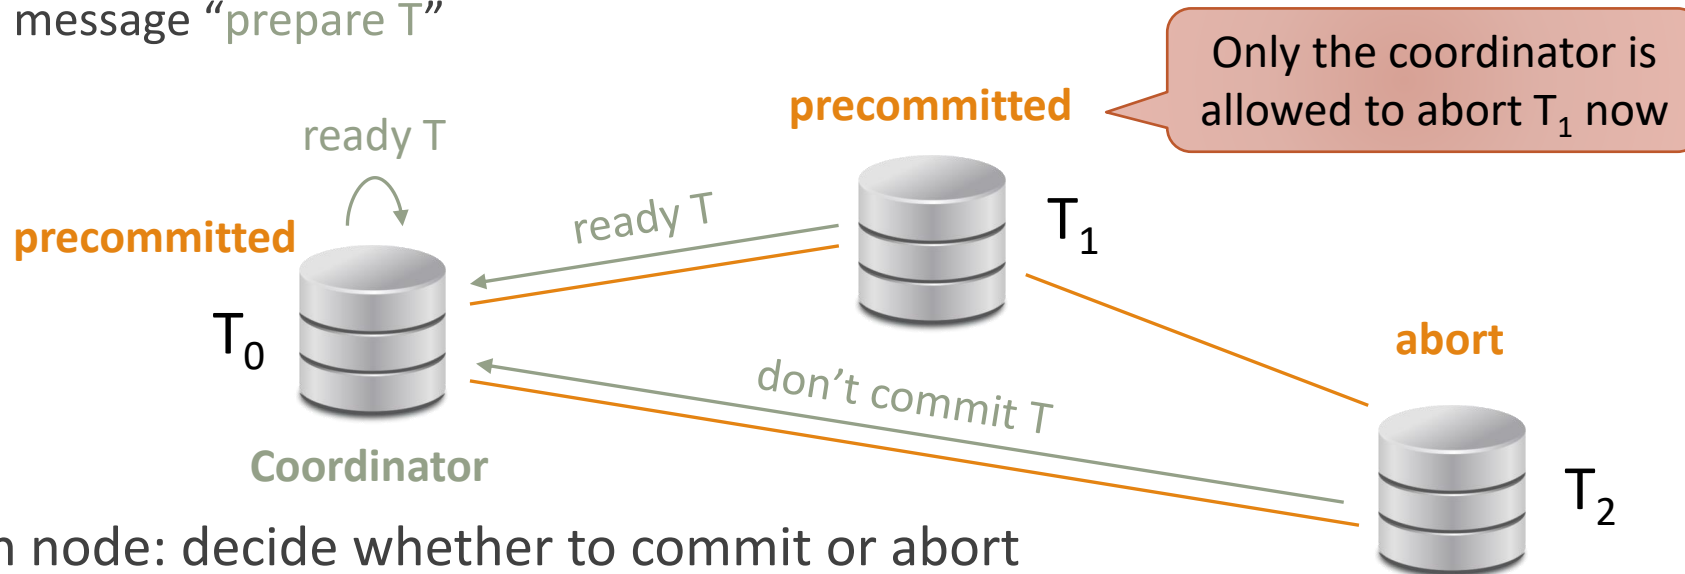
- Send message “prepare T”



Phase 1: Ready to Commit?

Coordinator: ask nodes if they want to commit

- Send message “prepare T”



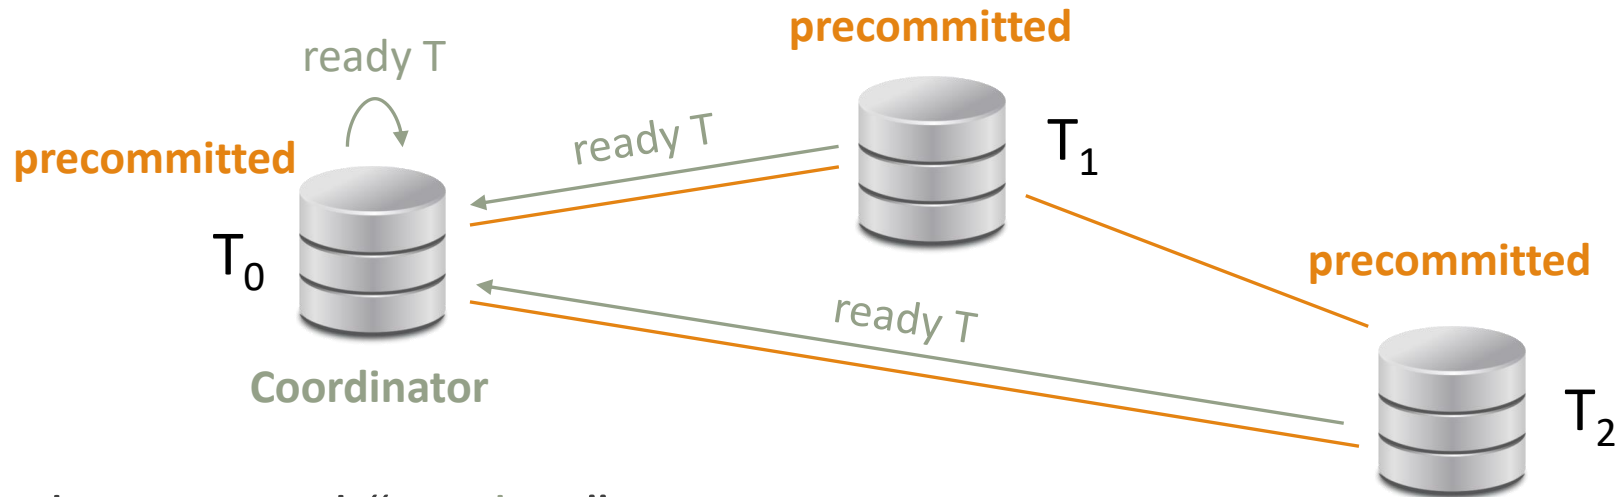
At each node: decide whether to commit or abort

- If commit → go into **precommitted** state & send back “ready T”
- If abort → send back “don't commit T” and abort local transaction
- Can delay, but must decide eventually

Phase 2: Let's Do It

Coordinator: waits for responses of nodes

- Assume nodes who don't reply before a given timeout wish to abort

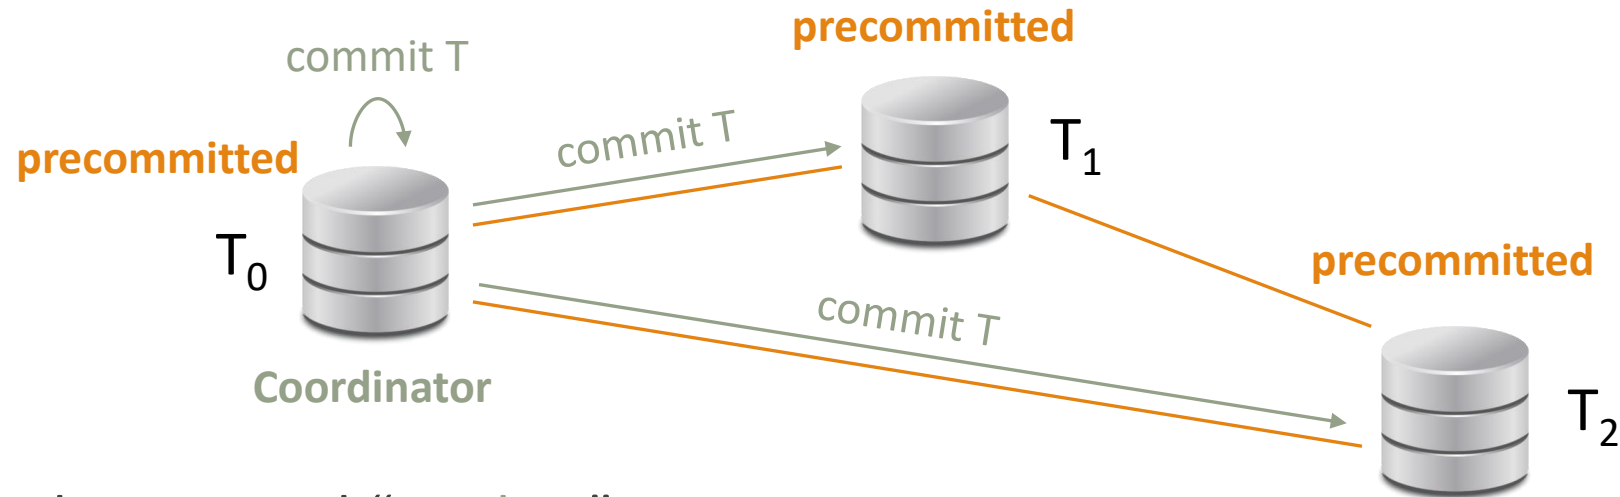


If all nodes respond "ready T"

Phase 2: Let's Do It

Coordinator: waits for responses of nodes

- Assume nodes who don't reply before a given timeout wish to abort



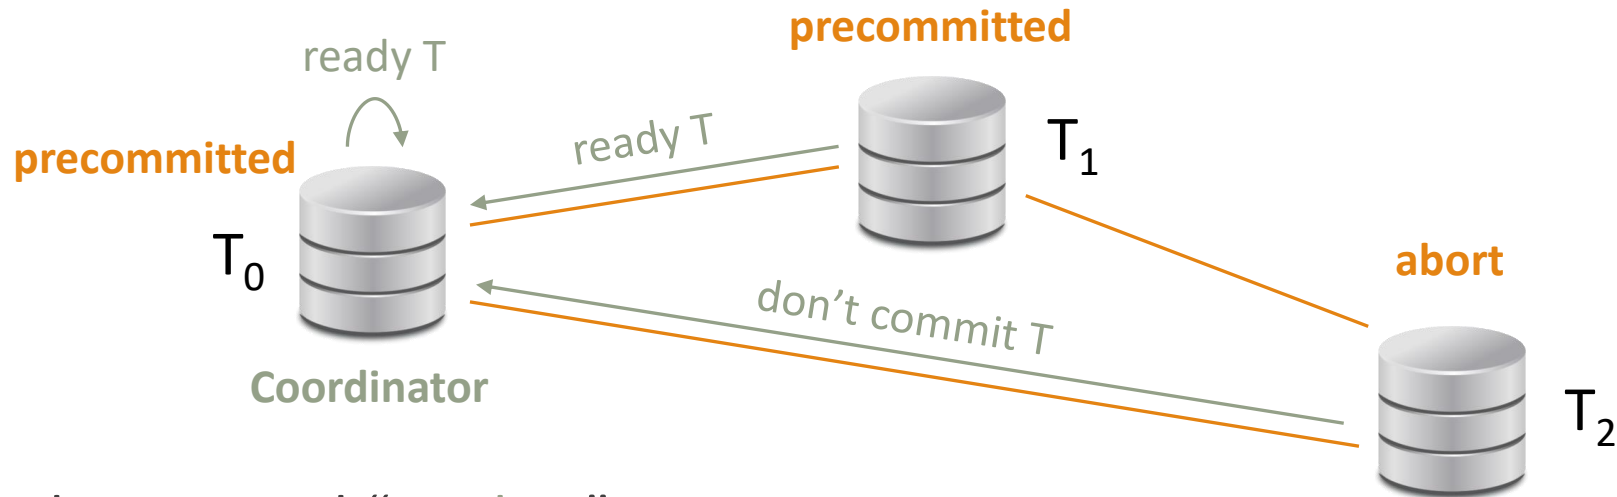
If all nodes respond "ready T"

- Send "commit T" to all nodes → nodes commit

Phase 2: Let's Do It

Coordinator: waits for responses of nodes

- Assume nodes who don't reply before a given timeout wish to abort



If all nodes respond "ready T"


- Send "commit T" to all nodes → nodes commit

If some node responds "don't commit T"

Phase 2: Let's Do It

Coordinator: waits for responses of nodes

- Assume nodes who don't reply before a given timeout wish to abort



If all nodes respond "ready T"

- Send "commit T" to all nodes → nodes commit


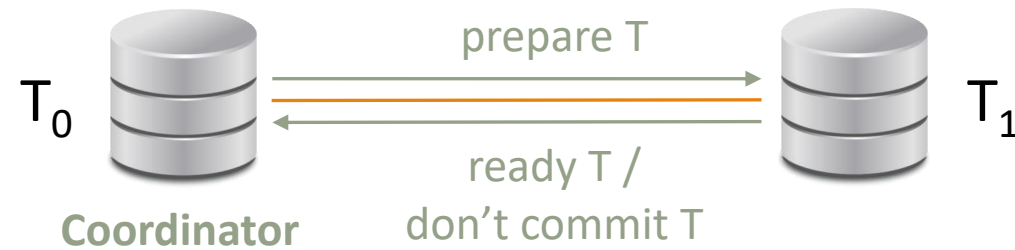
If some node responds "don't commit T"

- Send "abort T" to all nodes → nodes abort

Logging: Phase 1

Again, we have to be careful in which order to write to disk and to the log


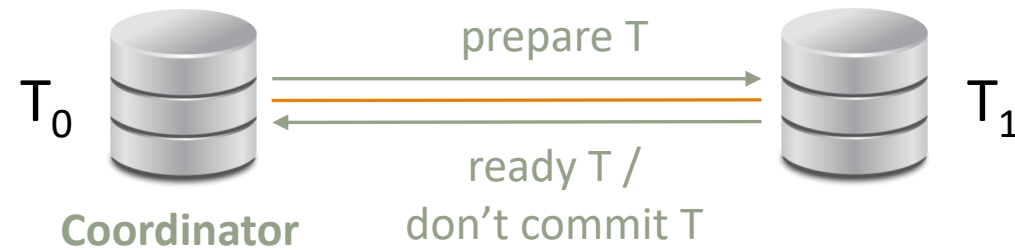
Phase 1:



Logging: Phase 1

Again, we have to be careful in which order to write to disk and to the log


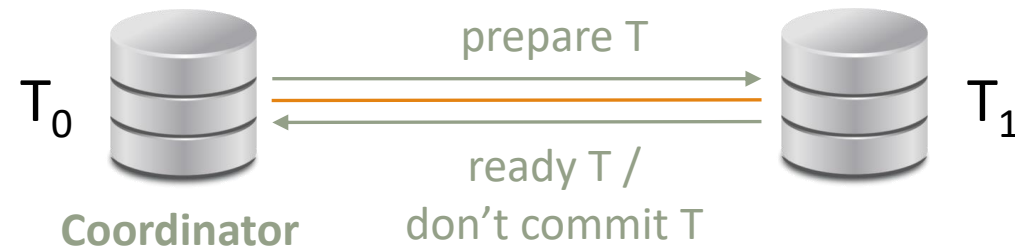
Phase 1:



Logging: Phase 1

Again, we have to be careful in which order to write to disk and to the log

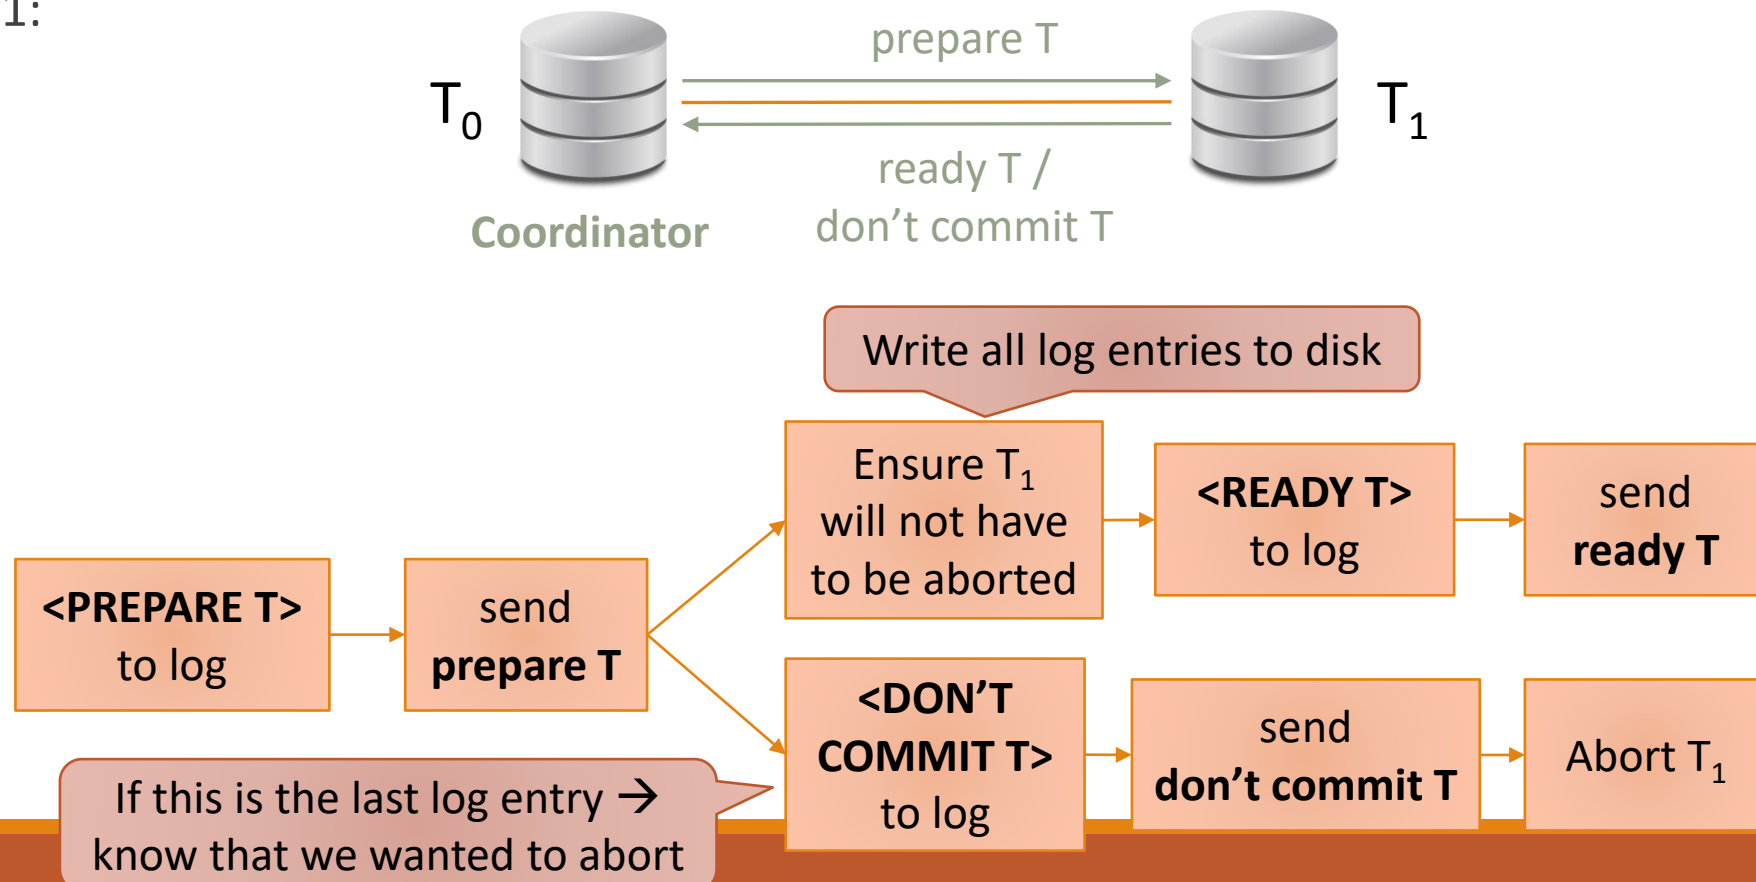
Phase 1:



Logging: Phase 1

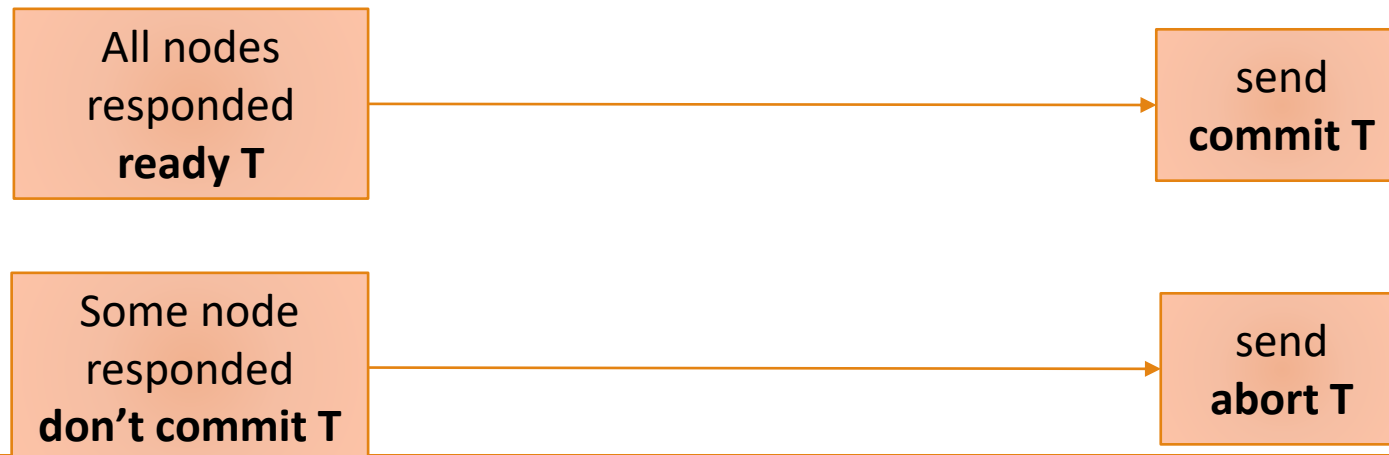
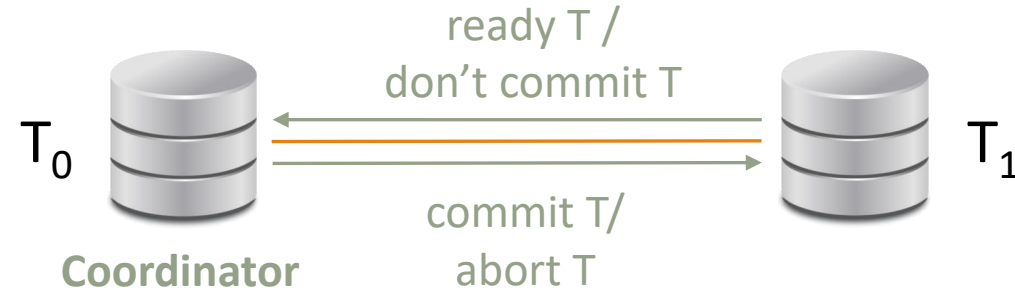
Again, we have to be careful in which order to write to disk and to the log

Phase 1:



Logging: Phase 2

Phase 2:



Logging: Phase 2

Phase 2:



If this is the last log entry \rightarrow decision was to commit

In case of failure, redo T_1

All nodes
responded
ready T

<COMMIT T>
to log

send
commit T

Similar to commit

Some node
responded
don't commit T

<ABORT T>
to log

send
abort T

Three-Phase Commit Protocol

Improvement of Two-Phase Commit

Can deal with the situation that in phase 2, the coordinator and some transaction crash, while everybody else are in precommitted state

Idea: divide phase 2 into two parts

- Phase 2(a): “Prepare to Commit”
 - Send the decision (commit/abort) to all nodes
 - Nodes go into prepare-to-commit state
- Phase 2(b): “Commit”
 - The old Phase 2

Advantage: if the coordinator fails, all nodes know if they should commit/abort a transaction

Summary

Saw a few ways to deal with concurrency control in DDMBs

- Two dimensions:
 - Different computers could be the primary for different database items or not
 - Primary have backups or not
- Voting as alternative

Saw 2PC and 3PC (two- and three-phase commit – no relation to 2PL, besides the words “two” and “phase”) for coordinating commits to deal with recovery

Did not deal with deadlocks...