



# Relational Data Model – Part 1

## Schema and State

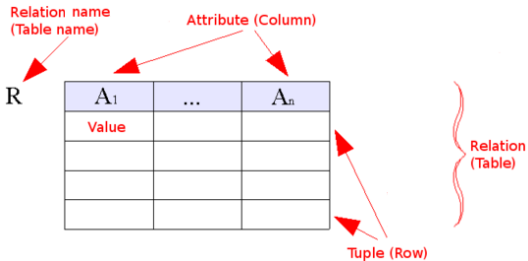
## What is the Relational Data Model?

- Introduced by Edgar F. Codd of IBM Research in 1970.  
*“A Relational Model for Large Shared Data Banks”, Communications of the ACM.*
- A database contains tables (called relations), and each table is made up of columns and rows.
- Humans have used tables for centuries to keep track of data.



- Used as the standard for relational DBMSs (e.g., Oracle, IBM DB2, Microsofts Access, Microsofts SQL Server, MySQL, postgresSQL, etc.).

## Relation



- Correspondence of informal and formal terms:

INFORMAL TERMS	FORMAL TERMS
Table	Relation
Column	Attribute
Data type	Domain
Row	Tuple
Table definition	Relation schema

## The Basics

- **Attributes** are used to describe the properties of information. In the relational model, they usually refer to atomic data.

**Example:** To capture the information of a person, we can use attributes like Name, Age, Gender, Address and PhoneNumber.

- **Domains** are the sets of all possible values for attributes.

- $\text{STRING} = \{A, B, CD, \dots\};$

**Example:** •  $\text{DATE} = \{01/01/2005, 03/07/1978, \dots\};$

- $\text{INT} = \{\dots, -1, 0, 1, 2, \dots\}.$

- Recall that, **Cartesian product**  $D_1 \times \dots \times D_n$  is the set of all possible combinations of values from the sets  $D_1, \dots, D_n$ .

**Example:** Let  $D_1 = \{\text{book}, \text{pen}\}$ ,  $D_2 = \{1, 2\}$  and  $D_3 = \{\text{red}\}$ . Then

- $D_1 \times D_2 \times D_3 = \{(\text{book}, 1, \text{red}), (\text{book}, 2, \text{red}), (\text{pen}, 1, \text{red}), (\text{pen}, 2, \text{red})\}$



## The Basics

- The attributes are StudentID, CourseNo, Semester, Status and EnrolDate.
- The domains of attributes are as follows.  
 $\text{dom}(\text{StudentID})=\text{INT};$   $\text{dom}(\text{CourseNo})=\text{STRING};$   
 $\text{dom}(\text{Semester})=\text{STRING};$   $\text{dom}(\text{Status})=\text{STRING};$   
 $\text{dom}(\text{EnrolDate})=\text{DATE}.$
- The whole table can be considered as a set  $\{(456, \text{COMP2400}, 2016 \text{ S2}, \text{active}, 25/05/2016), (458, \text{COMP1130}, 2016 \text{ S1}, \text{active}, 20/02/2016), (459, \text{COMP2400}, 2016 \text{ S2}, \text{active}, 11/06/2016)\}.$

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2400	2016 S2	active	11/06/2016

- Is the above set a subset of

$\text{INT} \times \text{STRING} \times \text{STRING} \times \text{STRING} \times \text{DATE}?$

**Answer: Yes.**

## The Basics

- A **relation schema** has a **relation name** and a list of **attributes**.
- Each attribute is associated with a **domain**.
- A relation schema can be expressed by
  - $R(A_1, \dots, A_n)$ , or
  - $R(A_1 : \text{dom}(A_1), \dots, A_n : \text{dom}(A_n))$ ,

where  $A_1, \dots, A_n$  are attributes of  $R$  and  $\text{dom}(A_i)$  is the domain of  $A_i$ .

**Example:** The relation schema in the previous example is

- $\text{ENROL}(\text{StudentID}, \text{CourseNo}, \text{Semester}, \text{Status}, \text{EnrolDate})$ , or
- $\text{ENROL}(\text{StudentID: INT}, \text{CourseNo: STRING}, \text{Semester: STRING}, \text{Status: STRING}, \text{EnrolDate: DATE})$ .

## The Basics

- Let  $R(A_1, \dots, A_n)$  be a relation schema.
- A **tuple** in  $R$  is a list  $t$  of values, i.e.,  $t \in \text{dom}(A_1) \times \dots \times \text{dom}(A_n)$ .

**Example:** The previous example has the following tuples:

- $(456, \text{COMP2400}, 2016 \text{ S2}, \text{active}, 25/05/2016) \in$   
 $\text{INT} \times \text{STRING} \times \text{STRING} \times \text{STRING} \times \text{DATE}.$
  - $(458, \text{COMP1130}, 2016 \text{ S1}, \text{active}, 20/02/2016) \in$   
 $\text{INT} \times \text{STRING} \times \text{STRING} \times \text{STRING} \times \text{DATE}.$
  - $(459, \text{COMP2400}, 2016 \text{ S2}, \text{active}, 11/06/2016) \in$   
 $\text{INT} \times \text{STRING} \times \text{STRING} \times \text{STRING} \times \text{DATE}.$
- A **relation**  $r(R)$  is a set of tuples  $r(R) \subseteq \text{dom}(A_1) \times \dots \times \text{dom}(A_n)$ .

**Example:** The previous example has the following relation:

- $r(\text{ENROL}) \subseteq \text{INT} \times \text{STRING} \times \text{STRING} \times \text{STRING} \times \text{DATE}.$



## The Basics

- A **relational database schema**  $S$  is
  - a set of relation schemas  $S = \{R_1, \dots, R_m\}$ , and
  - a set of integrity constraints  $IC$ .
- A **relational database state** of  $S$  is a set of relations such that
  - there is just one relation for each relation schema in  $S$ , and
  - all the relations satisfy the integrity constraints  $IC$ .





## The Basics

- Consider a relational database schema STUENROL that has three relation schemas:
  - STUDENT(StudentID, Name, DoB, Email).
  - COURSE(No, Cname, Unit);
  - ENROL(StudentID, CourseNo, Semester, Status, EnrolDate);

STUDENT			
StudentID	Name	DoB	Email

COURSE		
No	Cname	Unit

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate

- That is, STUENROL={STUDENT, COURSE, ENROL}.



## The Basics

- Relational Database State – Example

STUDENT			
StudentID	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com

COURSE		
No	Cname	Unit
COMP1130	Introduction to Advanced Computing I	6
COMP2400	Relational Databases	6

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2400	2016 S2	active	11/06/2016



# Relational Data Model – Part 2

## Integrity Constraints



## Integrity Constraints over Relations

- Constraints are **conditions** that must hold on *all* relations in a database state.
- The *main types* of constraints in the relational data model include:
  - 1 **Domain constraints;**
  - 2 **Key constraints;**
  - 3 **Entity integrity constraints;**
  - 4 **Referential integrity constraints.**



## (1) Domain Constraints

- Every value in a tuple must be from the **domain of its attribute**.
  - INT
  - VARCHAR
  - DATE
  - SMALLINT
  - NOT NULL



## (2) Key Constraints - Observation

- We observe that: **data does not occur independently from one another within individual relations.**
- No two students have the same student ID:

STUDENT			
StudentID	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com
...	...	...	...

- No two enrolments have the same student ID, the same course number in the same semester:

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2400	2016 S2	active	11/06/2016
...	...	...	...	...



## (2) Key Constraints - Definitions

- Let  $R(A_1, \dots, A_n)$  be a relation schema.
- A **superkey**  $SK$  of  $R$  is a subset of attributes of  $R$ , i.e.,  $SK \subseteq \{A_1, \dots, A_n\}$ , such that
  - no two distinct tuples in  $r(R)$  can have the same value for  $SK$ .
- A superkey  $SK$  of  $R$  is **minimal** if there is no other superkey  $SK' \subset SK$  held on  $R$ . A minimal superkey is also known as a **candidate key**.
- A **primary key**  $PK$  of  $R$  is a minimal superkey of  $R$ , (i.e., a primary key is one of the candidate keys). If a relation has only one candidate key then that would be the primary key.

## (2) Key Constraints - Example

STUDENT			
StudentID	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com
460	Tyrion	11/09/1987	tyrion@hotmail.com

- Is {DoB} a superkey of STUDENT? **No!**
- Is {StudentID, DoB} a superkey of STUDENT? **Yes!**
- Is {StudentID, DoB} a candidate key of STUDENT? **No!**
- Is {StudentID} a candidate key of STUDENT? **Yes!**
- Can {StudentID} be chosen as a primary key of STUDENT? **Yes!**
- Can {DoB} be chosen as a primary key of STUDENT? **No!**





## (2) Key Constraints - Example

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2400	2016 S2	active	11/06/2016
458	COMP1130	2015 S1	inactive	20/02/2015

- Is {CourseNo, Semester} a superkey of ENROL? **No!**
- Is {StudentID, CourseNo, Semester} a candidate key of ENROL? **Yes!**
- Can {StudentID, CourseNo} be chosen as a primary key of ENROL? **No!**



### (3) Entity Integrity Constraints

- Specifying a primary key also invokes the entity integrity constraint.
- **null** is a special value, which represents the value of an attribute that may be unknown or inapplicable.
- The **entity integrity constraint** states that **no primary key value can be NULL**.
  - This is because primary key values are used to *identify* individual tuples in a relation.
- **Note:** Other attributes of  $R$  may be constrained to disallow null values, even though they are not attributes in the primary key.



### (3) Entity Integrity Constraints – Example

- If STUDENTID is specified as the primary key of STUDENT, then the following relation violates the entity integrity constraint.

STUDENT			
StudentID	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
NULL	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com

- How about the case when EMAIL is the primary key of STUDENT?

**Answer: The relation does not violate the entity integrity constraint.**



## (4) Referential Integrity Constraints - Observation

- We observe that: **data does not occur independently from one another across relations.**
- Every course number appearing in ENROL must exist in COURSE:

STUDENT			
StudentID	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com

COURSE		
No	Cname	Unit
COMP1130	Introduction to Advanced Computing I	6
COMP2400	Relational Databases	6

ENROL				
StudentID	CourseNo	Semester	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2400	2016 S2	active	11/06/2016

## (4) Referential Integrity Constraints - Definition

- We use  $t[A]$  to denote the value of attribute  $A$  in tuple  $t$ .  
**Example:** For the tuple  $t=(459, \text{Fran}, 11/09/1987, \text{frankk@gmail.com})$ ,  $t[\text{Name}]=\text{Fran}$  and  $t[\text{DoB}]=11/09/1987$ .
- A **referential integrity constraint** specifies a reference between **two** relations, while the previous constraints involve **only one** relation.
- Let  $R_1$  and  $R_2$  be relation schemas in a database schema  $S$ , and  $R_2$  has the primary key  $\{B_1, \dots, B_n\}$ .
- A **foreign key** on  $R_1$  is a statement  $[A_1, \dots, A_n] \subseteq R_2[B_1, \dots, B_n]$  restricting states of  $S$  to satisfy the following property:
  - for each tuple  $t \in r(R_1)$  there exists a tuple  $t' \in r(R_2)$  with  $t[A_i] = t'[B_i]$  for  $i = 1, \dots, n$ .
- $R_1$  is called the **referencing relation** and  $R_2$  is called the **referenced relation**.



## (4) Referential Integrity Constraints – Example

- What foreign keys can be established in the database STUENROL?

STUDENT			
<u>StudentID</u>	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com

COURSE		
<u>No</u>	Cname	Unit
COMP1130	Introduction to Advanced Computing I	6
COMP2400	Relational Databases	6

ENROL				
<u>StudentID</u>	<u>CourseNo</u>	<u>Semester</u>	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2400	2016 S2	active	11/06/2016



## (4) Referential Integrity Constraints – Example

- In this case, we can establish the following foreign keys on ENROL:
  - 1 [CourseNo]  $\subseteq$  COURSE[No];
  - 2 [StudentID]  $\subseteq$  STUDENT[StudentID].
- This database state satisfies the above two foreign keys because
  - for each tuple  $t_1$  in ENROL, there is a tuple  $t_2$  in COURSE such that the CourseNo value in  $t_1$  is the same with the No value in  $t_2$ ;
  - for each tuple  $t'_1$  in ENROL, there is a tuple  $t'_2$  in STUDENT such that the StudentID value in  $t'_1$  is the same with the StudentID value in  $t'_2$ .

## (4) Referential Integrity Constraints – Question

- If the database STUENROL is slightly changed as follows, does this database still satisfy the foreign keys in the previous example?

STUDENT			
<u>StudentID</u>	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com

COURSE		
No	Cname	Unit
COMP1130	Introduction to Advanced Computing I	6
COMP2400	Relational Databases	6

ENROL				
<u>StudentID</u>	<u>CourseNo</u>	<u>Semester</u>	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2600	2016 S2	active	11/06/2016





## (4) Referential Integrity Constraints – Question

**Answer:** The following database does not satisfy the foreign key of  
**ENROL: [CourseNo]  $\subseteq$  COURSE[No].**

STUDENT			
<u>StudentID</u>	Name	DoB	Email
456	Tom	25/01/1988	tom@gmail.com
458	Peter	23/05/1993	peter@gmail.com
459	Fran	11/09/1987	frankk@gmail.com

COURSE		
No	Cname	Unit
COMP1130	Introduction to Advanced Computing I	6
COMP2400	Relational Databases	6

ENROL				
<u>StudentID</u>	<u>CourseNo</u>	<u>Semester</u>	Status	EnrolDate
456	COMP2400	2016 S2	active	25/05/2016
458	COMP1130	2016 S1	active	20/02/2016
459	COMP2600	2016 S2	active	11/06/2016



## Constraint Violations

- There are three basic operations that can change a database state:
  - **Insert**: insert one or more new tuples in a relation;
  - **Delete**: delete tuples in a relation;
  - **Update** (or **Modify**): change the values of attributes in existing tuples.
- Whenever these operations are applied, the integrity constraints specified in a database schema **should not be violated**.
- However,
  - Insert may violate ...
  - Delete may violate ...
  - Update may violate ...