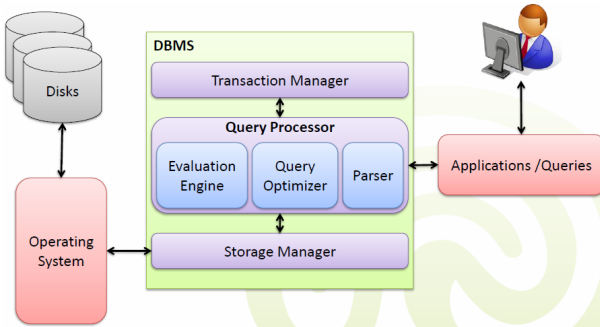




Query Processing

Query Processing – Overview

- 1 Users **submit** SQL queries to a DBMS.
- 2 The DBMS **processes and executes** them in a database.



- **Note:** SQL is a declarative language, so it is the task of DBMSs to decide how SQL queries should be executed.



Query Processing – Example

- **From:**

```
SELECT name FROM Person WHERE age<21;
```

- **To:**

name
Rickon Bran

- **Questions:**

- How does a relational DBMS process this?
- How can a relational DBMS process this efficiently?




Query Processing – Example

```
SELECT name FROM Person WHERE age<21;
```

High-level language
(SQL)

\Downarrow
 $\pi_{name}(\sigma_{age<21}(\text{Person}))$

Low-level language
(Relational Algebra)



Execution plan
(Query tree)

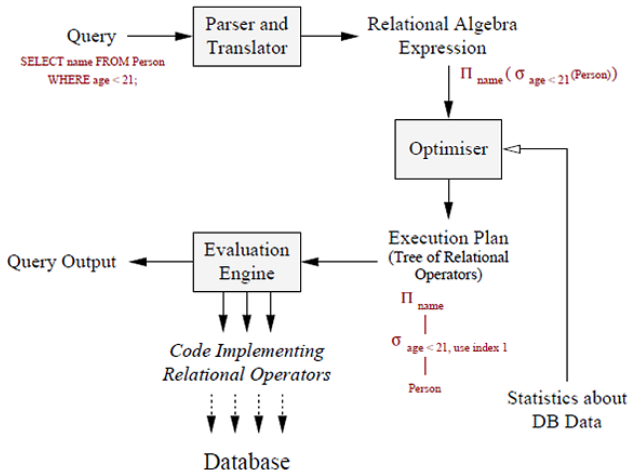
\Downarrow

name
Rickon
Bran

Query result



Query Processing – Example



Query Processing Steps

- **Query parser and translator**

- 1 Check the syntax of SQL queries
- 2 Verify that the relations do exist
- 3 Transform into relational algebra expressions

- **Query optimiser**

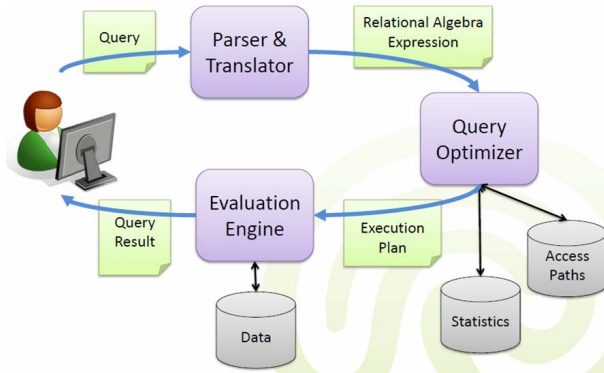
- 1 Transform into the best possible execution plan
- 2 Specify the implementation of each operator in the execution plan

- **Evaluation engine**

- 1 Evaluate the query execution plan
- 2 Return the result to the user

Query Processing – Parser

- The **parser** checks the syntax of the query:
 - Validation of table names, attributes, data types, access permission ...;
 - Either the query is executable or an error message is generated.





Query Processing – Parser

- Consider the relation schema:

Person(id:integer, name:string, age:integer, address:string)

- Note:** **System catalog** (also called **data dictionary**) is used at this stage, which contains the information about data managed by the DBMS.

Example:

attr_name	rel_name	type	position
id	Person	integer	1
name	Person	string	2
age	Person	integer	3
address	Person	string	4
...

- Question:** Can the following query be accepted by the parser?

```
SELECT fname, lname FROM Person WHERE address<21;
```




Query Processing – Parser

- Consider the relation schema:

`Person(id:integer, name:string, age:integer, address:string)`

- Question:** Can the following query be accepted by the parser?

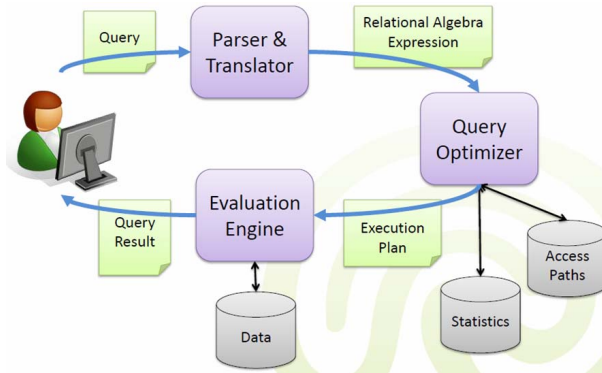
```
SELECT fname, lname FROM Person WHERE address<21;
```

- Answer:** The query **would be rejected** because

- The attributes `fname` and `lname` are not defined;
- The attribute `address` is not comparable with 21.

Query Processing – Translator

- The **translator** translates queries into RA expressions (not necessarily equivalent due to duplicates):
 - A query is first decomposed into **query blocks**.
 - Each query block is translated into an RA expression.





Recall: RA and SQL Queries

● RA operators

- **selection** σ_{φ}
- **projection** π_{A_1, \dots, A_n}
- **Cartesian product** $R_1 \times R_2$
- **join** $R_1 \bowtie_{\varphi} R_2$ and $R_1 \bowtie R_2$
- **renaming** $\rho_{R(A_1, \dots, A_n)}$
- **union** $R_1 \cup R_2$
- **intersection** $R_1 \cap R_2$
- **difference** $R_1 - R_2$

● SQL statement

```
SELECT attribute_list
  FROM table_list
  [WHERE condition]
  [GROUP BY attribute_list
  [HAVING group_condition]]
  [ORDER BY attribute_list];
```

$\sigma_{\varphi}(R) \Leftrightarrow \text{SELECT } * \text{ FROM } R \text{ WHERE } \varphi;$

$\pi_{A_1, \dots, A_n}(R) \Leftrightarrow \text{SELECT DISTINCT } A_1, \dots, A_n \text{ FROM } R;$

$R_1 \times R_2 \Leftrightarrow \text{SELECT DISTINCT } * \text{ FROM } R_1, R_2;$

...

- Aggregate operations in SQL require extended RA expressions.



Recall: RA and SQL Queries

- Nested subqueries are decomposed into separate query blocks.
- **Example:**

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary > (SELECT Salary
                FROM EMPLOYEE
                WHERE ssn=5);
```

Outer query block

```
SELECT Lname, Fname FROM EMPLOYEE WHERE
Salary > c
```

⇓ translated

$\pi_{Lname, Fname}(\sigma_{Salary > c}(EMPLOYEE))$

Inner query block

```
(SELECT Salary FROM EMPLOYEE WHERE
ssn=5)
```

⇓ translated

$\pi_{Salary}(\sigma_{ssn=5}(EMPLOYEE))$



Query Processing – Query Optimiser

- 1 Transform into the best possible execution plan

There are different possible relational algebra expressions for a single query!

(will be covered in this course)

- 2 Specify the implementation of each operator in the execution plan

There are different possible implementations for a relational algebra operator!

(will not be covered in this course)



Query Processing – Query Optimiser

- SQL queries only specify **what data to be retrieved** and **not how to retrieve data**.
- There are **many possible execution plans** for a SQL query.
- Query optimiser is responsible for identifying **an efficient execution plan**:
 - 1 enumerating alternative plans (typically, a subset of all possible plans);
 - 2 choosing the one with the least estimated cost.
- Query optimisation is one of the most important tasks of a relational DBMS.
A good DBMS must have a good query optimiser!



Equivalent RA Expressions

- Suppose that we have:

Students(matNr, firstName, lastName, email)

Exams(matNr, crsNr, result, semester)

Courses(crsNr, title, unit)

```
SELECT lastName, result, title
FROM STUDENTS, EXAMS, COURSES
WHERE STUDENTS.matNr=EXAMS.matNr AND
      EXAMS.crsNr=COURSES.crsNr AND result≤1.3;
```

- **Question:**

How many equivalent RA expressions for this SQL query can you find?



Equivalent RA Expressions

Students(matNr, firstName, lastName, email)

Exams(matNr, crsNr, result, semester)

Courses(crsNr, title, unit)

```
SELECT lastName, result, title
```

```
FROM STUDENTS, EXAMS, COURSES
```

```
WHERE STUDENTS.matNr=EXAMS.matNr AND
```


```
EXAMS.crsNr=COURSES.crsNr AND result≤1.3;
```

● Answer:

- 1 $\pi_{lastName, result, title}(\sigma_{result \leq 1.3}((\text{Students} \bowtie_{\text{Students.matNr}=\text{Exams.matNr}} \text{Exams}) \bowtie_{\text{Exams.crsNr}=\text{Courses.crsNr}} \text{Courses}))$
- 2 $\pi_{lastName, result, title}(\sigma_{result \leq 1.3}(\sigma_{\text{EXAMS.crsNr}=\text{Courses.crsNr}}(\sigma_{\text{STUDENTS.matNr}=\text{Exams.matNr}}(\text{Students} \times \text{Exams} \times \text{Courses}))))$
- 3 $\pi_{lastName, result, title}((\text{Students} \bowtie_{\text{Students.matNr}=\text{Exams.matNr}} (\sigma_{result \leq 1.3}(\text{Exams}))) \bowtie_{\text{Exams.crsNr}=\text{Courses.crsNr}} \text{Courses})$

Query Trees

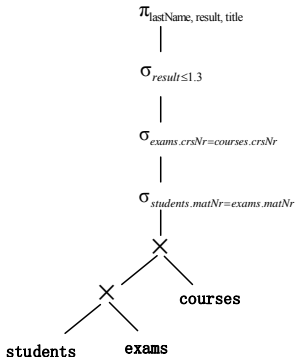
- Each RA expression can be represented as a **query tree**:
 - leaf nodes** represent the input relations;
 - internal nodes** represent the intermediate result;
 - the root node** represents the resulting relation.
- Example:**

$$\pi_{lastName, result, title}(\sigma_{result \leq 1.3}((Students \bowtie_{Students.matNr=Exams.matNr} Exams) \bowtie_{\sigma_{Exams.crsNr=Courses.crsNr}} Courses))$$


Query Trees


- Exercise:** Can you draw the query tree for the following RA expression?

$\pi_{lastName, result, title}(\sigma_{result \leq 1.3}(\sigma_{Exams.crsNr=Courses.crsNr}(\sigma_{Students.matNr=Exams.matNr}(Students \times Exams \times Courses))))$



Query Trees

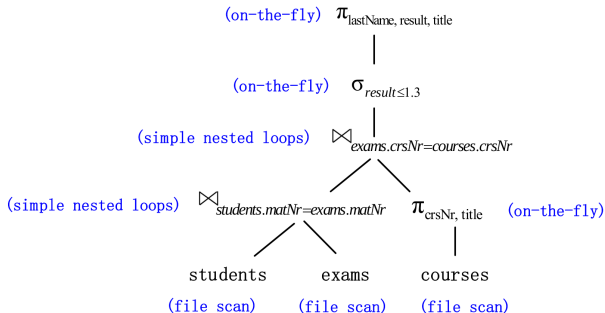
- For each query tree, computation proceeds **bottom-up**:
 - child nodes must be executed before their parent nodes;
 - but there can exist multiple methods of executing sibling nodes, e.g.,
 - process sequentially;
 - process in parallel.





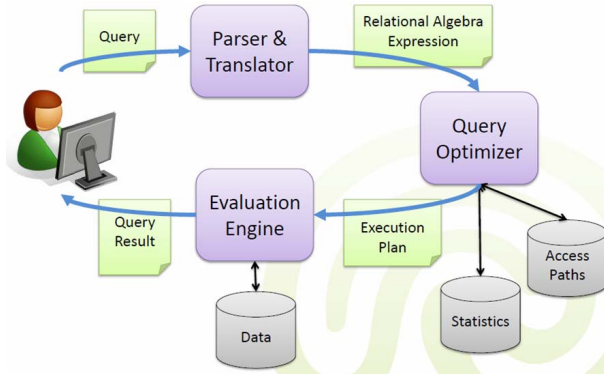
Execution Plan

- A **query execution plan** consists of an (extended) query tree with additional annotation at each node indicating:
 - (1) the *access method* to use for each table, and
 - (2) the *implementation method* for each RA operator.



Query Processing – Evaluation Engine

- The **evaluation engine** executes an execution plan, and returns the query answer to the user.





Query Optimisation



Query Optimisation

- In practice, query optimisers incorporate elements of the following three optimisation approaches:
 - **Semantic query optimisation**
Use application specific semantic knowledge to transform a query into the one with a lower cost (they return the same answer).
 - **Rule-based query optimisation**
Use heuristic rules to transform a relational algebra expression into an equivalent one with a possibly lower cost.
 - **Cost-based query optimisation**
Use a cost model to estimate the costs of plans, and then select the most cost-effective plan.



Semantic Query Optimisation

- Can we use **semantic information** stored in a database (such as integrity constraints) to optimise queries?
 - semantics: “meaning”.
- Recall that, **integrity constraints in the relational model** include:
 - key constraints
 - entity integrity constraints
 - referential integrity constraints
 - domain constraints
 - ...
 - user-defined integrity constraints
- **Key idea:** Integrity constraints may **not only be utilized to enforce consistency** of a database, but may **also optimise user queries**.



Semantic Query Optimisation

- **Example 1:**

Constraint: The relation `Employee` has the primary key `{ssn}`.

Query: `SELECT DISTINCT ssn FROM Employee;`

- We can avoid extra costs for duplicate elimination if the existing constraint tells us that tuples in the result will be unique.



Semantic Query Optimisation

- **Example 2:**

Constraint: No employee can earn more than 200000.

Query:

```
SELECT name  
FROM Employee  
WHERE salary > 300000;
```

- We do not need to execute a query if the existing constraint tells us that the result will be empty.



Semantic Query Optimisation

● Example 3:

Constraints: The relation WORKS_ON has the foreign keys:
[ssn] \subseteq EMPLOYEE[ssn] and [pno] \subseteq PROJECT[pnumber]

Query:

```
SELECT DISTINCT ssn
FROM Works_on INNER JOIN Project
on Works_on.pno=Project.pnumber;
```

- We can reduce the number of joins by executing the following query since both queries always return the same result.

```
SELECT DISTINCT ssn
FROM Works_on ;
```



Rule-based Query Optimisation

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.
- **Key ideas:** apply the most restrictive operation before other operations, which can reduce the size of intermediate results:
 - **Push-down selection:**
Apply as early as possible to reduce the number of tuples;
 - **Push-down projection:**
Apply as early as possible to reduce the number of attributes.
 - **Re-ordering joins:**
Apply restrictive joins first to reduce the size of the result.
- But we must ensure that the resulting query tree gives the same result as the original query tree, i.e., **the equivalence of RA expressions**.

Heuristic Rules

Staff(sid, fname, lname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

- There are **many heuristic rules for transforming** RA expressions, utilized by the **query optimiser**, such as:

$$(1) \sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R);$$

$$\sigma_{branchNo='1'}(\sigma_{salary>60000}(Staff)) = \sigma_{branchNo='1' \wedge salary>60000}(Staff)$$

$$(2) \pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y;$$

$$\pi_{salary}(\pi_{branchNo, salary}(Staff)) = \pi_{salary}(Staff)$$

$$(3) \sigma_{\varphi}(R_1 \times R_2) \equiv R_1 \bowtie_{\varphi} R_2$$

$$\sigma_{Staff.branchNo=Branch.branchNo}(Staff \times Branch) =$$

$$(Staff) \bowtie_{Staff.branchNo=Branch.branchNo} (Branch)$$



Heuristic Rules

Staff(sid, fname, lname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

$$(4) \sigma_{\varphi_1}(R_1 \bowtie_{\varphi_2} R_2) \equiv R_2 \bowtie_{\varphi_1 \wedge \varphi_2} R_1$$

$$\sigma_{\text{salary} > 60000}(\text{Staff} \bowtie_{\text{Staff.branchNo} = \text{Branch.branchNo}} (\text{Branch})) =$$

$$(\text{Staff}) \bowtie_{\text{Staff.branchNo} = \text{Branch.branchNo} \wedge \text{salary} > 60000} (\text{Branch})$$

$$(5) \sigma_{\varphi}(R_1 \bowtie R_2) \equiv \sigma_{\varphi}(R_1) \bowtie R_2, \text{ if } \varphi \text{ contains only attributes in } R_1$$

$$\sigma_{\text{salary} > 60000}(\text{Staff} \bowtie \text{Branch}) = \sigma_{\text{salary} > 60000}(\text{Staff}) \bowtie \text{Branch}$$

$$(6) \sigma_{\varphi_1 \wedge \varphi_2}(R_1 \bowtie R_2) \equiv \sigma_{\varphi_1}(R_1) \bowtie \sigma_{\varphi_2}(R_2) \text{ if } \varphi_1 \text{ contains only attributes in } R_1 \text{ and } \varphi_2 \text{ contains only attributes in } R_2.$$

$$\sigma_{\text{salary} > 60000 \wedge \text{city} = \text{'Canberra'}}(\text{Staff} \bowtie \text{Branch}) =$$

$$(\sigma_{\text{salary} > 60000}(\text{Staff})) \bowtie (\sigma_{\text{city} = \text{'Canberra'}}(\text{Branch}))$$

Heuristic Rules

Staff(sid, fname, lname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

(7) If the join condition involves only attributes in X , we have
 $\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, where X_i contains attributes in both R_1 and R_2 , and ones in both R_i and X , and

$$\pi_{branchNo, position, city}(Staff \bowtie Branch) =$$

$$\pi_{branchNo, position}(Staff) \bowtie (\pi_{branchNo, city}(Branch))$$

(8) If the join condition contains attributes not in X , we have
 $\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, where X_i contains attributes in both in R_1 and R_2 , and ones in both R_i and X

$$\pi_{position, city}(Staff \bowtie Branch) =$$

$$\pi_{position, city}(\pi_{branchNo, position}(Staff) \bowtie (\pi_{branchNo, city}(Branch)))$$



Push-down Selection – Example

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$$\pi_{first_name, last_name}(\sigma_{award_name='Oscar'}((PERSON \bowtie DIRECTOR) \bowtie MOVIE_AWARD))$$

- Question:** Can we apply the following rule to optimise the query?

$$\sigma_{\varphi}(R_1 \bowtie R_2) \equiv \sigma_{\varphi}(R_1) \bowtie R_2, \text{ if } \varphi \text{ contains only attributes in } R_1$$



Push-down Selection – Example

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}(\sigma_{award_name='Oscar'}((PERSON \bowtie DIRECTOR) \bowtie MOVIE_AWARD))$

- We would have:

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$

Push-down Projection – Example

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$$

- Question:** Can we apply the following rule to optimise the query?

$$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)),$$

where X_i contains attributes in both in R_1 and R_2 , and ones in both R_i and X



Push-down Projection – Example

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$$

- we would have:

$$\pi_{first_name, last_name}(\pi_{first_name, last_name, title, production_year}(PERSON \bowtie DIRECTOR) \bowtie \pi_{title, production_year}(\sigma_{award_name='Oscar'}(MOVIE_AWARD)))$$



A Common Query Pattern (Be Careful)

- A common query pattern is **join-select-project** involving three steps:
 - (1) **join** all the relevant relations,
 - (2) **select** the desired tuples, and
 - (3) **project** on the required attributes.
- This query pattern can be expressed as an RA expression

$$\pi_{A_1, \dots, A_n}(\sigma_{\varphi}(R_1 \times \dots \times R_k)),$$

or as an equivalent SQL statement


```
SELECT DISTINCT  $A_1, \dots, A_n$  FROM  $R_1, \dots, R_k$  WHERE  $\varphi$ ;
```

- Queries falling into this pattern can be **very inefficient**, which may yield huge intermediate result for the joined relations.

A Common Query Pattern (Be Careful)

push-down selection and **push-down projection**.

$$\pi_{A_1, \dots, A_n}(\sigma_{\varphi}(R_1 \times \dots \times R_k)),$$





Re-ordering Joins - Example

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

Suppose that it has **10000 tuples**.

DIRECTOR(id, title, production_year) with

$[title, production_year] \subseteq MOVIE_AWARD[title, production_year];$

$[id] \subseteq PERSON[id]$ and

Suppose that it has **100 tuples**.

MOVIE_AWARD(title, production_year, award_name, year_of_award)

Suppose that it has **1000 tuples**.

- Example:** Consider the following two RA queries. Which one is better?
 - PERSON \bowtie MOVIE_AWARD \bowtie DIRECTOR
 - PERSON \bowtie DIRECTOR \bowtie MOVIE_AWARD



Cost-based Query Optimisation

- A query optimiser does not depend solely on heuristic optimisation. It estimates and compares the costs of different plans.
- It estimates and compares the costs of executing a query using different execution strategies and chooses one with **the lowest cost estimate**.
- The query optimiser needs to **limit the number of execution strategies** to be considered for improving efficiency.



Summary

- In general, there are **many ways of executing a query** in a database.
- The user expects the result to be returned promptly, i.e., the query should be **processed as fast as possible**.
- But, the burden of optimising queries should not be put on the user's shoulder. **The DBMSs need to do the job!**
- Nonetheless, SQL is not a suitable query language in which queries can be optimised automatically.
- Instead, SQL queries are **transformed into their corresponding RA queries** and optimised subsequently.
- A major advantage of relational algebra is to **make alternative forms of a query easy to explore**.