



# Database Transactions – Part 1

## Introduction



## Transaction – A Classical Example

- **Scenario:** Suppose that Steve's account balance is **\$1000** and Bob's balance is **\$200**. Now Steve wants to transfer **\$500** into Bob's account.
- There are several steps involved in transferring the money:
  - 1 **Check** Steve's balance;
  - 2 **Update** Steve's balance;
  - 3 **Check** Bob's balance;
  - 4 **Update** Bob's balance.
- Steve later checked his balance (it was **\$500**), which looked good to Steve. However, Bob told Steve that he hadn't received his money yet (still **\$200** in Bob's account instead of **\$700**).

**Question:** What did happen?



## Transaction – A Classical Example

- **Reason:** Due to power outage, the system **stopped working just after updating** Steve's balance.
- **Task:** Transfer **\$500** from Steve's account to Bob's account

1 SELECT balance FROM ACCOUNT  
WHERE name = 'Steve';

2 UPDATE ACCOUNT  
SET balance = balance-500  
WHERE name='Steve';

3 SELECT balance FROM ACCOUNT  
WHERE name = 'Bob';

4 UPDATE ACCOUNT  
SET balance = balance+500  
WHERE name = 'Bob';

Operations	Steve	Bob
before 1	\$1000	\$200
after 1	\$1000	\$200
after 2	\$500	\$200
after 3	\$500	\$200
after 4	\$500	\$700

## Transaction – A Classical Example

- We need an approach to ensure that
  - either the balances of Steve and Bob remain unchanged **if the money transfer fails**
  - or Steve's balance is **\$500** and Bob's is **\$700** **if the money transfer succeeds.**

1 SELECT balance FROM ACCOUNT  
WHERE name = 'Steve';

2 UPDATE ACCOUNT  
SET balance = balance-500  
WHERE name='Steve';


3 SELECT balance FROM ACCOUNT  
WHERE name = 'Bob';

4 UPDATE ACCOUNT  
SET balance = balance+500

Operations	Steve	Bob
before 1	\$1000	\$200
after 1	\$1000	\$200
after 2	\$500	\$200
after 3	\$500	\$200
after 4	\$500	\$700

## What is a Transaction?

- DBMSs provide **transaction support** for solving this kind of problem.
- A **transaction** is a **sequence of database operations grouped together** for execution as a **logic unit** in a DBMS.
  - Different from an execution of a program outside the DBMS (e.g., a C program) in many ways!





## What is a Transaction?

- Database applications often access a database **by transactions** rather than individual operations.
  - e.g., large databases and hundreds of concurrent users: banking, supermarket checkout, airline reservation, online purchasing, etc.
- **Why transactions?** They can **enforce data integrity** in the following situations:
  - multiple users may modify and share data at the same time;
  - transaction, system, and media failures may happen from time to time.
- **What does a transaction look like?**
  - `INSERT`, `SELECT`, `UPDATE`, `DELETE`, `BEGIN`, `COMMIT`, `ABORT` (`ROLLBACK`), etc. from a high-level language perspective;
  - `read`, `write`, `begin`, `commit`, `abort` at the internal process level.



## Transaction – Language Level

- **Database operations** of a transaction (at the SQL language level) may include: **SELECT**, **INSERT**, **UPDATE**, **DELETE**.
- **Other operations**: **BEGIN**, **COMMIT**, **ABORT** (**ROLLBACK**)

### BEGIN TRANSACTION

- 1 `SELECT balance FROM ACCOUNT WHERE name = 'Steve';`
- 2 `UPDATE ACCOUNT`  
`SET balance = balance-500 WHERE name='Steve';`
- 3 `SELECT balance FROM ACCOUNT WHERE name = 'Bob';`
- 4 `UPDATE ACCOUNT`  
`SET balance = balance+500 WHERE name = 'Bob';`

### COMMIT



## Transactions - Internal Process Level

- **Basic operations** of a transaction (at the internal process level) are
  - **read(*X*)**: loads object *X* into main memory;
  - **write(*X*)**: modifies in-memory copy of object *X* (and writes it to disk later on);
- **Granularity of objects**: tables, rows, cells, or memory pages,
- **Other operations**:
  - **begin**: marks the beginning of a transaction;
  - **commit**: signals a successful end of the transaction - all changes can safely be applied to the database permanently;
  - **abort**: signals the transaction has ended unsuccessfully - undo all operations of the transaction.





## Transactions - Internal Process Level

```
T: BEGIN TRANSACTION
T: SELECT balance FROM ACCOUNT WHERE name = 'Steve';
T: UPDATE ACCOUNT SET balance = balance-500 WHERE name='Steve';
T: SELECT balance FROM ACCOUNT WHERE name = 'Bob';
T: UPDATE ACCOUNT SET balance = balance+500 WHERE name = 'Bob';
T: COMMIT;
```

### Objects:

- A - Steve's account balance;
- B - Bob's account balance.

Steps	<i>T</i>
1	read(A)
2	write(A) ( $A := A - 500$ )
3	read(B)
4	write(B) ( $B := B + 500$ )
5	commit



# Database Transactions – Part 2

## ACID Properties

## ACID Properties

- DBMSs ensure the following properties of transactions.
  - **Atomicity:**
    - The execution of each transaction is atomic, i.e., **either all operations are completed or not done at all.**
  - **Consistency:**
    - The states of a database are consistent (w.r.t. defined business rules) **before and after each transaction.**
  - **Isolation:**
    - Execution results of each transaction should be **unaffected by other concurrent executing transactions.**
  - **Durability:**
    - Once a transaction has been successfully completed, **its effects should persist in the database.**

**Note:** These properties are not independent from one another, but **atomicity is the central property.**




## Atomicity

- **Atomicity** requires that we execute a transaction to completion with only **two possibilities**:
    - **ALL**: all the operations are executed;
    - **NONE**: none of the operations are executed.
  - If a transaction fails to complete for some reason, it may leave database in an inconsistent state. Thus a DBMS **must remove effects of partial transactions** to ensure atomicity.
- Example:** The money can only be taken from Steve's account if the money has been transferred into Bob's account.

Operations	Steve	Bob	None are executed.
before 1	\$1000	\$200	
after 1	\$1000	\$200	
after 2	\$500	\$200	
after 3	\$500	\$200	All are executed.
after 4	\$500	\$700	

## Atomicity



## Consistency

- **Consistency** requires that, each transaction should **preserve the consistency of the database**.
- **Note:** Intermediate states may be inconsistent.

**Example:** Suppose that we have

**Steve's account balance + Bob's account balance = \$1200,**


Operations	Steve	Bob
before 1	\$1000	\$200
after 1	\$1000	\$200
after 2	\$500	\$200
after 3	\$500	\$200
after 4	\$500	\$700

$\$1000 + \$200 = \$1200$

**Not required to be consistent.**

$\$500 + \$700 = \$1200$

## Consistency<sup>1</sup>



- The database is in a consistent state before and after executing the transaction, but is not necessarily consistent in intermediate states.

<sup>1</sup> The figure is taken from <http://maxdb.sap.com>



## Isolation

- **Isolation** requires that transactions are **isolated from one another**.

**Example:** Other transactions can't see the changes on objects  $A$  (Steve's account balance) and  $B$  (Bob's account balance) until the transaction for the money transfer is completed.

$T_1$

---

read(A)  
write(A) ( $A := A - 500$ )  
read(B)  
write(B) ( $B := B + 500$ )  
commit

---

$T_2$


---

read(A)  
write(A) ( $A := A + 400$ )  
commit

---



## Isolation <sup>2</sup>




<sup>2</sup> The figure is taken from <http://michaeljswart.com/>



## Durability

- **Durability** requires that once the transaction is successfully completed, its changes to the database **must be persistent despite failures**.
- The decision is irrevocable: once committed, the transaction cannot revert to abort. **Changes are durable**.
- **Example:** Once Steve received the notification:  
    **"\$500 has been successfully transferred to Bob's account"**,  
the money can't go back to Steve's account and must appear in Bob's account.

## Durability<sup>3</sup>



<sup>3</sup> The figure is taken from <http://toyhouse.cc/profiles/blogs/the-acid-properties-of-transactions>




# Database Transactions – Part 3

## Concurrent Transactions


## Concurrent Transactions

- **Interleaved processing:** transactions are interleaved in a single CPU.



## Concurrent Transactions

- **Parallel processing:** transactions are executed in parallel in multiple CPUs.





## Concurrent Transactions

- Executing transactions concurrently will **improve database performance**
  - ↪ **Increase throughput** (*average number of completed transactions*)
    - For example, while one transaction is waiting for an object to be read from disk, the CPU can process another transaction (because I/O activity can be done in parallel with CPU activity).
  - ↪ **Reduce latency** (*average time to complete a transaction*)
    - For example, interleave execution of a short transaction with a long transaction usually allows the short one to be completed more quickly.
- But the DBMS has to guarantee that the interleaving of transactions **does not lead to inconsistencies**, i.e., **concurrency control**.



## Why is Concurrency Control Needed?

- Concurrency control is needed for preventing the following problems:
  - 1 The **lost update** problem
  - 2 The **dirty read** problem
  - 3 The **unrepeated read** problem
  - 4 The **phantom read** problem





## (1) - The Lost Update Problem

- Example:** Bob withdraws **\$100** from his account ( $T_1$ ) while Alice deposits **\$500** into Bob's account ( $T_2$ ).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_1$ : COMMIT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```

Steps	$T_1$	$T_2$
1	read(B)	read(B)
2		
3	write(B) ( $B := B - 100$ )	
4	commit	write(B) ( $B := B + 500$ )
5		
6		

Steps	B(Bob)
before 1	\$200
after 2	\$200
after 4	\$100
after 6	\$700

- Question:** What is the problem?



## (1) - The Lost Update Problem

- Example:** Bob withdraws **\$100** from his account ( $T_1$ ) while Alice deposits **\$500** into Bob's account ( $T_2$ ).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_1$ : COMMIT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```


Steps	$T_1$	$T_2$
1	read(B)	read(B)
2		
3	write(B) ( $B := B - 100$ )	
4	commit	write(B) ( $B := B + 500$ )
5		
6		

Steps	B(Bob)
before 1	\$200
after 2	\$200
after 4	\$100
after 6	\$700

- Answer:** Bob's balance should be **\$600**. The update by  $T_1$  is lost!

## (1) - The Lost Update Problem

- Occurs when two transactions update the same object, and one transaction could overwrite the value of the object which has already been updated by another transaction (**write-write conflicts**).
- Example:**



- $write(B)$  by  $T_2$  overwrites  $B$ , and the update by  $T_1$  is *lost*.



## (2) - The Dirty Read Problem

- Example:** Bob withdraws **\$100** from his account ( $T_1$ ) while Alice deposits **\$500** into Bob's account ( $T_2$ ).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : ABORT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```

Steps	$T_1$	$T_2$
1	read(B)	read(B)  write(B) (B:=B+500)  commit
2	write(B) (B:=B-100)	
3	abort	
4		
5		
6		

Steps	B(Bob)
before 1	\$200
after 1	\$200
after 2	\$100
after 4	\$200
after 6	\$600

- Question:** What is the problem?

## (2) - The Dirty Read Problem

- Example:** Bob withdraws **\$100** from his account ( $T_1$ ) while Alice deposits **\$500** into Bob's account ( $T_2$ ).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : UPDATE ACCOUNT SET balance=balance-100 WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_1$ : ABORT;  
 $T_2$ : UPDATE ACCOUNT SET balance=balance+500 WHERE name='Bob';  
 $T_2$ : COMMIT;
```


Steps	$T_1$	$T_2$
1	read(B)	read(B)
2	write(B) ( $B := B - 100$ )	
3		
4	abort	
5		write(B) ( $B := B + 500$ )
6		commit

Steps	B(Bob)
before 1	\$200
after 1	\$200
after 2	\$100
after 4	\$200
after 6	\$600

- Answer:** Bob's balance should be **\$700** since  $T_1$  was not completed.

## (2) - The Dirty Read Problem

- Occurs when one transaction could read the value of an object that has been updated by another transaction but has not yet committed (**write-read conflicts**).
- Example:**



- $T_1$  fails and must change the value of  $B$  back to **\$200**; but  $T_2$  has read the uncommitted ( $\cong$  *dirty*) value of  $B$  (**\$100**).



### (3) - The Unrepeatable Read Problem

- Example:** Bob checks his account ( $T_1$ ) twice (takes time to decide whether to withdraw **\$200**) while Alice withdraws **\$500** from Bob's account ( $T_2$ ).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : UPDATE ACCOUNT SET balance=balance-500 WHERE name='Bob';  
 $T_2$ : COMMIT;  
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';
```

Steps	$T_1$	$T_2$
1	read(B)	
2		read(B)
3		write(B) ( $B := B - 500$ )
4		commit
5	read(B)	

Steps	B(Bob)
before 1	\$500
after 2	\$500
after 3	\$0
after 4	\$0
after 5	\$0

- Question:** What is the problem?



### (3) - The Unrepeatable Read Problem

- Example:** Bob checks his account ( $T_1$ ) twice (takes time to decide whether to withdraw **\$200**) while Alice withdraws **\$500** from Bob's account ( $T_2$ ).

```
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : SELECT balance FROM ACCOUNT WHERE name='Bob';  
 $T_2$ : UPDATE ACCOUNT SET balance=balance-500 WHERE name='Bob';  
 $T_2$ : COMMIT;  
 $T_1$ : SELECT balance FROM ACCOUNT WHERE name='Bob';
```

Steps	$T_1$	$T_2$
1	read(B)	
2		read(B)
3		write(B) (B:=B-500)
4		commit
5	read(B)	

Steps	B(Bob)
<b>before 1</b>	<b>\$500</b>
<b>after 2</b>	\$500
<b>after 3</b>	\$0
<b>after 4</b>	<b>\$0</b>
<b>after 5</b>	<b>\$0</b>

- Answer:** Bob received two different account balances **\$500** and **\$0**, even though he hasn't withdrawn any money yet.






### (3) - The Unrepeatable Read Problem

- A transaction could change the value of an object that has been read by another transaction but is still in progress (could issue two read for the object, or a write after reading the object) (**read-write conflicts**).

- **Example:**





## (4) - The Phantom Read Problem

- **Example:** A query is submitted for finding all customers whose account balances are less than **\$300** ( $T_1$ ) while Alice is opening a new account with the balance **\$200** ( $T_2$ ).
- Assume that only Bob (B) has an account whose balance is less than **\$300** before Alice (A) opens his new account.

$T_1$ : SELECT name FROM ACCOUNT WHERE balance<300;

$T_2$ : INSERT INTO ACCOUNT(id, name, balance) VALUES(99, 'Alice', 250);

$T_2$ : COMMIT;

$T_1$ : SELECT name FROM ACCOUNT WHERE balance<300;

Steps	$T_1$	$T_2$
1	read(R)	
2		write(R)
3		commit
4	read(R)	

Steps	Query result
before 1	$R = \{B\}$
after 1	$R = \{B\}$
after 2	$R = \{A, B\}$
after 4	$R = \{A, B\}$

- **Question:** What is the problem?

## (4) - The Phantom Read Problem

- **Example:** A query is submitted for finding all customers whose account balances are less than **\$300** ( $T_1$ ) while Alice is opening a new account with the balance **\$200** ( $T_2$ ).
- Assume that only Bob (B) has an account whose balance is less than **\$300** before Alice (A) opens his new account.

```

T1: SELECT name FROM ACCOUNT WHERE balance<300;
T2: INSERT INTO ACCOUNT(id, name, balance) VALUES(99, 'Alice', 250);
T2: COMMIT;
T1: SELECT name FROM ACCOUNT WHERE balance<300;

```

Steps	$T_1$	$T_2$
1	read(R)	write(R) commit
2		
3		
4	read(R)	

Steps	Query result
before 1	$R = \{B\}$
after 1	$R = \{B\}$
after 2	$R = \{A, B\}$
after 4	$R = \{A, B\}$

- **Answer:**  $T_1$  reads Account based on the condition  $\text{balance} < 300$  twice but gets two different results  $\{B\}$  and  $\{A, B\}$ .

## (4) - The Phantom Read Problem

- Occurs when tuples updated by a transaction  $T_1$  satisfy the search conditions of another transaction so that, by the same search condition, the transaction obtains different results at different times.
- Example:**

