

CS 480

Introduction to Artificial Intelligence

April 21, 2022

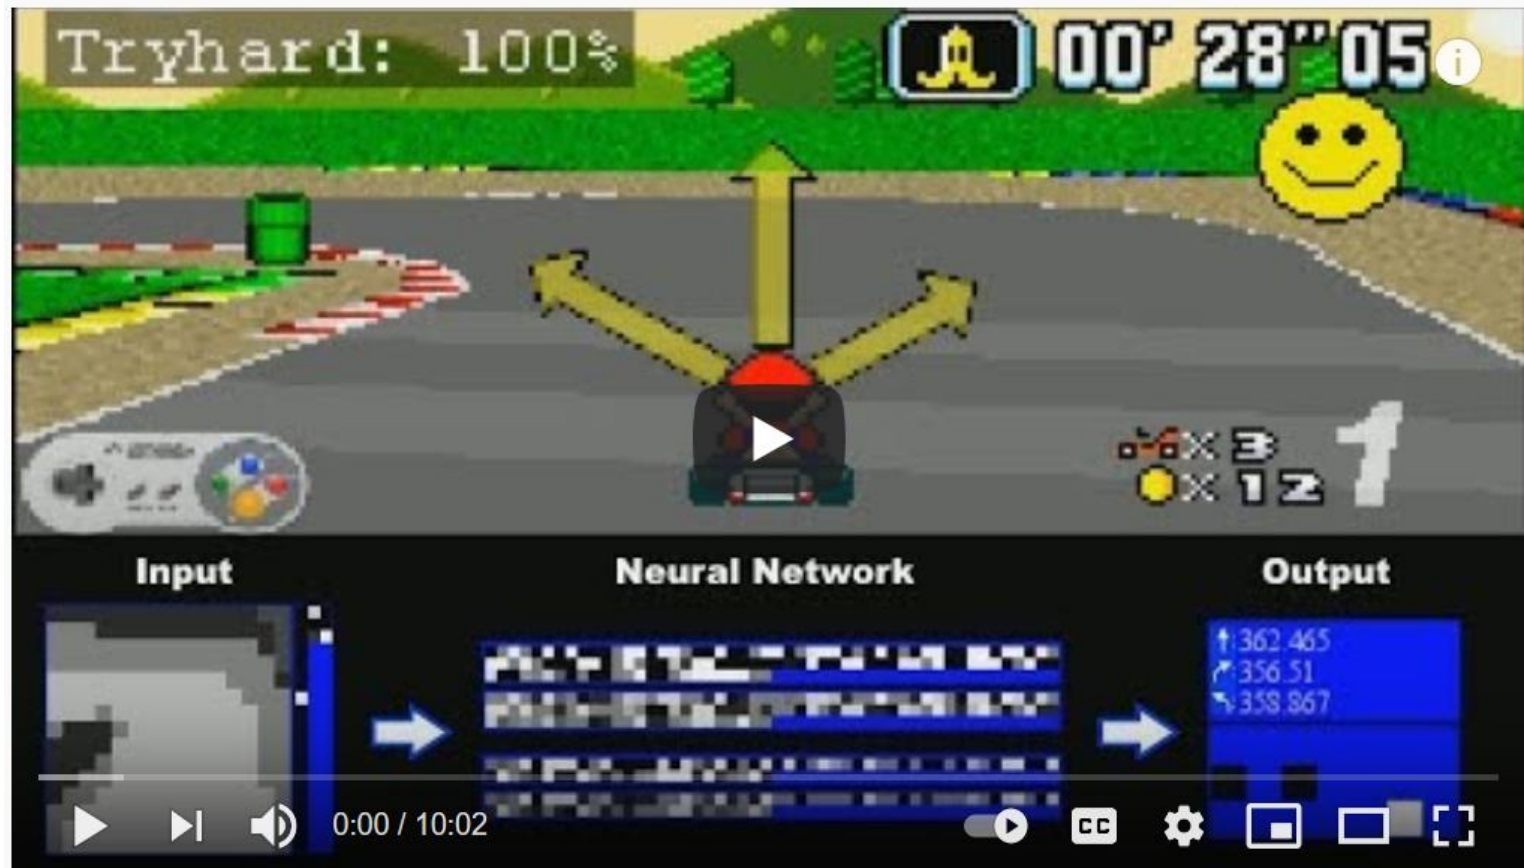
Announcements / Reminders

- **Final Exam: April 28th!**
 - Ignore Registrar date for CS 480
 - Online section: please contact Mr. Charles Scott (scott@iit.edu) to make arrangements if necessary
- End of semester course evaluation: open. Thank you!
- Programming Assignment #02: due **TOMORROW (04/22)**
 - **ADD COMMENTS!!!!**
- Written Assignment #04: due on Wednesday (04/27)
- Grading TA assignment:
https://docs.google.com/spreadsheets/d/1Cav_GBTGC7fLGzxuBCAUmEuJYPeF-HMLCYvwPbq8Fus/edit?usp=sharing

Plan for Today

- **Casual Introduction to Machine Learning**
 - **Deep Learning**
 - **Reinforcement Learning**

Reinforcement Learning in Action



MarlQ -- Q-Learning Neural Network for Mario Kart -- 2M Sub Special

330,560 views • Jun 29, 2019



18K



163



SHARE




SAVE




Source: https://www.youtube.com/watch?v=Tnu4O_xEmVk

ANN for Simple Game Playing



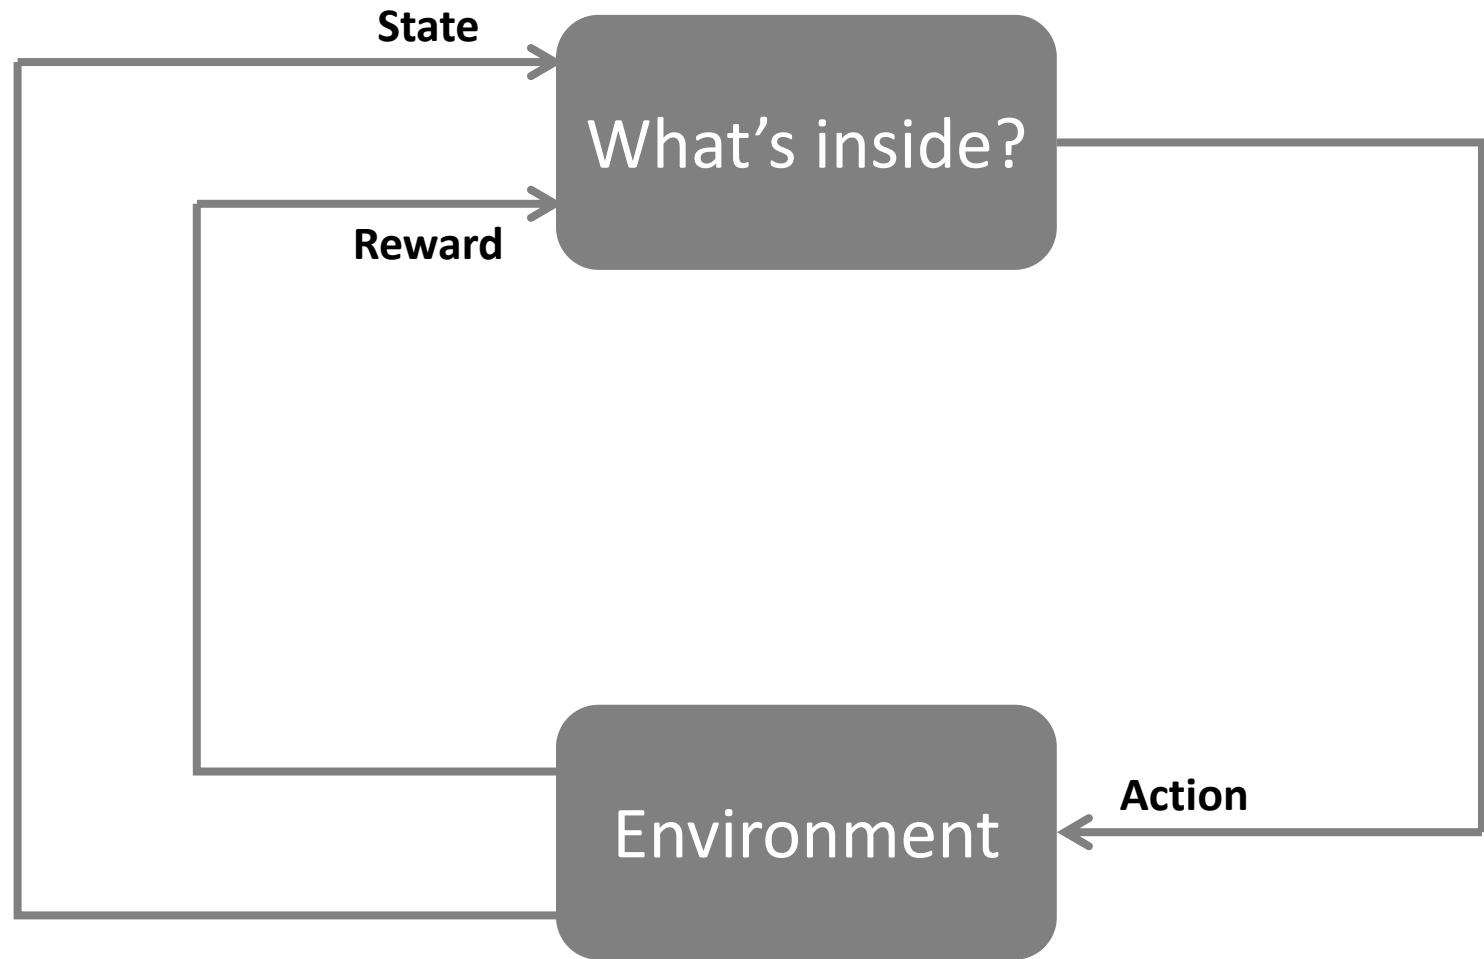
ANN for Simple Game Playing

Current game is an input. Decisions (UP/DOWN/JUMP) are rewarded/punished.




← **Correct all the weights** using Reinforcement Learning.

RL: Agents and Environments




RL: Agents and Environments




Deep Learning

Machine Learning vs. Deep Learning

Machine Learning




Deep Learning




Source: <https://www.quora.com/What-is-the-difference-between-deep-learning-and-usual-machine-learning>

Machine Learning vs. Deep Learning

Classic Machine Learning




Deep Learning



Source: <https://www.intel.com/content/www/us/en/artificial-intelligence/posts/difference-between-ai-machine-learning-deep-learning.html>

Deep Learning: Feature Extraction



Source: https://en.wikipedia.org/wiki/Deep_learning

Exercise: Object Recognition

<https://braneshop.com.au/object-detection-in-the-browser.html>

(you can try it on your smartphone)

Exercise: Image Colorizer

<https://deepai.org/machine-learning-model/colorizer>

Exercise: Deep Learning

<https://www.handwriting-generator.com/>

Main Machine Learning Categories

Supervised learning

Supervised learning is one of the most common techniques in machine learning. It is based on **known relationship(s) and patterns within data** (for example: relationship between inputs and outputs).

Frequently used types: **r e g r e s s i o n**, and **classification**.

Unsupervised learning

Unsupervised learning involves finding underlying patterns within data. Typically used in **clustering** data points (similar customers, etc.)

Reinforcement learning

Reinforcement learning is inspired by behavioral psychology. It is **based on a rewarding / punishing an algorithm**.

Rewards and punishments are based on algorithm's action within its environment.


What is Reinforcement Learning?

Idea:

Reinforcement learning is inspired by behavioral psychology. It is **based on a rewarding / punishing an algorithm.**

Rewards and punishments are based on algorithm's action within its environment.

RL: Agents and Environments



Reinforcement Learning in Action



Source: <https://www.youtube.com/watch?v=kopoLzvh5jY>

K-Armed Bandit Problem



K-Armed Bandit Problem

The K-armed bandit problem is a problem in which a **fixed limited set of resources** must be **allocated between competing (alternative) choices** in a way that **maximizes their expected gain**.

Each choice's **properties are only partially known** at the time of allocation, and **may become better understood as time passes** or by allocating resources to the choice.

K-Armed Bandit Problem

In the problem, **each machine provides a random reward from a probability distribution specific to that machine, that is not known a-priori.**

The objective of the gambler is to **maximize the sum of rewards earned through a sequence of lever pulls.**

K-Armed Bandit Problem

Bandit/Arm 1

33 %

current
success (win)
rate

Bandit/Arm 2

52 %

current
success (win)
rate


Bandit/Arm 3

78 %

current
success (win)
rate

Which bandit shall we play next?

K-Armed Bandit



Exploration vs. Exploitation

The crucial tradeoff the gambler faces at each trial is between "**exploitation**" of the machine that has the highest expected payoff and "**exploration**" to get more information about the expected payoffs of the other machines.

ϵ -greedy Algorithm

generate random number $p \in [0, 1]$

if ($p < \epsilon$) // explore

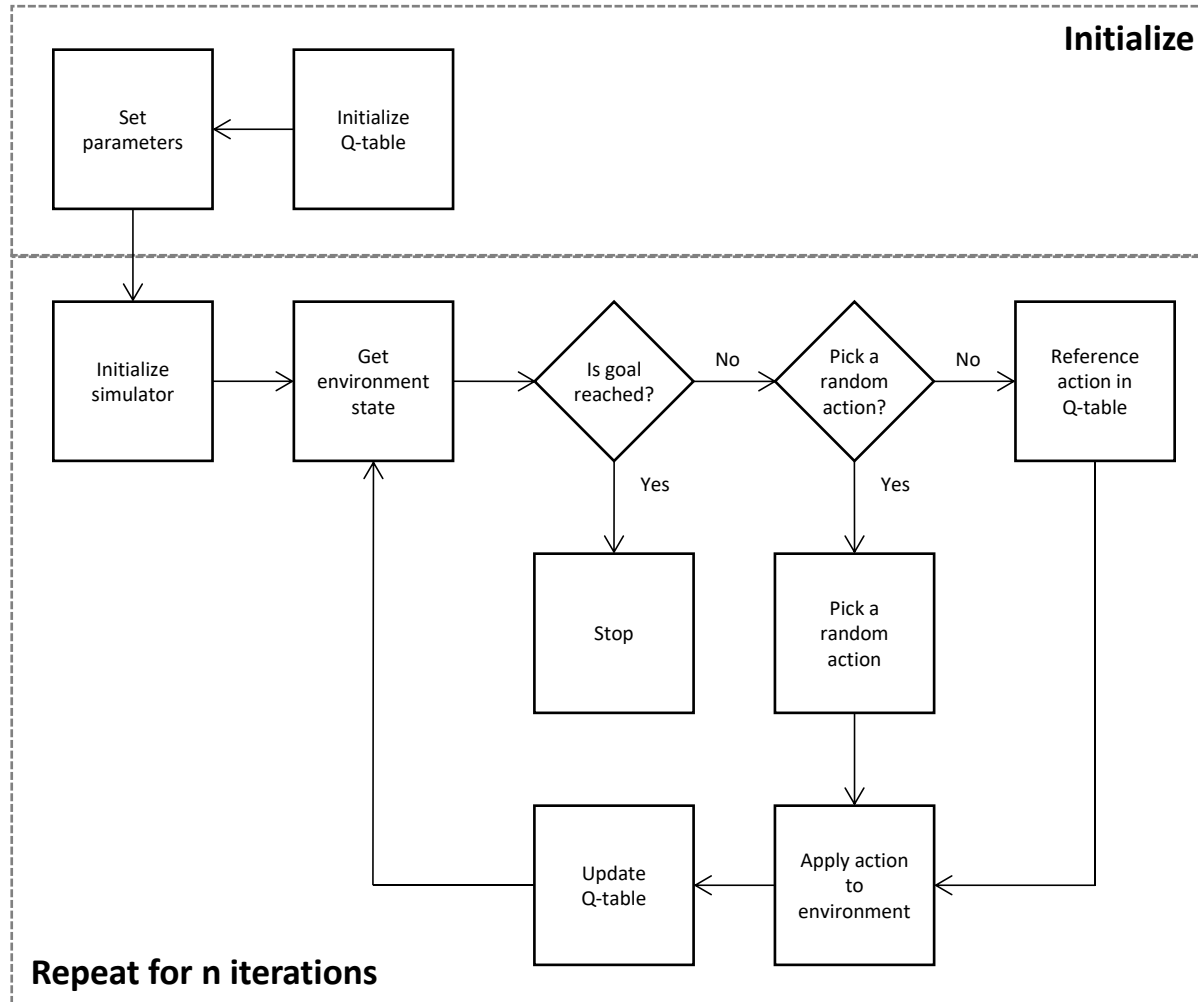
 select random arm

else // exploit

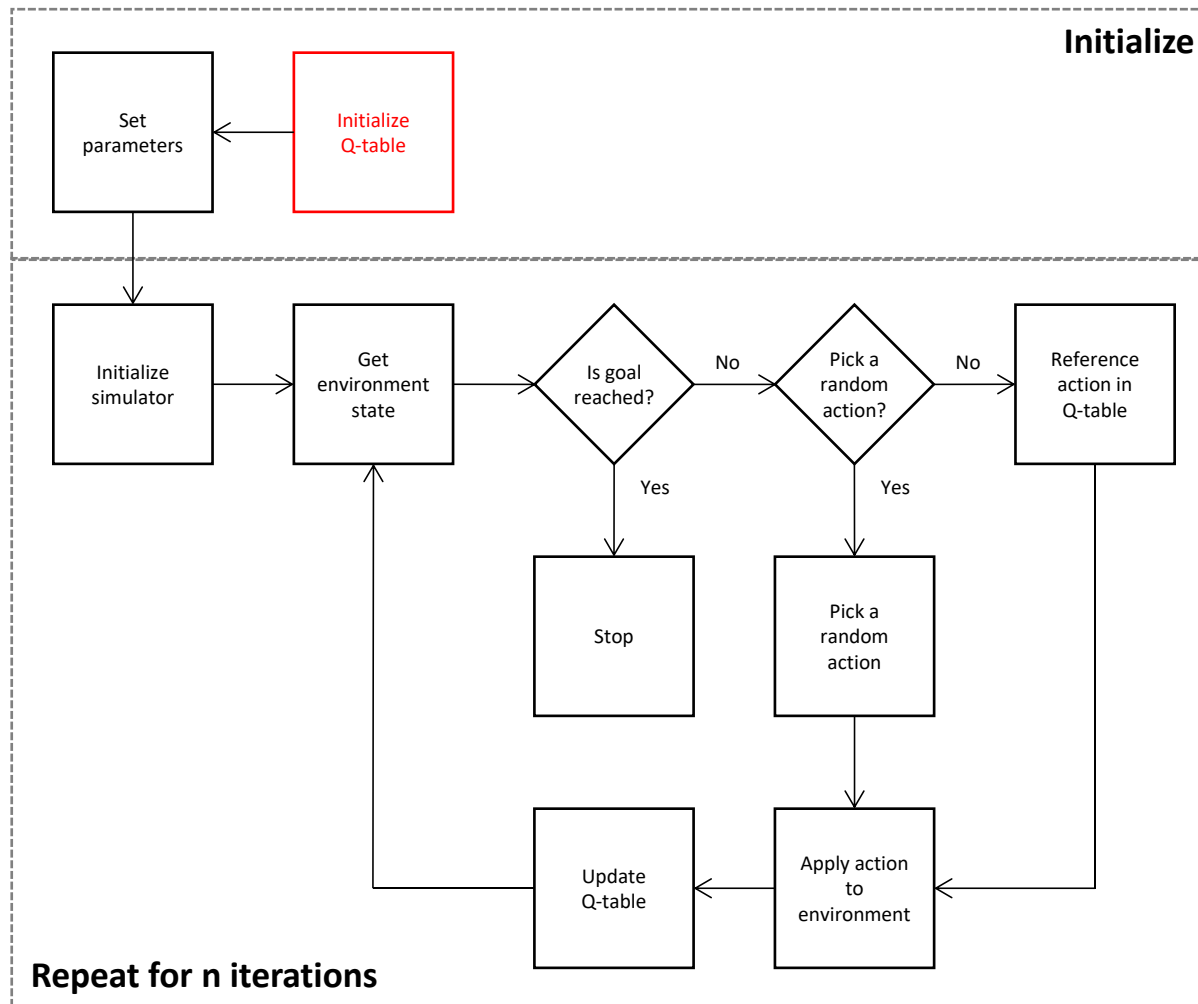
 select current best arm

end

Q-Learning Algorithm



Q-Learning Algorithm



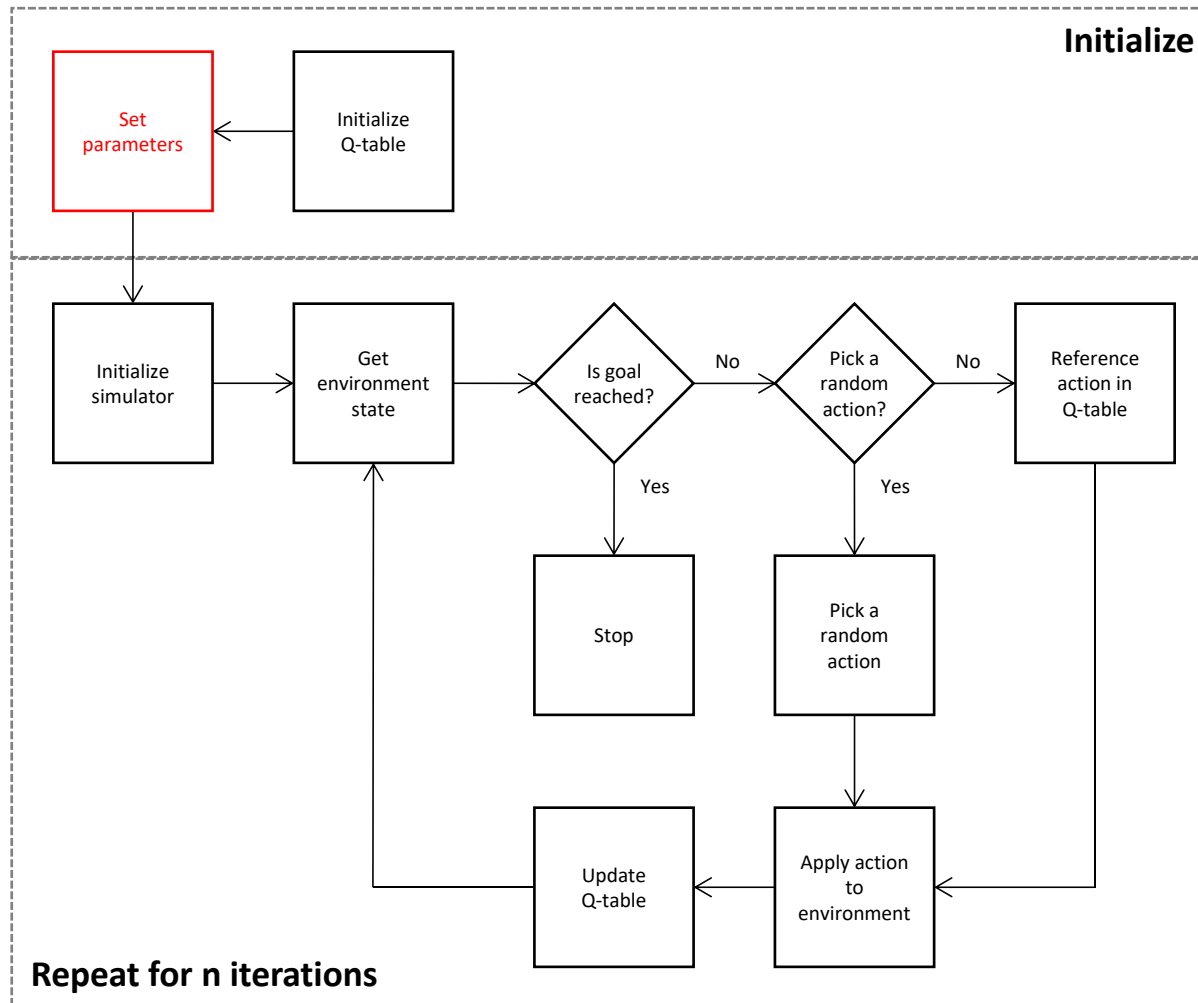
Initialize Q-table:

Set up and initialize (all values set to 0) a table where:

- rows represent **possible states**
- columns represent **actions**

Note that additional states can be added to the table when encountered.

Q-Learning Algorithm




Set parameters:

Set and initialize **hyperparameters** for the Q-learning process.

Hyperparameters include:

- **chance of choosing a random action:** a threshold for **choosing a random action over an action from the Q-table**
- **learning rate:** a parameter that describes **how quickly the algorithm should learn from rewards** in different states
 - high: faster learning with erratic Q-table changes
 - low: gradual learning with possibly more iterations
- **discount factor:** a parameter that **describes how valuable are future rewards**. It tells the algorithm whether it should seek “immediate gratification” (small) or “long-term reward” (large)

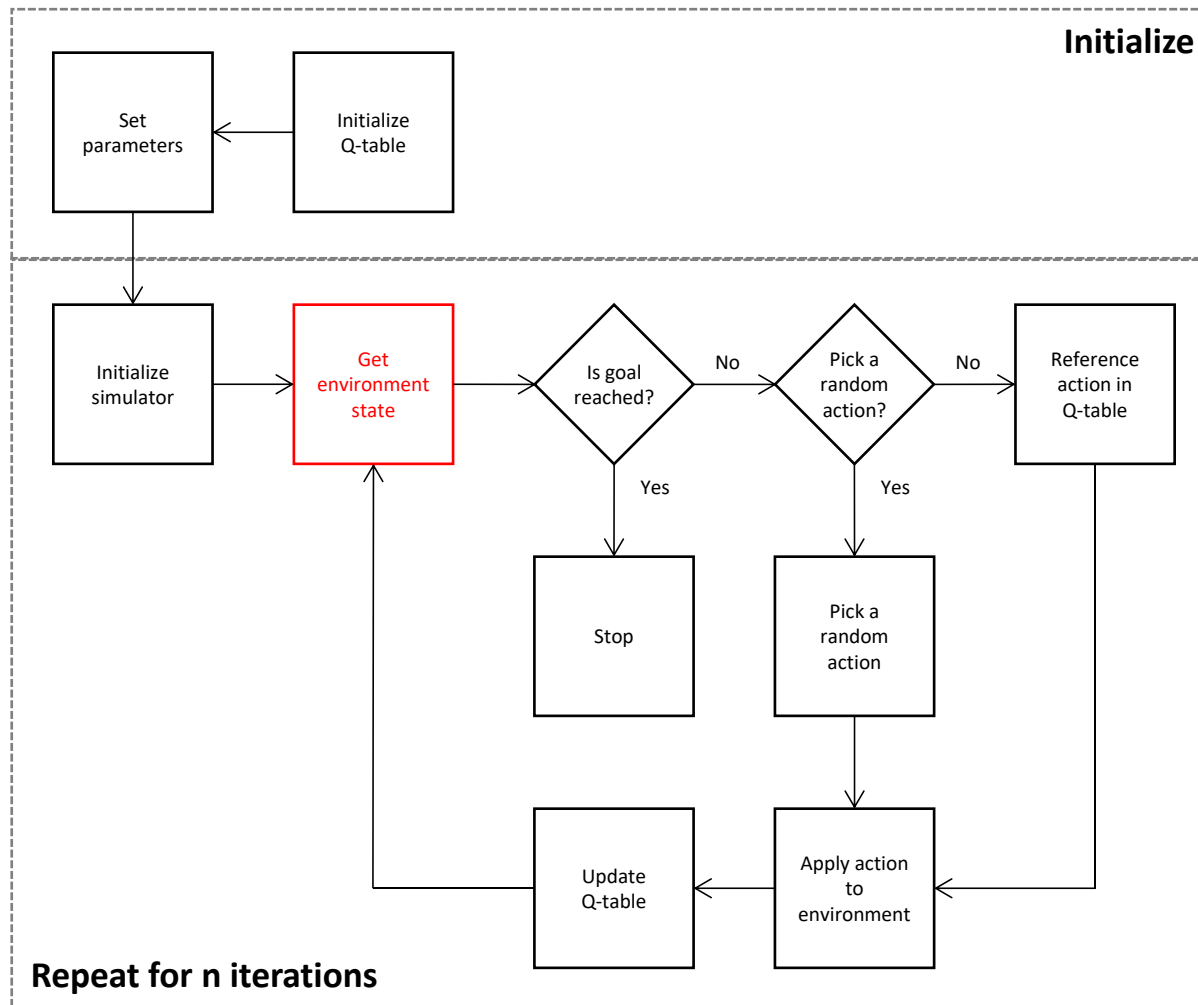
Q-Learning Algorithm



Initialize simulator:

Reset the simulated environment to its initial state and place the agent in a neutral state.

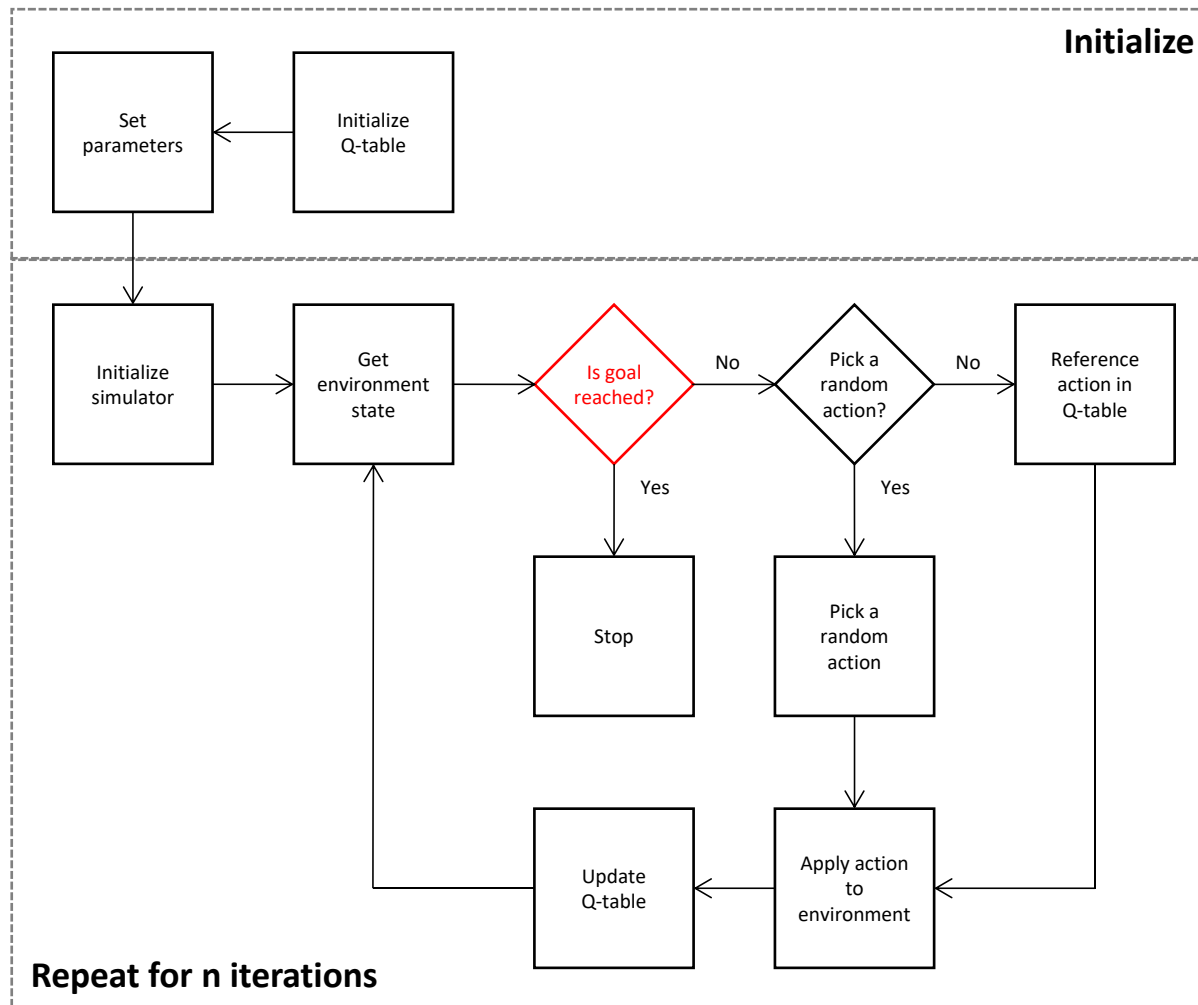
Q-Learning Algorithm



Get environment state:

Report the current state of the environment. Typically a vector of values representing all relevant variables.

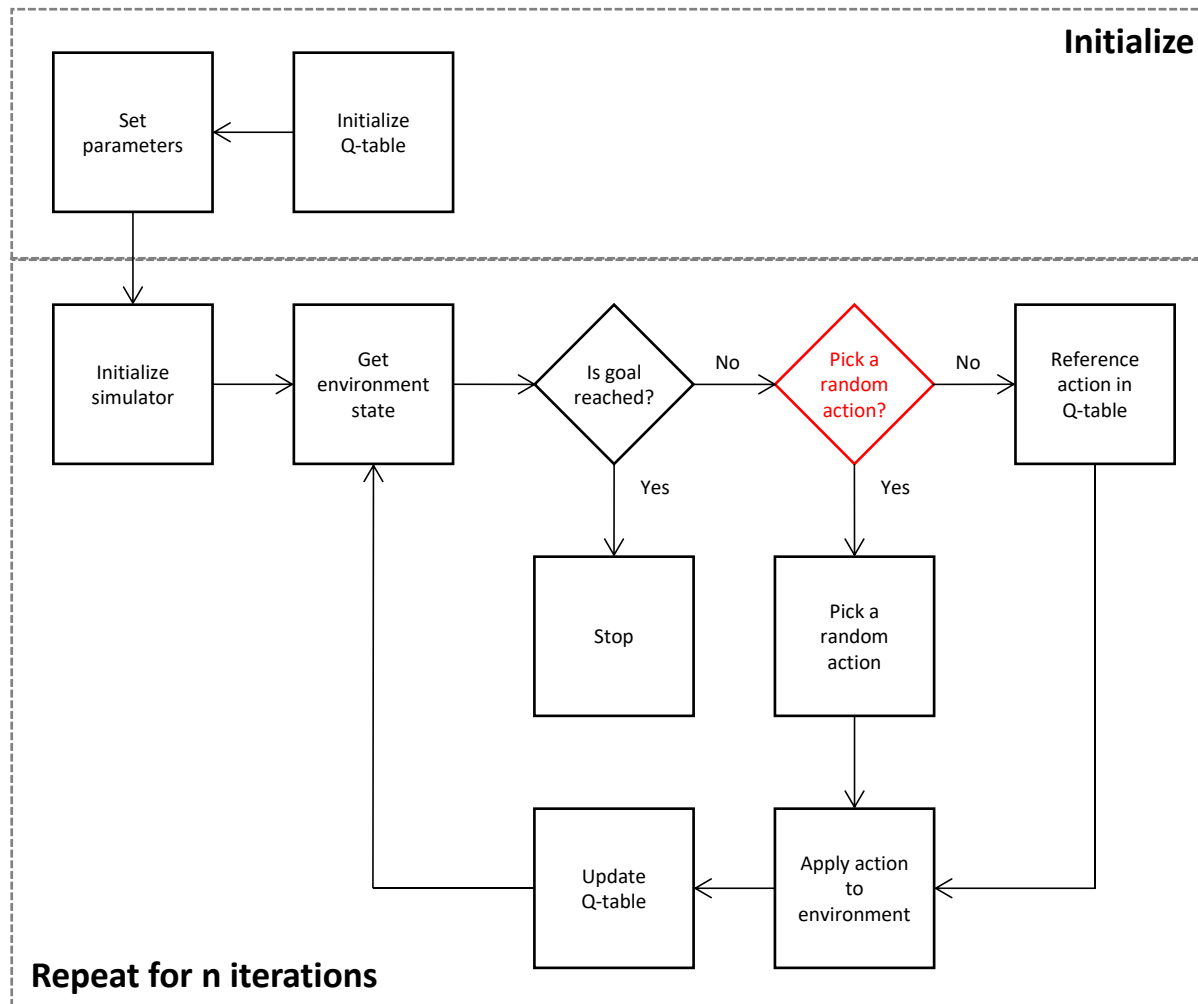
Q-Learning Algorithm



Is goal reached?:

Verify if the goal of the simulation has been achieved. It could be decided with the agent arriving in expected final state or by some simulation parameter.

Q-Learning Algorithm




Pick a random action?:

Decide whether next action should be picked at random or not (it will be selected based on Q-table data then).

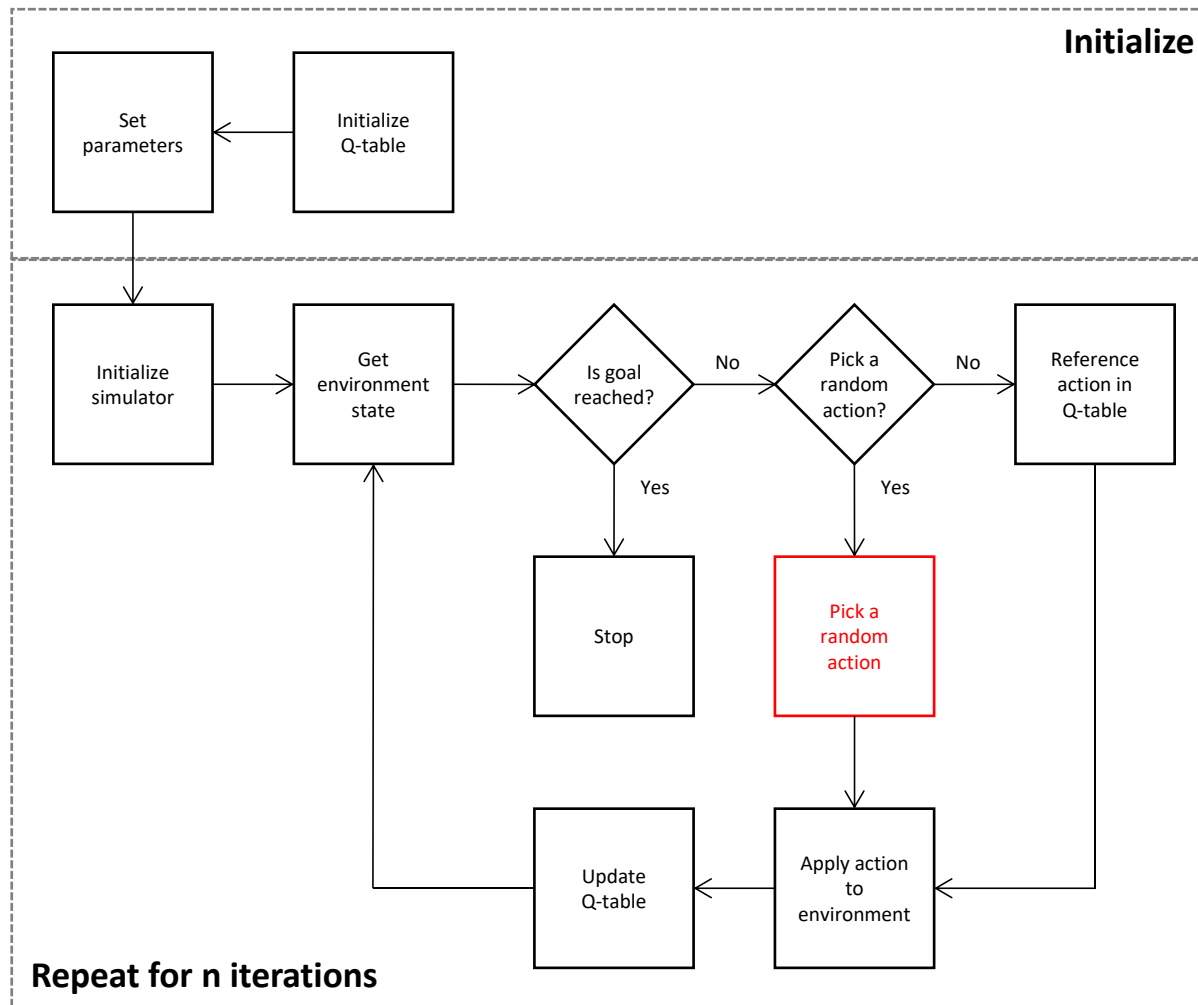
Use the **chance of choosing a random action hyperparameter** to decide.

Q-Learning Algorithm



Reference action in Q-table:
Next action decision will be based on data from the Q-table **given the current state of the environment**.

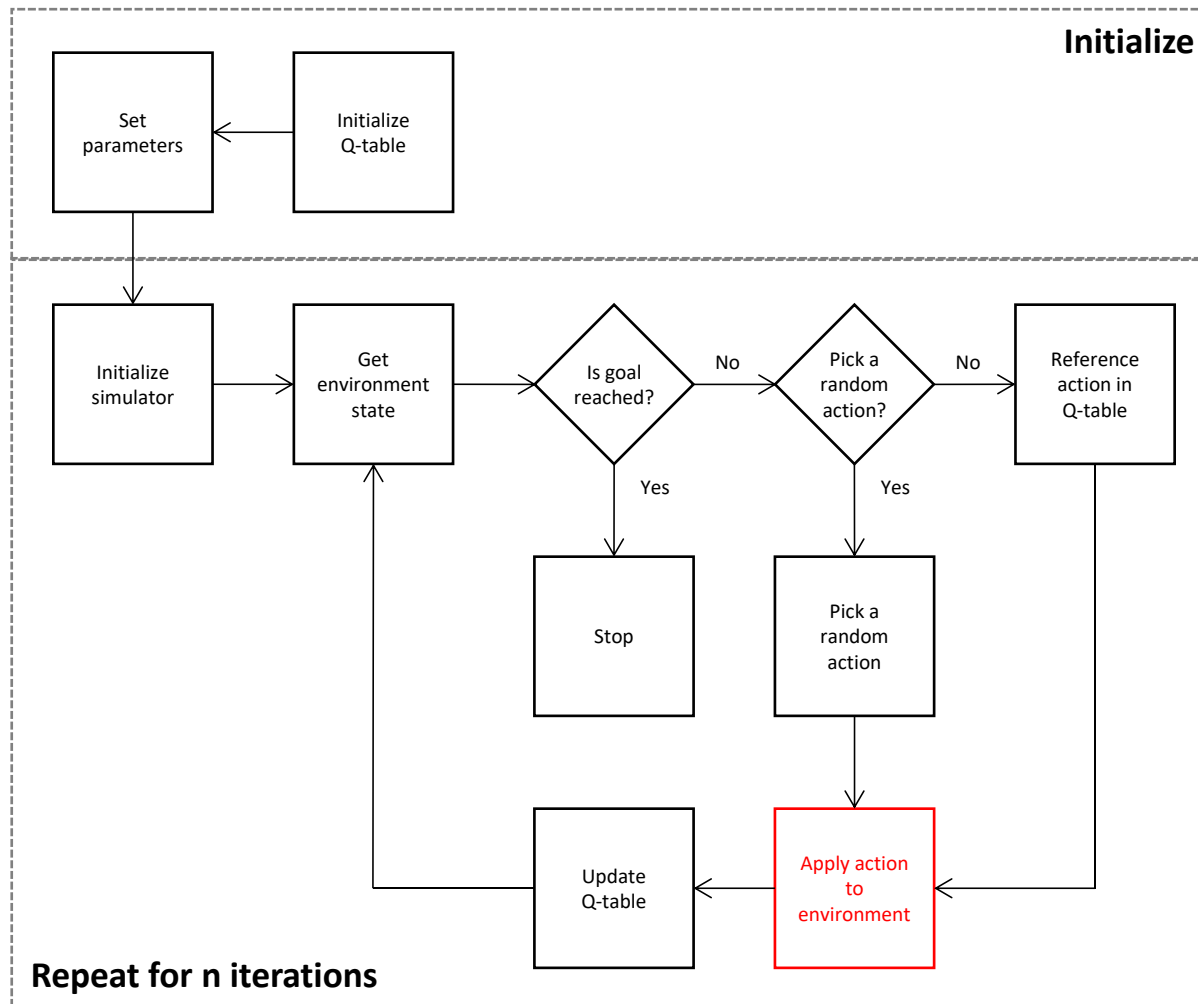
Q-Learning Algorithm



Pick a random action:

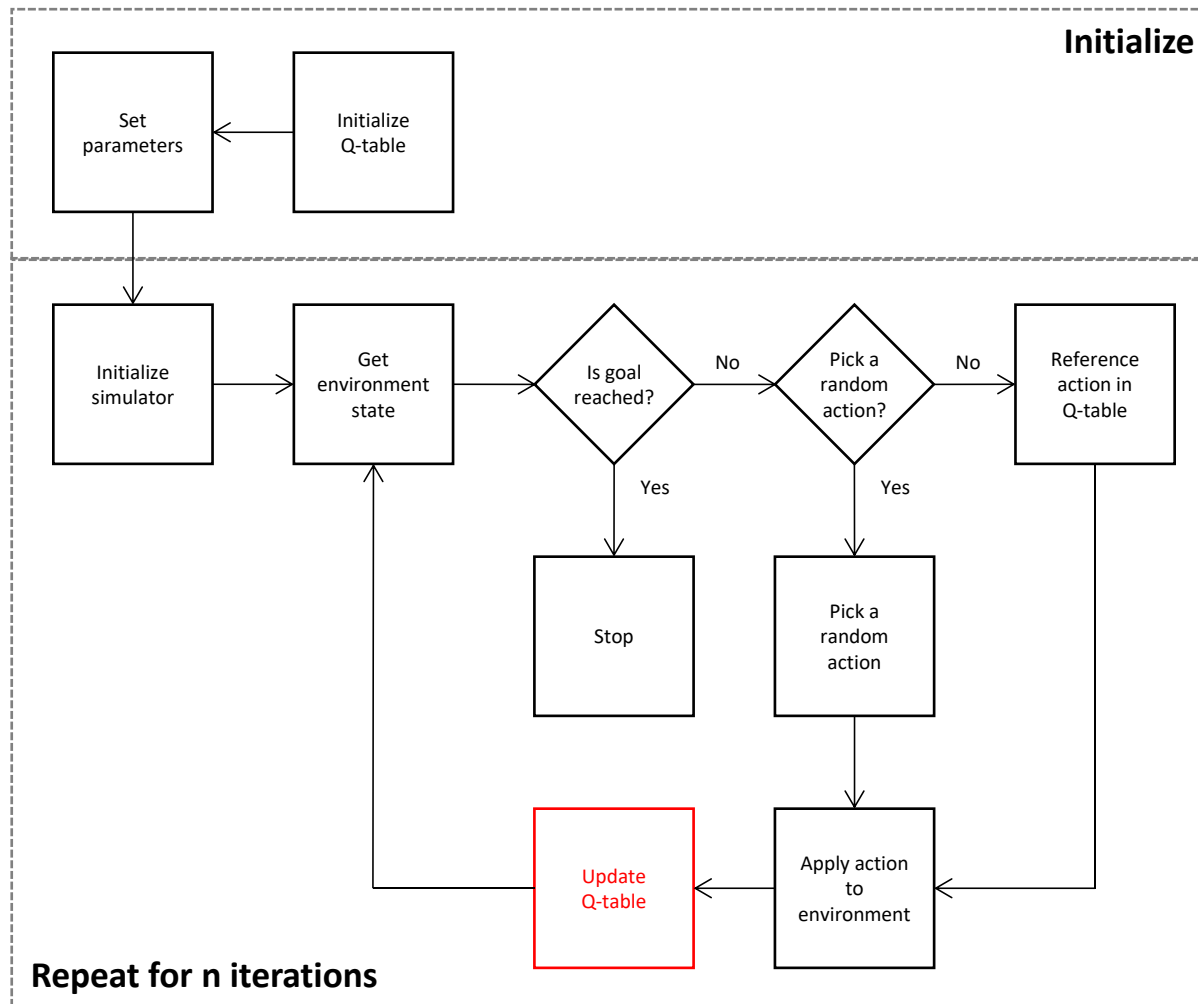
Pick any of the available actions at random. Helpful with exploration of the environment.

Q-Learning Algorithm



Apply action to environment:
Apply the action to the environment to change it. Each action will have its own reward.

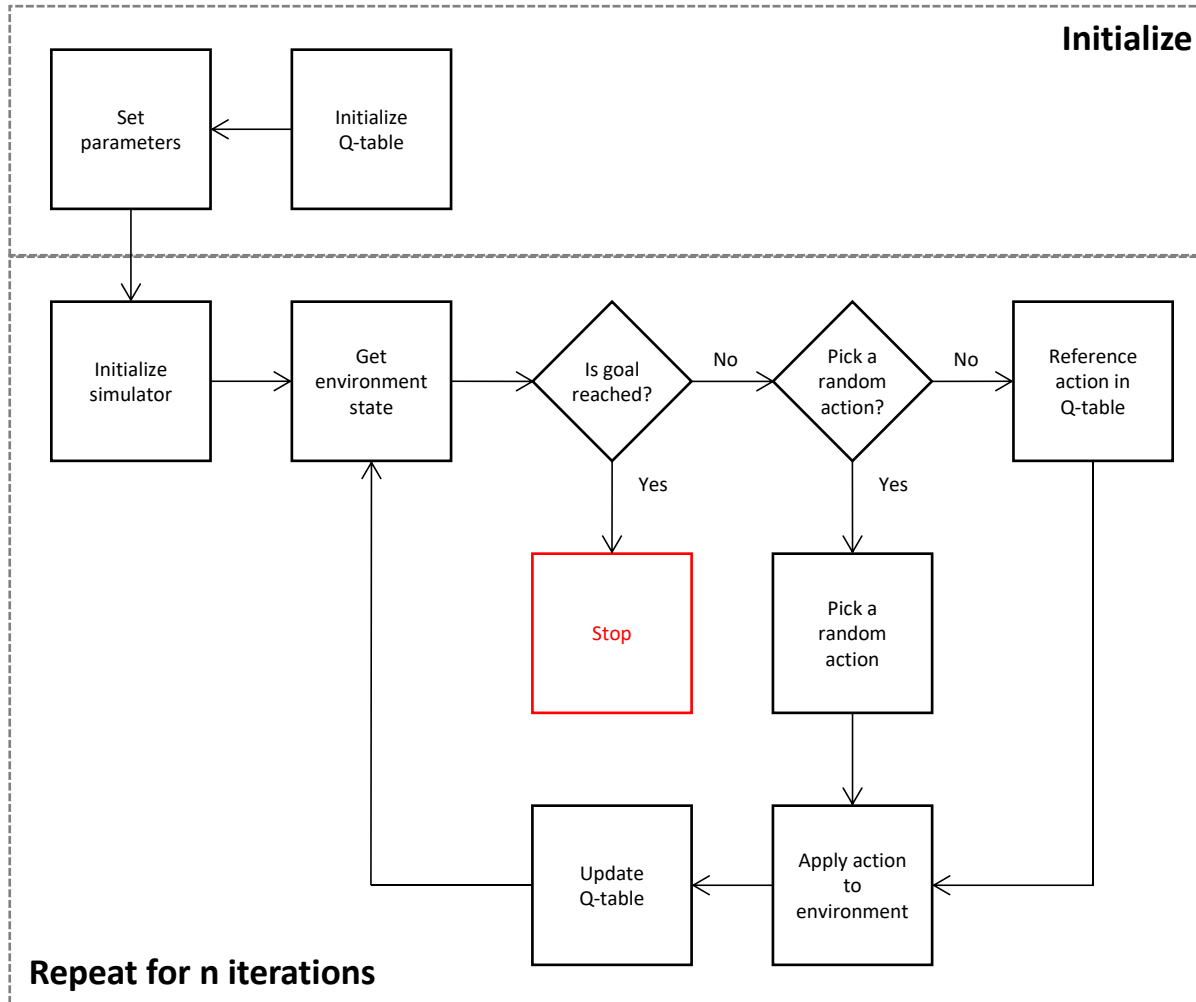
Q-Learning Algorithm



Update Q-table:


Update the Q-table given the reward resulting from recently applied action (feedback from the environment).

Q-Learning Algorithm



Stop:
Stop the learning process

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state →

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

- Move into car: -100
- Move into pedestrian: -1000
- Move into empty space: 100
- Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state →

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state →

Q-Learning Algorithm



Action: →

Reward: -100

Q-table value:

$$Q(1, \text{east}) = (1 - 0.1) * 0 + 0.1 * (-100 + 0.6 * \max \text{ of } Q(2, \text{all actions}))$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:


Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Q-Learning Algorithm



Action: →

Reward: -100

Q-table value:

$$Q(1, \text{east}) = (1 - 0.1) * 0 + 0.1 * (-100 + 0.6 * 0) = -10$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:














































Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Q-Learning Algorithm

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500


Action: →

Reward:   -1000

Q-table value:

$$Q(2, \text{south}) = (1 - 0.1) * 0 + 0.1 * (-1000 + 0.6 * \max \text{ of } Q(3, \text{all actions}))$$

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	-100	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action: →

Reward: -1000

Q-table value:

$$Q(2, \text{south}) = (1 - 0.1) * 0 + 0.1 * (-1000 + 0.6 * 0) = -100$$

Deep Reinforcement Learning

