

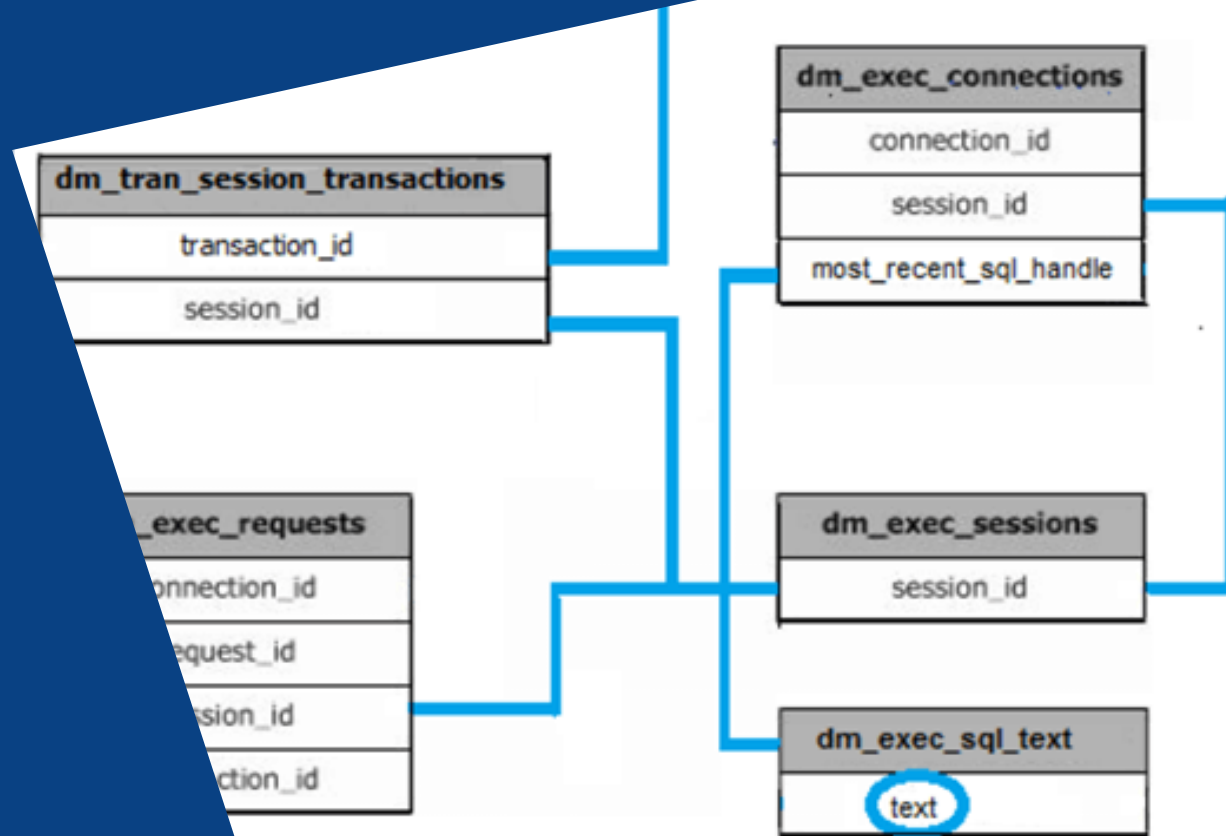


THE UNIVERSITY OF  
MELBOURNE

# Transactions and Concurrency

Database Systems & Information Modelling  
INFO90002

Week 7 – Database transactions  
Dr Tanya Linden  
David Eccles





# This Lecture Objectives

Why we need user-defined transactions

Properties of transactions

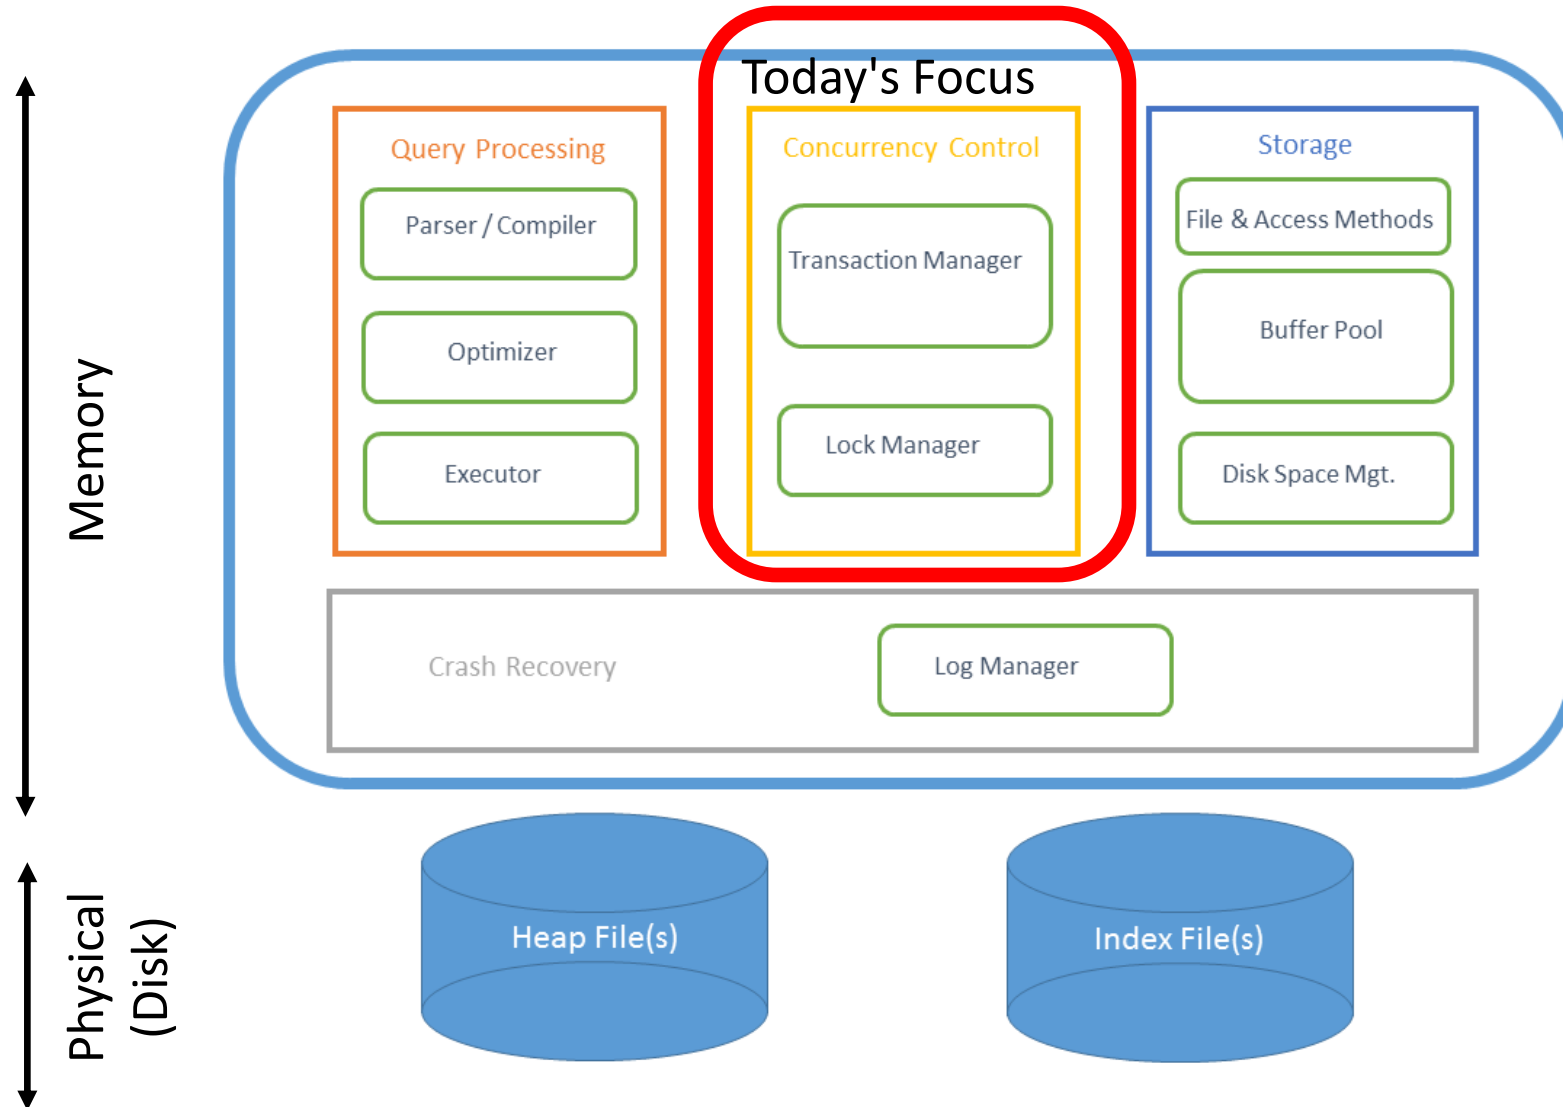
How to use transactions


Concurrent access to data

Locking and deadlocking

Transaction's role in database recovery

# Components of a DBMS





# Transactions

Why we need database transactions

# Business Transactions

Example “business transactions”:

- Placing order: Insert one row in *Order* table, then several in *OrderItem* table
- Withdrawal: Check amount < balance. If so, subtract amount from one row in bank account table, then add amount to another row
- Sending monthly statements: For all rows in Customer table, generate and send out monthly statements

Each requires several distinct database operations ...

- Example: move money from savings account to credit card account
  - Accept inputs from user (via ATM, internet banking or mobile app)
  - Select balance from savings account
  - Is there enough money to withdraw? If so:
    - Update savings account balance = balance – withdrawal
    - Update credit card balance = balance + withdrawal
  - If no errors encountered, end successfully

# Invoice example

## Bill To

John  
Synex Inc  
128 AA Juanita Ave  
Glendora  
CA 91740 US

## Ship To

John  
Synex Inc  
128 AA Juanita Ave  
Glendora  
CA 91740 US

Date	14-Aug-2009	Order No		Sales Person	Charles Wooten
Shipping Date	13-Aug-2009	Shipping Terms		Terms	COD
ID	SKU / Description	Unit Price (USD)	Qty	Amount (USD)	
PS.V880.005	AMD Athlon X2DC-7450, 2.4GHz/1GB/160GB/SMP-DVD/VB	580.00	6.00	3,480.00	
PS.V880.037	PDC-E5300 - 2.6GHz/1GB/320GB/SMP-DVD/FDD/VB	645.00	4.00	2,580.00	
LC.V890.002	LG 18.5" WLCD	230.00	10.00	2,300.00	
HP.Q754.071	HP LaserJet 5200	1,103.00	1.00	1,103.00	

# Normalised Relations and ER Diagram

- We can name the relations now

Customer (CustomerNumber, CustomerName, CustomerAddress)

Clerk (ClerkNumber, ClerkName)

Product (ProductNumber, ProductDescription)

Invoice (InvoiceNumber, Date, *CustomerNumber*, *ClerkNumber*)

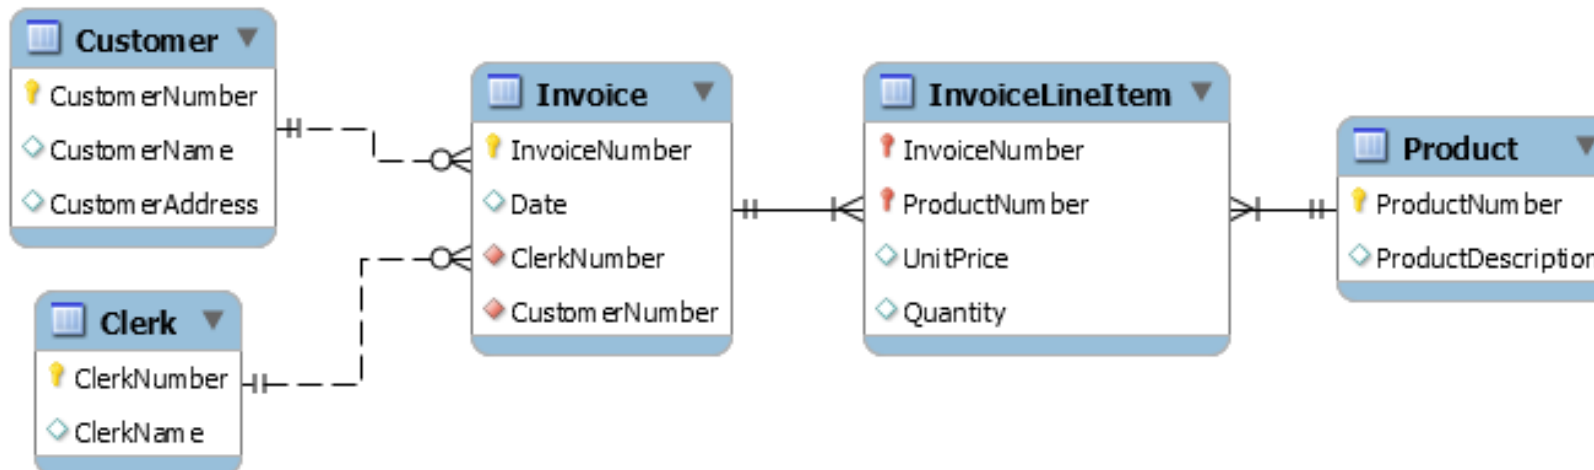
InvoiceLineItem (*InvoiceNumber*, *ProductNumber*, UnitPrice, Quantity)

KEY:

PK = Underline

FK = Italic

PFK = Underline + Italic





# What is a (database) Transaction?

A logical unit of work that must either be entirely completed or aborted (indivisible, atomic)

DML statements are already atomic in MySQL\*, SQL Server\*

RDBMS also allows for *user-defined* transactions

These are a sequence of DML statements, such as

- a series of UPDATE statements to change values
- a series of INSERT statements to add rows to tables
- DELETE statements to remove rows

Transactions will be treated as atomic

A successful transaction changes the database from one consistent state to another

- All data integrity constraints are satisfied

\* Database specific, e.g Oracle DML are not atomic transactions by default





# Why do we need Transactions?

Transactions solve TWO problems:

1. users need the ability to define a unit of work
2. concurrent access to data by >1 user or program

(also acts as an “undo” for manual database manipulation)

# Problem 1: Unit of work

- Single DML or DDL command (implicit transaction)
  - Example: update 700 records, but database crashes after 200 records processed
  - Restart server: you will find no changes to any records
  - Changes are “all or none”
- Multiple statements (user-defined transaction)
  - START TRANSACTION; (or, ‘BEGIN’)
    - SQL statement;
    - SQL statement;
    - SQL statement;
    - ...
  - COMMIT; (commits the whole transaction)
    - Or ROLLBACK (to undo everything)

SQL keywords: **BEGIN; START TRANSACTION; COMMIT, ROLLBACK**



# Business transactions as units of work

Each transaction consists of several SQL statements, embedded within a larger application program

Transaction needs to be an indivisible unit of work

- “**Indivisible**” means that either the whole job gets done, or none gets done:
  - if an error occurs, we don’t leave the database with the job half done, in an inconsistent state

In the case of an error:

- Any SQL statements already completed must be reversed
- Show an error message to the user
- When ready, the user can try the transaction again
- This is briefly annoying – but inconsistent data is disastrous



# Transaction Properties (ACID)

## Atomicity

- A transaction is treated as a single, indivisible, logical unit of work. All operations in a transaction must be completed; if not, then the transaction is aborted

## Consistency

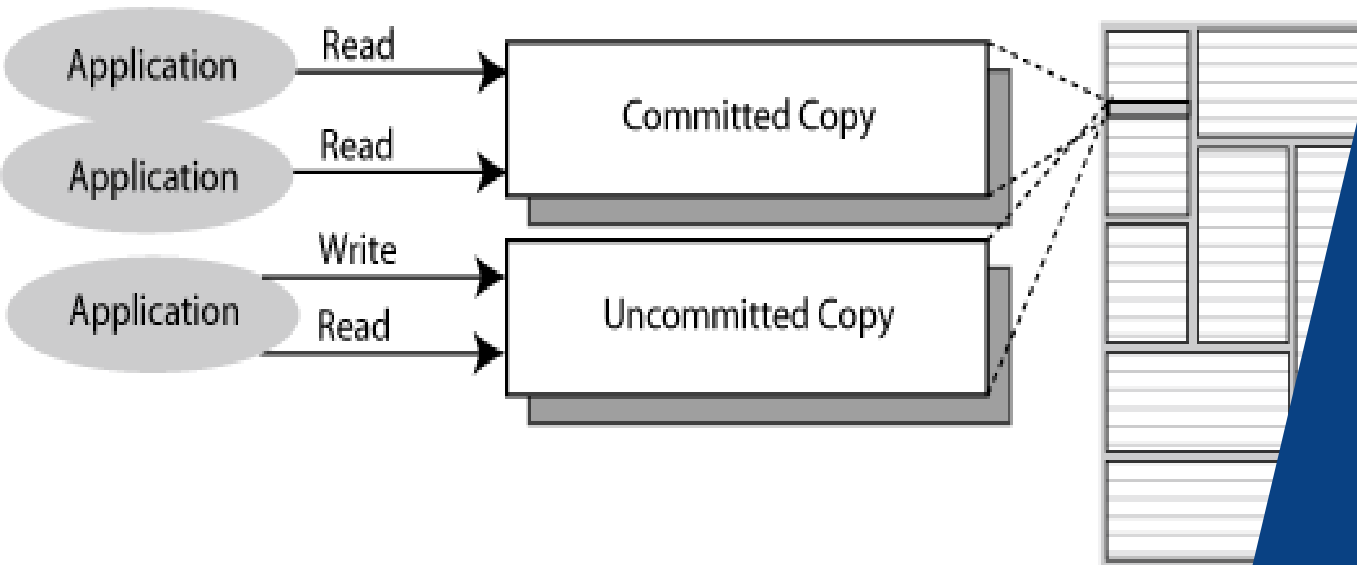
- Constraints that hold before a transaction must also hold after it
- (multiple users accessing the same data see the same value)

## Isolation

- Changes made during execution of a transaction cannot be seen by other transactions until this one is completed

## Durability

- When a transaction is complete, the changes made to the database are permanent, even if the system fails



# Concurrency

What happens if there is more than one user?



## Problem 2: Concurrent access

What happens if we have multiple users accessing the database at the same time...

Concurrent execution of DML against a shared database


Note that the sharing of data among multiple users is where much of the benefit of databases derives – users communicate and collaborate via shared data

But what could go wrong?

- lost updates
- uncommitted data
- inconsistent retrievals

# The Lost Update problem

Alice



Read account  
balance  
(balance = \$1000)

Withdraw \$100  
(balance = \$900)

Write balance  
**balance = \$900**

*t1a*

*t2a*

*t3a*


Time

*t1b*

*t2b*

*t3b*

Bob



Read account  
balance  
(balance = \$1000)

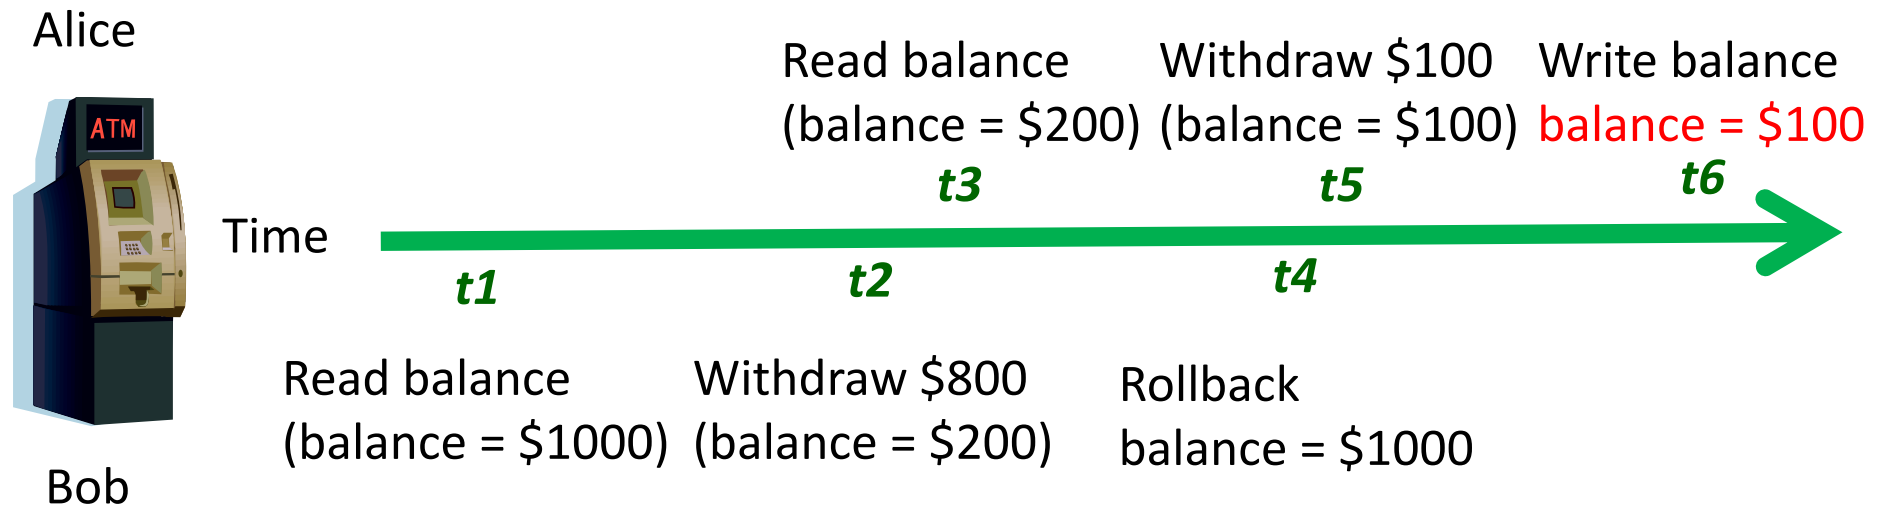
Withdraw \$800  
(balance = \$200)

Write balance  
balance = \$200

**Balance should be \$100**

# The Uncommitted Data problem

Uncommitted data occurs when two transactions execute concurrently and the first is rolled back after the second has already accessed the uncommitted data



**Balance should be \$900**



# The Inconsistent Retrieval problem

Occurs when one transaction calculates some aggregate functions over a set of data, while other transactions are updating the data

- Some data may be read after they are changed and some before they are changed, yielding inconsistent results

Alice	Bob
SELECT SUM(Salary) FROM Employee;	UPDATE Employee SET Salary = Salary * 1.01 WHERE EmpID = 33;
	UPDATE Employee SET Salary = Salary * 1.01 WHERE EmpID = 44;
(finishes calculating sum)	COMMIT;

# Example: Inconsistent Retrieval


Time	Trans- action	Action	Value	T1 SUM	Comment
1	T1	Read Salary for EmpID 11	10,000	10,000	
2	T1	Read Salary for EmpID 22	20,000	30,000	
3	T2	Read Salary for EmpID 33	30,000		
4	T2	Salary = Salary * 1.01			
5	T2	Write Salary for EmpID 33	30,300		
6	T1	Read Salary for EmpID 33	30,300	60,300	<i>after update</i>
7	T1	Read Salary for EmpID 44	40,000	100,300	<i>before update</i>
8	T2	Read Salary for EmpID 44	40,000		
9	T2	Salary = Salary * 1.01			
10	T2	Write Salary for EmpID 44	40,400		
11	T2	COMMIT			
12	T1	Read Salary for EmpID 55	50,000	150,300	
13	T1	Read Salary for EmpID 66	60,000	210,300	

we want either  
*before* \$210,000 or  
*after* \$210,700

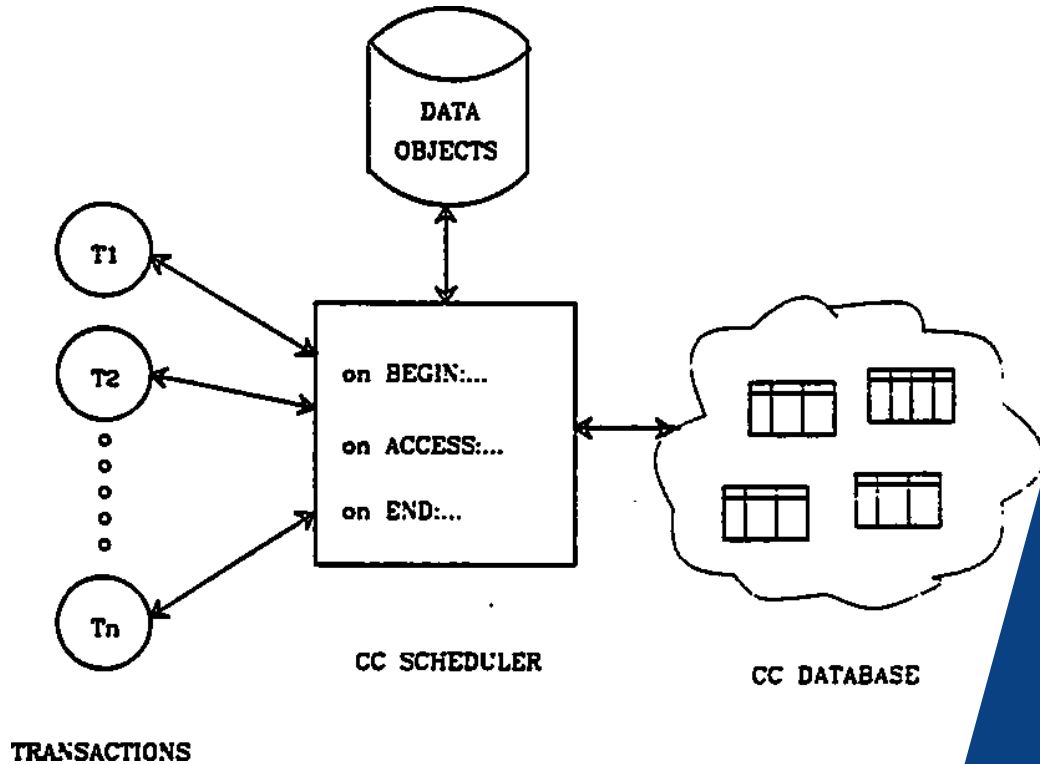
# Serializability

Transactions ideally are “serializable”

- Multiple, concurrent transactions *appear as if* they were executed one after another
- Ensures that the concurrent execution of several transactions yields consistent results



but true serial execution (i.e. no concurrency) is very expensive!



# Concurrency Control Methods

Locking  
Timestamps  
Optimistic Method



# Concurrency control methods

To achieve efficient execution of transactions, the DBMS creates a schedule of read and write operations for concurrent transactions

Interleaves the execution of operations, based on concurrency control algorithms such as locking and time stamping

Several methods of concurrency control

- *Locking* is the main method used
- Time Stamping
- Optimistic Methods



# Concurrency Control with Locking

## Lock

- Guarantees exclusive use of a data item to a current transaction
  - T1 acquires a lock prior to data access; the lock is released when the transaction is complete
  - T2 does not have access to data item currently being used by T1
  - T2 has to wait until T1 releases the lock
- Required to prevent another transaction from reading inconsistent data

## Lock manager

- Responsible for assigning and policing the locks used by the transactions

Question: at what granularity should we apply locks?



# Lock Granularity: options

## Database-level lock

- Entire database is locked
- Good for batch processing but unsuitable for multi-user DBMSs
- T1 and T2 can not access the same database concurrently even if they use different tables
- (SQLite, Access)

## Table-level lock

- Entire table is locked - as above but not quite as bad
- T1 and T2 can access the same database concurrently as long as they use different tables
- Can cause bottlenecks, even if transactions want to access different parts of the table and would not interfere with each other
- Not suitable for highly multi-user DBMSs



# Lock Granularity: options

## Page-level lock

- An entire disk page is locked (a table can span several pages and each page can contain several rows of one or more tables)
- Not commonly used now

## Row-level lock

- Allows concurrent transactions to access different rows of the same table, even if the rows are located on the same page
- Improves data availability but with high overhead (each row has a lock that must be read and written to)
- Currently the most popular approach (MySQL, Oracle)

## Field-level lock

- Allows concurrent transactions to access the same row, as long as they access different attributes within that row
- Most flexible lock but requires an extremely high level of overhead
- Not commonly used





# Types of Locks

## Binary Locks

- has only two states: locked (1) or unlocked (0)
- eliminates “Lost Update” problem
  - the lock is not released until the statement is completed
- considered too restrictive to yield optimal concurrency, as it locks even for two READs (when no update is being done)

The alternative is to allow both *Exclusive* and *Shared* locks

- often called Read and Write locks

Writers never block Readers

- INSERT UPDATE DELETE should never block a SELECT

Readers never block Writers

- SELECT should never block INSERT UPDATE DELETE

Read Consistent; Read Committed are 2 options



# Shared and Exclusive Locks

## Exclusive lock

- Access is reserved for the transaction that locked the object
- Must be used when transaction intends to WRITE
- Granted if and only if no other locks are held on the data item
- In MySQL: “select ... for update”

## Shared lock

- Other transactions are also granted Read access
- Issued when a transaction wants to READ data, and no Exclusive lock is held on that data item
  - multiple transactions can each have a shared lock on the same data item if they are all just reading it
- In MySQL: “select ... lock in share mode”

# Strict Two Phase Locking (Strict 2PL)

## Two Rules

1. If a transaction  $T$  wants to read an object  $O$ , it first requests a *shared* lock on the object and then obtains an *exclusive* lock
2. All locks held by transaction  $T$  are released when the transaction is completed.

In effect the locking protocol only allows 'safe' interleaving of transactions  $T_i$ .

If two transactions are accessing completely different parts of the db they concurrently obtain the locks and proceed.

If two transactions  $T_1, T_2$  are accessing the same object and one wants to modify it, their actions are ordered serially – all actions of one transaction  $T_1$  must complete before the other transaction  $T_2$  can proceed.

# Deadlock


Condition that occurs when two transactions wait for each other to unlock data

- T1 locks data item X, then wants Y
- T2 locks data item Y, then wants X
- Each waits to get a data item which the other transaction is already holding
- Could wait forever if not dealt with

Only happens with **exclusive** locks

Deadlocks are dealt with by:

- prevention
- detection
- (we won't go into the details of how in this course)



# Deadlock demo

- Two separate sessions
  - In order
  - Tx1 Update row 3 (Green)
  - Tx2 Update row 2 (White)
  - Tx3 Update row 2 (Green)
  - Tx4 Update row 3 (White)
- 
- Note: Only the session which detects the deadlock rolls back the transaction. The Green session still holds locks on row 2 and 3

```
deccles2 — mysql -u root -p — 80x24
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> SELECT * FROM DEPT;
+-----+-----+-----+
| deptno | name      | location |
+-----+-----+-----+
| 10     | ACCOUNTING | NEW YORK |
| 20     | RESEARCH  | DALLAS   |
| 30     | SALES     | CHICAGO  |
| 40     | OPERATIONS | BOSTON   |
+-----+-----+-----+
4 rows in set (0.00 sec)

1 mysql> UPDATE DEPT set location = 'SYDNEY' where deptno = 30;
Query OK, 1 row affected (0.08 sec)
Rows matched: 1  Changed: 1  Warnings: 0

3 mysql> UPDATE DEPT set location = 'MELBOURNE' where deptno = 20;
Query OK, 1 row affected (13.07 sec)
Rows matched: 1  Changed: 1  Warnings: 0

mysql>

deccles2 — mysql -u root -p — 71x24
mysql> USE SCOTT;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> SELECT * FROM DEPT;
+-----+-----+-----+
| deptno | name      | location |
+-----+-----+-----+
| 10     | ACCOUNTING | NEW YORK |
| 20     | RESEARCH  | DALLAS   |
| 30     | SALES     | CHICAGO  |
| 40     | OPERATIONS | BOSTON   |
+-----+-----+-----+
4 rows in set (0.00 sec)

2 mysql> UPDATE DEPT set location = 'AUCKLAND' WHERE deptno = 20;
Query OK, 1 row affected (0.00 sec)
Rows matched: 1  Changed: 1  Warnings: 0

4 mysql> Update dept set location = 'WELLINGTON' where deptno = 30;
ERROR 1213 (40001): Deadlock found when trying to get lock; try restart
ing transaction
mysql>
```



# Transaction Isolation Levels

Isolation levels provide the level any given transaction is exposed to the actions of other executing transactions

Isolation levels are *READ UNCOMMITTED*, *READ COMMITTED*, *REPEATABLE READ* and *SERIALIZABLE*

*Serializable* is the highest degree of isolation

- Locks are obtained for both reads (R) and writes (W)
- Transaction *T* only reads committed transactions

*Repeatable Read*

- T1 reads the values established by the first read. This means that if you issue several SELECT statements within the same transaction (T1), these SELECT statements are consistent also with respect to each other – even if the values have been modified by other transacts (T2,T3,...Tn) until a commit is issued.

# Transaction Isolation Levels

## *READ COMMITTED*

- Only reads committed changes by other transactions
- No value written by  $T1$  is changed by another transaction until  $T1$  completes (commits)
- However a value read by  $T1$  can be modified by other transactions  $Tn$  while  $T1$  is in progress.
- Read committed holds exclusive locks (x) before writing objects and shared locks (s) before reading objects.

## *READ UNCOMMITTED*

- Does not obtain shared locks before reading
- Prohibits any transaction from making changes (i.e. write)
- Access mode is *read only*

SET TRANSACTION ISOLATION LEVEL <level>;

# Alternative concurrency control methods

## Timestamp

- Assigns a global unique timestamp to each transaction
- Each data item accessed by the transaction gets the timestamp
- Thus for every data item, the DBMS knows which transaction performed the last read or write on it
- When a transaction wants to read or write, the DBMS compares its timestamp with the timestamps already attached to the item and decides whether to allow access

## Optimistic

- Based on the assumption that the majority of database operations do not conflict
- Transaction is executed without restrictions or checking
- Then when it is ready to commit, the DBMS checks whether any of the data it read has been altered – if so, rollback



# Logging transactions

If transaction cannot be completed, it must be aborted and any changes rolled back

To enable this, DBMS tracks all updates to data


This *transaction log* contains:

- a record for the beginning of the transaction
- for each SQL statement
  - operation being performed (update, delete, insert)
  - objects affected by the transaction
  - “before” and “after” values for updated fields
  - pointers to previous and next transaction log entries
- the ending (COMMIT) of the transaction

# Transaction log

Also provides the ability to restore a corrupted database

If a system failure occurs, the DBMS will examine the log for all uncommitted or incomplete transactions and it will restore the database to its previous state



# Example transaction log

TRL ID	TRX NUM	PREV PTR	NEXT PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	****Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, ...
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, ...
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423	CHECKPOINT								
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, ...
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	****Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



# What's examinable

**Why we need user-defined transactions**

**Properties of transactions**

**ACID**

**How to use transactions**

**BEGIN; START TRANSACTION; COMMIT; ROLLBACK;**

**Concurrent access to data**

**Concurrent access strategies**

**Locking and deadlocking**

**Types of Locks (Binary; Shared)**

**Database recovery**

**Fundamentals of transaction recovery**

**\* All material is examinable – these are the suggested key skills you would need to demonstrate in an exam scenario**



THE UNIVERSITY OF  
MELBOURNE

# Thank you