

COMP9120 Database Management Systems

Tutorial Week 6: Integrity Constraints

This week's exercises are formulated such that they can be carried out in your own time and are important in giving you practice with the Integrity Constraint features of Relational Database Management Systems.

Exercise 1. Data Integrity Constraints

Consider the University relational database schema from Week 5, which is available in Canvas.

If you have not done so already, create this schema by running all the SQL statements in the downloaded file on PostgreSQL.

a) Key Constraints

The schema already contains some appropriate key and **NOT NULL** constraints. All key constraints are named so that you can drop them easily. You can add more constraints with the following syntax:

ALTER TABLE table **ADD CONSTRAINT** name ...

Here table is the name of the table you wish to modify, and name is the name for the constraint. Supplying a name for the constraint is recommended to make it easier to drop them later, as well as making it easier to trace back to which constraint caused a violation.

One example of a business rule we can enforce using a key constraint is that:

“Two units-of-study cannot be taught in the same room at the same time in a given semester;”

ALTER TABLE Lecture **ADD CONSTRAINT** ClassroomConflict

UNIQUE(classroomId,semester,year,classTime);

After adding this uniqueness constraint, check what happens if you try to add a row in the lecture table that violates this constraint.

b) Referential Integrity

The schema already contains some appropriate foreign key constraints. To *change* a constraint you must drop it before adding a replacement one (of the same name):

ALTER TABLE table **DROP CONSTRAINT** name;

Looking through the current schema, you will notice:

1. WhenOffered is missing a foreign key constraint on uosCode which references the UnitOfStudy table. Go ahead and add this using:

ALTER TABLE WhenOffered **ADD CONSTRAINT** WhenOffered_fk_uos **FOREIGN KEY**

(uosCode) **REFERENCES** UnitOfStudy **ON DELETE CASCADE**;

You will notice that there is an **ON DELETE CASCADE** option specified on the foreign key constraint. This means that if a row with a uosCode existing in WhenOffered is deleted from UnitOfStudy then rows in WhenOffered having this uosCode are also deleted. Try deleting a row from UnitOfStudy for a unit of study referenced in WhenOffered.

2. Consider whether the default of '**ON DELETE NO ACTION**' is appropriate for foreign key constraints in the current schema. For instance, if a unit of study has prerequisites in the Requires table, then are we allowed to delete the unit of study from the UnitOfStudy table?

We can modify the existing ON DELETE option as follows:

ALTER TABLE Requires **DROP CONSTRAINT** Requires_uoSCode_fk_UOS;

ALTER TABLE Requires **ADD CONSTRAINT** Requires_uoSCode_fk_UOS **FOREIGN KEY** (uosCode) **REFERENCES** UnitOfStudy(uoSCode) **ON DELETE CASCADE**;

c) Check Constraints

You can add check constraints with the following syntax:

ALTER TABLE table **ADD CONSTRAINT** name **CHECK** (condition)

Here condition is the condition that should hold true for each tuple in the table.

Add check constraints for the following:

1. The system shall contain the transcript records of each student that show each unit of study the student has completed, the semester the student took the course, and the grade the student received (all grades are in the set {'F','P','CR','D','HD','W'});

ALTER TABLE Transcript **ADD CONSTRAINT** valid_grades **CHECK**(grade IN ('F','P','CR','D','HD','W'));

What happens if you try and modify a grade to 'M'?

2. The system shall contain information about the unit-of-studies offered, and for each unit of study the system shall contain whether the unit is offered in semester S1, or in semester S2, or in winter semester (WS) or summer semester (SS);

ALTER TABLE WhenOffered **ADD CONSTRAINT** valid_semesters **CHECK**(semester IN ('S1','S2','WS','SS'));

What happens if you try to modify a semester to 'S3'?

d) Assertions

Assertions are not implemented in most mainstream RDBMS, but you should be able to express them with the standard syntax. Below are some examples for how to write assertions to enforce some business rules:

1. The number of students enrolled in a unit-of-study must equal the current enrolment;

```
CREATE ASSERTION EnrollmentAssert CHECK (
  NOT EXISTS (
    SELECT 1
    FROM UoSOffering o
    WHERE enrollment != (SELECT COUNT(*)
      FROM Transcript t
      WHERE t.uosCode=o.uosCode AND
      t.semester=o.semester AND
      t.year=o.year)
    )
);
```

Try executing the underlined code followed by a semicolon. Will this assertion be violated?

2. The room assigned to a unit-of-study must have at least as many seats as the maximum allowed enrolment for the unit;

```
CREATE ASSERTION RoomCapacityAssert CHECK (
  NOT EXISTS (
    SELECT 1
    FROM (UoSOffering NATURAL JOIN Lecture)
      NATURAL JOIN Classroom
    WHERE seats < maxEnrollment
    )
);
```

Try executing the underlined code followed by a semicolon. Will this assertion be violated?

3. A student cannot be registered for more than 24 credit points in a given semester;

```
CREATE ASSERTION MaxCreditsAssert CHECK (
  NOT EXISTS (
    SELECT 1
    FROM Transcript NATURAL JOIN UnitOfStudy
    GROUP BY studId, semester, year
    HAVING SUM(credits) > 24
    )
);
```

Try executing the underlined code followed by a semicolon. Will this assertion be violated?

e) Triggers

Some of the behaviour intended from assertions (which are not implemented by most mainstream RDBMS) can be achieved through triggers. Check documentation at

<https://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

1. We can write a trigger to update the enrolment number for a unit of study offering *when a student is added* (in Transcript), as shown below:

```
CREATE OR REPLACE FUNCTION updateEnrol() RETURNS trigger AS $$
BEGIN
    UPDATE UoSOffering U
    SET enrollment = enrollment+1
    WHERE U.uosCOde = NEW.uosCode AND semester = NEW.semester
           AND year = NEW.year;
    RETURN NEW;
END; $$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER UpdateEnrol AFTER INSERT ON Transcript
FOR EACH ROW EXECUTE PROCEDURE updateEnrol();
```

2. This trigger will not guarantee that the first assertion in Exercise 1d will always hold true. What other triggers are needed?

Consider:

- What happens if enrolment exceeds maxenrolment? (should update be attempted before insert?)
- What happens if Transcript is later updated to change the course enrolled?
- What happens if Transcript entry is deleted?
- What happens if UoSOffering.maxenrolment is updated?

Triggers should be added to deal with such situations.