

# COMP9414: Artificial Intelligence

## Lecture 9b: Reinforcement Learning

Wayne Wobcke

e-mail: w.wobcke@unsw.edu.au

### This Lecture

- Reinforcement Learning vs Supervised Learning
- Models of Optimality
- Exploration vs Exploitation
- Temporal Difference Learning
- Q-Learning

### Types of Learning

- **Supervised Learning**
  - ▶ Agent is presented with examples of inputs and their target outputs, and must learn a function from inputs to outputs that agrees with the training examples and generalizes to new examples
- **Reinforcement Learning**
  - ▶ Agent is not presented with target outputs for each input, but is periodically given a reward, and must learn to maximize (expected) rewards over time
- **Unsupervised Learning**
  - ▶ Agent is only presented with a series of inputs, and must find useful patterns in these inputs

### Supervised Learning

- Given a **training set** and a **test set**, each consisting of a set of items for each item in the training set, a set of features and a target output
- Learner must learn a **model** that can **predict** the target output for **any** given item (characterized by its set of features)
- Learner is given the input features and target output for each item in the training set
  - ▶ Items may be presented all at once (batch) or in sequence (online)
  - ▶ Items may be presented at random or in time order (stream)
  - ▶ Learner **cannot** use the test set **at all** in defining the model
- Model is evaluated by its performance on predicting the output for each item in the **test set**

## Learning Actions

Supervised learning can be used to learn actions from a training set of situation-action pairs (called Behavioural Cloning)

However, there are many applications for which it is difficult, inappropriate, or even impossible to provide a “training set”

- Optimal control
  - ▶ Mobile robots, pole balancing, flying a helicopter
- Resource allocation
  - ▶ Job shop scheduling, mobile phone channel allocation
- Mix of allocation and control
  - ▶ Elevator control, Backgammon

## Models of Optimality

Is a fast nickel worth a slow dime?

Finite horizon reward  $\sum_{i=0}^h r_{t+i}$

Average reward  $\lim_{h \rightarrow \infty} \frac{1}{h} \sum_{i=0}^{h-1} r_{t+i}$


Infinite discounted reward  $\sum_{i=0}^{\infty} \gamma^i r_{t+i}, \quad 0 \leq \gamma < 1$

- Finite horizon reward is simple computationally
- Infinite discounted reward is easier for proving theorems
- Average reward is hard to deal with, because can't sensibly choose between small reward soon and large reward very far in the future

## Reinforcement Learning Framework

- Agent interacts with its environment
- There is a set  $S$  of *states* and a set  $A$  of *actions*
- At each time step  $t$ , the agent is in some state  $s_t$  and must choose an action  $a_t$ , whereupon it goes into state  $s_{t+1} = \delta(s_t, a_t)$  and receives reward  $r(s_t, a_t)$
- In general,  $r()$  and  $\delta()$  can be multi-valued, with a random element
- The aim is to find an optimal *policy*  $\pi : S \rightarrow A$  which maximizes the **cumulative** reward

## Comparing Models of Optimality



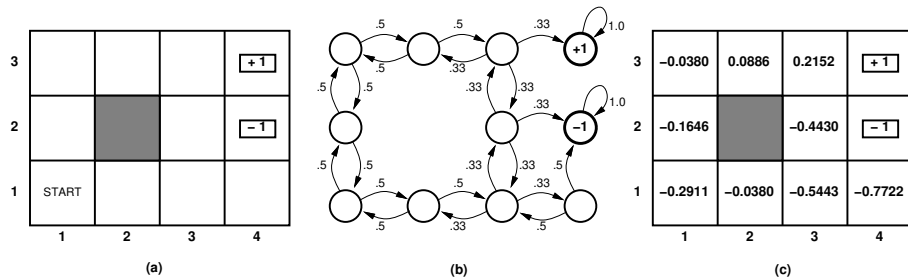
## Environment Types

Environments can be

- Passive and stochastic
- Active and deterministic (Chess)
- Active and stochastic (Backgammon)


## Value Function

For each state  $s \in S$ , let  $V^*(s)$  be the maximum discounted reward obtainable from  $s$



The optimal value function determines the optimal policy

## Example: Delayed Rewards



## Calculation

**Theorem:** In a deterministic environment, for an optimal policy, the value function  $V^*$  satisfies the Bellman equations:  $V^*(s) = r(s, a) + \gamma V^*(\delta(s, a))$  where  $a = \pi^*(s)$  is the optimal action at state  $s$ .

Let  $\delta^*(s)$  be the transition function for  $\pi^*(s)$  and suppose  $\gamma = 0.9$

1. Suppose  $\delta^*(s_1) = s_1$ . Then  $V^*(s_1) = 5 + 0.9V^*(s_1)$  so  $V^*(s_1) = 50$   
 Suppose  $\delta^*(s_2) = s_2$ . Then  $V^*(s_2) = 10 + 0.9V^*(s_2)$  so  $V^*(s_2) = 100$
2. Suppose  $\delta^*(s_1) = s_2$ . Then  $V^*(s_1) = 2 + 0.9V^*(s_2)$  so  $V^*(s_1) = 92$   
 Suppose  $\delta^*(s_2) = s_2$ . Then  $V^*(s_2) = 10 + 0.9V^*(s_2)$  so  $V^*(s_2) = 100$
3. Suppose  $\delta^*(s_1) = s_2$ . Then  $V^*(s_1) = 2 + 0.9V^*(s_2)$  so  $V^*(s_1) = 81.6$   
 Suppose  $\delta^*(s_2) = s_1$ . Then  $V^*(s_2) = 15 + 0.9V^*(s_1)$  so  $V^*(s_2) = 88.4$

So 2 is the optimal policy

## Exploration/Exploitation Tradeoff

Most of the time, the agent should choose the “best” action

However, in order to ensure the optimal strategy can be **learned**, the agent must occasionally choose a different action, e.g.

- Choose a random action 5% of the time, or
- Use a Boltzmann distribution to choose the next action

$$P(a) = \frac{e^{\hat{V}(a)/T}}{\sum_{b \in A} e^{\hat{V}(b)/T}}$$

## K-Armed Bandit Problem



The special case of an active stochastic environment with only one state is called a **K-Armed Bandit Problem**, because it is like being in a room with several (friendly) slot machines, for a limited time, and trying to collect as much money as possible

Each **action** (slot machine) provides a different average reward

## Temporal Difference Learning

TD(0) [also called AHC, or Widrow-Hoff Rule]

$$\hat{V}(s) \leftarrow \hat{V}(s) + \eta [r(s, a) + \gamma \hat{V}(\delta(s, a)) - \hat{V}(s)]$$

( $\eta$  = learning rate)

The (discounted) value of the next state, plus the immediate reward, is used as the target value for the current state

A more sophisticated version, called TD( $\lambda$ ), uses a weighted average of future states

## Q-Learning

For each  $s \in S$ , let  $V^*(s)$  be the maximum discounted reward obtainable from  $s$ , and let  $Q(s, a)$  be the discounted reward available by first doing action  $a$  and then acting optimally

Then the optimal policy is

$$\pi^*(s) = \arg \max_a Q(s, a)$$

where

$$Q(s, a) = r(s, a) + \gamma V^*(\delta(s, a))$$

Then

$$V^*(s) = \max_a Q(s, a)$$

so

$$Q(s, a) = r(s, a) + \gamma \max_b Q(\delta(s, a), b)$$

This allows iterative approximation of  $Q$  by

$$\hat{Q}(s, a) \leftarrow r(s, a) + \gamma \max_b \hat{Q}(\delta(s, a), b)$$

## Theoretical Results

---

**Theorem:** Q-learning will eventually converge to the optimal policy, for any deterministic Markov Decision Process, assuming an appropriately randomized strategy.

(Watkins & Dayan 1992)

**Theorem:** TD-learning will also converge, with probability 1.

(Sutton 1988, Dayan 1992, Dayan & Sejnowski 1994)

## Summary

---

- Reinforcement Learning is an active area of research
- Mathematical results (more than in other areas of AI)
- Need to have an appropriate representation
- Future algorithms which choose their own representations?
- Many practical applications

## Limitations of Theoretical Results

---

- Delayed reinforcement
  - ▶ Reward resulting from an action may not be received until several time steps later, which also slows down the learning
- Search space must be finite
  - ▶ Convergence is slow if the search space is large
  - ▶ Relies on visiting every state infinitely often
- For “real world” problems, can’t rely on a lookup table
  - ▶ Need to have some kind of [generalization](#) (e.g. TD-Gammon)